

CERTIFIED BELLMAN OPERATORS FOR COSTLY POLICY REVISION

QIAN QI
Peking University

We study the minimum implementation cost of revising an installed policy while preserving operating performance after every state and date. A common randomized Markov continuation must serve all restarts. We construct a backward feasibility repair with an explicit cost modulus. A witness-driven form avoids an exact operating solve. It yields convergent global policy search for finite-state problems and an end-to-end approximation bound for continuous-state problems admitting certified total-variation aggregation. Exact rational spatial search and a verified McCormick comparator target the same constrained objective. A prospectively fixed eight-environment experiment reports all budget stops. On the original 42 continuous-state development configurations, whole-domain lotteries certify strict improvement over every deterministic policy in 27 cases. Nonlinear two-state lower bounds become positive, but longer-horizon gaps remain material. Independent checks distinguish arithmetic validity from approximation assumptions. The results identify when local operating guarantees can be converted into global revision certificates without changing the policy class.

KEYWORDS: Dynamic programming, Policy revision, Constrained Markov control, Global optimization, Verified computation.

1. INTRODUCTION

An organization may wish to improve an installed rule without replacing every decision it makes. Implementation cost then depends both on the decisions changed today and on the states visited under the revised rule tomorrow. A feasible revision must preserve the operating objective, whereas a cost-effective revision must economize on intervention along its own induced state distribution. These two requirements create a constrained dynamic problem. Pointwise selection of inexpensive changes need not solve it.

The relevant continuation must also be specified. Two histories that meet at the same date and state use the same continuation under a Markov rule. Allowing a different continuation for each initial condition or each remembered randomization can reduce costs, but changes the economic object. Conversely, requiring a deterministic rule can increase implementation cost even when randomization is admissible. We therefore take a common randomized Markov rule as the central choice variable. Performance is protected after every date and state; implementation cost is integrated under a stated initial distribution. A performance guarantee at every restart is not a claim of cost optimality at every restart.

The main contribution is a constructive connection between these requirements and a complete global search. A policy that violates every operating constraint by at most δ can be repaired backward without adding a persistent random seed or a history variable. The repair improves operating value at every state. Its additional implementation cost is bounded explicitly by a modulus that vanishes with δ . Strict discounting and a positive operating allowance provide the slack; no global duality assumption is used. This repair gives an error bound for rational policy lattices. For binary actions it also gives a stopping theorem for the interval Bellman branch-and-bound procedure used in the experiments. Unlike convergence to a necessary-action relaxation, the limiting object is the original common-continuation constrained optimum.

The witness-driven extension separates operating approximation error from policy-search error in one computable global bound. It supplies a positive-margin check and a constructive

policy without passing the exact optimum to the repair routine. A further perturbation result converts model error into operating and cost allowances; strict randomized separations also yield break-even administration charges. These statements extend the numerical method rather than changing its economic objective.

A second result transfers finite-state certificates to a continuous state space. Under a certified total-variation approximation of the transition kernel by a within-cell reset kernel, cell averaging maps an arbitrary measurable Markov policy to a finite-state policy. Lifting reverses this construction. Uniform primitive errors give an explicit operating-tolerance adjustment and a cost adjustment. The repair modulus controls the remaining difference between the relaxed and tightened finite-state optima. The theorem makes the approximation assumption operational and visible. It does not treat weak convergence of an atomic transition law as total-variation convergence, and it does not silently replace discrete shocks by continuous noise.

The numerical analysis separates three questions. First, does complete search return sound intervals for the common Markov problem, and how do those intervals compare with a direct global mathematical-programming formulation? Second, does randomization change the conclusion in the original continuous-state maintenance configurations? Third, can a nontrivial revision lower bound survive nonlinear, two-state dynamics and nonzero operating-witness error? Each experiment retains the same objective on the two sides of its interval. The eight new finite-state environments and both solver implementations were fixed before those environments were generated. The original 42 configurations are explicitly a development cohort, not a prospective test.

The resulting distinctions are economically consequential. In the development cohort, 27 whole-domain randomized deployments cost strictly less than a certified lower bound for every feasible deterministic policy under the same uniform initial law. This is a statement about the value of admissible randomization, not about a neural label attached to the installed rule. The nonlinear experiment changes the zero lower bound into a positive bound; its finest horizon-32 interval is $[8.8718, 11.8492]$. That interval is informative but not tight. In the prospective finite-state comparison, both methods face the same node allowance and wall-clock limit. The stronger linear relaxation often improves the lower endpoint, but longer horizons retain substantial residual intervals. The theory supplies completeness, not a claim that complete search is inexpensive.

1.1. *Relation to the literature*

Constrained Markov decision processes provide the natural starting point (Altman, 1999). Occupation measures are particularly effective when the constraints are expectations under a designated initial law. Our formulation additionally requires every restarted problem to meet its operating guarantee using one shared continuation. A separate occupation measure for each restart would need compatibility restrictions: at a shared date and state, its action probabilities cannot depend on which restart generated the occupancy. Dropping that restriction produces a different feasible class. Our finite formulation imposes it directly through a single policy variable for each date and state. This observation does not preclude occupation-measure methods; it specifies the coupling they must preserve in this problem.

Finite approximations to constrained control are not new. Saldi (2018) proves value convergence for discounted and average-cost constrained problems, and develops approximately optimal policies and rates under additional continuity conditions. Saldi, Linder, and Yüksel (2018) provides a systematic treatment of quantized stochastic control. Our transfer result concerns finite-horizon, all-restart constraints and combines the within-cell reset identity with an explicit Markov feasibility repair. It states the total-variation condition needed to average unrestricted measurable policies rather than assuming that a convenient policy sieve is automatically

dense in the constrained optimum. The theorem is not a replacement for those approximation theories; its additional object is the common-continuation restart constraint and its computable tolerance-to-cost modulus.

Primal–dual policy methods also solve constrained expected-return problems. Ding, Zhang, Basar, and Jovanovic (2020) establishes global convergence guarantees for natural policy-gradient methods in discounted constrained MDPs. Ni and Kamgarpour (2025) studies safe exploration under cumulative constraints. Those learning problems have information and feasibility requirements different from our known-model, all-restart certification problem. Safe baseline improvement (Laroche, Trichelair, and Tachet des Combes, 2019) and approximation-error analysis for fitted dynamic programming (Munos and Szepesvári, 2008, Scherrer, 2013) further distinguish model uncertainty, policy approximation, and optimization error. We claim no sample-efficiency comparison against methods given less information.

The global optimization ingredients are classical. Convex envelopes of factorable products originate in McCormick (1976); the numerical LP proposals here use a modern simplex implementation associated with Huangfu and Hall (2018). We do not attribute novelty to bilinear envelopes, best-bound search, or weak duality. The contribution is a common-policy economic formulation whose local approximation errors can be repaired into an admissible deployment, with an explicit global stopping bound. We use the direct McCormick formulation as a serious same-object comparator, not as a deliberately weak price portfolio. A rational residual correction turns proposed LP multipliers into lower bounds even when floating-point dual feasibility is imperfect. This separation of numerical proposal and proof follows the broader program of verified computation (Rump, 2010).

Semi-infinite optimization emphasizes the difference between checking finitely many constraints and controlling them on a region (Hettich and Kortanek, 1993). Our continuous-state certificates therefore include whole-cell inequalities and isolated switching points. Information-relaxation bounds supply another route to dynamic-programming lower and upper bounds (Brown, Smith, and Sun, 2010, Brown and Smith, 2014); they are related to, but not identical with, the restart support functions used below. The support functions remain useful accelerators even though no finite price set is asserted to exhaust a global dual.

The economic motivation is adjustment rather than unconstrained operating optimization. Numerical economics has long emphasized model structure, approximation error, and state-space complexity (Judd, 1998, Brumm and Scheidegger, 2017). Empirical replacement models such as Rust (1987) discipline operating primitives with data. Our revision cost is instead the cost of departing from an installed decision rule, not a cost paid whenever the previous physical action changes. Adding physical switching costs would require including the previous action in the state. The present maintenance environments are designed, known-model economies. Their normalized operating and implementation units are not estimated monetary welfare, and their certificates are not estimates of empirical prevalence. The general convergence results, rather than a claimed calibration, carry the methodological argument.

2. THE ECONOMIC PROBLEM

Let X be a standard Borel space, $A = \{1, \dots, m\}$ a finite action set, and $t = 0, \dots, T$. Rewards $r_t(x, a)$, implementation costs $k_t(x, a)$, and terminal operating payoff $g(x)$ are bounded Borel functions. The transition kernels P_t^a are known Borel probability kernels. Write $|r_t| \leq R$, $|g| \leq G$, and $0 \leq k_t \leq K$. Unless otherwise indicated, $0 < \beta < 1$ and $\varepsilon > 0$. Define

$$T_t^a v(x) = r_t(x, a) + \beta P_t^a v(x), \quad V_T = g, \quad V_t = \max_{a \in A} T_t^a V_{t+1}. \quad (1)$$

Finiteness of A makes V_t Borel and permits a Borel maximizing selector, for example the lowest indexed maximizing action.

A randomized Markov policy is a sequence of Borel kernels $p_t(\cdot | x)$ on A . Its values satisfy

$$J_T^p = g, \quad J_t^p = \sum_a p_t(a | x) T_t^a J_{t+1}^p, \quad (2)$$

$$C_T^p = 0, \quad C_t^p = \sum_a p_t(a | x) \{k_t(x, a) + \beta P_t^a C_{t+1}^p(x)\}. \quad (3)$$

The same p_t is used by all histories reaching (t, x) . Randomization is newly drawn conditional on the current date and state; a persistent hidden coin indexing different future policies is not introduced.

For an initial probability law ν , the central object is

$$K_R(\nu, \varepsilon) = \inf_{p \in \mathcal{F}_R(\varepsilon)} \int C_0^p d\nu, \quad \mathcal{F}_R(\varepsilon) = \{p : J_t^p(x) \geq V_t(x) - \varepsilon \forall t, x\}. \quad (4)$$

A deterministic policy restricts p_t to vertices of the simplex and defines K_D . An operating-optimal policy is feasible, so both classes are nonempty. On a finite state space, the product of policy simplexes is compact, J^p and C^p are continuous, and the feasible set is closed. The infimum in (4) is then attained. We use infimum, not an unproved function-space minimum, in the Borel-state formulation.

For an installed rule $a_t^0(x)$, the principal cost specification is

$$k_t(x, a) = w_t(x) \mathbf{1}\{a \neq a_t^0(x)\}, \quad w_t \geq 0. \quad (5)$$

The cost is charged for the action actually drawn. Its expectation under a lottery is not an indicator that the lottery's most probable action differs from the installed action. The unchanged installed policy has zero revision cost even when it is operating-infeasible. Consequently, dividing an optimality interval by its revision cost is undefined, not an informative relative-gap statistic. We report $(U - L)/U$ when $U > 0$, and report zero-cost exactness separately.

The uniform law in the development economy weights all normalized initial conditions equally. It is a numerical reference distribution, not an estimated stationary law. We also integrate exactly or outwardly under densities $2x$ and $2(1 - x)$. These calculations change the initial objective while leaving the all-state operating constraints and installed rule unchanged. A pointwise cost certificate would cover arbitrary initial laws, including atoms; an almost-everywhere cost enclosure need not.

3. FEASIBILITY REPAIR AND COMPLETE POLICY SEARCH

Set

$$D_C = K \sum_{s=0}^{T-1} (s+1) \beta^s, \quad (6)$$

$$D_J = \max_{1 \leq h \leq T} \left\{ 2R \sum_{s=0}^{h-1} (s+1) \beta^s + 2Gh\beta^h \right\}. \quad (7)$$

The maximum over remaining horizons in D_J matters: the terminal term $h\beta^h$ need not increase with h . Total variation between action distributions is one half of their ℓ_1 distance.

LEMMA 1—Policy perturbations: *If $\sup_{t,x} \text{TV}(p_t(\cdot | x), q_t(\cdot | x)) \leq \theta$, then*

$$\sup_{t,x} |J_t^p - J_t^q| \leq D_J \theta, \quad \sup_{t,x} |C_t^p - C_t^q| \leq D_C \theta.$$

PROOF: For costs, an action's remaining cost lies in $[0, K \sum_{s=0}^{h-1} \beta^s]$ at a state with h periods remaining. Changing its distribution changes the conditional expectation by at most the length of this interval times θ . Adding β times the next-date value difference and iterating gives $K \sum_{s=0}^{h-1} (s+1) \beta^s \theta$. For operating rewards use the interval length $2R$ for each stage and $2G$ for the terminal reward. A reward at distance s can be affected by $s+1$ action changes, and the terminal reward by h changes. Taking the maximum over h gives (7). The same conclusions follow by a finite-horizon coupling, but do not require one. *Q.E.D.*

THEOREM 1—Common-continuation repair: *Suppose $p \in \mathcal{F}_R(\varepsilon + \delta)$, $\delta \geq 0$. There is a Borel randomized Markov policy $\hat{p} \in \mathcal{F}_R(\varepsilon)$ with $J_t^{\hat{p}} \geq J_t^p$ for every (t, x) and*

$$\sup_{t,x} \text{TV}(\hat{p}_t, p_t) \leq \rho_\varepsilon(\delta) := \frac{\delta}{(1-\beta)\varepsilon + \delta}, \quad \int C_0^{\hat{p}} d\nu \leq \int C_0^p d\nu + D_C \rho_\varepsilon(\delta). \quad (8)$$

In particular,

$$0 \leq K_R(\nu, \varepsilon) - K_R(\nu, \varepsilon + \delta) \leq D_C \rho_\varepsilon(\delta). \quad (9)$$

PROOF: Proceed backward, with $J_T^{\hat{p}} = g$. Given a repaired continuation, put $q_a = T_t^a J_{t+1}^{\hat{p}}$, $q_p = \sum_a p_t(a | x) q_a$, $q_* = \max_a q_a$, and $b = V_t - \varepsilon$. By the induction hypothesis, $q_p \geq J_t^p \geq V_t - \varepsilon - \delta$. Also $J_{t+1}^{\hat{p}} \geq V_{t+1} - \varepsilon$ implies $q_* \geq V_t - \beta\varepsilon$. If $q_p \geq b$, retain p_t . Otherwise choose a measurable maximizing action a_* and set

$$\hat{p}_t = (1 - \alpha)p_t + \alpha e_{a_*}, \quad \alpha = \frac{b - q_p}{q_* - q_p}.$$

Writing $d = b - q_p \in (0, \delta]$, the denominator is at least $(1 - \beta)\varepsilon + d$. Thus $\alpha \leq \rho_\varepsilon(\delta)$ and $J_t^{\hat{p}} = b \geq q_p \geq J_t^p$. All operations are Borel; the division is used only where the denominator is strictly positive. The construction uses only t and x and the already fixed next-date continuation. Lemma 1 bounds cost. Apply the construction to arbitrarily near-minimizers at tolerance $\varepsilon + \delta$ and take their infimum to obtain (9). *Q.E.D.*

The repair is not a convex combination of whole policies. Such a combination, if selected once and remembered, would enlarge the information state. Here the mixture is selected anew at each state using one repaired continuation. The slack $(1 - \beta)\varepsilon$ explains why the theorem assumes strict discounting and a positive tolerance. No vanishing modulus at $\beta = 1$ or $\varepsilon = 0$ is asserted by this argument.

3.1. Certified search without an exact operating solve

The exact value in the preceding statements is not an oracle that the method must possess. A verified upper witness supplies the slack needed by the repair. This matters when operating approximation, rather than policy search alone, is expensive. Let $L_t \leq V_t \leq U_t$ be bounded

Borel witnesses with $U_T = g$. Define checked scalar bounds

$$\begin{aligned} e &\geq \max_{t < T} \|U_t - L_t\|_\infty, & \sigma &= (1 - \beta)\varepsilon - d, \\ d &\geq \max \left\{ 0, \sup_{t < T, x} \left[U_t(x) - \max_a T_t^a U_{t+1}(x) \right] \right\}. \end{aligned} \quad (10)$$

For the finite-state statement, let $m = |A|$, let $M \geq 1$ be an integer, and set $\theta_M = \min\{1, (m - 1)/M\}$. Lattice probabilities are multiples of $1/M$. The same primitive transition and reward operators occur in the witness check and the repair. Error in a different operator must first be included in its enclosure. A small residual is not, by itself, a certificate that $U \geq V$; the witness inequalities and terminal condition are also required.

THEOREM 2—Witness-driven repair and global error: *Suppose $\sigma > 0$ and $\delta \geq 0$. If a common randomized Markov policy satisfies $J_t^p \geq U_t - \varepsilon - \delta$ for all $t < T$, the backward construction with target $U_t - \varepsilon$ produces a common policy $\hat{p}$ such that*

$$J_t^{\hat{p}} \geq U_t - \varepsilon \geq V_t - \varepsilon, \quad J_t^{\hat{p}} \geq J_t^p, \quad \sup_{t, x} \text{TV}(\hat{p}_t, p_t) \leq \frac{\delta}{\sigma + \delta}. \quad (11)$$

Its additional initial implementation cost is at most $D_C \delta / (\sigma + \delta)$. On a finite state space, let $v_{M,L}$ be the smallest cost among rational-lattice policies satisfying $J_t^p \geq L_t - \varepsilon - D_J \theta_M$. Repair a minimizer using U , and let its cost be $U_{M,L,U}$. Then

$$\max\{0, v_{M,L} - D_C \theta_M\} \leq K_R(\nu, \varepsilon) \leq U_{M,L,U}, \quad (12)$$

$$U_{M,L,U} - \max\{0, v_{M,L} - D_C \theta_M\} \leq D_C \left\{ \theta_M + \frac{e + D_J \theta_M}{\sigma + e + D_J \theta_M} \right\}. \quad (13)$$

Thus $M \rightarrow \infty$, $e \rightarrow 0$, and $d \rightarrow 0$ give convergence to the actual constrained optimum using checked witnesses rather than an exact operating solve. All displayed errors and the stopping test are computable for rational represented finite-state primitives.

PROOF: Fix a repaired continuation and put $q_a = T_t^a J_{t+1}^{\hat{p}}$, $q_p = \sum_a p_t(a | x) q_a$, and $b = U_t - \varepsilon$. At the final decision date $J_T^{\hat{p}} = U_T$; at every earlier date the induction hypothesis gives $J_{t+1}^{\hat{p}} \geq U_{t+1} - \varepsilon$. Positivity and (10) imply

$$q_* := \max_a q_a \geq \max_a T_t^a U_{t+1} - \beta \varepsilon \geq U_t - d - \beta \varepsilon = b + \sigma.$$

The inductive improvement gives $q_p \geq J_t^p \geq b - \delta$. Where $q_p < b$, set $\alpha = (b - q_p)/(q_* - q_p)$ and mix p_t with the least-index maximizing action. For $z = b - q_p \in (0, \delta]$, one has $\alpha \leq z/(\sigma + z) \leq \delta/(\sigma + \delta)$. Otherwise leave p_t unchanged. This proves (11); the cost bound follows from Lemma 1. Finite action maximization and division on a strictly positive-denominator Borel set preserve measurability.

For the lattice statement, round any genuinely feasible policy as in Theorem 3. Its rounded value exceeds $V - \varepsilon - D_J \theta_M \geq L - \varepsilon - D_J \theta_M$, and its cost increases by at most $D_C \theta_M$. This proves the lower endpoint. A lattice minimizer is at worst $e + D_J \theta_M$ below the upper-witness target. Apply the first assertion with that value of δ to obtain the upper endpoint and (13). Nonnegative costs justify clipping the lower endpoint at zero. Q.E.D.

Failure of $\sigma > 0$ means that this particular witness pair has not certified the repair margin; it does not establish infeasibility of the original problem. Witness refinement and complete search have different roles in (13). Refining the witnesses alone does not close the remaining policy-search gap. Conversely, an increasingly fine policy lattice does not erase a fixed operating-witness error. The exact-value theorem is recovered by $L = U = V$ and $d = e = 0$.

3.2. A complete finite-state construction

Let X have n states. For an integer $M \geq 1$, let $\mathcal{P}_M$ be the policies whose action probabilities are multiples of $1/M$. Put $\theta_M = \min\{1, (m-1)/M\}$ and $\delta_M = D_J \theta_M$. Evaluate every candidate by the common Bellman equations and let

$$v_M = \min \left\{ \int C_0^p d\nu : p \in \mathcal{P}_M, J_t^p \geq V_t - \varepsilon - \delta_M \forall t \right\}.$$

For rational primitives these comparisons, values, and policy probabilities use rational arithmetic.

THEOREM 3—Convergence to the constrained optimum: *Repair a minimizer defining v_M and denote its feasible cost by U_M . Then*

$$L_M := \max\{0, v_M - D_C \theta_M\} \leq K_R(\nu, \varepsilon) \leq U_M, \quad U_M - L_M \leq D_C \{\theta_M + \rho_\varepsilon(\delta_M)\}. \quad (14)$$

Thus a prescribed positive absolute accuracy is reached at a finite M . At most $\binom{M+m-1}{m-1} nT$ policies need be evaluated at that level. Running maxima of lower bounds and running minima of upper bounds give nested certificates without changing their validity.

PROOF: Round the first $m-1$ probabilities of any feasible policy down to multiples of $1/M$ and assign the remainder to the last action. The total-variation error is at most $(m-1)/M$, capped by one. Lemma 1 shows that the rounded policy belongs to the relaxed finite set and has cost at most the original cost plus $D_C \theta_M$. Hence $v_M - D_C \theta_M \leq K_R$. Theorem 1 applied to the minimizing grid policy gives $U_M \leq v_M + D_C \rho_\varepsilon(\delta_M)$. The weak-composition count gives the number of simplex points at each date and state. The displayed error tends to zero. *Q.E.D.*

The theorem controls distance to (4), not distance to a one-step relaxation. Its count is exponential in $nT(m-1)$; it is a completeness construction, not a polynomial-time claim. The deterministic finite-state problem can instead be solved by evaluating its m^{nT} vertices with the original tolerance. The randomized repair is neither required nor valid as a deterministic-preserving operation.

3.3. The executed spatial search

For binary actions, write p_{ti} for the probability of action 1. A policy box $Q = \prod_{ti} [\ell_{ti}, u_{ti}]$ retains the same coordinates across all restarts. Backward interval Bellman evaluation gives operating upper values $\bar{J}$ and unconstrained implementation lower values $\underline{C}$ for that box. Necessary operating tests can contract its coordinate intervals: if q_0^+, q_1^+ are next-date operating upper values, an admissible p must satisfy

$$(1-p)q_0^+ + pq_1^+ \geq V_t(i) - \varepsilon. \quad (15)$$

Intersect this half-interval with $[\ell_{ti}, u_{ti}]$. Empty intersections or inconsistent directed operating enclosures remove the box. Every removal is justified by a necessary inequality. Feasible policies are never removed because a sufficient candidate check fails.

The remaining cost lower bound is $\int \underline{C}_0 d\nu$. Repairing the midpoint of the contracted box produces a feasible upper policy. The algorithm selects a live box with the smallest lower bound and bisects a coordinate maximizing $\beta^t(u_{ti} - \ell_{ti})$. It preserves both children and every live leaf. A global lower endpoint is the minimum of all live or cost-pruned leaf lower bounds and the incumbent upper cost. It stops when their difference is no greater than the requested tolerance, or reports the remaining interval when the preassigned budget is exhausted.

PROPOSITION 1—Stopping bound for interval Bellman search: *Let $w = \max_{ti}(u_{ti} - \ell_{ti})$ for a surviving policy box. Its midpoint-repair cost minus its interval Bellman lower bound is at most*

$$H(w) = D_C\{w + \rho_\varepsilon(D_J w)\}. \quad (16)$$

Without a node or time cap, the best-bound bisection procedure reaches every prescribed positive gap η in finitely many steps. In particular, choose $w_ > 0$ with $H(w_*) \leq \eta$. A sufficient worst-case bisection depth is*

$$d \left\lceil \log_2 \frac{1}{\beta^{T-1} w_*} \right\rceil, \quad d = nT, \quad (17)$$

with the ceiling replaced by zero when its argument is at most one. The associated full binary-tree node bound is exponential in this depth. Adding valid McCormick lower bounds cannot invalidate the stopping guarantee.

PROOF: With only rectangular probability restrictions, ordinary Bellman maximization produces the maximum operating value simultaneously at every state. It is a common Markov policy, not a family of incompatible histories. A surviving contracted box has its operating upper value at least $V_t - \varepsilon$ at every date and state. Each such maximizing policy differs from the midpoint policy by total variation at most w . Lemma 1 therefore makes the midpoint $D_J w$ -approximately feasible. The same rectangular Bellman argument for minimum cost gives a lower bound no smaller than midpoint cost minus $D_C w$. Repair and (8) imply (16). A box satisfying $H(w) \leq \eta$ cannot need refinement while the global gap exceeds η : its repaired cost is already an incumbent candidate, so its lower bound is within η of the global incumbent or higher. On any still-refined path, weighted longest-side bisection makes every physical width at most w_* by the depth in (17). There are finitely many nodes up to that depth. Valid stronger lower bounds only remove or improve boxes. The supplement details the effect of the necessary contractions; they shrink the rectangles and preserve every feasible point. *Q.E.D.*

This is a convergence theorem for the implemented search rule, not an inference from favorable outcomes. It also explains why a 60-second run can end with a large gap: a finite worst-case bound can be prohibitively large. The numerical results report that distinction directly.

4. A SINGLE ERROR BUDGET FOR CONTINUOUS-STATE APPROXIMATION

A finite-state theorem alone does not justify replacing a continuous model by a grid. The following sufficient conditions permit such a transfer without restricting the original policy to be piecewise constant.

ASSUMPTION 1—Certified reset approximation: *For a finite Borel partition $(E_i)_{i=1}^{n_h}$, let μ_i be probability measures supported on E_i and write $\nu = \sum_i \nu_i \mu_i$. There are cellwise rewards $r_{h,t}$, costs $k_{h,t} \in [0, K]$, and terminal values g_h , and a kernel*

$$Q_{h,t}^a(x, dy) = \sum_j P_{h,t}^a(i, j) \mu_j(dy), \quad x \in E_i,$$

with certified uniform errors d_r, d_k, d_g and α_h satisfying

$$|r - r_h| \leq d_r, \quad |k - k_h| \leq d_k, \quad |g - g_h| \leq d_g, \quad |(P_t^a - Q_{h,t}^a)f(x)| \leq \alpha_h \|f\|_\infty$$

for every bounded Borel f . The approximation rewards obey the same bounds R, G as the original rewards, or R, G are enlarged to cover both models. All finite transition entries and error bounds are computable with verified directions.

The last inequality uses the operator convention for total variation; it is not the half- ℓ_1 convention for action probabilities above. Define $M_{J,t} = R \sum_{s=0}^{T-t-1} \beta^s + G \beta^{T-t}$ and $M_{C,t} = K \sum_{s=0}^{T-t-1} \beta^s$. Set

$$a_T = d_g, \quad a_t = d_r + \beta a_{t+1} + \beta \alpha_h M_{J,t+1}, \quad (18)$$

$$b_T = 0, \quad b_t = d_k + \beta b_{t+1} + \beta \alpha_h M_{C,t+1}, \quad (19)$$

$$a_h = \max_t a_t, \quad b_h = b_0.$$

Let $v_h(s)$ be the optimum of the finite common-Markov model at operating allowance s and initial masses (ν_i) .

THEOREM 4—End-to-end transfer: *Under Assumption 1, if $\varepsilon > 2a_h$, then*

$$v_h(\varepsilon + 2a_h) - b_h \leq K_R(\nu, \varepsilon) \leq v_h(\varepsilon - 2a_h) + b_h. \quad (20)$$

Suppose finite computation returns intervals $[L_+, U_+]$ and $[L_-, U_-]$ for these relaxed and tightened optima, of widths at most η_+ and η_- . A valid global interval is

$$[\max\{0, L_+ - b_h\}, U_- + b_h], \quad (21)$$

and its width is at most

$$\eta_+ + \eta_- + 2b_h + D_C \frac{4a_h}{(1 - \beta)(\varepsilon - 2a_h) + 4a_h}. \quad (22)$$

If all certified primitive errors tend to zero and both finite-solver widths tend to zero, this interval converges to the original measurable-policy optimum. Its measured width supplies a termination test for the actual economic objective.

PROOF: Fix any measurable Markov policy p , not necessarily cellwise constant. Backward subtraction of its original and reset-model Bellman equations gives (18)–(19). Bellman comparison gives the same operating error for optimal values. In the reset model, averaging the policy within E_i , $\bar{p}_t(a | i) = \int p_t(a | x) \mu_i(dx)$, averages its value exactly; after every transition the within-cell law is again μ_j , and stage rewards and costs are cellwise constant. This identity follows backward, including the common continuation. An original feasible policy has reset

operating value at least the reset optimum minus $\varepsilon + 2a_h$ pointwise, hence also after cell averaging; its finite-model cost differs from the original initial cost by at most b_h . Taking an infimum gives the lower inequality. Conversely, lift a finite policy feasible at $\varepsilon - 2a_h$ by using its action probabilities throughout each cell. The reset value is constant in each cell; the two operating errors make the lifted original policy feasible, and its cost differs by at most b_h . This proves the upper inequality. Finally apply Theorem 1 in the finite model with target allowance $\varepsilon - 2a_h$ and violation $4a_h$, then add the two finite-computation widths and the cost errors. *Q.E.D.*

4.1. Where the transfer theorem applies

On a compact domain, a transition density with computable uniform continuity in the source variable and an L_1 -approximable density in the destination variable can furnish the reset approximation. The cell reference measures must also represent the initial law as in Assumption 1. For example, with Lebesgue-uniform cell measures and a uniform initial law, a uniformly Lipschitz transition density gives a computable $\alpha_h = O(h)$ on a shape-regular mesh. Cellwise operating and cost approximations give $d_r, d_k, d_g = O(h)$ when their moduli are Lipschitz. Aligned discontinuities in an installed rule can be retained as exact partition boundaries. In this case (22) is $O(h + \eta_+ + \eta_-)$ at fixed positive ε and $\beta < 1$, with explicitly displayed constants. The number of cells still scales as $h^{-\dim X}$; the theorem does not remove the curse of dimensionality.

The primary affine-shock model and the nonlinear two-state model below have atomic transition laws. Smearing their successor atoms into cells generally leaves the total-variation error bounded away from zero. We therefore do *not* apply Theorem 4 to those models by asserting a small mesh. Their reported intervals are direct whole-domain bounds for the unchanged atomic models. They remain valid without a convergence claim for that particular spatial representation. This distinction prevents the approximation theorem from being used as an unsupported explanation for closing an atomic-model gap.

5. CERTIFICATES THAT ACCELERATE THE GLOBAL PROBLEM

The complete search and the transfer theorem provide the convergence argument. The inherited one-sided certificates are useful accelerators and diagnostic tools, but are not substituted for that argument.

5.1. Deterministic outer restrictions and measurability

Let $L_t \leq V_t \leq U_t$ be bounded Borel operating witnesses, including the terminal date. Every all-restart feasible deterministic policy must select an action in

$$\mathcal{N}_t^{L,U}(x) = \{a : L_t(x) - T_t^a U_{t+1}(x) \leq \varepsilon\}. \quad (23)$$

Indeed, its operating continuation cannot exceed $V_{t+1} \leq U_{t+1}$. The set is nonempty because an operating-optimal action satisfies the test. Its graph is Borel, since it is a finite union of Borel sublevel sets. Backward minimization of $k_t + \beta P_t^a B_{t+1}$ over this set yields a Borel lower function; the lowest-index minimizer is a Borel selector. Thus the general-state recursion does not rely on an implicit measurable-selection assertion.

This is a necessary-action outer class, not the original constrained problem. A selected continuation can accumulate losses that violate later operating requirements. Equality between its integrated cost lower bound and a separately feasible policy's cost certifies $K_D(\nu, \varepsilon)$. Equality of the full functions certifies more. Tightening L, U alone does not force equality; complete search is what controls the remaining structural gap.

5.2. Randomized support floors and a connected local program

For $\lambda \geq 0$, a verified supersolution W_t^λ of

$$W_T^\lambda \geq \lambda g, \quad W_t^\lambda \geq \max_a \{\lambda r_t - k_t + \beta P_t^a W_{t+1}^\lambda\}$$

implies $C_t^p \geq \lambda(V_t - \varepsilon) - W_t^\lambda$ for every feasible policy, including randomized Markov policies. Define a floor F_t by the maximum of zero and finitely many such terms, using L_t in place of V_t when necessary. Given a valid next-date lower function B_{t+1} , define

$$\mathcal{D}_t(x) = \left\{ p \in \Delta(A) : \sum_a p_a T_t^a U_{t+1}(x) \geq L_t(x) - \varepsilon \right\}.$$

$$B_t(x) = \max \left\{ F_t(x), \min_{p \in \mathcal{D}_t(x)} \sum_a p_a \{k_t(x, a) + \beta P_t^a B_{t+1}(x)\} \right\}, \quad B_T = 0. \quad (24)$$

The feasible action distribution of any admissible policy belongs to the local constraint, and its continuation cost dominates B_{t+1} . Induction proves that (24) is a global randomized cost lower bound. It does not prove that the local optimizer can be combined into a feasible policy or that finite support prices close a global duality gap.

For three actions and one local operating inequality, the local polytope has only pure feasible vertices and two-action boundary mixtures. Our new primary calculation evaluates these vertices with exact rational arithmetic on directed whole-cell bounds. Unlike the earlier diagnostic with a zero global floor, it uses the nonzero F_t at every date and is paired with a feasible policy for the same uniform-initial randomized objective. It improves that lower endpoint in 23 of the 42 configurations. The prior adaptive local-LP theorem and all of its experimental records remain in the retained supplement; their target is still the local relaxation, not (4).

5.3. Whole-domain randomized deployments

A simple backward construction supplies feasible upper policies without changing the atomic model. Given a feasible next-date value J_{t+1} , write $q_0 = T_t^{a_0} J_{t+1}$, $q_* = \max_a T_t^a J_{t+1}$, and $b = V_t - \varepsilon$. If $q_0 \geq b$, retain the installed action. Otherwise mix it with a lowest-index maximizing action with probability

$$p_t^{\text{change}}(x) = \frac{b(x) - q_0(x)}{q_*(x) - q_0(x)}. \quad (25)$$

The denominator is positive wherever this expression is used. The resulting operating value is $J_t = \max\{q_0, b\}$ and is feasible at every state. In the piecewise-affine model, J_t remains piecewise affine while the mixing probability can be a ratio of affine functions. Cost is therefore enclosed rather than falsely represented as an exact piecewise-affine function.

The verifier includes isolated switching points in every operating identity. Directed cost boxes integrate against the three absolutely continuous initial laws. All affine transition slopes are strictly positive, so a finite set of switching points and its finite-horizon preimages have zero probability under these laws. That fact justifies almost-everywhere cost enclosures; it does not justify ignoring knots in a restart operating constraint or using the same cost enclosure for a Dirac initial law.

5.4. A verified direct mathematical-programming comparator

The direct finite-state formulation minimizes $\sum_i \nu_i C_{0i}$ subject to (2)–(3), $J_{ti} \geq V_{ti} - \varepsilon$, and the common policy simplexes. The constraints are bilinear in policy probabilities and continuation values. A McCormick relaxation on each probability box provides a classical global lower program. The baseline uses the same incumbent repair, contraction rule, box selection, node allowance, initial law, and all-restart constraints as interval Bellman search. The additional work is the stronger LP bound and an LP-proposed feasible candidate after repair.

Floating-point LP status is not a certificate. For a minimization relaxation $c'z$ subject to $Az \leq b$ and $\ell \leq z \leq u$, every rational vector $\lambda \geq 0$ gives

$$\mathcal{L}(\lambda) = -\lambda'b + \sum_j \min\{(c + A'\lambda)_j \ell_j, (c + A'\lambda)_j u_j\} \leq \min c'z. \quad (26)$$

This follows by adding $\lambda'(Az - b) \leq 0$ and minimizing the resulting affine function over the box. Rounded nonnegative solver marginals are evaluated in (26) using exact rational arithmetic. A failed LP falls back to the interval Bellman lower bound; a floating infeasibility declaration never removes the policy box. A separate verifier rebuilds every relaxation and every search-tree partition from the primitive model.

6. NUMERICAL EVIDENCE

The experiments keep model class, information, initial distribution, and operating allowance explicit. The supplement reports every configuration; no failed or budget-stopped case is removed. The computation report records proof-object sizes, memory, wall time, the exact environment, and the boundary between independent arithmetic checks and shared-code validation.

6.1. Prospectively fixed common-Markov problems

Before generating eight new environments, we committed their seeds, joint parameter changes, two solver hashes, node cap 4,096, wall-clock cap 60 seconds per solver and case, and absolute stopping allowance 10^{-3} . The freeze is `commit 13fd3ee`; it follows development seed 390001 and precedes all eight evaluated seeds. Both solvers see the same known model. The two-state models represent low and high condition; four-state models use a pair of condition indicators with interaction terms in their transition probabilities. Horizons are 3, 4, 8, and 12; discount factors are 0.8 or 0.95; allowances are 0.01 or 0.05; operating prices, transition persistence, intervention effects, and initial masses vary jointly. The exact formulas and resulting rational matrices are in the supplement and proof objects.

Table I contains all eight pairs, including unfinished searches. A positive absolute-gap certificate is not necessarily a small relative-gap certificate when the optimum is small. We therefore report both. The linear relaxation often improves the lower endpoint substantially, but costs more per node. The four-state longer-horizon cases expose the combinatorial burden: the remaining interval is not hidden by a success count. The comparison does not establish that interval Bellman search dominates direct global programming. Instead it identifies the value and cost of the stronger relaxation under matched information and caps.

6.2. The unchanged continuous-state development cohort

The original maintenance model has $x \in [0, 1]$, three actions, two shocks with probabilities $3/5$ and $2/5$, $\beta = 19/20$, stage reward $x - c_a$ with $c = (0, 3/20, 9/20)$, terminal payoff $x/2$,

TABLE I
SAME-OBJECT COMPARISON ON THE PROSPECTIVELY FIXED ENVIRONMENTS

States	T	Method	Lower	Upper	Gap/upper (%)	Seconds
2	3	I	2.8414	2.8425	0.03	0.99
2	3	M	2.8414	2.8425	0.03	1.77
2	4	I	0.0000	0.0000	0.00	0.00
2	4	M	0.0000	0.0000	0.00	0.38
2	8	I	7.2534	9.0046	19.45	17.19
2	8	M	8.5290	9.0046	5.28	60.41
2	12	I	7.1419	7.5553	5.47	16.32
2	12	M	7.4446	7.5553	1.46	60.40
4	3	I	2.6157	2.6217	0.23	10.82
4	3	M	2.6206	2.6217	0.04	35.99
4	4	I	0.0233	0.0244	4.00	0.47
4	4	M	0.0235	0.0244	3.13	2.60
4	8	I	4.6631	7.8481	40.58	60.04
4	8	M	6.1990	7.8481	21.01	60.41
4	12	I	4.8936	6.6439	26.34	60.04
4	12	M	5.8606	6.6439	11.79	60.43

I: interval Bellman search. M: verified McCormick search. Bounds concern the identical common randomized Markov objective for each pair. Reported times include model construction and the operating solve; the search's 60-second cap is checked between iterations, not by interrupting a proof-producing node. A zero upper cost is an exact zero-cost case.

Full stop reasons, node counts, and memory appear in the supplement.

and revision weight $1 + x$. Its two affine successors for actions 0, 1, and 2 respectively are

$$(3x/4, 3x/4 + 1/20), \quad (17x/20 + 1/10, 17x/20 + 7/50), \quad (x/10 + 4/5, x/10 + 17/20).$$

The seven installed rules, horizons 4, 8, and 12, and allowances 0.01, 0.05 are unchanged. Neural and spline rules are fixed computational incumbents, not estimated economic policies. Exact operating values are inexpensive in this model and their computation is charged. We do not use this cohort to claim performance when operating approximation itself is difficult.

The inherited deterministic comparison has 12 zero-cost and 12 positive-cost integrated equalities. Ten of the positive-cost equalities occur at horizon 4, two at horizon 8, and none at horizon 12. The maximum deterministic relative gaps at each horizon are reported alongside absolute gaps in the supplement. In the horizon-8 neural31003 case at $\varepsilon = 0.05$, the two initial cost functions are identical pointwise, but the arrays differ at later restart dates. Thus initial equality does not imply equality at every restart. The new diagnostic records the full rational disagreement locus.

The lotteries in (25) certify strict improvement over *every* feasible deterministic policy in 27 of the 42 uniform-initial configurations (Table II). The horizon-12 defer case at allowance 0.01, for example, has a randomized interval $[15.3172, 15.4670]$ and a positive deterministic separation. This addresses the economically relevant policy-class difference directly, without substituting a fixed-restart history-conditioned frontier. The 12 zero-cost cases are already optimal by feasibility and nonnegativity and are not counted as nontrivial search successes. All 30 positive-cost randomized intervals retain a reported nonzero width; the narrowest and widest relative intervals are distinguishable in the full table.

Changing the initial density to $2x$ yields 28 strict deterministic separations; changing it to $2(1 - x)$ yields 27. These are sensitivity results on the same developed environments, not independent observations. The connected local program improves 23 uniform lower endpoints,

TABLE II
COMMON RANDOMIZED MARKOV CERTIFICATES ON THE DEVELOPMENT COHORT

T	Cases	Zero-cost exact	Strict gain over K_D	Improved lower	Maximum gap/upper (%)
4	14	4	10	8	14.61
8	14	4	8	7	32.29
12	14	4	9	8	34.23

Strict gain means that a verified randomized upper cost is below the deterministic lower bound for the same initial law and constraints. It does not mean that the randomized optimum is attained by that candidate. Improved lower compares the connected local-support bound with the prior randomized support lower bound. All 42 uniform-initial rows and both nonuniform sensitivity tables are retained.

TABLE III
SELECTED NONLINEAR COMMON-MARKOV INTERVALS

T	Mesh per axis	Lower	Upper	Gap/upper (%)	Independent check
8	64	0.9438	3.9348	76.01	full
16	128	4.1908	7.6644	45.32	not independent
32	64	6.5389	16.7884	61.05	full
32	128	8.1264	12.5226	35.11	full R40
32	256	8.8718	11.8492	25.13	not independent

Uniform initial law on the square and operating allowance $1/2$. Inherited full independent scalar checks cover $(T, N) = (8, 64)$ and $(32, 64)$. R40 additionally replays the complete $(32, 128)$ object, including the candidate upper cost and all-restart operating bound. The $(32, 256)$ result retains directed construction and structural checks, not the independent label of adjacent rows.

including the longer-horizon cases; it thus affects the primary global interval rather than a disconnected local diagnostic. It still does not establish global strong duality or close every interval.

6.3. Nonlinear two-state bounds

The retained nonlinear model has $(x, y) \in [0, 1]^2$, four actions, bilinear transitions, two shocks, reward $(x + y)/2 - c_a$, terminal payoff $(x + y)/4$, and revision weight $1 + (x + y)/2$. Its exact rational primitive specification is reproduced in the supplement. Whole-box operating witnesses and support supersolutions use outward arithmetic. The new lower construction uses 11 global support prices and five nonnegative local multipliers, including zero, at every date. It is paired with exactly the retained feasible upper policies. We run no new upper-portfolio competition and do not compare unequal proposal success fractions as evidence of optimality.

The old zero lower endpoint was not forced by the definition of revision cost: the operating constraints can force departures from the installed rule. Nonzero support supersolutions now detect that restriction. At $T = 32, N = 256$, the interval is $[8.8718, 11.8492]$, with relative width 25.13%. This is a genuine improvement over $[0, 11.8492]$, but it is not a near-exact optimum or a high-dimensional scalability result. Both endpoints and the all-restart operating inequality are independently rebuilt at $T = 32, N = 64$, and the R40 audit extends this to the complete $T = 32, N = 128$ object. Memory for the retained support arrays grows as $8N^2(T + 1)(11 + 4)$ bytes before temporary arrays and interpreter overhead; the machine record reports the actual payload and process high-water mark.

The allowance $1/2$ belongs to the inherited exploratory nonlinear design. It is not made comparable to 0.01 simply by normalizing state variables: the operating scale and remaining

horizon matter. We additionally compute lower-bound diagnostics at $\varepsilon = 0.05$ on four stated configurations. None is paired with a policy certified only at $1/2$. Without a newly feasible upper policy those tighter-tolerance rows remain lower-bound-only results. A failed sufficient box certificate is labeled uncertified, not infeasible. A box crossing an installed-rule threshold can charge the upper revision weight conservatively; it does not assert that every realized state in that box changes action.

7. VERIFICATION, ECONOMIC INTERPRETATION, AND LIMITS

A certificate has three distinct burdens. Its mathematical reduction must preserve the target policy class. Its finite arithmetic must check the asserted inequalities with the correct direction. Its publication must correspond to the checked objects. Hashes address the third burden, not the first two. Independent arithmetic addresses part of the second burden and can expose a model-contract error, but is not a proof that every economic assumption is appropriate.

The finite-state verifier uses SymPy rationals and imports neither search implementation nor the optimizer. It rebuilds the common Markov equations, primitive contracts, every policy-box split, every necessary contraction, the LP residual bounds, and the global leaf minimum. The continuous-state verifier independently checks the new lotteries, whole-cell costs, support supersolutions, and local primal vertices, using the prior independent primitive parser rather than the constructor's piecewise-affine operations. The nonlinear verifier uses symbolic rational corner geometry and arbitrary-precision scalar integer induction instead of the constructor's vectorized box operations. It checks two complete representative objects, not merely selected rows. The verification report specifies precisely which objects passed; it never substitutes cross-run equality or a mutation test for an inequality proof.

The synthesis preserves all prior scientific material. Deterministic bounds, history-conditioned frontiers, the two-period analytic randomization example, adaptive local-LP analysis, 66 economic sensitivities, failed runs, and the stopped-control derivations remain available in full. The supplement attaches the prior article and technical supplement verbatim, with a content map. The older 349-page historical annex remains unchanged as an archival object rather than a prerequisite for reading this article. The stopped consumption–portfolio problem is not relabeled solved: its inherited full-domain bound of approximately 7.1818 does not meet its 0.01 target. A killed Bellman comparison supports witness propagation but supplies neither missing action coverage nor a diffusion-discretization error bound (Kushner and Dupuis, 2001).

The economic conclusion established by the new primary calculation is precise. When randomized action assignment is admissible and costs attach to realized departures, common Markov lotteries can strictly reduce the minimum implementation cost relative to every deterministic feasible rule. This conclusion holds on the certified configurations and initial laws. It does not estimate how frequently organizations permit such lotteries, how costly they are to administer, or the monetary welfare effect of a real deployment. Fixed randomization overhead, parameter uncertainty, and empirical calibration would require additional primitives and certificates. The joint-perturbation holdout study tests a fixed algorithm beyond its development seed, but is not a statistical generalization theorem.

7.1. Model uncertainty and the economic margin for randomization

The certificates can be translated into a robustness statement without treating designed numerical primitives as estimated data. Fix a nominal model, initial law, and policy class. Suppose a nonempty set of alternative models has verified uniform bounds

$$|J_t^{p,M} - J_t^{p,0}| \leq a, \quad |V_t^M - V_t^0| \leq a, \quad |C_0^{p,M} - C_0^{p,0}| \leq b \quad (27)$$

for every common policy p , state, and date, and every M in the set. Superscript 0 denotes the nominal model. The error recursion in Section 4 supplies such bounds from uniform reward, cost, terminal, and kernel enclosures on a common state space; no reset aggregation is needed for this fixed-space perturbation statement. Uniformity over policies is essential. A fitted residual checked only under the proposed policy does not establish (27).

PROPOSITION 2—Portable operating and cost intervals: *Let $\varepsilon > 2a$. A nominal policy feasible at allowance $\varepsilon - 2a$ is feasible at allowance ε in every alternative model. If L_+ is a nominal lower bound at allowance $\varepsilon + 2a$ and U_- is the nominal cost of such a tightened feasible policy, then every model-specific common-policy optimum belongs to*

$$[\max\{0, L_+ - b\}, U_- + b].$$

The same interval encloses the optimum that minimizes worst-case implementation cost over policies feasible in all the alternative models. In the latter assertion the upper endpoint uses a single policy, not a different policy for each model.

PROOF: For the proposed policy, $J^{p,M} \geq J^{p,0} - a \geq V^0 - \varepsilon + a \geq V^M - \varepsilon$. Conversely a policy feasible in a model M satisfies $J^{p,0} \geq J^{p,M} - a \geq V^M - \varepsilon - a \geq V^0 - \varepsilon - 2a$. Its nominal cost is at least L_+ and its cost in M at least $L_+ - b$. Take infima for the model-specific statement. For the robust problem, the first implication supplies one policy for all models and the second applies to every robustly feasible policy in each model. Taking the worst-case cost preserves the bounds. *Q.E.D.*

A strict randomized separation also has an economically interpretable administration-cost allowance. Let U_R be the certified cost of a feasible common lottery and L_D a lower bound for every feasible deterministic policy, for the same operating constraint and initial law. Suppose operating payoffs and dynamics are unchanged by administering the lottery. A nonnegative extra administration charge at date t bounded by h_t increases its expected discounted cost by at most $H = \sum_{t < T} \beta^t h_t$. If

$$U_R + H < L_D, \tag{28}$$

then the lottery remains cheaper than every deterministic feasible policy even after this additional charge. This conclusion does not require administration at every visited state; charging the bound everywhere is conservative. Deterministic rules are assumed not to pay the randomization-specific charge.

For a common per-date ceiling h , the certified break-even threshold is

$$h_* = (L_D - U_R)_+ \bigg/ \sum_{t=0}^{T-1} \beta^t. \tag{29}$$

Every $h < h_*$ preserves the strict ranking. Equality at h_* yields no strict conclusion. For model uncertainty, use a deterministic nominal lower bound at $\varepsilon + 2a$ and a nominal lottery feasible at $\varepsilon - 2a$; then replace the numerator by $(L_{D,+} - U_{R,-} - 2b)_+$. The current development results report $a = b = 0$: they calculate economic margins under known designed primitives, not empirical uncertainty radii.

TABLE IV
CERTIFIED ALLOWANCE FOR A PER-DATE RANDOMIZATION ADMINISTRATION CHARGE

T	Installed rule	ε	Det. lower	Lottery upper	h_* lower bound
4	defer	0.01	4.93668	4.73011	0.055684
4	defer	0.05	4.93206	3.89012	0.280857
4	neural31001	0.01	0.11117	0.08928	0.005902
8	defer	0.01	10.81072	10.63794	0.025667
8	defer	0.05	10.75071	9.85440	0.133151
8	neural31001	0.01	0.46864	0.40232	0.009853
12	defer	0.01	15.61281	15.46691	0.015872
12	defer	0.05	15.47129	14.72413	0.081278
12	neural31001	0.01	3.49393	3.36015	0.014553

Uniform initial law, unchanged development economy. Endpoints and the threshold are rounded outward. A positive threshold means that every smaller constant per-date overhead preserves a strict cost advantage over every feasible deterministic policy. All 126 rule/horizon/tolerance/initial-law margins are retained in the machine-readable results, including zeros.

8. CONCLUSION

A local operating requirement becomes a global revision method only when one can control both continuation compatibility and the cost of restoring feasibility. The backward repair provides that control for positive tolerance and strict discounting. Combined with complete policy search, it produces a convergent certificate for the actual common randomized Markov objective. Combined with a verified reset approximation, it yields an explicit end-to-end error budget for an unrestricted measurable-policy problem. These results retain the optimization target rather than replacing it by a necessary-action or finite-price relaxation.

The computations show both the benefit and the burden of this construction. Whole-domain lotteries resolve a substantive deterministic-versus-randomized comparison on the original development cohort. Connected support bounds improve the same primary objective, and nonzero nonlinear lower bounds replace feasibility-only intervals. A matched direct global-programming comparator is often stronger per explored box. Longer horizons and four-state joint perturbations retain material gaps under fixed budgets. The stopping theorem explains how complete refinement can close them; the reported execution does not claim to have done so. The remaining numerical challenge is to exploit economic structure strongly enough that this complete certification becomes practical at larger state and horizon dimensions without changing its policy class or silently weakening its operating constraints.

REFERENCES

- Altman, E. (1999): *Constrained Markov Decision Processes*. Chapman and Hall/CRC. [2]
Brown, D. B., and J. E. Smith (2014): “Information Relaxations, Duality, and Convex Stochastic Dynamic Programs,” *Operations Research*, 62(6), 1394–1415. doi:10.1287/opre.2014.1322. [3]
Brown, D. B., J. E. Smith, and P. Sun (2010): “Information Relaxations and Duality in Stochastic Dynamic Programs,” *Operations Research*, 58(4, part 1), 785–801. doi:10.1287/opre.1090.0796. [3]
Brumm, J., and S. Scheidegger (2017): “Using Adaptive Sparse Grids to Solve High-Dimensional Dynamic Models,” *Econometrica*, 85(5), 1575–1612. [3]
Ding, D., K. Zhang, T. Basar, and M. R. Jovanovic (2020): “Natural Policy Gradient Primal-Dual Method for Constrained Markov Decision Processes,” *Advances in Neural Information Processing Systems*, 33. [3]
Hettich, R., and K. O. Kortanek (1993): “Semi-Infinite Programming: Theory, Methods, and Applications,” *SIAM Review*, 35, 380–429. doi:10.1137/1035089. [3]
Huangfu, Q., and J. A. J. Hall (2018): “Parallelizing the Dual Revised Simplex Method,” *Mathematical Programming Computation*, 10, 119–142. doi:10.1007/s12532-017-0130-5. [3]

- Judd, K. L. (1998): *Numerical Methods in Economics*. Cambridge, MA: MIT Press. [3]
- Kushner, H. J., and P. G. Dupuis (2001): *Numerical Methods for Stochastic Control Problems in Continuous Time*, second edition. New York: Springer. [15]
- Laroche, R., P. Trichelair, and R. Tachet des Combes (2019): “Safe Policy Improvement with Baseline Bootstrapping,” *Proceedings of Machine Learning Research*, 97, 3652–3661. [3]
- McCormick, G. P. (1976): “Computability of Global Solutions to Factorable Nonconvex Programs: Part I—Convex Underestimating Problems,” *Mathematical Programming*, 10, 147–175. doi:10.1007/BF01580665. [3]
- Munos, R., and C. Szepesvári (2008): “Finite-Time Bounds for Fitted Value Iteration,” *Journal of Machine Learning Research*, 9(27), 815–857. [3]
- Ni, T., and M. Kamgarpour (2025): “A Safe Exploration Approach to Constrained Markov Decision Processes,” *Proceedings of Machine Learning Research*, 258, 3592–3600. [3]
- Rump, S. M. (2010): “Verification Methods: Rigorous Results Using Floating-Point Arithmetic,” *Acta Numerica*, 19, 287–449. [3]
- Rust, J. (1987): “Optimal Replacement of GMC Bus Engines: An Empirical Model of Harold Zurcher,” *Econometrica*, 55, 999–1033. doi:10.2307/1911259. [3]
- Saldi, N. (2018): “Finite-State Approximations to Discounted and Average Cost Constrained Markov Decision Processes,” arXiv:1807.02994. [2]
- Saldi, N., T. Linder, and S. Yüksel (2018): *Finite Approximations in Discrete-Time Stochastic Control: Quantized Models and Asymptotic Optimality*. Birkhäuser. doi:10.1007/978-3-319-79033-6. [2]
- Scherrer, B. (2013): “Performance Bounds for λ Policy Iteration and Application to the Game of Tetris,” *Journal of Machine Learning Research*, 14(36), 1181–1227. [3]