

NEURAL BELLMAN OPERATORS

QIAN QI

Peking University

This paper develops Neural Bellman Operators for policy evaluation and feasible improvement in controlled economies. Centered continuation errors link economic transport to full-policy loss. A constructive neural backend retains the feasible actions attaining its labels and supports implementation under acquired states and endogenous capacity. We give an exact continuous-state compiler that preserves the original policy, including ties, while removing a redundant state-cover factor from its evaluation work. The construction combines classical distance transforms with witness-preserving off-grid evaluation and numerical certification. A priced observation technology determines precision from a certified economic account. Direct constrained-policy comparisons distinguish signed cost differences from the ability to recoup a controller replacement fee. Matched services include uniform and error-driven conventional methods, failed refinements, and three- and four-dimensional nonlinear stress cases. The results separate representation savings, complete work, and economic decisions without attributing a native-envelope capability exclusively to neural encoding. Broader applications retain their own comparison conditions.

KEYWORDS: Dynamic programming, Neural approximation, Policy certification, Recursive utility, Computational economics.

CURRENT READING NOTICE

The R61 article, technical supplement and response form the current submission. This complete companion retains all preceding statements and adds the current exact-search, stability, fixed-policy and economic evidence. Earlier reading notices are historical.

The R59 article and response form the current submission. This companion retains the preceding NBO development and adds lossless trained-neural action recovery

and its matched prospective work study. Historical claims refer to their own cohorts and hypotheses.

R57 READING NOTICE

The R57 active article and response are the current submission. This companion preserves all earlier statements and adds the new fitted-witness transfer and executed action-search comparison. Historical uses of current refer to their own editions.

CURRENT READING MAP

The active R56 article is the current exposition of Neural Bellman Operators. This companion preserves the preceding body and adds the current results without deleting its earlier assumptions, proofs, applications or adverse observations. Historical uses of the word current below identify their own editions.

READING NOTICE FOR THE COMPLETE DEVELOPMENT EDITION

The current additions are the directed full-sweep certificate and the complete policy-cost study. The historical body remains available without deleting its results or applications.

This volume retains the baseline development verbatim, including historical descriptions, followed by the current R54 argument. Its historical reference to a formerly authoritative revision is not a second current scientific claim. The active article, technical supplement and response govern the new evidence and numerical corrections. Every baseline result and its original assumptions remain available for examination.

1. INTRODUCTION

An economic decision uses a continuation value because its consequences extend beyond the current period. A learned continuation is useful only to the extent that the decisions made with it are accurate. Prediction loss, a Bellman residual at sampled states, and the value of the policy actually implemented are different objects. Their distinction is particularly important when a counterfactual changes

future policies or when a recursive preference evaluates the continuation generated by those policies rather than the optimum.

This paper develops Neural Bellman Operators (NBOs) as a policy-evaluation and feasible-improvement method. An operator takes economic primitives, a current policy, and computational resources as inputs. It returns a continuation, a feasible policy, and an economic error account, or it records that the requested account was not obtained within its cap. The continuation is evaluated against its own fitted future. Current actions, policy-specific future values, numerical execution, and the cost of verification enter the same account. The construction is not a proposal to rename an arbitrary neural predictor as a Bellman operator.

There are four parts to the argument. First, only action-dependent continuation errors affect a current decision. Their oscillation, rather than their absolute level, controls loss. A transport bound identifies when an anchor continuation remains valid after the future economy changes. These results connect reuse and refresh decisions to economic tolerances. They also preserve the original formulation in controlled diffusions, recursive preferences, endogenous preferences, and games, subject to the distinct hypotheses of those applications.

Second, a signed residual certificate controls a complete returned policy. Its upper bound contains the widths of the residual enclosures, the suboptimality of the selected actions, and the terminal enclosure. It is invariant to additive date-specific shifts of the fitted continuation. This leads to a useful separation: a policy can be accurate even though a coarse verifier does not establish its accuracy. We prove that residual verification can be refined while retaining the original network, the original deployed actor, and its original allowance. The resulting bound measures improved knowledge of the same policy, not a policy improvement disguised as additional verification.

Third, certification of each policy relative to the optimum does not rank two policies. We construct their cost difference directly. Under expected-cost preferences, an own-continuation Bellman telescoping identity produces an unbiased policy-cost score with a support width governed by the signed certificate. Coupling innovations across policies and enclosing every simulated path gives simultaneous finite-sample intervals for the difference in the original continuous-innovation economy. The

statistical assertion is conditional on frozen policies and the stated independent simulation model; outward arithmetic and unresolved action-cell boundaries have their own explicit treatment.

Fourth, a deterministic neural construction connects primitive moduli to a complete policy tolerance. A minimum of Lipschitz cones, realized exactly by ReLU gates, fits each date’s own-future Bellman labels without inflating its Lipschitz bound when those labels are inexact. The state, action, and arithmetic budgets are specified before construction. This yields finite termination for a defined neural backend rather than a claim about arbitrary nonlinear training. A new full matched catalogue executes the same resolution ladder for this backend and a conventional spline, including every unsuccessful rung and three isolated timing repetitions.

The action that attains a Bellman label supplies additional information. We retain its index through the ReLU minimum circuit and deploy that action when its cone is selected. A one-sided policy sandwich then avoids charging the separate actor and critic accounts twice. Both the nearest-node neural ablation and the spline receive a strengthened certificate. Their complete matched frontiers isolate certificate sharpness from first-target work and actual policy cost.

A feasible policy must also survive the information and arithmetic used to implement it. We give a uniform acquisition-aware theorem that permits discontinuous witness selection, uncertain measured states and capacity-dependent repair. A coupled two-state investment economy realizes all covering, integration and repair operations explicitly, with a conventional fitted-value comparator. The original scalar witness policies receive direct expected-cost comparisons at every common first crossing, under a tolerance and simultaneous error account fixed before execution. This establishes proximity within the declared decision band, not a signed neural ranking. A catalogue-choice theorem makes the result usable when implementation charges differ. The complete resource curves and an identical native-envelope control isolate what the construction, the certificate and the neural representation each contribute.

Three further results connect this construction to computation and economic implementation. First, a literal cone network repeats a scan over every state label whenever a Bellman query evaluates its future continuation. On a tensor cover, we

compile this network once and answer an arbitrary continuous-state query from the cell's corner data. The compiler carries the original action witness and the original tie rule. Consequently the improvement in evaluation work does not come from changing the policy being evaluated. Classical separable distance transforms supply the nodal computation; the additional policy argument establishes the continuous off-grid identity and feasible witness realization. In the two-state construction this removes the redundant factor that makes dense cone evaluation scale as N^6 rather than N^4 . It does not remove the state cover or give a capability unavailable to the identical native envelope.

Second, we make observation accuracy a priced economic input. A finite-bit sensor reports a state cell. Robust repair selects an action feasible throughout that cell. The policy theorem then gives an augmented loss account consisting of the full-information construction bound, an acquisition term decreasing geometrically in the number of bits, and a linear information charge. Exact neighboring-bit comparisons identify the precision minimizing this guarantee. This is an information-design result for the stated technology, not an assertion that minimizing a bound also minimizes the unknown actual loss.

Third, we evaluate actual constrained-policy cost differences, independently of their all-state regret certificates. The two installed controllers face common continuous innovations and the same observation technology. A prospective replacement charge turns simultaneous expected-cost intervals into a decision about whether either replacement recoups its cost. That question differs both from the sign of the policy difference and from exact equality. A complete matched construction study, an error-driven conventional comparator, and higher-dimensional nonlinear stress executions accompany the decision analysis. We charge compilation, adaptive pilots, unsuccessful resolutions, repair and output, and report the joint policy-comparison service separately. These distinctions permit a positive computational contribution to the original NBO program without confusing a sharper certificate, a cheaper representation, and a better economic policy.

The numerical implementation starts from economic primitives. In the scalar investment economy, it fits a fresh ReLU continuation backward against targets using its own previously fitted future network. In the two-capital economy, it trains

hidden locations in a convex squared-ReLU continuation, with continuous simplex actions and continuous uniform innovations. The second experiment includes a flexible nonnegative ridge dictionary as well as a quadratic projection. Each comparator constructs its own future targets and uses the same final economic criterion. Neither conventional values nor inherited neural weights are used as labels for the neural services.

We report the complete 150-service catalogue, not only successful candidates. In the original capped scalar services the neural procedure reaches the all-state target 0.06 in all twelve cases and 0.04 in six, whereas the spline reaches both targets in every case. In the coupled services the neural procedure reaches 0.25 in six cases and 0.10 in five; the flexible ridge reaches both in all six. These statements describe the original service, including its original verification cap.

A subsequent diagnostic holds both the neural weights and the deployed policies fixed. Refining only residual verification certifies all six previously unresolved scalar policies at 0.04 and the previously unresolved coupled policy at 0.10. The original failures are not relabeled. The diagnostic shows that they do not, by themselves, identify a failure to construct an accurate neural policy. It does not identify the population success probability of the optimizer, establish a new end-to-end runtime, or change any policy's cost.

The comparative findings are equally integral to the paper. The original scalar signed policy comparisons frequently favor the spline. The new coupled direct intervals all contain zero and therefore do not order the policies. The precision controller genuinely escalates after rejected low-precision proposals, but its complete recorded work does not improve on fixed binary64 in this experiment. These findings distinguish the certified capability of the construction from a comparative advantage that would require additional evidence.

The paper retains its original economic subject and full theory. The exposition now follows the economic problem, the operator, the policy theorem, and the evidence. The complete previous article, proof supplement, and economic applications are materialized unchanged within the current source tree as preserved companions; the preservation map records their exact identities. No historical theorem or application

is deleted to obtain a shorter active exposition. The present article and new technical supplement state and prove the additional results used by the new evidence.

1.1. *Relation to the literature*

Neural Bellman Operators sit within approximate dynamic programming, fitted-value methods, and policy iteration, not outside them. Regression-based continuation methods and simulation-based dynamic programming supply important conventional benchmarks (Judd, 1998, Powell, 2011, Longstaff and Schwartz, 2001, Ernst et al., 2005). Actor-critic and approximate policy-iteration methods separate evaluation and improvement (Konda and Tsitsiklis, 1999, Scherrer et al., 2015); safe-improvement arguments likewise distinguish a proposed action from a verified gain (Kakade and Langford, 2002, Pirota et al., 2013). We use these distinctions to organize the economic object returned by the algorithm.

Neural methods for differential equations provide ways to approximate values and their derivatives (Sirignano and Spiliopoulos, 2018, Han et al., 2018, Hornik et al., 1990). Such approximation does not on its own verify a returned dynamic policy or its numerical execution. Our certificates can be applied to neural or conventional continuations. The neural-specific construction lies in fresh own-future training and, in the coupled experiment, learned hidden locations. Its comparative value must be isolated against flexible conventional methods, not inferred from the certificate's validity.

The general economic formulation admits recursive utility (Epstein and Zin, 1989, Duffie and Epstein, 1992). The compact-domain policy theorem uses monotonicity and cash invariance, while the original Gaussian entropic construction has separate moment and spectral conditions. These are different bridges, not interchangeable assumptions. The direct paired-cost identity below applies to conditional expectation; it is not asserted for a nonlinear certainty equivalent. The confidence step uses the empirical Bernstein inequality of Maurer and Pontil (2009). Our additional ingredients are the policy-specific telescoping score, its certificate-derived support, and the numerical enclosure of the original continuous-law comparison.

2. THE CONTROLLED ECONOMY AND THE OPERATOR

2.1. Policy-specific continuation

The original continuous-time economy has an admissible control a and state

$$dS_t = b(t, S_t, a_t) dt + \sigma(t, S_t, a_t) dW_t. \quad (1)$$

The feasible correspondence, invariant or stopped domain, and exit and terminal rewards are economic primitives. For recursive evaluation, the continuation of a fixed policy is represented, on its specified horizon, by

$$Y_s^a = g(\tau, S_\tau) + \int_s^\tau f(v, S_v, a_v, Y_v^a) dv - \int_s^\tau Z_v^a dW_v.$$

The value is $V^a(t, s) = Y_t^a$ under the well-posedness and comparison hypotheses in the retained controlled-economy analysis. In particular, the recursive argument is the policy's own continuation. Replacing it by the optimal continuation would change the policy-evaluation problem.

The applications companion retains the full boundary specifications, infinite-horizon transversality conditions, utility-domain restrictions, temporal-self equations, and game conditions. A finite-domain residual calculation does not establish any of these by itself. In a game, certifying a specified deviation problem is not the same as establishing the equilibrium fixed point. In continuous time, the controlled generator contains state derivatives; freezing a critic during an actor update does not suppress legitimate derivatives of the represented state-dependent function.

For the constructive implementation and the main policy theorem, consider dates $t = 0, \dots, T$, compact state space X , and feasible correspondence $A_t(x)$. States evolve as

$$X_{t+1} = F_t(X_t, a_t, Z_{t+1}). \quad (2)$$

Let c_t be cost, g terminal cost, $\beta_t > 0$, $B_0 = 1$, and $B_t = \prod_{j < t} \beta_j$. We use a cost-minimization convention in the discrete implementation. It is equivalent to changing signs in an additive reward formulation, but does not replace the original recursive-utility assumptions by a sign convention.

Write

$$Q_t(v; x, a) = c_t(x, a) + \beta_t \rho_t(v(F_t(x, a, Z_{t+1}))), \quad (3)$$

$$\mathcal{T}_t v(x) = \inf_{a \in A_t(x)} Q_t(v; x, a), \quad V_T^* = g.$$

The conditional evaluations ρ_t are monotone and cash invariant on a domain containing every value and comparison function used below. Feasible measurable selectors exist; compactness, continuity of primitives, and continuity of the evaluated integrals are sufficient conditions in the experiments. The Bellman recursion is valid against all admissible adapted policies. For expectation, J^π is expected discounted cost. For a recursive evaluation it denotes the value generated by that evaluation's own recursion.

2.2. An executable NBO service

An implemented operator maps economic data and a current policy to

$$(\hat{f}, \hat{\pi}, \mathcal{C}, W) \quad \text{or} \quad (\text{unresolved}, W, \text{diagnostic}),$$

where $\mathcal{C}$ is an economic certificate and W records the work used. Its stages are target construction, fitting, feasible action selection, independent certification, and durable storage. Critic fitting freezes its future inputs; action selection freezes the resulting critic. In the fresh backward construction, date t uses the already fitted $\hat{f}_{t+1}$, not an externally solved optimal continuation. The terminal target is the primitive g .

A finite optimizer schedule is an algorithm, not a convergence theorem. A service that has not met its target at its cap records failure, with the complete cost of its attempts. The retained factor-training analysis provides a finite constructive theorem for a separate structured class. The nonlinear schedules below instead have an explicit, auditable capped-service estimand.

For a fixed catalogue $\mathcal{K}$ and target ε , define

$$H_m(k, \varepsilon) = \mathbf{1}\{\text{method } m \text{ stores a valid certificate at } \varepsilon \text{ within its cap on object } k\}. \quad (4)$$

The empirical object is the full array of H_m , costs to the first qualifying checkpoint, and capped costs for failures. Its arithmetic average is a catalogue fraction. A probability over random initializations would require a specified law and a sampling design for that different estimand. A post-review refinement of a failed object's verifier is a new diagnostic and does not change (4) for the original service.

3. CONTINUATION ACCURACY, TRANSPORT, AND POLICY LOSS

3.1. Centered errors and economic reuse

Consider a current objective $r_\theta(s, a) + S_\theta(T_\theta(s, a))$ to be maximized. Let $e = \hat{S}_0 - S_0$ be anchor approximation error and $\Delta_\theta = S_\theta - S_0$ the change in the actual future continuation. Suppose the implemented action is δ -optimal for $r_\theta + \hat{S}_0 \circ T_\theta$.

PROPOSITION 1—Centered decision and transport bounds: *The current decision loss satisfies*

$$L_\theta(s, \hat{a}) \leq \delta + \text{osc}_{A_\theta(s)}(e \circ T_\theta) + \text{osc}_{A_\theta(s)}(\Delta_\theta \circ T_\theta). \quad (5)$$

For a common finite action menu with positive weights p_j , set $R_p = \sum_j p_j(e_j - \sum_i p_i e_i)^2$ and $C_p = \max_{i \neq j} (p_i^{-1} + p_j^{-1})$. If $|\Delta_\theta| \leq b_\theta$ on the queried postdecision set, then $L_\theta \leq \delta + \sqrt{C_p R_p} + 2b_\theta$.

If common-innovation future transitions have Lipschitz constants L_k and perturbation bounds $\epsilon_{\theta,k}$, and future rewards have Lipschitz constants H_k and perturbations $\eta_{\theta,k}$ on an invariant domain, put $d_0 = 0$, $d_{k+1} = L_k d_k + \epsilon_{\theta,k}$. For nonnegative reward weights w_k ,

$$|S_\theta - S_0| \leq b_\theta := \sum_k w_k (\eta_{\theta,k} + H_k d_k). \quad (6)$$

The proof compares the true and fitted objective at the two selected actions and couples the future states. The weighted finite-menu bound is Cauchy–Schwarz applied to $e_i - e_j$. Details are in the technical supplement. An additive continuation error common to all actions is irrelevant to the decision. A finite menu, however, does not certify a continuous action set merely because it includes both observed actions.

For recursive maps with sup-norm Lipschitz constants β_k on a common invariant utility domain, a one-step perturbation v_k propagates as $b_k \leq v_k + \beta_k b_{k+1}$. This is

the recursive counterpart of (6). Its use requires the common domain; it cannot cross a singular utility boundary by extending a local constant. These transport bounds retain the original refresh logic: use an anchor while its economic allowance covers approximation and future change, and recertify or recompute when that condition is exhausted.

3.2. A direct all-state policy theorem

For continuous fitted functions f_t and a feasible measurable policy π , suppose the verifier supplies, uniformly in x ,

$$\ell_t \leq f_t(x) - \mathcal{T}_t f_{t+1}(x) \leq u_t, \quad t < T, \quad (7)$$

$$Q_t(f_{t+1}; x, \pi_t(x)) - \mathcal{T}_t f_{t+1}(x) \leq \eta_t,$$

$$\ell_T \leq f_T(x) - g(x) \leq u_T.$$

These are enclosures for the actual coefficients and actual returned actions. Neither an optimizer's success flag nor a residual on sampled states replaces them.

THEOREM 2—Centered full-policy certificate: *Under the conditions of Section 2 and (7), every initial state satisfies*

$$0 \leq J_0^\pi(x) - V_0^*(x) \leq G(f, \pi) := \sum_{t < T} B_t(u_t - \ell_t + \eta_t) + B_T(u_T - \ell_T). \quad (8)$$

More specifically,

$$V_0^*(x) \geq f_0(x) - \sum_{t < T} B_t u_t - B_T u_T, \quad (9)$$

$$J_0^\pi(x) \leq f_0(x) + \sum_{t < T} B_t(\eta_t - \ell_t) - B_T \ell_T.$$

Replacing f_t by $f_t + d_t$, with date-specific constants d_t , leaves the bound unchanged after the residual endpoints are shifted by $d_t - \beta_t d_{t+1}$ and the terminal endpoints by d_T .

PROOF: Monotonicity and cash invariance imply $\mathcal{T}_t(v + c) = \mathcal{T}_t v + \beta_t c$ and the analogous identity for the fixed-policy map. Starting from $g \geq f_T - u_T$, backward induction gives the first inequality in (9). Starting from $g \leq f_T - \ell_T$ and using

$Q_t(f_{t+1}; x, \pi_t(x)) \leq f_t(x) - \ell_t + \eta_t$ gives the second. Subtraction proves the upper bound; optimality proves the lower bound. The shifted residuals have unchanged widths and unchanged action gaps. *Q.E.D.*

The theorem certifies a full adapted policy, not a finite action menu or a grid of initial states. With expected cost its proof is a Bellman telescoping argument; for a monotone cash-invariant recursive evaluation it is an order argument. The numerical experiment is not needed for its validity. Its numerical usefulness depends on producing small, valid enclosures.

COROLLARY 3—Retained-policy recertification: *Fix the functions f_t , the deployed policy π , valid actor allowances η_t^{old} , and a terminal enclosure. Let a second verifier return valid residual enclosures $[\ell_t^{\text{new}}, u_t^{\text{new}}]$ for those same functions. Then*

$$J_0^\pi - V_0^* \leq \sum_{t < T} B_t(u_t^{\text{new}} - \ell_t^{\text{new}} + \eta_t^{\text{old}}) + B_T(u_T^{\text{old}} - \ell_T^{\text{old}}). \quad (10)$$

No retraining or change in deployed actions is required. If desired, intersecting two valid residual enclosures makes the new bound no larger than the old one.

The original actor allowance is valid on the entire state domain, so a finer critic verification mesh does not invalidate it. Conversely, replacing it by an allowance for a newly optimized fine-grid actor would certify a different policy. The implementation records the original actor hash and explicitly discards any fine-grid action proposals generated internally by the verifier. Those proposals' work and transient storage are still charged.

The corollary does not promise that arbitrary precision is attainable by mesh refinement. Nodal residual variation, the retained actor allowance, and the terminal enclosure can leave a positive floor. Nor is the corollary a claim that additional verification improves welfare. It distinguishes the accuracy of a fixed economic decision rule from the sharpness of our knowledge of that accuracy.

3.3. Construction beyond post-training verification

The original trainable-factor construction remains a constructive part of NBO. For a quadratic own-policy continuation $x'P_tx + c_t$, a square-activation network

$\|W_t x\|^2/d + \hat{c}_t$ trains every entry of W_t . Its population derivative loss and gradient are

$$\mathcal{L}_t(W; P) = \frac{1}{4} \|W'W - dP\|_F^2, \quad \nabla_W \mathcal{L}_t = W(W'W - dP).$$

If $\sigma_{\min}(W) \geq s > 0$, then $\|W'W/d - P\|_F \leq \|\nabla_W \mathcal{L}_t\|_F/(ds)$. Together with the full-vector actor defect, this gives a policy-error account for a jointly trained hidden factor. The own-policy coefficient is re-evaluated for the returned policy, not silently borrowed from an earlier training policy.

The primitive-budget analysis in the retained article chooses a feasibility witness, coefficient envelopes, rank margins, and update budgets before training. In the Gaussian entropic extension it also verifies positive exponential-moment margins. Under its stated spectral hypotheses, a scaled orthogonal initialization and bounded gradient steps give the finite training construction proved there. The finite matrix and Gaussian formulas are not a nonlinear dimension-scaling experiment. They also do not imply that Adam or L-BFGS has a deterministic residual cap for a general squared-ReLU network. We preserve both results and keep the two logical statements separate.

4. FROM ECONOMIC PRIMITIVES TO NONLINEAR CERTIFICATES

4.1. *Backward fitting and selected policies*

For each method, initialize a fresh terminal fit from g . At date t , construct targets from $\mathcal{T}_t f_{t+1}$ using that method's own already fitted future. Fit the current continuation under the frozen schedule. Then propose feasible actions, enclose the Bellman infimum and the selected-action gap, and validate residual coverage. Store a checkpoint before reporting its first target crossing. Continue through the declared ladder only as specified by the original study. The record includes all attempted resolutions, unsuccessful target queries, target-generation work, and storage operations.

The scalar network has 63 ReLU features. Its training uses 513 states, a 257-point target action construction, 400 Adam steps at learning rate 5×10^{-5} , and a final output fit. The scalar verifier uses its own finer action and state ladders. The

conventional spline is reconstructed from the same primitives with its own ladder and targets. A conventional target solver is not used to label a neural continuation.

The coupled neural continuation is

$$f_t(x) = a_t + b'_t x + \sum_{j=1}^{48} c_{tj} (w'_{tj} x + v_{tj})_+^2, \quad c_{tj} \geq 0. \quad (11)$$

Hidden locations and weights are trained. Nonnegative coefficient initialization and a finite L-BFGS schedule of 300 iterations act on fresh own-future targets. The ridge comparator instead fixes a 96-feature squared-hinge dictionary and fits nonnegative coefficients to its own targets; zero coefficients may be pruned. Its dictionary seed is fixed by date, so repeated training-seed labels do not create independent random ridge representations. A quadratic projection is retained as a restricted comparator, not treated as the strongest conventional alternative.

4.2. All-state and continuous-action verification

At each state node, outward arithmetic encloses the actual network value and the action objective. In the coupled economy, convexity in actions and a positive strong-convexity constant yield a continuous-simplex primal–dual gap. A projected optimizer merely proposes the action; its stopping flag is not the certificate. State interpolation then extends the signed residual bounds off-grid. The nearest-node deployed actor receives a separate, uniform allowance for using its node action at the actual state.

Uniform innovations are integrated through positive-part moments, rather than replaced by a sampled innovation law. The elementary identity for a uniform variable U on $[-h, h]$ is

$$\mathbb{E}(y + U)_+^2 = \frac{(y + h)_+^3 - (y - h)_+^3}{6h}, \quad h > 0.$$

Successive integration gives the corresponding formula for independent coordinates and affine neuron arguments; zero coefficients are handled by the lower-dimensional formula. Interval evaluation encloses cancellation and rounding. The supplement specifies the mesh and action bounds used in the inherited verifier.

The residual account distinguishes a nodal enclosure from state coverage, actor error, and terminal error. A nodal enclosure contains the effects of representation, fitting, Bellman optimization, and arithmetic; it is not an identified decomposition into a pure approximation error and a pure optimizer error. In particular, subtracting a training loss from the residual does not estimate either component. The diagnostic reports every separately available allowance without counting nested subcomponents twice.

4.3. Work, dimension, and execution

For tensor state, action, and innovation covers of sizes K_x, K_a, m , a direct verifier pays work of order

$$\sum_{t < T} K_x K_a \{m(C_F + P_{t+1}) + C_\rho\} + \sum_{t \leq T} K_x P_t + W_{\text{storage}}, \quad (12)$$

where P_t counts network evaluation work. Actor storage is of order $TK_x d_a$. Analytic integration replaces m by the work of the moment formula, not by zero. With tensor covers, K_x scales as $h_x^{-d_x}$; analogous factors arise for action and innovation discretizations. A dimension-explicit certificate is therefore not a dimension-free algorithm.

If a Lipschitz innovation approximation contributes at most ζ_t to Q_t , widening residual endpoints by ζ_t and the actor allowance by $2\zeta_t$ adds $4 \sum_t B_t \zeta_t$ to (8). This is a discretization account, not permission to replace a continuous law without charging its error. Independent policy queries may have their own scenario-tree or simulation cost and are recorded separately.

The complete service includes own-future targets, fitting, unsuccessful refinement, action extraction, verification, applicable policy queries, serialization, and durable checkpoint storage. Initialization and outer-process time are retained separately. New local diagnostic clocks are not added to old GitHub-hosted clocks to manufacture an end-to-end estimate. All reported arithmetic certificates concern stored binary coefficients with outward enclosures. The inherited execution account treats rounded state acquisition and transition/cost evaluation separately from ideal policy loss.

5. DIRECT COMPARISONS OF ECONOMIC POLICY COSTS

A smaller upper bound on $J^\pi - V^*$ need not mean a smaller J^π . We therefore compare the neural and ridge policies directly at common states, using common innovations and the expectation version of (3).

For method m , let X^m follow its own fixed policy and define

$$Z_m = f_0^m(x_0) + \sum_{t < T} B_t \{Q_t(f_{t+1}^m; X_t^m, \pi_t^m(X_t^m)) - f_t^m(X_t^m)\} + B_T \{g(X_T^m) - f_T^m(X_T^m)\}. \quad (13)$$

This is a control-variate score for actual policy cost, not the difference of two regret bounds.

THEOREM 4—Paired economic score and support: *For fixed policies and integrable continuations under conditional expectation, $\mathbb{E}Z_m = J_0^{\pi^m}(x_0)$. If method m satisfies (7), then every trajectory satisfies*

$$L_m \leq Z_m \leq U_m, \quad (14)$$

$$L_m = f_0^m(x_0) - \sum_{t < T} B_t u_t^m - B_T u_T^m,$$

$$U_m = f_0^m(x_0) + \sum_{t < T} B_t (\eta_t^m - \ell_t^m) - B_T \ell_T^m.$$

Consequently $Z_N - Z_R$ has expectation $J^{\pi^N} - J^{\pi^R}$ and support $[L_N - U_R, U_N - L_R]$, whose width is $G_N + G_R$. These statements hold under any coupling of the two methods that preserves each policy's innovation law.

PROOF: Conditional expectation gives $\mathbb{E}Q_t(f_{t+1}^m; X_t^m, \pi_t^m) = \mathbb{E}c_t + \beta_t \mathbb{E}f_{t+1}^m(X_{t+1}^m)$. Summing telescopes all continuation terms. The selected-action gap is between zero and η_t^m , so each centered summand in (13) is between $-u_t^m$ and $\eta_t^m - \ell_t^m$. The terminal summand is between $-u_T^m$ and $-\ell_T^m$. Summing proves the support statement and its width. Q.E.D.

To evaluate the original continuous law, simulate uniform innovation bins and enclose all within-bin innovations. Interval propagation includes every actor that can be selected when a state enclosure crosses a nearest-node boundary. Suppose the

resulting endpoint pair satisfies $A_i \leq Z_{N,i} - Z_{R,i} \leq D_i$ and lies in a known support interval of width B . For independent draws of bins, the lower and upper endpoints are bounded independent observations. Let $\bar{A}, \bar{D}$ and s_A^2, s_D^2 be their sample means and variances. For q predeclared comparisons, set

$$r(s^2) = \sqrt{\frac{2s^2 \log(4q/\alpha)}{n}} + \frac{7B \log(4q/\alpha)}{3(n-1)}. \quad (15)$$

Applying the one-sided empirical Bernstein inequality to the lower and upper endpoints separately, then taking a union bound, gives simultaneous coverage at least $1 - \alpha$ for

$$[\bar{A} - r(s_A^2), \bar{D} + r(s_D^2)] \quad (16)$$

over all q comparisons. Mean, variance, logarithm, and square-root calculations are outward enclosed in the implementation; intersecting with the deterministic support preserves coverage. The proof and numerical details are in the supplement.

The coverage statement is conditional on frozen training outputs and an independent simulation model. Fixed PCG64 streams make this implementation reproducible; a fixed stream is not itself a proof of probabilistic independence. No policy is selected using these simulations. Intervals crossing zero are reported as unresolved, not as equivalence or as evidence of a superior neural policy.

6. INVESTMENT ECONOMIES AND THE COMPLETE EVIDENCE

6.1. Design and provenance

The original science design has twelve scalar neural services, twelve scalar spline services, six coupled neural services, six coupled ridge services, six coupled quadratic services, and 36 services for each of three precision rules. The successful source-bound execution contains all 150 records. The original first run, with four nonfinite-network exceptions, is retained. Its numerical inverse-softplus initialization was subsequently replaced by an algebraically equivalent stable expression; economic primitives, seeds, targets, budgets, and comparators were unchanged. We do not treat the failed first run as nonexistent or combine its clocks with the successful run.

The R44 protocol was committed before the additional diagnostics were executed. It is a post-review protocol, not an original preregistration. It selects the seven unresolved tight-target neural objects by their already observed outcomes and freezes their weights and policies. The paired comparison uses all six coupled pairs and all four originally declared states. A source manifest binds the original inputs, the diagnostic code, outcomes, and publication sources. Every original record and all five original science-source hashes were replayed.

6.2. Scalar and coupled primitives

The scalar experiment retains the original bounded investment state, price and risk cells, three declared training seeds, continuous feasible action interval, and the independent own-policy evaluator. Prices are one and four. Both methods construct their own continuations from the same primitive problem. The retained evidence and code specify every scalar primitive; the full economic derivation remains in the original investment section and applications.

In the coupled economy, $x \in [0, 1]^2$, $T = 4$, $\beta = 15/16$, and $A = \{a \geq 0 : a_1 + a_2 \leq 1/8\}$. With independent $z_j \sim U[-1/32, 1/32]$,

$$F_j(x, a, z) = \frac{1}{8} + \frac{1}{2}x_j + \frac{1}{8}x_{3-j} + \frac{1}{16}x_j(1 - x_{3-j}) + a_j + z_j.$$

This transition maps the stated state-action-innovation domain into the state domain. Define

$$s(x) = \sum_{j=1}^2 (x_j - 11/16)^2 + 2(1/2 - x_1 - x_2)_+^2 + \frac{1}{4}(x_1 - x_2)^4.$$

The stage cost is $s(x) + p \sum_j a_j^2 + \frac{1}{2}a_1a_2 + 4 \sum_j a_j^4$; terminal cost doubles the squared-target term in s and retains its remaining terms. Prices $p \in \{1, 4\}$ penalize costly investment. The shared simplex represents a common resource constraint, the transition couples the two capital stocks, and the state cost penalizes aggregate shortage and imbalance. This is a stylized controlled economy, not an empirically calibrated welfare exercise.

For both coupled methods the final all-state verification ladder is $N = 16, 32, 64, 128$. The action objective has strong-convexity lower bound

$2p - 1/2 > 0$. Uniform positive-part moments integrate continuous uncertainty analytically. The finest original actor has $4(129)^2 \cdot 2 = 133,128$ stored scalar actions. This actor cost is distinct from continuation-network storage.

6.3. Original attainment and comparative outcomes

TABLE I

ORIGINAL CAPPED-SERVICE TARGET ATTAINMENT

Method	All-state target	Certified
Scalar neural	0.06	12/12
Scalar neural	0.04	6/12
Scalar spline	0.06	12/12
Scalar spline	0.04	12/12
Coupled neural	0.25	6/6
Coupled neural	0.10	5/6
Coupled ridge	0.25	6/6
Coupled ridge	0.10	6/6
Coupled quadratic	0.25	0/6
Coupled quadratic	0.10	0/6

Table I reports the original capped-service outcomes. The original scalar neural bound is larger than its matched spline bound in every one of the twelve pairs. At the first common 0.06 crossing, 29 of 48 signed neural-minus-spline intervals have positive lower endpoints, none has a negative upper endpoint, and 19 contain zero. At 0.04, the six common pairs give 17 positive intervals, none negative, and seven overlaps. All lie within the prespecified descriptive margin ± 0.005 . This establishes numerical closeness at those states but is not an equivalence design; where signs are identified, they favor the spline.

The coupled ridge reaches both targets in all six services. It has the smaller final gap in four pairs and the neural method in two. At 0.25 the ridge reaches the target faster in all six original observations. At 0.10 it is faster in four of the five pairs in which both pass, while the neural method is faster in one and fails in the

remaining pair. These singleton nonlinear timing observations are descriptive. The ridge continuation has storage comparable to the neural continuation, so its strength cannot be dismissed as a comparison with an artificially restricted quadratic.

6.4. Diagnosing tight-target failures without changing policies

TABLE II
ADDITIONAL VERIFICATION OF UNCHANGED NEURAL POLICIES

Service	Original bound	Refined bound	Target	Attained
041	0.04092128	0.03068472	0.04	Yes
089	0.04118465	0.03090917	0.04	Yes
120	0.04112077	0.03085353	0.04	Yes
124	0.04093160	0.03070697	0.04	Yes
128	0.15223570	0.09775755	0.10	Yes
145	0.04136567	0.03108538	0.04	Yes
147	0.04101560	0.03077643	0.04	Yes

For the six unresolved scalar policies, only the residual state grid is refined from $N = 2048$ to $N = 4096$, retaining the original action verification resolution $A = 1024$. For the coupled policy, residual verification uses $N = 256$ instead of 128. The original actor allowances and terminal enclosures remain in (10); all original action arrays and networks are hash-checked unchanged. Fine-grid actor proposals generated internally are discarded, not deployed.

Table II gives all seven results. The retained scalar policy bounds become approximately 0.03068–0.03109, below 0.04. The retained coupled policy bound becomes 0.09775756, below 0.10. There are zero training updates and zero newly stored deployed-action scalars. The coupled diagnostic still pays for a fine verification mesh with 66,049 states, including transient fine-grid proposals; retaining the original actor does not make verification free.

The original coupled decomposition is informative. Its discounted state-cover contribution to nonterminal residual widths is about 0.07268224, its actor allowance is 0.03287249, its remaining nodal residual enclosure is 0.04574738, and its termi-

nal contribution is 0.00093360. These sum to the original bound up to displayed rounding. The extra mesh quarters the principal second-order coverage term while retaining the actor and terminal terms. The resulting crossing does not establish that representation and optimization errors are separately small; it establishes the required full-policy bound for the unchanged policy.

The diagnostic answers a specific method question. The original capped service was less reliable at its tight target, but the saved neural objects were not necessarily inadequate policies. All seven selected policies admit the tighter certificates when verification is extended under a fixed rule. This finding does not alter the original attainment table or justify a comparison between a selectively refined neural certificate and an unrefined comparator as evidence of a work advantage.

6.5. *Direct coupled economic differences*

TABLE III
DIRECT NEURAL-MINUS-RIDGE COST COMPARISONS, SIMULTANEOUS 99 PERCENT INTERVALS

Price	Comparisons	Positive	Negative	Unresolved
1	12	0	0	12
4	12	0	0	12

For each of the six original neural/ridge pairs, select each method's first 0.25 crossing. Compare at $(1/8, 1/8)$, $(1/4, 1/2)$, $(1/2, 1/4)$, and $(3/4, 3/4)$. Each comparison uses 65,536 common innovation-bin trajectories, 40-bit uniform bins, and the own-policy score in (13). The total family error is 0.01 across 48 one-sided bounds. The complete 24 intervals, endpoint moments, seeds, ambiguity counts, and support bounds appear in the supplement and machine-readable records.

All 24 direct confidence intervals contain zero. Thus the coupled evidence now addresses the actual neural-minus-ridge policy-cost difference, but does not rank the methods. This is stronger information than subtracting regret bounds and different information from a smaller all-state certificate. The simultaneous interpretation pertains to the declared states and fixed policies under the simulation model, not

to every state in the compact domain. The all-state policy guarantees remain the separate deterministic statements in Table I and Theorem 2.

6.6. Precision and complete recorded work

TABLE IV
ORIGINAL REPEATED PRECISION CATALOGUE

Method	Certified	Accepted32	Rejected32	Accepted64
Fixed32	6/36	147	30	0
Fixed64	36/36	0	0	474
Adaptive	36/36	348	126	126

The precision experiment repeats twelve dimension/conditioning/target objects three times. Fixed binary64 and adaptive each certify all 36 services; fixed binary32 certifies six. Fixed64 and adaptive each make 474 accepted updates. Adaptive makes 348 accepted binary32 updates, 126 rejected binary32 proposals, and 126 accepted binary64 updates. Every rejection is charged.

Adaptive is slower in 26 of the 36 matched recorded services and faster in ten; aggregate recorded time is about 0.98% larger. The controller’s activation therefore demonstrates its numerical mechanism, not a complete-work saving in this implementation. This result does not invalidate the conditional numerical acceptance theorem. It prevents the theorem from being mistaken for an empirical efficiency conclusion.

6.7. What the evidence identifies

The evidence identifies a finite from-primitives nonlinear construction, all-state certificates for its returned policies, a useful retained-policy verification refinement, and direct economic comparisons under the declared state and innovation design. It also identifies the continued strength of conventional comparators. It does not identify a general neural speed advantage, a nonlinear high-dimensional frontier, or a population probability of optimizer success.

An adaptive sparse-grid or adaptive-partition comparator has not been executed for the original neural optimizers. Section 15 adds a separate error-driven coordinate-FVI comparator and three- and four-dimensional stress catalogue for the constructive backend; neither retroactively changes the optimizer study. Neither has a nonlinear dimension frontier or repeated isolated timing of every original scalar and coupled service. Section 7.3 adds a separate predeclared horizon/price catalogue with repeated isolated timing for the deterministic constructive backend and spline; it does not replace the original experiments. The retained dimension-explicit analysis exposes the tensor factors those studies would have to confront, rather than treating a two-state experiment as their substitute. The general economic framework and conditional constructive results remain in force; the missing experiments concern comparative evidence for broader implementation claims.

7. A CONSTRUCTIVE NEURAL SERVICE WITH A PRIMITIVE ERROR BUDGET

The policy theorem is deliberately independent of the representation of the continuation. A construction theorem must do more: it must produce a representation, feasible actions, and verified residuals from economic primitives. This section supplies such a theorem for a deterministic ReLU backend. It does not assume convergence of the nonlinear optimizers used in the earlier experiments. Its distinguishing property is that the continuation has a prescribed Lipschitz constant even when its Bellman labels are inexact.

7.1. Primitive conditions and the neural architecture

Consider the compact Bellman economy of Section 2, with a common action set $A \subset \mathbb{R}^a$ and state set $X \subset \mathbb{R}^d$. All distances in this section use the ℓ^1 norm. The innovation law and the monotone, cash-invariant functional ρ_t do not depend on the state or action. The transition remains in X . Suppose known nonnegative constants satisfy, uniformly in the innovation,

$$|c_t(x, a) - c_t(y, b)| \leq C_{x,t} \|x - y\|_1 + C_{a,t} \|a - b\|_1, \quad (17)$$

$$\|F_t(x, a, z) - F_t(y, b, z)\|_1 \leq F_{x,t} \|x - y\|_1 + F_{a,t} \|a - b\|_1. \quad (18)$$

The terminal cost is L_g -Lipschitz. These assumptions are additional conditions on the constructive backend, not a replacement for the more general feasible correspondences in the original economic formulation. A state-dependent feasible correspondence requires its own feasible-cover and transport constants.

We require an effective numerical description of the primitives: certified finite state and action covers can be obtained at every positive radius; their nodes are feasible; and a numerical evaluator can enclose each Bellman query to any requested positive tolerance in finite work. The last condition includes evaluation of ρ_t , not just evaluation of the network. For expectation with a compact innovation domain, known cell probabilities and a certified innovation cover provide one such evaluator. If F_t is $F_{z,t}$ -Lipschitz in the innovation, replacing each cell by a representative at distance at most h_z contributes at most $\beta_t L_{t+1} F_{z,t} h_z$. Arithmetic, probability-mass approximation, and any remaining integration error must be added. The experiment below instead integrates a piecewise-linear network analytically under the original continuous uniform law.

Set

$$L_T = L_g, \quad L_t = C_{x,t} + \beta_t F_{x,t} L_{t+1}, \quad L_t^a = C_{a,t} + \beta_t F_{a,t} L_{t+1}. \quad (19)$$

Choose state nodes $x_{t,i}$ of covering radius h_t and action nodes $a_{t,j}$ of covering radius k_t . Given the continuation already constructed at date $t + 1$, compute numbers $\hat{q}_{t,ij}$ satisfying

$$|\hat{q}_{t,ij} - Q_t(f_{t+1}; x_{t,i}, a_{t,j})| \leq e_t. \quad (20)$$

Define $y_{t,i} = \min_j \hat{q}_{t,ij}$, choose a minimizing action index $j(t, i)$ with a fixed tie rule, and construct

$$f_t(x) = \min_i \{y_{t,i} + L_t \|x - x_{t,i}\|_1\}. \quad (21)$$

At date T , use labels within e_T of g and the same formula with constant L_g . At any other date, deploy the action $a_{t,j(t,i(x))}$ attached to a nearest state node, again with a fixed tie rule. The action is feasible by construction. No future optimal value or conventional comparator is used as a label.

Formula (21) is an exact ReLU construction, not an appeal to a universal approximation theorem. Absolute values use $|u| = \max\{u, 0\} + \max\{-u, 0\}$; minima use $\min\{u, v\} = u - \max\{u - v, 0\}$. A balanced tree of these gates implements the finite minimum. With K_t state nodes, the resulting sparse affine-ReLU circuit has $O(dK_t)$ coefficients and $O(1 + \log K_t)$ depth with unrestricted affine fan-in. Signed intermediate values can be carried by their positive and negative parts. The construction is related to the classical Lipschitz extension formula of McShane (1934); neither that formula nor its elementary ReLU realization is claimed as new. Here its role is to close an own-future dynamic construction, including the deployed actor and every numerical allowance.

THEOREM 5—Constructive own-future neural policy: *Under the preceding conditions, the network (21) is L_t -Lipschitz at every date, independently of the numerical label errors. Its residual and its deployed actor satisfy*

$$-e_t \leq f_t - T_t f_{t+1} \leq L_t^a k_t + e_t + 2L_t h_t, \quad (22)$$

$$Q_t(f_{t+1}; x, \pi_t(x)) - T_t f_{t+1}(x) \leq L_t^a k_t + 2e_t + 2L_t h_t. \quad (23)$$

At the terminal date, $-e_T \leq f_T - g \leq e_T + 2L_g h_T$. Consequently the cost of the returned policy, relative to the optimum over all admissible adapted policies, is uniformly bounded by

$$\mathcal{G}_{\text{con}} = \sum_{t < T} B_t \{2L_t^a k_t + 4e_t + 4L_t h_t\} + B_T \{2e_T + 2L_g h_T\}. \quad (24)$$

For fixed nodes and L_t , changing every label by at most δ changes the represented function by at most δ in the uniform norm; it does not change the Lipschitz bound.

PROOF: Monotonicity and cash invariance imply nonexpansiveness in the uniform norm. Thus $Q_t(f_{t+1}; \cdot, a)$ and $T_t f_{t+1}$ are L_t -Lipschitz, while the action modulus is L_t^a . At a node, minimization of the numerical queries gives $-e_t \leq y_{t,i} - T_t f_{t+1}(x_{t,i}) \leq L_t^a k_t + e_t$. Each cone in (21) is L_t -Lipschitz, as is their finite minimum. The lower residual bound follows from the triangle inequality for every cone; the upper bound uses a nearest node twice, once for the cone and once for the Bellman value. At that node the selected action is within $L_t^a k_t + 2e_t$ of the

continuous-action infimum. Transport to x costs at most $2L_th_t$. The terminal argument is identical without action minimization. Substitution of these widths and actor allowances into Theorem 2 gives (24). Finally, minima of two arrays whose entries differ by at most δ differ by at most δ . The supplement gives the construction, measurability, and work details. Q.E.D.

COROLLARY 6—Deterministic termination at an economic tolerance: *Let $\varepsilon > 0$, $H = \sum_{t=0}^T B_t$, and $s = \varepsilon/H$. At nonterminal dates choose*

$$e_t \leq s/8, \quad L_t^a k_t \leq s/8, \quad L_t h_t \leq s/16. \quad (25)$$

At the terminal date choose $e_T \leq s/4$ and $L_g h_T \leq s/4$. Terms with zero coefficient impose no restriction. Then the finite constructive service returns a feasible neural policy with loss at most ε . No probability law over initializations and no optimizer stopping criterion is needed.

The result follows because each discounted bracket in (24) is at most s . Effective compactness and the numerical evaluator give finite work for these choices. Arbitrarily small tolerances may require arbitrary precision; the corollary does not promise termination in fixed binary64. A finite implementation ladder can still return a recorded failure, even though the unbounded constructive algorithm terminates under the stated hypotheses.

The earlier scalar constructive result in the retained paper implemented piecewise-linear interpolation as a ReLU network. Its sufficient allocation controlled a label-error-to-mesh ratio in the propagated slope. Formula (21) works in arbitrary finite state dimension and enforces the slope bound directly. In particular, label and action allowances can be allocated at order ε without a separate slope-preservation requirement that those allowances be quadratic in the state spacing. This is an improvement of that sufficient proof budget, not a lower bound for all spline methods or evidence that conventional Lipschitz projection is unavailable.

7.2. What the construction costs

The construction does not eliminate state coverage. With K_t state nodes, J_t action nodes, M_t innovation representatives, and network evaluation cost P_{t+1} , direct target formation requires

$$\sum_{t < T} K_t J_t \{M_t (C_{F,t} + P_{t+1}) + C_{\rho,t} + C_{c,t}\} + O\left(\sum_{t=0}^T d K_t\right) \quad (26)$$

arithmetic operations under the indicated evaluation model. $C_{\rho,t}$ includes the finite-risk functional and probability processing. Certified integration can replace the innovation sum but its own cost must be charged. Bit lengths, adaptive arithmetic, and nearest-node policy queries are additional explicit costs, not constants suppressed by a dimension claim. A rectangular state cover with radius h_t can require order h_t^{-d} nodes; a rectangular action cover can require order k_t^{-a} actions. The theorem is therefore a dimension-explicit construction, not an assertion that a nonlinear high-dimensional frontier has been executed or that the curse of dimensionality has been removed.

In one state dimension, the finite envelope can be compiled exactly into a piecewise-linear representation. Two passes compute the greatest L -Lipschitz minorant of the node labels. Between neighboring nodes the envelope is the minimum of the two adjacent affine cones. At most one additional intersection is needed per interval. Rational knots and values permit exact antiderivatives; an interval evaluator encloses their numerical evaluation. This compilation evaluates the defining network exactly and does not substitute labels from a separately solved Bellman problem. The spline comparator receives the same exact-integration acceleration, so it is not charged for an artificial dense neural evaluation loop.

7.3. A predeclared matched service catalogue

The new experiment tests complete services, not only previously failed neural candidates. It uses a further nonlinear investment laboratory with $x \in [0, 1]$, $a \in [0, 1/4]$, $\beta = 15/16$, and independent z uniform on $[-1/32, 1/32]$:

$$F(x, a, z) = \frac{1}{32} + \frac{11}{16}x + \frac{1}{16}x(1 - x) + a + z, \quad (27)$$

$$c_p(x, a) = (x - \frac{11}{16})^2 + pa^2 + 4a^4 + 2(\frac{3}{8} - x)_+^2, \quad (28)$$

$$g(x) = 2(x - \frac{11}{16})^2 + 2(\frac{3}{8} - x)_+^2. \quad (29)$$

The invariant domain follows from monotonicity in x, a, z and the two extreme transitions 0 and 1. Valid primitive state moduli are $C_x = 23/8$, $F_x = 3/4$, and $L_g = 17/4$; action moduli are $C_a = p/2 + 1/4$ and $F_a = 1$. The downside term penalizes a low capital stock, and the quartic investment cost discourages rapid adjustment. This is a numerical economic laboratory, not an estimated or calibrated quantitative exercise.

The complete design was fixed before execution: horizons $\{2, 4\}$, prices $\{1, 4\}$, two methods, three isolated repetitions, and the full common state/action ladder $N = A \in \{16, 32, 64, 128, 256, 512\}$. The targets are 0.5, 0.25, and 0.125 in units of discounted cost. These are policy-loss tolerances, not equivalence margins for a neural-minus-spline contrast. Every rung restarts from terminal primitives and constructs its own future backward. The neural method uses (21); the spline uses linear interpolation of its own labels and propagates its actual exact slope bound. Both evaluate the same continuous innovation law analytically and round labels to dyadic multiples of 2^{-40} . All numerical query errors are outwardly enclosed and enter the certificate.

All rungs are retained, including those preceding the first target crossing. The clock to a target includes every preceding rung, target formation, rational compilation, actor construction, verification, and durable checkpoint output. Fresh processes run sequentially on one selected CPU with one numerical-library thread and an identical non-economic warm-up. Imports and warm-up are excluded from the internal clock but included in a separate process clock. CPU frequency is not fixed. Independent policy-cost comparison is not performed in this catalogue and is not assigned a zero cost. Three timing repetitions describe local dispersion, not a population law for neural optimization.

The neural backend certifies all twelve services at 0.5 and 0.25, and nine of twelve at 0.125. The spline certifies all twelve at every target. The three neural failures are repetitions of the same horizon-four, price-four economy: the final gap is approximately 0.12503, just above 0.125. They remain failures rather than being rounded

TABLE V

FIRST CERTIFIED STATE / ACTION RESOLUTION IN THE COMPLETE CONSTRUCTIVE CATALOGUE

Generator	T	p	0.5	0.25	0.125	Final gap
ReLU envelope	2	1	64	128	256	0.06132
Spline	2	1	64	128	256	0.05619
ReLU envelope	2	4	128	256	512	0.06274
Spline	2	4	64	128	256	0.05799
ReLU envelope	4	1	128	256	512	0.12236
Spline	4	1	128	256	512	0.09711
ReLU envelope	4	4	256	512	–	0.12503
Spline	4	4	128	256	512	0.10221

Entries are first qualifying $N = A$ on the fixed ladder. Each row represents three identical deterministic policy/certificate arrays, not three independent economic tasks. A dash is a retained failure at the declared cap $N = 512$. Final gaps use the cap for both methods.

TABLE VI

COMPLETE PREFIX WORK AT TARGET 0.125, OR AT THE FAILED CAP

Generator	T, p	N	Queries	Median s	Range s	Exit
ReLU envelope	2,1	256	176,586	0.3009	[0.2982, 0.3094]	target
Spline	2,1	256	176,586	0.2144	[0.2136, 0.2365]	target
ReLU envelope	2,4	512	702,924	0.8450	[0.8354, 0.8472]	target
Spline	2,4	256	176,586	0.2151	[0.2119, 0.2153]	target
ReLU envelope	4,1	512	1,405,848	1.5540	[1.5497, 1.5626]	target
Spline	4,1	512	1,405,848	1.2654	[1.2341, 1.2745]	target
ReLU envelope	4,4	512	1,405,848	1.5495	[1.4602, 1.5506]	cap
Spline	4,4	512	1,405,848	1.2705	[1.2625, 1.2810]	target

Queries sum all Bellman evaluations on earlier failed rungs and the reported rung. Clocks include construction, verification and durable checkpoints. Ranges contain the three isolated observations. A cap row is not time to successful certification. Independent economic comparison is not included.

into attainment. Across the 33 common successful method/target/repetition comparisons, the spline reaches the target earlier in every recorded observation. Thus the new execution establishes reproducible complete construction and a transparent frontier, not neural cost superiority. The deterministic termination corollary concerns its specified sufficient budget; the predeclared finite ladder need not contain

that budget. The finite catalogue measures the deterministic backend just specified. It neither overwrites the original neural optimizer’s attainment array nor turns the seven retained-policy diagnostics into prospective observations. The theorem supplies a method-level termination statement for this backend; the recorded services establish only the achieved targets, costs, and storage on the declared economies. A smaller certificate is not a lower realized policy cost. The new numerical comparison and the earlier direct policy-cost comparisons therefore remain distinct throughout.

8. NEURAL CONTINUATION WITH AN ACTION WITNESS

A constructed continuation contains information about the action that produced each Bellman label. Discarding that information and subsequently controlling an unrelated actor can charge the same approximation twice. We now retain the label’s feasible action through the minimum circuit. This changes the returned policy, not merely the description of an old certificate. The construction continues to use the original economic primitives and its own fitted future.

8.1. A one-sided policy sandwich

Write $\mathcal{T}_t^\pi f_{t+1}(x) = Q_t(f_{t+1}; x, \pi_t(x))$. Suppose that, uniformly on the invariant domain,

$$f_t - \mathcal{T}_t f_{t+1} \leq u_t, \quad \mathcal{T}_t^\pi f_{t+1} - f_t \leq d_t, \quad \ell_T \leq f_T - g \leq u_T. \quad (30)$$

The quantities u_t and d_t are signed upper bounds. Their separate absolute values need not enter a policy-loss account.

PROPOSITION 7—One-sided policy sandwich: *Under the monotonicity, cash-invariance and Bellman comparison conditions of Section 2, inequalities (30) imply*

$$0 \leq J_0^\pi(x) - V_0^*(x) \leq \sum_{t < T} B_t(u_t + d_t) + B_T(u_T - \ell_T). \quad (31)$$

The bound is invariant to additive date-specific shifts of the continuation. It holds for the stated recursive evaluations as well as expectation; no simulation identity is required.

PROOF: Backward order comparison gives $V_0^* \geq f_0 - \sum_{t < T} B_t u_t - B_T u_T$ and $J_0^\pi \leq f_0 + \sum_{t < T} B_t d_t - B_T \ell_T$. Subtraction proves the result. A shift $f_t \mapsto f_t + b_t$ adds $b_t - \beta_t b_{t+1}$ to u_t and subtracts it from d_t , leaving their sum unchanged. *Q.E.D.*

The centered theorem remains valid: its choice $d_t = \eta_t - \ell_t$ recovers (8). The new proposition avoids that indirect bound when the selected-policy residual can be certified jointly with the continuation. It is an order comparison, not a new claim of neural superiority.

8.2. Retaining the minimizing cone's action

Use the primitive conditions, grids, query tolerances and labels of Section 7. Let $a_{t,i}$ denote the action whose numerical query attains $y_{t,i}$. In place of the nearest-node actor, return

$$i_t(x) \in \arg \min_i \{y_{t,i} + L_t \|x - x_{t,i}\|_1\}, \quad \pi_t^w(x) = a_{t,i_t(x)}. \quad (32)$$

A fixed tie convention makes this finite selector measurable. The common action set makes each selected action feasible, even when the attaining cone is not centered at a nearest node. The critic remains the exact affine-ReLU circuit. The actor additionally uses the circuit's comparison index; a discontinuous selector is not asserted to be a continuous ReLU output.

THEOREM 8—Witness-preserving neural construction: *Under the conditions of Theorem 5, the policy (32) satisfies*

$$Q_t(f_{t+1}; x, \pi_t^w(x)) - f_t(x) \leq e_t. \quad (33)$$

Consequently its all-state loss relative to the optimum over all admissible adapted policies is at most

$$\mathcal{G}_w = \sum_{t < T} B_t \{L_t^a k_t + 2e_t + 2L_t h_t\} + B_T \{2e_T + 2L_g h_T\}. \quad (34)$$

The nonterminal allowance is half the separated sufficient allowance in (24); the terminal account is unchanged. This comparison of certified bounds does not assert that the two policies have ordered realized costs.

PROOF: For the attaining index $i = i_t(x)$, the numerical label and its action obey $Q_t(f_{t+1}; x_{t,i}, a_{t,i}) \leq y_{t,i} + e_t$. Transporting this same action to x costs at most $L_t \|x - x_{t,i}\|_1$. Hence its objective is at most $f_t(x) + e_t$. The unchanged optimal-residual upper bound is $u_t = L_t^a k_t + e_t + 2L_t h_t$. Apply Proposition 7 with $d_t = e_t$ and the unchanged terminal enclosure. Q.E.D.

The label, the cone and the selected action are one mathematical object in this argument. Retaining only a Lipschitz-regularized label value without an attaining original action index would not justify (33). The supplement proves that the scalar compiler preserves precisely this information, including dominated labels, rational intersections and ties.

COROLLARY 9—Resource and execution allowances: Let $s = \varepsilon / \sum_{t=0}^T B_t$. At non-terminal dates it suffices to choose

$$e_t \leq s/4, \quad L_t^a k_t \leq s/4, \quad L_t h_t \leq s/8,$$

and retain the terminal allocation of Corollary 6. Under the same effective-computability hypotheses the construction terminates with loss at most ε .

If the chosen cone exceeds the minimum by at most v_t , state acquisition error is at most r_t in the stated norm, and a feasible executed action is within α_t of that cone's action, replace e_t in (33) by

$$e_t + v_t + 2L_t r_t + L_t^a \alpha_t.$$

Add the discounted execution terms to (34).

The relaxed allocation is a sufficient-budget improvement relative to the earlier proof, not a complexity lower bound for conventional methods. A tensor cover still incurs its dimension factors. Exact state acquisition is the reference experiment; unchecked lookup using rounded switchpoints does not inherit an exact-index guarantee.

8.3. *A strengthened conventional comparison*

For either the original nearest-node cone actor or the spline actor, let C_i be the node's closed nearest-neighbor cell and let f_t be that method's own continuation. The same selected-label argument gives the computable bound

$$d_t = e_t + D_t, \quad D_t = \max_i \sup_{x \in C_i} \{y_{t,i} + L_t \|x - x_{t,i}\|_1 - f_t(x)\}. \quad (35)$$

On the scalar domain, all terms are piecewise affine. Their maximum is attained among the continuation knots, state nodes and cell boundaries; both adjacent actors are included at a boundary. The exact rational maximum supplies a valid all-state allowance without an additional verification mesh. Both conventional interpolation and the nearest-node neural ablation receive this improved account. Thus the experiment does not contrast the new bound only against a deliberately loose baseline. The old separated bound remains recorded, and the smaller of two bounds for the same returned policy may be used.

8.4. *Matched construction, policy and certificate frontiers*

The following matched catalogue is the frozen R46 experiment. Its statements about direct policy cost refer to the evidence then available; Section 13 supplies the new actual-policy comparison, and Section 12 executes the state-dependent construction. The design was fixed in the repository before execution. It retains the nonlinear continuous-innovation economy (27), horizons 2, 4, prices 1, 4, and all six resolutions 16, 32, 64, 128, 256, 512. Three methods and three isolated timing repetitions produce 36 services and 216 rungs. Each rung restarts from primitives. The two cone methods have identical critics and node-action arrays; their deployed selectors differ. The spline constructs its own labels. All methods receive the same exact integration and outward numerical arithmetic.

The declared targets are $1/4$, $1/8$ and $1/16$. Every unsuccessful rung and every capped failure is retained. Internal prefix clocks include target construction, witness or nearest-cell compilation, certificates and durable checkpoints; process clocks additionally include imports and the common non-economic warm-up. Runs use one selected CPU and one numerical-library thread, with method order rotated

across repetitions and economic cells. Frequency is not controlled. These are local measurements, not a hardware-independent work theorem. Primitive counts and rational storage are recorded separately from clocks, and are not called complete floating-point or bit-operation counts.

TABLE VII

FINAL ALL-STATE POLICY-LOSS BOUNDS AT $N = 512$

T	p	Witness neural	Nearest neural	Spline
2	1	0.034308	0.050092	0.039641
2	4	0.035017	0.051637	0.040957
4	1	0.064387	0.093593	0.066369
4	4	0.065720	0.097039	0.069979

Bounds are rounded upward in the records; displayed decimals are summaries, not the target test. All three repetitions have identical checkpoints. Both nearest-node comparators use the strengthened one-sided certificate.

All three methods certify 12/12 services at 1/4 and at 1/8, and 6/12 at 1/16. The latter failures are all horizon-four services and are retained. The witness bound is strictly tighter than the nearest-node neural ablation in all 24 distinct economy/resolution comparisons, and tighter than the strengthened spline in 23 of 24. The exception is $T = 4$, $p = 1$, $N = 16$, where the witness and spline bounds are approximately 2.06039 and 2.03515. The declared dyadic targets have the same first-crossing resolution for all three methods in every economic cell. Thus the certificate improvement does not yield an earlier first crossing on this particular ladder. The spline is faster in all 30 common successful target/repetition clock comparisons; the witness construction is faster than the nearest-node neural ablation in all 30. These local clocks and tighter regret bounds establish neither policy-cost superiority nor a general neural speed advantage. The comparisons isolate the effect of carrying the action witness and of using a joint policy certificate. They do not reclassify earlier capped services or claim that the new selector was previously deployed. In particular, an improvement shared by the strengthened nearest-node ablation cannot be attributed solely to witness selection. A tighter regret certificate does not

identify the sign of a direct policy-cost contrast. The existing direct expected-cost intervals and adverse precision findings remain unchanged.

9. STATE-DEPENDENT CONSTRAINTS AND FEASIBLE WITNESSES

Investment opportunities often depend on current capital, collateral or remaining capacity. The common-action construction above does not by itself authorize transporting a node's action to a state where that action is inadmissible. We now incorporate a feasible repair into the same neural Bellman construction. The original controlled economy, rather than a relaxed control problem, remains the comparison object.

Use the invariant state domain, numerical-query conditions and common evaluation functional of Section 7. Replace its common action set by nonempty compact sets $A_t(x)$ satisfying

$$d_H(A_t(x), A_t(z)) \leq K_t^A \|x - z\|_1, \quad (36)$$

where d_H is Hausdorff distance in the action norm. The primitive inequalities (17) are now required for pairs of feasible state-action points; neither costs nor dynamics need to be evaluated at infeasible actions. Assume effective finite k_t -nets of $A_t(x_{t,i})$ and a feasible Borel repair $\Gamma_{t,i}(x) \in A_t(x)$ of each stored action $a_{t,i}$, with

$$\|\Gamma_{t,i}(x) - a_{t,i}\|_1 \leq K_t^A \|x - x_{t,i}\|_1 + \zeta_t. \quad (37)$$

The nonnegative ζ_t is a certified displacement allowance, not permission to violate feasibility. Exact nearest-point selection gives $\zeta_t = 0$ when available. The effective repair and feasible-net procedures are additional computational inputs, whose work must be charged.

Starting with $L_T = L_g$, define

$$A_t^Q = C_{x,t} + \beta_t F_{x,t} L_{t+1}, \quad (38)$$

$$D_t^Q = C_{a,t} + \beta_t F_{a,t} L_{t+1}, \quad (39)$$

$$L_t = A_t^Q + D_t^Q K_t^A.$$

At each state node, minimize the certified queries over its own feasible action net. Retain the original minimizing action with its numerical label $y_{t,i}$. Form the same cone network $f_t(x) = \min_i \{y_{t,i} + L_t \|x - x_{t,i}\|_1\}$. Let $i_t(x)$ attain this minimum and return $\pi_t(x) = \Gamma_{t,i_t(x)}(x)$.

THEOREM 10—Feasible-witness transport: *Under these conditions and the Bellman comparison hypotheses of Section 2, the constructed continuation is L_t -Lipschitz, and the returned policy is feasible and measurable. Uniformly on the original state domain,*

$$-e_t \leq f_t - \mathcal{T}_t f_{t+1} \leq D_t^Q k_t + e_t + 2L_t h_t, \quad (40)$$

$$\mathcal{T}_t^\pi f_{t+1} - f_t \leq e_t + D_t^Q \zeta_t.$$

Consequently its loss relative to the optimum over all admissible adapted policies is at most

$$\sum_{t < T} B_t \{D_t^Q k_t + 2e_t + 2L_t h_t + D_t^Q \zeta_t\} + B_T (2e_T + 2L_T h_T). \quad (41)$$

No optimal continuation is supplied to the label generator. The bound concerns the policy after feasible repair, not the possibly infeasible array of unmodified node actions.

PROOF: The primitive inequalities imply a modulus A_t^Q in the state and D_t^Q in the action for feasible pairs. Given an action at x , the Hausdorff condition supplies an action at z within $K_t^A \|x - z\|_1$. Comparing the two objectives and reversing the roles of x, z shows that $\mathcal{T}_t f_{t+1}$ is L_t -Lipschitz. The feasible node net therefore gives $-e_t \leq y_{t,i} - \mathcal{T}_t f_{t+1}(x_{t,i}) \leq D_t^Q k_t + e_t$. The cone argument proves the first enclosure. For the attaining index, transport between the two feasible pairs and (37) give

$$Q_t(f_{t+1}; x, \Gamma_{t,i}(x)) \leq y_{t,i} + e_t + L_t \|x - x_{t,i}\|_1 + D_t^Q \zeta_t.$$

The cone equals $f_t(x)$. Apply Proposition 7 and the unchanged terminal enclosure.

Q.E.D.

This result closes a constructive route for a class of endogenous constraints within the original model. It does not assume that all feasible correspondences have a finite Hausdorff modulus. With $s = \varepsilon / \sum_{t=0}^T B_t$, sufficient nonterminal allocations are $e_t \leq s/8$, $D_t^Q k_t \leq s/4$, $L_t h_t \leq s/8$ and $D_t^Q \zeta_t \leq s/4$. Together with the earlier

terminal allocation, they imply finite termination at tolerance ε when all listed operations are effective. Zero coefficients impose no restriction. When exact repair is available, the larger query allowance of Corollary 9 is again sufficient. Setting $K_t^A = \zeta_t = 0$ recovers the common-action witness theorem.

A capacity-constrained investment rule. For two nonnegative state coordinates on $[0, 1]^2$, let $A(x) = [0, (x_1 + x_2)/2]$. This is a state-dependent capital-capacity constraint with $K^A = 1/2$ in the ℓ^1 state norm. The repair $\Gamma_i(x) = \min\{a_i, (x_1 + x_2)/2\}$ is feasible, satisfies (37) with $\zeta = 0$, and has an elementary affine-ReLU representation for each fixed witness. Thus changing capacity changes the actor and the constructive state modulus through an explicit term $D^Q/2$. The index selector remains a finite comparison procedure. The supplement provides exact rational checks for this example; these are checks of the theorem and admissibility contract, not an additional multidimensional performance study.

Neither feasible-net construction nor repair is free. With $J_{t,i}$ node actions, target-query work is summed over $\sum_i J_{t,i}$ rather than a common $K_t J_t$. Each deployed query also incurs the selected repair's work and storage. Under imperfect state acquisition, repair must be certified feasible at the true state, or throughout the acquisition set for a robust implementation. An error allowance for the objective cannot substitute for that feasibility condition.

10. FROM ACQUIRED STATES TO FEASIBLE POLICIES

A policy is implemented at an acquired state, not at an exact real number available without cost. This distinction matters for a witness selector: a small change in the state can change its index discontinuously. The appropriate object is therefore an objective allowance for the selected witness together with feasibility at the true state. A Lipschitz assumption on the actor would replace, rather than solve, this problem.

At date t , an observation $\hat{x}$ specifies a nonempty acquisition set $\mathcal{D}_t(\hat{x}) \subset X$ containing the true state, with $\|x - \hat{x}\|_1 \leq r_t$ for every $x \in \mathcal{D}_t(\hat{x})$. Use the feasible-witness data of Section 9. Suppose the numerical selector returns an index i such that

$$y_{t,i} + L_t \|\hat{x} - x_{t,i}\|_1 \leq f_t(\hat{x}) + v_t. \quad (42)$$

The implemented action $\hat{a}_t(\hat{x})$ must belong to $A_t(x)$ for every x in the acquisition set. Assume, on that entire set,

$$\|\hat{a}_t(\hat{x}) - a_{t,i}\|_1 \leq K_t^A \|x - x_{t,i}\|_1 + \zeta_t^{\text{acq}}. \quad (43)$$

The observation and selection rules are measurable and causal. Enlarging the history by these observations leaves the stipulated conditional innovation evaluation unchanged; otherwise the full-information comparator must be defined on the enlarged information structure. The allowances are uniform, not fitted to a favorable sample of observation errors.

THEOREM 11—Acquisition-aware feasible realization: *Under the hypotheses of Theorem 10, the implemented action satisfies, uniformly over every permitted observation and true state,*

$$Q_t(f_{t+1}; x, \hat{a}_t(\hat{x})) - f_t(x) \leq e_t + v_t + 2L_t r_t + D_t^Q \zeta_t^{\text{acq}}. \quad (44)$$

Consequently the original full-information optimum bounds the implemented policy's value below, and its loss is bounded above by

$$\sum_{t < T} B_t \{D_t^Q k_t + 2e_t + 2L_t h_t + v_t + 2L_t r_t + D_t^Q \zeta_t^{\text{acq}}\} + B_T(2e_T + 2L_g h_T). \quad (45)$$

The result requires no continuity of the index or actor and no probabilistic independence of measurement error. It applies to the common monotone, cash-invariant recursive evaluation under its original comparison assumptions.

PROOF: Transport from the original feasible node action to the implemented feasible action gives

$$Q_t(f_{t+1}; x, \hat{a}_t) \leq y_{t,i} + e_t + L_t \|x - x_{t,i}\|_1 + D_t^Q \zeta_t^{\text{acq}}.$$

The cone at x is at most its value at $\hat{x}$ plus $L_t r_t$. Apply (42), then use $f_t(\hat{x}) \leq f_t(x) + L_t r_t$. This proves (44). Backward order comparison, conditional also on the observation history when necessary, proves (45). Feasibility permits comparison with the optimum over all admissible adapted policies. The proof does not compare distances between the two selected actions at nearby states. *Q.E.D.*

COROLLARY 12—Capacity, measurement, and action quantization: Let $A_t(x) = [0, \kappa_t(x)]$, where κ_t is nonnegative and K_t^A -Lipschitz. Suppose a computed lower capacity satisfies

$$0 \leq \underline{\kappa}_t(\hat{x}) \leq \kappa_t(x), \quad \kappa_t(x) - \underline{\kappa}_t(\hat{x}) \leq 2K_t^A r_t + \chi_t \quad (x \in \mathcal{D}_t(\hat{x})).$$

For an action spacing $\Delta_t > 0$, implement

$$\hat{a}_t = \Delta_t \left\lfloor \frac{\min\{a_{t,i}, \underline{\kappa}_t(\hat{x})\}}{\Delta_t} \right\rfloor.$$

It is feasible on the whole acquisition set and satisfies (43) with $\zeta_t^{\text{acq}} = 2K_t^A r_t + \chi_t + \Delta_t$. Exact, unquantized robust repair uses $\Delta_t = 0$ and the minimum itself.

PROOF: The downward quantization preserves nonnegativity and the lower capacity. Since $a_{t,i} \leq \kappa_t(x_{t,i})$, exact repair at x displaces the node action by at most $K_t^A \|x - x_{t,i}\|_1$. Replacing true capacity by its certified lower bound adds at most $2K_t^A r_t + \chi_t$; downward quantization adds less than Δ_t . Summing proves the displacement inequality. Q.E.D.

10.1. Ordinary numerical comparisons

The execution in Section 12 uses binary64 scores at dyadic acquired states. Stored labels and the coordinate distances are exactly representable at the declared input resolutions. Put $u = 2^{-53}$, $\gamma_2 = 2u/(1 - 2u)$, $Y_t = \max_i |y_{t,i}|$, and let $\tilde{L}_t$ be the represented slope. In state dimension d , each computed cone score differs from its exact score by at most

$$\tau_t = d|\tilde{L}_t - L_t| + \gamma_2\{Y_t + d|\tilde{L}_t|\}. \quad (46)$$

The first term charges coefficient conversion; the second charges multiplication and addition. A minimum of the computed scores satisfies (42) with $\nu_t = 2\tau_t$, including ties. The supplement states the representability and arithmetic conditions and proves the error account. The general theorem also permits other certified score enclosures.

This implementation does not round a rational switchpoint and assume that its original cell is unchanged. For the scalar comparisons, all cells intersecting a state

enclosure are retained. For acquired constrained deployment, a numerical index has the objective allowance (46) and its action is repaired against the entire acquisition set. Both routes address discontinuity directly.

10.2. The identical non-neural envelope

The envelope construction must be separated from the language used to encode it. The classical Lipschitz extension (McShane, 1934), a native minimum computation, and an exact ReLU realization can represent the same function.

PROPOSITION 13—Representation control: *Fix the node locations, labels, slopes, original action witnesses, and selector tie rule. A conventional min-plus implementation and the affine-ReLU implementation of the same envelope return the same mathematical continuation and policy. They have the same exact policy certificate and direct policy cost. A balanced realization uses $O(dK)$ arithmetic and comparison gates for K cones and depth $O(\log K + \log d)$. Merely changing the representation does not improve the covering or Bellman-query factors.*

PROOF: For every pair of real numbers, $\min\{a, b\} = a - (a - b)_+$. Applying this identity through a balanced reduction tree proves equality of the continuation. Carrying the same attaining index, with the prescribed tie rule, proves equality of the action witness and repaired action. The certificates and costs concern these mathematical objects. The gate count includes the coordinate absolute values and their sums. The supplement separately bounds floating-point gate errors; exact equality is not asserted for unchecked rounded evaluations. *Q.E.D.*

We execute this non-neural control on all 24 distinct scalar witness checkpoints. It shares the defining data rather than masquerading as an independently trained competitor. The native selector and its ReLU encoding are thus one construction with two representations. The learned hidden-location methods studied earlier are different candidate generators. Their assessment, and that of the deterministic backend, remains a comparison of actual policies and complete computational work.

11. WITNESS-PRESERVING COMPILATION OF THE NEURAL OPERATOR

The cost of a Bellman construction depends on how its continuation is evaluated. A literal evaluation of the cone network in Theorem 10 visits every state label at every action–innovation query. Repeating that scan is unnecessary on a tensor cover. The following result compiles the same continuation and original action witnesses; it neither changes the fitted function nor substitutes a newly optimized policy.

Let $\mathcal{X} = \prod_{j=1}^d \{x_{j,0}, \dots, x_{j,n_j}\}$, with strictly increasing coordinates and $S = \prod_j (n_j + 1)$. The rectangular domain is the product of the coordinate intervals. Index the nodes lexicographically. For arbitrary real labels y_i and $L \geq 0$, write

$$f(x) = \min_{i \leq S} \{y_i + L\|x - x_i\|_1\}, \quad i_*(x) = \min_{i \leq S} \argmin \{y_i + L\|x - x_i\|_1\}. \quad (47)$$

Each original index carries a feasible node action a_i . Labels need not be Lipschitz-consistent: the numerical errors allowed in the earlier construction are permitted here without alteration. Define the transformed label and its original owner at a grid node v by

$$z_v = \min_i \{y_i + L\|v - x_i\|_1\}, \quad o_v = \min_i \argmin \{y_i + L\|v - x_i\|_1\}. \quad (48)$$

The owner is an original label index, not the index of the node at which its transformed value is stored.

THEOREM 14—Continuous compilation with original witnesses: *For every closed tensor cell C and every $x \in C$, let $\mathcal{V}(C)$ be its corner set. Then*

$$f(x) = \min_{v \in \mathcal{V}(C)} \{z_v + L\|x - v\|_1\}, \quad (49)$$

$$(f(x), i_*(x)) = \min_{\text{lex}} \{(y_{o_v} + L\|x - x_{o_v}\|_1, o_v) : v \in \mathcal{V}(C)\}. \quad (50)$$

Forward and backward relaxations along each coordinate line compute all pairs (z_v, o_v) using $O(dS)$ additions and comparisons and $O(S)$ stored values and indices. After cell location, a value or exact witness query uses at most 2^d corner pairs and $O(d2^d)$ arithmetic operations. A binary search locates the cell in $O(\sum_j \log(n_j + 1))$ comparisons; a uniform grid permits direct location.

When node coordinates, labels and L are rational, exact rational preprocessing and comparisons implement these identities, including zero slopes, inconsistent labels and ties. Applying the same repair map to the original returned action gives exactly the same feasible policy and policy cost as the uncompiled construction. The residual and acquisition bounds of Theorems 10 and 11 are unchanged.

PROOF: The triangle inequality gives $f(x) \leq z_v + L\|x - v\|_1$. Choose an attaining original site x_i . In each coordinate choose the corner coordinate lying between x_j and $x_{i,j}$. The resulting $v \in \mathcal{V}(C)$ satisfies $\|x - x_i\|_1 = \|x - v\|_1 + \|v - x_i\|_1$, proving (49). Choose now $i = i_*(x)$. At this corner an owner with a strictly smaller transformed value would improve the objective at x ; an owner with the same value and a smaller original index would contradict the prescribed tie rule. Hence $o_v = i_*(x)$ for a suitable corner. Every candidate in (50) is an original cone, so its lexicographic minimum is the desired pair. Coordinatewise forward and backward minimizations evaluate the separable ℓ^1 transform; keeping original indices in each lexicographic comparison preserves ties. The supplement gives the line invariant, operation account and numerical enclosure details. Q.E.D.

The nodal distance transform is classical; the linear-time coordinate algorithm is described by Felzenszwalb and Huttenlocher (2012). We use it for a different computational obligation: continuous off-grid Bellman queries must retain the original feasible action and the specified tie rule. Equation (49) alone does not authorize using the action attached to v ; it is (50) that licenses the policy. The result belongs to the cone-based NBO construction, but an identical native min-plus implementation can use the same compiler. It does not establish a computational property exclusive to a neural encoding. Proposition 13 remains an equality statement, not a comparison between independently trained methods.

11.1. Complete construction work

Let $J_{t,i}$ denote the size of the feasible action net at state node i , and let M_t be the number of integration points in a certified query. Put $Q_t = M_t \sum_i J_{t,i}$. For a tensor implementation, the representation part of a backward construction changes from

$O(\sum_t dS_{t+1}Q_t)$ to

$$O\left(\sum_t \left[dS_{t+1} + Q_t \left\{ d2^d + \sum_j \log(n_{t+1,j} + 1) \right\} \right] \right). \quad (51)$$

Primitive evaluation, feasible-net construction, integration, error accounting, actor formation, repair and output are additional terms. Exact arithmetic has a bit cost; the count in (51) is not a claim that rational comparisons have unit physical cost. Our records therefore report coefficient bit lengths, storage and complete process and CPU times alongside the component counts.

In the two-state investment study, $S = (N + 1)^2$, $J = N + 1$ and $M = N$. At fixed dimension and horizon, compiling removes the extra state-cover factor from the dense representation workload: $O(TN^6)$ becomes $O(TN^4)$, plus the explicitly charged setup and output. This does not remove the tensor cover, the action cover, or the innovation factor. In dimension d , both S and the 2^d corner factor remain visible. The four-dimensional stress cases below are executions of this finite-dimensional construction, not evidence that the curse of dimensionality has disappeared.

COROLLARY 15—A resource allocation for implemented policy accuracy: *In the two-state economy of Section 12, let $C_{T,p}$ be the explicit primitive constant in (54), $S_T = \sum_{t < T} \beta^t$, and choose a dyadic N with*

$$N \geq 4C_{T,p}/\varepsilon, \quad \xi, \xi_T \leq \frac{\varepsilon}{8(S_T + \beta^T)}.$$

Let A be the acquisition coefficient in (59) and $K = \sum_{t < T} \beta^t D_t^Q$. With exact witness selection, a b -bit sensor and a common downward action spacing Δ satisfying

$$A2^{-b} \leq \varepsilon/4, \quad K\Delta \leq \varepsilon/4,$$

the constructed implemented policy has loss at most ε against the full-information optimum. A zero coefficient imposes no corresponding precision restriction. An observation fee adds its explicit discounted charge to this loss account.

For fixed primitive parameters and horizon, choosing the smallest sufficient dyadic resolution gives $N = O(\varepsilon^{-1})$. The compiled algebraic construction uses $O(TN^4)$ operations,

versus $O(TN^6)$ for the literal dense cone evaluation, and stores $O(TN^2)$ labels, owners and actions. Executing the full dyadic ladder to this resolution has the same orders, with all preceding attempts charged. Numerical precision, coefficient bit lengths, uncertainty-box work and physical output latency remain separate costs; these orders are not a precision-free or machine-independent timing claim.

PROOF: The first two allocations make the right side of (54) at most $\varepsilon/2$. The acquisition and quantization terms add at most another $\varepsilon/2$ by (59). Compilation preserves the selected policy, so it does not change this account. On the dyadic ladder, the sum of every positive power of the earlier resolutions is bounded by a constant times that power of its largest resolution. Summing the complete algebraic operation counts and retained arrays proves the resource assertions. Effective certified queries at the prescribed accuracy, rather than unchecked fixed-precision evaluations, are part of the hypotheses. *Q.E.D.*

The corollary turns the compiler into a sufficient work-to-economic-accuracy allocation. A finite experimental ladder can still end before the sufficient resolution; its failure must remain a failure. Nor does more precise observation eliminate construction error: the two portions of the budget are distinct.

11.2. Certified numerical queries and selector ambiguity

The executable compiler computes (z_v, o_v) in rational arithmetic. A value query encloses each corner expression outward and takes the minima of lower and upper endpoints. For a state box contained in one cell, the possible original actors are among its corner owners; for a box crossing cells, the union of their corner owners is used. An owner is discarded only when its lower objective enclosure exceeds another candidate's upper enclosure. Exact ambiguous point queries use rational comparisons and the original lexicographic rule. Uncertain-state queries retain the hull of all remaining actions and propagate that hull through repair and the dynamics.

Thus a floating-point tie is not silently assigned to a convenient policy. A numerical implementation that deliberately returns a nonattaining index instead has to pay

the v_t allowance in (42). The present frozen-object checks separate exact identities from outward-rounding diagnostics. Overlapping machine intervals are useful checks; the identity theorem, rather than that overlap, proves that the compiler represents the original policy.

12. AN ORACLE-EXPLICIT CONSTRAINED INVESTMENT ECONOMY

The following evidence is the frozen R47 study. Its results and limitations are retained; Section 15 gives the new execution and decision evidence. We now execute the feasible-witness construction in a coupled two-state economy. Capital affects both future production conditions and current investment capacity. This makes feasible repair part of the economic policy rather than an algebraic example appended to a common-action calculation.

Let $x \in [0, 1]^2$, $\beta = 15/16$, and $0 \leq a \leq \kappa(x) = 1/8 + (x_1 + x_2)/16$. A common innovation z is continuously uniform on $[-1/32, 1/32]$, with opposite exposures in the two transitions:

$$F_1(x, a, z) = 1/16 + x_1/2 + x_2/8 + x_1(1 - x_2)/16 + a/2 + z, \quad (52)$$

$$F_2(x, a, z) = 1/16 + x_2/2 + x_1/8 + x_2(1 - x_1)/16 + a/4 - z.$$

Define the shortage and imbalance cost

$$b(x) = (x_1 - x_2)^2/4 + 2(1/2 - x_1 - x_2)_+^2.$$

Stage cost is $\sum_{j=1}^2 (x_j - 5/8)^2 + b(x) + pa^2 + 4a^4$; terminal cost doubles the squared-target terms and retains b . The predeclared cells have $T \in \{2, 3\}$ and $p \in \{1, 4\}$. These are normalized theoretical economies, not estimated parameters or calibrated welfare magnitudes.

PROPOSITION 16—Finite primitives and continuous-law queries: *The state domain is invariant under (52). In the ℓ^1 state norm, valid primitive constants are*

$$C_x = 13/4, \quad L_g = 9/2, \quad F_x = 11/16, \quad F_a = 3/4, \quad C_a = p/2 + 1/4, \quad K^A = 1/16. \quad (53)$$

The uniform tensor state grid with N subdivisions has covering radius $h = 1/N$. At each node, $N + 1$ equally spaced feasible actions have covering radius at most $k = 1/(8N)$. The repair is $\min\{a_i, \kappa(x)\}$ and has zero extra displacement allowance. If a continuation is L -Lipschitz, M uniform midpoint cells enclose its original continuous-law expectation after adding the deterministic remainder $L/(32M)$ to the numerical integration error.

PROOF: The primitive formulas give $F_1 \in [1/32, 29/32]$ and $F_2 \in [1/32, 27/32]$ by interval bounds. The sums of absolute derivatives in each state column are at most $11/16$, and the action derivative has ℓ^1 norm $3/4$. Differentiating the costs on the two shortage regions gives (53). Capacity is $1/16$ -Lipschitz and at most $1/4$, proving the grid and repair assertions. Within a uniform shock cell, the mean distance to its midpoint is $1/(64M)$. Since the transition has shock modulus two in ℓ^1 , the expectation error is at most $L/(32M)$. Thus quadrature does not change the innovation law to a discrete one. Q.E.D.

The two generators are the feasible cone witness and conventional fitted-value iteration in the tensor bilinear basis. Each works backward against its own fitted future, uses the same feasible action net, and evaluates the same continuous law with the preceding remainder. No optimal value or other method's continuation is supplied as a label. Rounded labels, their valid query errors, slopes, node actions and all policy bounds are stored.

The conventional actor selects a nearest state node and repairs its action at the current capacity. It receives the one-sided account rather than only the older separated certificate. For the bilinear interpolant, the Bellman-residual enclosure is

$$-e_t - L_t h \leq f_t - \mathcal{T}_t f_{t+1} \leq D_t^Q k + e_t + L_t h.$$

The interpolation weights are nonnegative and reproduce coordinates. Their weighted mean ℓ^1 distance to the surrounding vertices is at most $1/N$, which proves the stated extension from the nodal enclosure.

LEMMA 17—Complete cell-corner verification: *For the bilinear actor, the selected-policy allowance is $e_t + D_t$, where*

$$D_t = \max_i \sup_{x \in C_i} \{y_{t,i} + L_t \|x - x_i\|_1 - f_t(x)\}.$$

Here C_i is the closed nearest-node cell and both adjacent choices are admitted on a boundary. It suffices to evaluate the at most nine vertices obtained by splitting each such cell along its node coordinates. Outward endpoint evaluation supplies a valid upper bound on D_t without an additional verification mesh.

PROOF: On each resulting subrectangle the distance term is affine and the continuation is bilinear. An affine function minus a bilinear function attains its maximum at a vertex: first maximize over one coordinate, then over the other. The label's feasible action, repaired at the queried state, has objective at most $y_{t,i} + e_t + L_t \|x - x_i\|_1$ by feasible transport. Taking the complete maximum proves the allowance. *Q.E.D.*

12.1. A cap obtained from economic accuracy

In this class the previously abstract covering, integration and repair operations are finite routines with explicit costs. Let $S_T = \sum_{t < T} \beta^t$, $a_p = C_x + K^A C_a$ and $b = \beta(F_x + K^A F_a) = 705/1024$. Define

$$\bar{L} = \max\{L_g, a_p / (1 - b)\}, \quad \bar{D} = C_a + \beta F_a \bar{L}.$$

The primitive recursion gives $L_t \leq \bar{L}$ and $D_t^Q \leq \bar{D}$ at every horizon. Suppose that, apart from the quadrature remainder, each nonterminal numerical query has error at most ξ and terminal labels have error at most ξ_T . With $M = N$, Theorem 10 implies

$$\mathcal{G}_N \leq C_{T,p} / N + 2S_T \xi + 2\beta^T \xi_T, \tag{54}$$

$$C_{T,p} = S_T \{\bar{D} / 8 + 2\bar{L} + \beta \bar{L} / 16\} + 2\beta^T L_g.$$

Thus a dyadic $N \geq 2C_{T,p} / \varepsilon$ and $\xi, \xi_T \leq \varepsilon / [4(S_T + \beta^T)]$ suffice for loss at most ε . Acquisition and action quantization add the explicit terms of (45); these must

receive part of the same tolerance. Fixed binary64, a fixed label lattice and a finite three-rung study are not asserted to achieve every positive ε .

For this two-state, one-action, one-shock implementation, Bellman queries number $T(N + 1)^3$. Each has N midpoint evaluations. Direct cone evaluation examines $(N + 1)^2$ cones, giving $O(TN^6)$ primitive arithmetic work. Bilinear future evaluation uses a local patch, giving $O(TN^4)$ work for that implementation. Both store $O(TN^2)$ critic labels and actions. The present batched program additionally holds $O(N^3)$ state-action arrays and $O(BN^2)$ temporary cone entries for fixed batch size B . Feasible acquisition and repair add their recorded deployment work, rather than being included at zero cost.

With p_b -bit arithmetic and multiplication cost $M(p_b)$, the corresponding arithmetic upper bounds acquire a factor $M(p_b + T + \log N)$, with coefficient conversion, comparison and serialization charged separately. The precision must also meet the query enclosure requirement. These are algorithmic upper bounds, not lower bounds for all conventional methods. Proposition 13 shows that the identical native envelope has the same leading factors as its ReLU realization. No representation label removes the tensor dependence.

12.2. Execution and acquired deployment

The fixed ladder is $N = M = 4, 8, 16$, with three isolated processes per method and economic cell. The 24 services store all 72 rungs, including failed attempts. Targets are 4, 2, 1, 1/2 in the normalized cost units. They test the construction-to-certificate service; a mathematically valid coarse bound is not automatically an economically adequate tolerance. The full attained bounds and work therefore accompany the target counts.

Both methods attain the target 4 in all twelve services, at the last rung, and neither attains 2, 1 or 1/2 within this cap. The witness's final bound is smaller in each economic cell, but bilinear fitted-value iteration is faster in all twelve common successful comparisons. These coarse all-state certificates establish an executed feasible construction, not an economically precise solution at the unachieved tolerances. The

TABLE VIII

COUPLED CONSTRAINED ECONOMY: FULL-PREFIX WORK AT $N = 16$

T	p	Method	Bound	Queries	Seconds	Time range
2	1	Witness	2.3148	11,534	7.045	[7.003,7.070]
2	1	Bilinear FVI	2.7960	11,534	0.326	[0.325,0.352]
2	4	Witness	2.3691	11,534	6.992	[6.975,7.188]
2	4	Bilinear FVI	2.8288	11,534	0.324	[0.323,0.325]
3	1	Witness	3.3212	17,301	10.519	[10.480,10.552]
3	1	Bilinear FVI	3.8915	17,301	0.490	[0.487,0.500]
3	4	Witness	3.4107	17,301	10.482	[10.425,10.538]
3	4	Bilinear FVI	3.9321	17,301	0.494	[0.489,0.520]

Query counts and clocks include all three rungs, including failed attempts. Seconds are medians over three isolated process observations; the range reports their minimum and maximum. CPU affinity and library thread counts are fixed, but CPU frequency is not controlled.

sufficient-cap formula states what further resources the mathematical guarantee requires without relabeling these failures.

For each of the twelve distinct saved witness policies we also execute the selector on a 33-by-33 measurement grid, with coordinate acquisition radii 0, 2^{-12} and 2^{-8} . Exact robust repair and downward action spacing 2^{-12} give 72 deployment cases. Capacity is bounded on the entire acquisition box; it is not checked only at its midpoint. The binary64 score allowance is uniform by (46); the grid records exercise implementation and repair activity rather than replacing that proof.

Repair is active rather than a vacuous operation, and downward action quantization is charged explicitly. At exact acquired input the small nonzero numerical selection allowance remains in the exact records even when it rounds to zero in this table. At a nonzero radius the acquisition term is positive; the total bound is not reported as the exact-state certificate.

13. POLICY COSTS AND ECONOMIC METHOD CHOICE

The following evidence is the frozen R47 study. Its results and limitations are retained; Section 14 gives the new execution and decision evidence. We next evaluate the actual witness policy, not its distance bound relative to an unknown optimum.

TABLE IX

ACQUIRED-STATE AND QUANTIZED-ACTION DEPLOYMENT

Coordinate radius	Action spacing	Active repairs	Quantized	Largest added bound
0	0	4,817	0	0.000000
0	1/4096	4,817	0	0.004619
1/4096	0	6,609	0	0.022513
1/4096	1/4096	6,609	6,609	0.027132
1/256	0	6,609	0	0.360204
1/256	1/4096	6,609	0	0.364823

Each row covers all twelve distinct constrained witness policies and 32,670 measured state–date queries. Feasibility holds on the entire acquisition box. The additional loss is a uniform theoretical allowance for the executed arithmetic, acquisition and action rule, not an estimated average policy cost.

The policies are the frozen objects at every common first crossing in the original witness catalogue. There are ten such economic-cell and target combinations. Three timing repetitions of an identical checkpoint are not three independent policies.

The primary initial-state law is uniform on $[0,1]$. The additional fixed states $1/8$, $1/2$ and $7/8$ test particular economic starting positions. At each of these forty groups, the three pairwise contrasts concern witness, nearest-cone and spline policies. A fixed design specifies 131,072 common paths, forty-bit innovation bins and family error 0.01 for all 120 contrasts. The economic comparison uses the same continuous innovations as the original economy.

Let $[\underline{z}_m, \bar{z}_m]$ enclose $g - f_T^m$ and let the initial continuation difference have range $[a_0, b_0]$ over the chosen initial-state law. With the one-sided residuals of (30), the score difference of two policies has deterministic support

$$\begin{aligned}
 a &= a_0 + \sum_{t < T} B_t(-u_t^m - d_t^n) + B_T(\underline{z}_m - \bar{z}_n), \\
 b &= b_0 + \sum_{t < T} B_t(d_t^m + u_t^n) + B_T(\bar{z}_m - \underline{z}_n).
 \end{aligned} \tag{55}$$

For a common random initial state, one bounds the difference $f_0^m - f_0^n$, not the two separate ranges. Its extrema are obtained exactly from the union of the rational

1 piecewise-affine knots. Conditional expectation and then integration over the initial 1
2 law extend Theorem 4 to this estimand. 2

3 Every simulated bin generates lower and upper score endpoints. When a state 3
4 enclosure meets a selector boundary, the calculation encloses all possible actors and 4
5 subsequent states. It neither drops the path nor substitutes an unverified midpoint 5
6 actor. The mean and variance computations have rational roundoff accounts. For 6
7 the declared family, $\log(4 \cdot 120/0.01) < 11$; a finite positive Taylor sum for $\exp(11)$ 7
8 certifies this inequality. Using 11 in (15) is conservative and avoids an unchecked 8
9 logarithm evaluation. The supplement gives the complete endpoint construction 9
10 and arithmetic audit. 10

11
12 THEOREM 18—Simultaneous policy comparison and adoption: *Conditional on the 12
13 frozen policies and the independent uniform-bin model, the endpoint intervals constructed 13
14 with (55) and (15) cover all declared expected policy-cost differences simultaneously with 14
15 probability at least 0.99. If an interval $[L_{mn}, U_{mn}]$ has $U_{mn} < 0$, policy m has smaller 15
16 expected cost than n on this event. For a previously specified tolerance η , $U_{mn} \leq \eta$ certifies 16
17 noninferiority within that tolerance, whereas $[L_{mn}, U_{mn}] \subset [-\eta, \eta]$ certifies two-sided 17
18 proximity. 18*

19 More generally, fix a reference policy b and known implementation charges k_m . Suppose 19
20 simultaneous intervals $[L_m, U_m]$ enclose $J_m - J_b$ for a finite candidate catalogue, including 20
21 $[0, 0]$ for b . Select $\hat{m}$ to minimize $U_m + k_m$. On the coverage event, 21

$$J_{\hat{m}} + k_{\hat{m}} \leq \min_m \{J_m + k_m + (U_m - L_m)\}. \quad (56)$$

22
23
24 Thus an economic choice made after seeing all intervals retains its declared uncertainty 24
25 account. The catalogue comparison is within one common economic task and initial-state 25
26 law, not across different prices, horizons or evaluations. 26
27

28 PROOF: The score identity identifies the expected difference. Its interval end- 28
29 points are bounded independent observations within each group. Apply the one- 29
30 sided empirical Bernstein inequality of Maurer and Pontil (2009) to each lower 30
31 and upper endpoint and take a union bound over 240 tails. The numerical enclo- 31
32 sures only enlarge the valid intervals. Independence between different contrasts is 32

unnecessary. On the resulting event, $J_{\hat{m}} - J_b \leq U_{\hat{m}}$ and $U_m \leq J_m - J_b + (U_m - L_m)$.
Minimize the former upper bound plus the known charge to obtain (56). *Q.E.D.*

The study fixes $\eta = 1/1024$ in its normalized cost units before execution. It is a stated decision tolerance, not an estimated welfare threshold. Zero remains the superiority threshold. An interval merely intersecting zero meets neither criterion automatically. The fixed pseudorandom implementation supplies repeatable inputs; it does not prove the independent sampling model. The theorem states that model explicitly. No claim about initialization risk or the distribution of newly trained policies follows from these frozen-policy comparisons.

TABLE X
DIRECT EXPECTED-COST COMPARISONS UNDER A PREDECLARED DECISION TOLERANCE

Contrast	Groups	Signed	Within 1/1024	Largest endpoint
Witness – Cone nearest	40	0	40	0.0001161
Witness – Spline	40	0	40	0.0001975
Cone nearest – Spline	40	0	40	0.0002078

Signed counts exclude intervals containing zero. “Within” requires the entire interval inside $[-1/1024, 1/1024]$. Largest endpoint is the largest absolute endpoint over the 40 groups. All 120 intervals use one 0.01 family error account, conditional on the frozen policies and independent-bin model. Each group uses 131,072 common paths.

Every direct interval contains zero, so the experiment identifies no signed policy ranking. Every interval is also contained in the predeclared decision band. It therefore certifies two-sided expected-cost proximity for the specified initial laws at that tolerance, rather than inferring equivalence from a failure to reject equality. The uniform initial-state law is a distributional conclusion, not an all-state assertion. The individual initial states are separate declared estimands.

13.1. Work as a function of the certified bound

A coarse target list can hide a genuine difference between certificate frontiers. For the original saved rungs j , let $G_{m,j}$ be the valid bound and let $C_{m,j}$ include all work

through rung j . The complete first-crossing curve is

$$W_m(\varepsilon) = C_{m,j_m(\varepsilon)}, \quad j_m(\varepsilon) = \min\{j : G_{m,j} \leq \varepsilon\}, \quad (57)$$

with $W_m = +\infty$ when the cap is exhausted. Its breakpoints are the recorded bounds; both methods' breakpoints partition the entire positive tolerance axis. We reconstruct every interval of this partition, not only intervals favorable to the witness. This is an exact descriptive account of frozen evidence, not a retrospective change to the original targets.

TABLE XI
COMPLETE TOLERANCE-AXIS PARTITION OF THE ORIGINAL SCALAR FRONTIERS

Relation on a partition interval	Number
Both attain; witness uses fewer queries	19
Both attain; spline uses fewer queries	1
Both attain; same prefix queries	20
Only witness attains within the cap	4
Only spline attains within the cap	0
Both exhaust the cap	4

The complete curves contain 19 nonempty tolerance intervals on which both methods attain and the witness uses fewer Bellman queries, one favoring the spline, and twenty with identical counts. Four further intervals are attainable only by the witness; four below both final bounds are unattained. Interval counts are not probabilities or measures of economic importance. Above the largest breakpoint in each cell both first-rung counts coincide. The original targets remain unchanged and still have identical crossing resolutions. These descriptive intervals identify where a sharper certificate changes a resource decision without manufacturing new prospective targets.

13.2. *Economic interpretation and the NBO program*

The new evidence connects a feasible implementation to two distinct economic statements. First, the acquisition-aware theorem bounds the loss of a policy that can

actually be executed under uncertain state measurement and capacity constraints. Second, the direct experiment certifies proximity of the original scalar policy costs within an explicitly chosen decision tolerance, so that a user of that catalogue may consider implementation charges without treating separate regret bounds as values. These statements are stronger than a certificate comparison alone.

Neither statement implies that a neural representation is indispensable. The identical native envelope is a mathematical control, and the conventional fitted-value and spline implementations remain strong alternatives. The constrained experiment demonstrates two-state continuous-uncertainty construction and nontrivial repair, not a high-dimensional scaling result or an estimated economic application. Its tight targets remain unachieved. The direct cost experiment does not cover the new constrained policies or nonlinear recursive utility; its expectation-specific conclusion is not transferred to those settings. The controlled-diffusion, recursive-preference, endogenous-preference, temporal-self and game developments below retain their own conditions and original role in the NBO program. Their preservation is not counted as a new comparative execution.

14. PRICED OBSERVATION AND DIRECT CONSTRAINED-POLICY DECISIONS

A state-error radius is an economic input when acquiring more precise information is costly. We specify a finite-bit observation technology and a controller replacement charge. These are primitives of a theoretical decision problem in resource units, not estimates from a calibrated economy. They permit a decision interpretation without treating an interval containing zero as equality or treating an all-state regret bound as a ranking of two policies.

14.1. *Choosing the observation precision*

On $[0, 1]^d$, a b -bit sensor reports, in each coordinate, the containing dyadic interval and its midpoint. The right endpoint 1 belongs to the last interval. The resulting acquisition set has ℓ^1 radius

$$r(b) = d2^{-(b+1)}. \quad (58)$$

For $A_t(x) = [0, \kappa_t(x)]$, use the lowest capacity on the reported box, then round the selected action downward to spacing Δ_t . Feasibility holds throughout that box. In the affine capacity economy this lower capacity is computed exactly from the lower box coordinates. The acquisition repair therefore has $\chi_t = 0$ in Corollary 12.

Suppose the same b is chosen before date zero and each coordinate-bit costs $\lambda > 0$ per date. The observation fee is paid in the same additive resource units as the objective. With $B_t = \prod_{s < t} \beta_s$, it is Cb , where $C = \lambda d \sum_{t < T} B_t$. Write G for the full-information witness loss certificate with exact selector and repair. The acquisition theorem implies an augmented loss bound, relative to the full-information optimum without the observation fee,

$$J^{\pi, b}(x) + Cb - V^*(x) \leq G + K_\Delta + A2^{-b} + Cb, \quad (59)$$

$$A = d \sum_{t < T} B_t (L_t + D_t^Q K_t^A), \quad K_\Delta = \sum_{t < T} B_t D_t^Q \Delta_t.$$

A known selector allowance is added to G when necessary. The fee is a deterministic date cost, so the same argument also applies to a common cash-invariant recursive evaluation for which the earlier comparison theorem holds. Direct sampling below is, separately, expectation-specific.

PROPOSITION 19—Optimal precision for the certified economic account: *On a declared integer range $b_{\min} \leq b \leq b_{\max}$, the smallest minimizer of the right side of (59) is characterized by the neighboring inequalities*

$$C \leq A2^{-b} \quad \text{if } b > b_{\min}, \quad C \geq A2^{-(b+1)} \quad \text{if } b < b_{\max}, \quad (60)$$

with the smallest minimizer selected when equality gives two. The criterion can be decided by rational comparisons, without an unverified logarithm. It minimizes a proved upper account; it need not minimize the unknown actual policy cost as a function of precision.

PROOF: The difference between the accounts at $b + 1$ and b is $C - A2^{-(b+1)}$, which is nondecreasing in b . The two neighboring comparisons are therefore necessary and sufficient, with the indicated endpoint convention. Equation (59) follows by substituting (58) and $\zeta_t^{\text{acq}} = 2K_t^A r(b) + \Delta_t$ into (45). Q.E.D.

This calculation provides an explicit information-price comparative static: an increase in λ weakly decreases the smallest minimizing precision, while an increase in the sensitivity account A weakly increases it. The result concerns the guarantee, not a fitted distribution of measurement errors. The executed curves use every $b = 4, \dots, 16$ and every $\lambda \in \{1/16384, 1/4096, 1/1024\}$. All curves and endpoint choices are retained, rather than reporting only an attractive radius.

14.2. *The actual constrained-policy cost*

Fix the two-state witness and bilinear-FVI policies stored at the finest original R47 resolution, $N = 16$. In each economic cell, the two policies are evaluated with the same observation contract: exact state, six bits per coordinate, or ten bits per coordinate. The finite-bit policies use robust capacity repair and action spacing $1/4096$; the exact-state policies are unquantized. Thus there are three distinct pairs of implemented policies, not three observations of an assumed continuous actor.

For an initial law μ , the estimand is

$$D_{\mu,b} = \mathbb{E}_{\mu} \left[\sum_{t < T} \beta^t \{c(X_t^W, A_t^W) - c(X_t^F, A_t^F)\} + \beta^T \{g(X_T^W) - g(X_T^F)\} \right]. \quad (61)$$

The observation fees cancel within a pair because its two controllers use the same sensor. The paired policies share innovations but follow their own endogenous state paths. This is their actual cost difference, not a difference of regret certificates. A raw path score is used here; the earlier own-continuation score remains a valid alternative when its signed support is sharper.

For the two-state primitives, $0 \leq c \leq (99 + 4p)/64$ and $0 \leq g \leq 37/16$. Consequently the path difference belongs to $[-H, H]$, with

$$H = \frac{99 + 4p}{64} \sum_{t < T} \beta^t + \frac{37}{16} \beta^T. \quad (62)$$

This uniform support is conservative. It is independent of the observed contrast and is fixed before simulation.

THEOREM 20—Interval simulation and replacement decisions: *For each of K fixed policy-pair tasks, partition the underlying initial and innovation uniforms into equal bins*

and draw independent uniform bin indices. Suppose a numerical path routine returns bounded random endpoints L_j, U_j enclosing the exact path difference for every continuous realization inside the selected bins, including every possible actor, observation and repair branch. Let $[a, b]$ be a deterministic support enclosure and clip endpoints to it. For $n \geq 2$, let $\underline{m}_L, \bar{m}_U$ enclose their sample means and s_L^2, s_U^2 upper-bound their unbiased sample variances. Put

$$R(s^2) = \sqrt{\frac{2s^2\ell}{n}} + \frac{7(b-a)\ell}{3(n-1)}, \quad \ell \geq \log(4K/\alpha). \quad (63)$$

Then, under the independent-bin sampling model, simultaneously for all K tasks with probability at least $1 - \alpha$,

$$D_{\mu,b} \in [\max\{a, \underline{m}_L - R(s_L^2)\}, \min\{b, \bar{m}_U + R(s_U^2)\}]. \quad (64)$$

If replacing either installed controller costs $k > 0$ once at date zero, an interval strictly inside $(-k, k)$ proves that neither replacement recoups that fee in expected resource cost for that task. An interval with upper endpoint below zero ranks the witness policy below FVI; a lower endpoint above zero ranks it above FVI. Containing zero does not establish equality.

PROOF: Conditional on a bin sequence, integrating the continuous uniform variables gives a conditional path expectation between its enclosing endpoints. Averaging over uniform bin indices reproduces the original continuous law, so $\mathbb{E}L_j \leq D_{\mu,b} \leq \mathbb{E}U_j$. Apply the empirical-Bernstein inequality of [Maurer and Pontil \(2009\)](#) to the lower and upper endpoint samples, with a union bound across $2K$ one-sided assertions; certified moment and arithmetic allowances enlarge, rather than shrink, the intervals. Finally, switching from FVI to witness changes cost by $D_{\mu,b} + k$, while switching in the opposite direction changes it by $k - D_{\mu,b}$. Both are positive when $|D_{\mu,b}| < k$. Q.E.D.

The experiment fixes $K = 48$, $n = 262144$ and $\alpha = 0.01$ before execution. An exact rational exponential bound verifies that $\ell = 10$ is sufficient. Forty-bit bins cover the continuous initial and innovation laws. The initial laws are independent uniform coordinates and the three fixed vectors $(1/8, 1/8)$, $(1/2, 1/2)$ and $(7/8, 7/8)$, crossed with all four horizon-price cells and three sensors. A fixed pseudorandom stream

1 makes the records reproducible; it is not a proof that those indices are independent 1
 2 random variables. 2

3 The replacement fee is $k = 1/64$, one quarter of the maximal quadratic investment 3
 4 expenditure at $p = 1$. It is an explicit theoretical transaction charge. We do not 4
 5 calibrate it to welfare data, choose it after inspecting the intervals, or infer an all- 5
 6 state ranking from the four initial laws. Reporting the zero threshold separately 6
 7 keeps the decision statement distinct from a claim of neural superiority. 7

8 15. COMPLETE COMPILED-POLICY SERVICES AND ECONOMIC COMPARISONS 9

10 The new study separates three questions: the cost of representing a fixed NBO 10
 11 policy, the complete cost of constructing and certifying a policy, and the actual cost 11
 12 of using one controller rather than another. A representation benchmark cannot by 12
 13 itself answer either of the latter questions. Conversely, a conventional comparator 13
 14 receives the same innovation law, feasibility conditions and all-state comparison 14
 15 principle, not a deliberately inefficient implementation chosen to favor the neural 15
 16 operator. 16

17 15.1. *Design, stopping rules and conventional baselines* 18

19 The design was fixed before the complete execution. The principal catalogue uses 19
 20 the original constrained two-state economy of Section 12, horizons $T = 2, 3$, prices 20
 21 $p = 1, 4$, and $N = M \in \{4, 8, 16, 32, 64\}$. The three methods are the compiled cone 21
 22 witness, uniform tensor FVI and error-driven nonuniform tensor FVI. There are 22
 23 three isolated process repetitions for each method and economic cell: 36 services and 23
 24 180 rungs. Every rung starts from primitives and uses its own future continuation. 24
 25 All rungs are executed, including those after a target crossing; first-target work 25
 26 charges the complete prefix, including every prior failure. 26
 27 27

28 The conventional adaptive method forms an own-future pilot on a 5×5 grid, 28
 29 uses its coordinatewise second differences to allocate nonuniform dyadic knots, and 29
 30 certifies with the actual largest coordinate gaps. Its pilot, allocation and repeated 30
 31 target evaluations are charged. The section on nonuniform fitted-value comparison 31
 32 in the supplement specifies the fixed heuristic and proves the strengthened nearest- 32

node certificate on these unequal grids. It is a conventional adaptive-coordinate FVI comparator, not a sparse-grid method or an unrestricted optimal adaptive partition.

All methods share outward arithmetic, feasible action fractions and midpoint integration with its deterministic remainder under the continuous uniform law. Each service includes a finite acquired-state deployment checksum and durable checkpoint output. Complete prefix wall and CPU times, method-specific evaluation counts, coefficient bit lengths, stored actor scalars, checkpoint bytes and peak resident memory are recorded. The internal clock excludes startup and the common non-economic warm-up; a separate process clock includes them. No independent cost comparison is hidden inside that boundary: its joint service is reported separately below.

TABLE XII

COMPLETE TWO-STATE CONSTRUCTION SERVICES AT THE DECLARED CAP

T	p	Method	Loss bound	Median (s)	Range (s)	Peak MiB
2	1	C	0.579	48.001	[47.981, 49.481]	92.0
2	1	U	0.719	36.308	[36.280, 36.426]	74.5
2	1	A	0.719	36.689	[36.685, 36.696]	74.6
2	4	C	0.592	48.077	[47.790, 48.600]	91.3
2	4	U	0.727	36.391	[36.389, 37.257]	73.5
2	4	A	0.727	37.025	[36.669, 37.110]	73.7
3	1	C	0.830	71.933	[71.549, 72.464]	93.1
3	1	U	1.003	54.649	[54.146, 54.650]	74.8
3	1	A	1.003	55.677	[55.256, 57.886]	75.5
3	4	C	0.853	71.963	[71.798, 72.082]	93.5
3	4	U	1.013	54.151	[53.822, 54.812]	74.9
3	4	A	1.013	55.026	[54.930, 55.110]	75.3

C denotes compiled witness, U uniform tensor FVI, and A error-driven coordinate FVI. Each entry ends at $N = M = 64$ and includes the complete five-rung prefix through durable record output. The range is over three isolated repetitions with identical checkpoints, not three economic draws. Peak resident memory includes the interpreter and warm-up. Bounds are displayed to three decimals; exact rational tests determine attainment.

Against uniform FVI, compiled witness is earlier in 3 of 30 common successful target/repetition comparisons, later in 27, and tied in 0. Against error-driven FVI,

TABLE XIII
PROSPECTIVE TARGET ATTAINMENT IN THE TWO-STATE CATALOGUE

Method	4	2	1	1/2	1/4
C	12/12	12/12	12/12	0/12	0/12
U	12/12	12/12	6/12	0/12	0/12
A	12/12	12/12	6/12	0/12	0/12

All failures and every prior rung are retained. Counts concern twelve deterministic services per method, crossing four economic cells and three timing repetitions. A repeated failure is not an independent estimate of optimizer reliability. No historical failure is reclassified.

compiled witness is earlier in 15 of 30 common successful target/repetition comparisons, later in 15, and tied in 0. These are observed complete-prefix wall times, not a hardware-general speed law. The comparison uses the prospective targets rather than retrospectively selecting a favorable breakpoint. The full positive-tolerance partition, including one-sided and joint nonattainment, is reported in the supplement and machine-readable frontier. The error-driven rule returns uniform grids in all 210 recorded date models; all 60 adaptive checkpoints have exactly the same continuations and actors as their uniform-FVI counterparts. Thus its additional observed cost is pilot and allocation overhead, not evidence of a realized nonuniform-grid improvement. The nonuniform verification theorem is proved and tested, but an effective spatial-adaptation advantage is not established by this catalogue. One prospective tolerance illustrates the distinction between common-resolution speed and target work. At horizon three, price four and target two, witness crosses at $N=32$ while uniform FVI first crosses at $N=64$. Median complete-prefix times are 5.498 and 54.146 seconds, respectively. This is one economic cell with three timing repetitions. At target one, witness attains all twelve services and each FVI implementation attains six; none attains one half or one quarter at its cap.

15.2. Representation identity and higher-dimensional execution

The fixed-object diagnostic uses all twelve distinct original constrained witness checkpoints, crossing $T = 2, 3$, $p = 1, 4$ and $N = 4, 8, 16$. Both the original dense evaluator and the compiler evaluate the same fixed point catalogue. Compiler setup

is charged. Exact rational off-grid checks verify values and original winning indices, including the deterministic tie convention; outward machine intervals are also checked. These clocks do not include new target generation and are not added to historical training times.

TABLE XIV

FIXED-POLICY REPRESENTATION COSTS, INCLUDING COMPILER SETUP

N	Observations	Median ratio	Minimum	Maximum	Compiled (s)
4	12	1.93	1.87	2.37	0.0074
8	12	2.92	2.80	4.10	0.0137
16	12	4.48	4.42	4.52	0.0385

The ratio is dense load-plus-query time divided by compiled setup-plus-query time, over the same original continuation and points. Each resolution crosses four economic cells and three repetitions. This is a same-policy representation experiment, not a construction-to-certificate speedup or an independently trained native comparator. Exact witness-index checks accompany the numerical interval checks.

To expose, rather than suppress, dimension factors, a cycle-coupled extension executes $d = 3, 4$ with continuous common uncertainty, state-dependent capacity and a nontrivial scalar action. Its normalized stage and terminal costs recover the two-state economy exactly at $d = 2$. The full primitive formulas and valid moduli appear in the supplement. The stress catalogue fixes $p = 1$, $T = 2, 3$, and the compiled witness and uniform multilinear FVI methods. Dimension three uses $N = M = 4, 8, 16$; dimension four uses $N = M = 2, 4, 8$. This adds eight services and 24 rungs. The smaller four-dimensional cap is explicit. These are memory-limited stress designs with one observation per cell, not repeated fixed-accuracy scaling estimates. They test an executed nonlinear implementation in dimensions above two without turning the finite-dimensional theorem into a dimension-free performance claim. The retained tensor state count, action factors, innovation work and 2^d corner terms determine what this implementation can afford.

TABLE XV

EXECUTED NONLINEAR DIMENSION STRESS CASES

d	T	Method	N	States	Loss bound	Seconds	MiB
3	2	C	16	4913	2.330	9.505	64.9
3	2	U	16	4913	2.814	6.677	52.3
3	3	C	16	4913	3.349	13.906	66.5
3	3	U	16	4913	3.915	9.905	52.6
4	2	C	8	6561	4.630	7.733	69.7
4	2	U	8	6561	5.456	7.368	50.3
4	3	C	8	6561	6.642	11.550	71.2
4	3	U	8	6561	7.589	10.982	51.2

Price is one. Each service includes all three declared rungs and has one repetition. The maximum resolution differs by dimension and is not a matched-accuracy scaling design. The common innovation remains continuous; its quadrature remainder, state-dependent capacity and all-state loss account are retained.

15.3. *Direct constrained-policy decisions and information prices*

The direct comparison freezes the original $N = 16$ witness and FVI policies. It does not select the new R48 policies after examining their costs. Forty-eight policy-pair tasks cross the four economic cells, four initial laws and three observation contracts of Section 14. The full study contains $48 \times 262144 = 12582912$ paired paths. Numerical path widths, every unresolved actor or sensor branch, endpoint moments, simultaneous intervals and decisions are deposited. Table XVI summarizes interval hulls across the four declared initial laws; its rows are not all-state confidence bands. The supplement reports every interval separately. The zero threshold and the replacement-fee threshold $1/64$ are different decision questions. A no-recoupment result is conditional on that predeclared theoretical fee. It neither establishes exact policy equality nor measures empirical welfare equivalence.

The sensor-price experiment evaluates all thirteen precisions $b = 4, \dots, 16$ at each of the three declared information prices, for every original witness checkpoint. Table XVII gives the finest original resolution; the evidence includes all coarser checkpoints and every candidate precision, not only the minimizing choice. The optimum minimizes the uniform augmented upper account. Since the full-information

TABLE XVI

DIRECT WITNESS-MINUS-FVI CONSTRAINED-POLICY COSTS

T	p	Sensor	Interval hull over four laws	Signed	No fee recovery
2	1	exact	$[-0.001247, 0.002717]$	1/4	4/4
2	1	6	$[-0.001113, 0.002716]$	1/4	4/4
2	1	10	$[-0.001220, 0.002717]$	1/4	4/4
2	4	exact	$[-0.001142, 0.001840]$	0/4	4/4
2	4	6	$[-0.001014, 0.001602]$	0/4	4/4
2	4	10	$[-0.001151, 0.001832]$	0/4	4/4
3	1	exact	$[-0.001501, 0.002627]$	1/4	4/4
3	1	6	$[-0.001364, 0.002528]$	1/4	4/4
3	1	10	$[-0.001455, 0.002639]$	1/4	4/4
3	4	exact	$[-0.001356, 0.001806]$	0/4	4/4
3	4	6	$[-0.001342, 0.001783]$	0/4	4/4
3	4	10	$[-0.001365, 0.001808]$	0/4	4/4

Each hull contains the four separate simultaneous intervals for the declared initial laws; it is not an all-state band. There are 48 contrasts, each with 262,144 common interval paths, and family error at most 0.01 under the independent-bin model. Signed counts identify intervals excluding zero. The final column counts intervals strictly inside $(-1/64, 1/64)$, ruling out recouping either replacement fee. Display endpoints are rounded outward. The observation fee cancels when the two controllers use the same sensor.

certificate G is unchanged across observation precisions, the integer choice responds to the objective sensitivity and the price of information, not to an arbitrary selection of an attractive state-error radius. A coarse state cover can still make the total account large; paying for more bits cannot repair that separate construction error.

15.4. *What changes in the evidence*

The compiler theorem supplies a constructive work reduction for a defined NBO backend without changing its policy. The complete catalogue measures that backend against uniform and error-driven conventional FVI with all failed resolutions included. Direct constrained costs and priced observation support explicit economic decisions for fixed policies under declared information contracts. These results extend the original program, rather than replace it with a different subject.

TABLE XVII

PRICED PRECISION FOR THE ORIGINAL FINEST WITNESS POLICIES

T	p	Bit price	b_*	Acquisition	Fee	Total account
2	1	$\frac{1}{16384}$	16	0.0027	0.0038	2.3212
2	1	$\frac{1}{4096}$	14	0.0040	0.0133	2.3320
2	1	$\frac{1}{1024}$	12	0.0092	0.0455	2.3694
2	4	$\frac{1}{16384}$	16	0.0034	0.0038	2.3763
2	4	$\frac{1}{4096}$	14	0.0047	0.0133	2.3871
2	4	$\frac{1}{1024}$	12	0.0101	0.0455	2.4247
3	1	$\frac{1}{16384}$	16	0.0043	0.0056	3.3310
3	1	$\frac{1}{4096}$	14	0.0063	0.0193	3.3468
3	1	$\frac{1}{1024}$	12	0.0145	0.0661	3.4017
3	4	$\frac{1}{16384}$	16	0.0054	0.0056	3.4215
3	4	$\frac{1}{4096}$	15	0.0061	0.0207	3.4374
3	4	$\frac{1}{1024}$	13	0.0103	0.0716	3.4925

The original $N = 16$ checkpoint is fixed. Prices are per coordinate-bit per date; b ranges from four to sixteen and action spacing is $1/4096$. The selected bit count minimizes the exact augmented uniform upper account, not measured policy cost. All thirteen candidate precisions and the coarser original checkpoints remain in the evidence. Displayed nonnegative accounts are rounded upward.

The historical adverse findings remain: the original nonlinear optimizer catalogue, selected residual-only recertifications, unresolved neural-ridge signs, and mixed-precision timing records are not reclassified. Identical native and ReLU envelopes remain one mathematical construction. The new deterministic repetitions do not estimate a population success probability for nonconvex optimization; the direct confidence statements concern their fixed policies under the sampling model; the stress clocks do not identify hardware-general scaling. The comparative claims in this article are restricted to the objects actually proved or executed, while the controlled-economy and application-specific theoretical scope remains intact.

16. ECONOMIC APPLICATIONS AND CONCLUSION

The original applications concern policy-dependent recursive utility, endogenous preference specifications, temporal selves, controlled diffusion problems, and dynamic games. Their complete economic statements and proofs are retained without

alteration. The centered and transport bounds explain how an NBO can support economic counterfactuals when changes in future primitives are controlled. The policy theorem explains what must be verified for the actual policy. Neither statement supplies an equilibrium or transversality condition absent from the application.

The nonlinear investment evidence adds an executable link from primitives to learned continuations and directly certified policies. Its retained-policy diagnostic demonstrates that a failed tight certificate can reflect the resolution of verification rather than the quality of the deployed policy. Its direct coupled comparison answers a different question—the relative cost of two policies—without confusing separate regret bounds with a method contrast.

The acquired-state result makes that returned policy implementable under a stated information and arithmetic contract. The constrained economy supplies finite covering, integration and repair routines; the direct scalar experiment supplies simultaneous tolerance-based cost comparisons and a catalogue-adoption rule. Their distinct assumptions and work accounts remain visible.

The compiler adds a policy-preserving computational step to this program. It connects the explicit neural continuation to an executable representation whose per-query work no longer scans the full state cover. The economic implementation adds two further choices: how precisely to acquire the state, and whether replacing an installed controller pays for itself. The first is governed by a priced, certified information account; the second by direct expected costs with an explicit replacement charge. They are not inferred by ranking regret bounds.

The empirical content is correspondingly separated. Frozen-object experiments identify representation costs for the same policy; complete services compare fresh own-future constructions against conventional FVI; direct intervals concern fixed policies and declared initial laws. The higher-dimensional executions expose the remaining tensor and innovation costs. They do not constitute a dimension-free approximation theorem, a calibration to welfare data, or a reliability law for unrestricted neural training. Within these boundaries, the operator supplies a constructive route from economic primitives through a feasible implemented policy to a verified resource decision.

Neural Bellman Operators thus organize construction, certification, economic reuse, and resource accounting around the policy actually returned. Their scientific assessment requires all of these objects. Flexible conventional methods remain important competitors; numerical activation is not a work saving, and a fixed catalogue is not a general reliability theorem. Theorem 5 and Corollary 6 now supply a separate deterministic reliability statement for an explicitly constructed neural backend. They do not retroactively prove convergence of the earlier optimizers. Preserving those distinctions makes the original economic program testable while retaining its full theory and applications.

APPENDIX A: READING AND PRESERVATION MAP

The main article is the authoritative R48 exposition of the existing Neural Bellman Operators paper. It includes the controlled-economy formulation, centered transport and policy comparison, constructive continuation and original witnesses, acquired-state feasibility, exact compilation, and the economic decision analysis. The technical supplement supplies the proofs and numerical accounts needed for those results, as well as every new direct interval and the complete first-target time partitions. All R47 main and supplement labels, including their inherited results, remain in the corresponding active documents. The current response addresses the R47 report item by item.

The experimental sections distinguish frozen historical evidence from new executions. Earlier optimizer, precision, scalar and coupled policy records are not regenerated or relabeled. The full historical theory and application companions are retained as ordinary files in the revision tree; preservation is not asserted to be a new validation of every application. The main finite-horizon certificate needs the monotonicity, cash-invariance, feasibility and comparison assumptions stated here. The direct sampling result additionally uses conditional expectation and its sampling model. Controlled diffusions, infinite horizons, recursive utility domains and equilibrium applications retain their separate analytic conditions.

The source tree contains the ordinary manuscript, proofs, bibliography, table generator, scientific sources, complete frozen records and build instructions. A clean

1 rebuild audits the records and regenerates the documents without a new scientific 1
2 run or an expiring artifact download. The publication audit binds the review anchor, 2
3 scientific source, evidence, tests and compiled documents. These checks establish 3
4 a reproducible submitted object; they are not a substitute for mathematical peer 4
5 review. 5

6 6
7 7
8 8
9 9
10 10
11 11
12 12
13 13
14 14
15 15
16 16
17 17
18 18
19 19
20 20
21 21
22 22
23 23
24 24
25 25
26 26
27 27
28 28
29 29
30 30
31 31
32 32

APPENDIX: CURRENT R52 DEVELOPMENT

APPENDIX B: INTRODUCTION

A continuation value is useful to an economist because it changes a decision. Its prediction error is not the loss of the policy actually implemented. Nor is a smaller upper bound on that loss evidence that one policy has lower cost than another. When observation, action selection and verification are costly, the numerical procedure and the economic decision have to be evaluated together.

We develop Neural Bellman Operators as a policy-evaluation and feasible-improvement method. The economic object remains a controlled economy with policy-specific continuations. The operator constructs each date's continuation using its own fitted future, returns a feasible policy, and accounts for its loss against the Bellman optimum. A general signed policy sandwich permits arbitrary fitted candidates. A constructive affine-ReLU backend additionally produces the candidate from primitive moduli, without assuming that a nonconvex training algorithm eventually succeeds. Exact compilation retains the original action witness through off-grid evaluation and state-dependent feasibility repair.

The principal question in this revision is how that construction produces a certified economic answer with finite resources. State coverage, action search and integration have different error coefficients. Tying them to one grid size conceals the allocation problem. We derive a separated primitive budget and an optimal continuous allocation for its leading work proxy, including lower resource bounds and dyadic implementation. The empirical comparison then records the complete positive-tolerance frontier of a fixed common resource ladder, not only an isolated favorable target. All failed earlier attempts belong to the work required to return a policy.

The next question is what the returned policy does. We evaluate the actual policies at the declared first crossings, even when two generators cross at different resolutions. Independent interval simulation preserves the original continuous initial and innovation laws and encloses discontinuous actors and robust capacity repair. Conditioning on construction information permits simultaneous cost inference after this selection. The result does not infer a policy ranking from two certificates and

1 does not extrapolate three deterministic timing repetitions into an optimizer-success 1
 2 probability. 2

3 Two economic decisions make this distinction operational. First, a controller 3
 4 replacement is worthwhile only if its actual gain exceeds the replacement fee. We 4
 5 identify the directional gain intervals and the full break-even fee account, retaining 5
 6 but no longer relying on a single generous charge. Second, observation precision is 6
 7 chosen on the same finite bit grid on which the implemented policies are evaluated. 7
 8 A simultaneous upper-net-cost selector has a regret bound relative to the best 8
 9 precision in that catalogue. The bound holds for every nonnegative information 9
 10 price on one coverage event. This is a result about actual net policy cost, different 10
 11 from minimizing a conservative all-state loss account. 11

12 These advances do not depend on calling a classical envelope neural. A native 12
 13 min-plus envelope and its exactly equivalent ReLU circuit are one method with 13
 14 two encodings. The original-owner compiler is available to both. We compare that 14
 15 method with uniform multilinear fitted-value iteration (FVI) and two separately 15
 16 recorded conventional refinement rules. The original curvature-bisection block 16
 17 remained uniform; the later graded block was nonuniform. A new residual-driven 17
 18 comparator below records its realized geometry and all pilot work. Each generates 18
 19 its own future labels and receives the same feasible action menu, continuous-law 19
 20 integration and one-sided economic criterion. The evidence must determine the cost 20
 21 of the returned policy as well as certificate sharpness. No uniform neural dominance 21
 22 is built into the design. 22

23
 24 *Related methods.* Lipschitz extension supplies the cone construction ([McShane](#), 24
 25 [1934](#)), and separable distance transforms supply its linear-work nodal engine 25
 26 ([Felzenszwalb and Huttenlocher, 2012](#)). Approximation in min-plus spaces has 26
 27 an established dynamic-programming literature: [Lakshminarayanan and Bhatnagar](#) 27
 28 ([2014](#)) study min-plus projected Bellman equations and supremum-norm errors. 28
 29 Bellman-inequality methods likewise construct lower value bounds and policies; 29
 30 [Wang et al. \(2015\)](#) use iterated inequalities with linear or semidefinite optimization. 30
 31 Our argument does not claim priority for these principles. Its subject is their con- 31
 32 nection to an own-future feasible policy, original-owner continuous compilation, 32

acquired-state implementation, complete resource allocation and direct economic comparison of the returned objects. The empirical-Bernstein component uses the bounded-sample inequality of [Maurer and Pontil \(2009\)](#); conditional selection and numerical path enclosures specify the estimand to which it is applied.

The broader NBO program includes controlled diffusions, recursive utility, endogenous preferences, temporal selves and games. Those results and their assumptions are retained in the complete development edition and its full proof supplement, together with the original training, matrix-factor and precision experiments. The present article supplies a self-contained constructive theorem–evidence chain. Its direct-cost statistical theorems concern expected discounted costs, not an undeclared extension to every nonlinear recursive application. This organization preserves the paper’s economic subject while making the current contribution and the actual comparison burden explicit.

The additional improvement result addresses the economically important case where the certified construction returns a more costly policy than a conventional competitor. We preserve that incumbent and permit only actions whose actual reference advantage is certified nonpositive after the numerical and observation allowances. At the final date of the original investment model, exact continuous-law integration makes this gate executable without assuming a successful policy-value fit. Both generators receive the same repair. The resulting experiment isolates removable final-action loss from the policy gap that remains; it does not rename the compiler as a uniquely neural advantage.

APPENDIX C: THE ECONOMIC OPERATOR AND ITS POLICY ACCOUNT

Consider dates $t = 0, \dots, T$ and an invariant compact state domain $X \subset \mathbb{R}^d$. At date t , the action belongs to a nonempty compact set $A_t(x)$. The transition is $F_t(x, a, Z_{t+1})$, the current cost is $c_t(x, a)$, and the terminal cost is $g(x)$. Write $B_0 = 1$ and $B_t = \prod_{j < t} \beta_j$, with $\beta_j > 0$. An implemented policy includes its observation rule, action selection and feasible numerical realization, not just its stored parameters.

For a continuation v , define

$$Q_t(v; x, a) = c_t(x, a) + \beta_t \rho_t(v(F_t(x, a, Z))), \quad \mathcal{T}_t v(x) = \inf_{a \in A_t(x)} Q_t(v; x, a). \quad (65)$$

Throughout the deterministic comparison, the evaluations ρ_t are monotone and cash invariant on a common domain containing all functions and constant translates used below. They are therefore nonexpansive in the supremum norm. Continuous objectives, effective compact feasible correspondences and measurable minimizers, or measurable arbitrarily near-minimizers with their error charged, are assumed. The stipulated conditional recursion gives Bellman comparison against all admissible adapted controls. In the experiments ρ_t is conditional expectation, costs are bounded and innovations have their original continuous uniform law. The direct-cost inference in Section J is specific to expectation; it is not automatically an inference theorem for nonlinear recursive utility or games.

A Neural Bellman Operator is a construction and evaluation procedure for (65). It forms targets using its *own* fitted f_{t+1} , returns a neural continuation and a feasible policy, and records an economic error allowance. It does not receive the optimal continuation as a training label. Arbitrary trained networks can be evaluated by the first result below. The subsequent constructive backend specifies how a network and its actor are obtained without assuming successful nonconvex optimization.

C.1. Action-dependent errors and complete-policy loss

Let $H(a)$ be a true current cost and $\hat{H}(a)$ its continuation approximation. If $\hat{a}$ is δ -optimal for $\hat{H}$ on the same feasible action set, then

$$H(\hat{a}) - \inf_a H(a) \leq \delta + \text{osc}_a \{\hat{H}(a) - H(a)\}. \quad (66)$$

Indeed, compare $\hat{a}$ with an arbitrarily near-optimal true action and cancel the approximate objectives. The infimum limit gives the result. Additive continuation levels cancel; action-dependent discrepancies do not. Under a change in future policies, the continuation must be transported to the new economic problem before this inequality licenses reuse. The full development edition retains the original quantitative transport, recursive-domain and economic-application results. No transport bound is presumed valid after an unaccounted change of constraints or preferences.

PROPOSITION 21—One-sided policy sandwich: *Suppose, uniformly in state,*

$$f_t - \mathcal{T}_t f_{t+1} \leq u_t, \quad \mathcal{T}_t^\pi f_{t+1} - f_t \leq v_t, \quad \ell_T \leq f_T - g \leq u_T.$$

Then the actual feasible policy satisfies

$$0 \leq J_0^\pi - V_0^* \leq \sum_{t < T} B_t(u_t + v_t) + B_T(u_T - \ell_T). \quad (67)$$

The bound is invariant to date-specific constant shifts of the fitted continuations, provided the associated residual accounts are shifted consistently.

PROOF: Backward comparison and cash invariance give

$$V_t^* \geq f_t - \sum_{k=t}^{T-1} B_{t,k} u_k - B_{t,T} u_T, \quad J_t^\pi \leq f_t + \sum_{k=t}^{T-1} B_{t,k} v_k - B_{t,T} \ell_T,$$

where $B_{t,k} = \prod_{j=t}^{k-1} \beta_j$. The terminal claim is the stipulated enclosure; applying the Bellman or fixed-policy operator proves each inductive step. Subtraction gives (67). A shift $f_t \mapsto f_t + c_t^0$ adds $c_t^0 - \beta_t c_{t+1}^0$ to u_t and subtracts it from v_t . The terminal width is unchanged. *Q.E.D.*

A sharper valid residual account for the same continuation and the same deployed actor sharpens knowledge of the policy. It does not retrain that policy or change its realized cost. Conversely, subtracting two upper bounds in (67) does not identify the difference in actual costs. These distinctions govern both the historical comparisons and the new experiments.

APPENDIX D: CONSTRUCTING AND COMPILING A FEASIBLE NEURAL POLICY

Assume the primitive objective has, for feasible state–action pairs, state modulus $A_t^Q = C_{x,t} + \beta_t F_{x,t} L_{t+1}$ and action modulus $D_t = C_{a,t} + \beta_t F_{a,t} L_{t+1}$. Let

$$d_H(A_t(x), A_t(z)) \leq K_t^A \|x - z\|_1, \quad L_t = A_t^Q + D_t K_t^A, \quad L_T = L_g.$$

The Hausdorff condition and the primitive moduli imply that $\mathcal{T}_t f_{t+1}$ is L_t -Lipschitz whenever f_{t+1} is L_{t+1} -Lipschitz. All moduli and numerical integration procedures are supplied by the economic primitives, not fitted from successful outcomes.

At state nodes $x_{t,i}$ with covering radius h_t , form effective feasible k_t -nets. Evaluate their own-future objectives with certified error at most e_t . Minimize the computed labels and retain the corresponding *original* action $a_{t,i}$. A Borel repair $\Gamma_{t,i}(x)$ is feasible at x and satisfies

$$\|\Gamma_{t,i}(x) - a_{t,i}\|_1 \leq K_t^A \|x - x_{t,i}\|_1 + \zeta_t. \quad (68)$$

The allowance ζ_t concerns displacement; feasibility itself may not fail. Define

$$f_t(x) = \min_i \{y_{t,i} + L_t \|x - x_{t,i}\|_1\}, \quad i_t(x) = \min \arg \min_i \{y_{t,i} + L_t \|x - x_{t,i}\|_1\}. \quad (69)$$

Return $\pi_t(x) = \Gamma_{t,i_t(x)}(x)$. Terminal labels approximate g with error e_T and use the same cone construction with modulus L_g .

THEOREM 22—Own-future construction with feasible witnesses: *The functions in (69) have exact affine-ReLU realizations and Lipschitz moduli L_t independent of label error. The returned policy is feasible and measurable. It satisfies*

$$-e_t \leq f_t - \mathcal{T}_t f_{t+1} \leq D_t k_t + e_t + 2L_t h_t, \quad (70)$$

$$\mathcal{T}_t^\pi f_{t+1} - f_t \leq e_t + D_t \zeta_t,$$

and hence

$$J_0^\pi - V_0^* \leq \sum_{t < T} B_t(D_t k_t + 2e_t + 2L_t h_t + D_t \zeta_t) + B_T(2e_T + 2L_g h_T). \quad (71)$$

With effective covers, repairs and certified queries at the prescribed accuracy, allocating each nonnegative term a share of a positive economic tolerance gives a terminating finite construction. A fixed numerical precision or a capped finite experimental ladder need not contain that allocation.

PROOF: The selected node label lies between the node Bellman infimum minus e_t and that infimum plus $D_t k_t + e_t$. The Lipschitz property of the infimum implies that every cone lies above it minus e_t . At a nearest node its cone exceeds the infimum by at most $D_t k_t + e_t + 2L_t h_t$. At the attaining original index, the two feasible pairs and

(68) give

$$Q_t(f_{t+1}; x, \Gamma_{t,i}(x)) \leq y_{t,i} + e_t + L_t \|x - x_{t,i}\|_1 + D_t \zeta_t.$$

This proves (70). The terminal residual belongs to $[-e_T, e_T + 2L_g h_T]$; Proposition 21 gives the loss bound. A finite minimum of equally Lipschitz cones is equally Lipschitz. The identities $|z| = \text{ReLU}(z) + \text{ReLU}(-z)$ and $\min(u, v) = u - \text{ReLU}(u - v)$ give the exact network. Finite lexicographic comparison and Borel repair give measurability. Effective refinement of the finitely many positively weighted terms yields the termination assertion. *Q.E.D.*

The Lipschitz-extension principle is classical (McShane, 1934). A native min-plus envelope and its exact ReLU circuit, using the same labels, witnesses and ties, are *one method with two encodings*. They return the same function and policy. Neither Theorem 22 nor a faster representation proves a uniquely neural capability or a training-success probability for Adam or L-BFGS. The mathematical increment is the complete feasible, own-future construction and implementation account; the empirical question is its work and policy value relative to conventional generators.

D.1. Original-owner compilation on a tensor cover

On a tensor cover $\mathcal{X} = \prod_{j=1}^d \{x_{j,0}, \dots, x_{j,n_j}\}$, put $S = \prod_j (n_j + 1)$. For arbitrary labels y_i , including inconsistent inexact labels, compute at each grid node v

$$(z_v, o_v) = \min_{\text{lex}, i} (y_i + L \|v - x_i\|_1, i).$$

The index o_v is the original label owner, not the transformed node v .

THEOREM 23—Continuous original-owner identity: *For a closed tensor cell C with corners $\mathcal{V}(C)$, and $x \in C$,*

$$f(x) = \min_{v \in \mathcal{V}(C)} (z_v + L \|x - v\|_1), \tag{72}$$

$$(f(x), i_*(x)) = \min_{\text{lex}, v \in \mathcal{V}(C)} (y_{o_v} + L \|x - x_{o_v}\|_1, o_v).$$

Coordinatewise forward and backward relaxations compute (z, o) in $O(dS)$ additions and comparisons and $O(S)$ storage. After locating the cell, a query uses $O(d2^d)$ arithmetic

operations. Exact rational comparisons preserve all ties, including $L = 0$. Using the original owner's repaired action preserves the uncompiled policy, its cost and its certificate exactly.

PROOF: The triangle inequality bounds $f(x)$ above by each corner expression. Choose the lexicographically first attaining site $i_*(x)$. In every coordinate choose the cell endpoint between x_j and $x_{i_*,j}$. At this corner distances add exactly. An owner of strictly smaller transformed value would improve the objective at x ; a tied smaller owner would contradict the original tie rule. Thus this corner has owner $i_*(x)$, proving both identities. Along a coordinate line, forward and backward relaxations respectively minimize over sites on either side, always comparing value–original-index pairs. Iterating over coordinates gives the separable distance transform and its stated work. Rational arithmetic and unchanged repair complete the claim. *Q.E.D.*

The nodal algorithm is the classical distance transform ([Felzenszwalb and Huttenlocher, 2012](#)); preserving the original action through continuous off-grid queries is the policy obligation. If there are $J_{t,i}$ node actions and M_t integration representatives, write $Q_t = M_t \sum_i J_{t,i}$. Literal dense cone evaluation costs $O(\sum_t dS_{t+1}Q_t)$ in representation work. Compilation replaces this by $O(\sum_t [dS_{t+1} + Q_t d2^d])$, plus cell location, primitive evaluation, integration, feasible repair, certificate and output. Rational bit lengths and state boxes crossing many cells incur additional costs. There is no dimension-free complexity assertion.

D.2. Acquired states and conventional actors

A b -bit coordinate sensor reports a dyadic cell and its midpoint $\hat{x}$. On the unit cube, $\|x - \hat{x}\|_1 \leq r(b) = d2^{-b-1}$. For the investment capacity $\kappa(x)$ below, clip the original action to the smallest capacity throughout that cell and round down to an action quantum Δ_t . This is feasible at every state in the cell. Relative to an exact true-state repair, its displacement is at most $2K_t^A r(b) + \Delta_t$. Transporting the witness cone from $\hat{x}$ to x adds at most $2L_t r(b)$, plus any certified selector objective error v_t . Thus the implemented-policy account adds

$$\sum_{t < T} B_t \{ (2L_t + 2D_t K_t^A) r(b) + D_t \Delta_t + v_t \}. \quad (73)$$

No continuity of the discrete owner selector is assumed. At uncertain numerical comparisons the simulator retains all possible actors, rather than choosing a favorable one.

For conventional FVI the continuation is multilinear on its realized tensor grid. Its actual modulus is bounded by the largest coordinate edge slope. A nearest-node actor uses a feasible node minimizer and true-state repair. Let h_t be the largest half-cell ℓ^1 radius, L_t^{gr} the Bellman graph modulus, and e_t the node query allowance. Its optimal residual upper bound is $D_t k_t + e_t + L_t^{\text{gr}} h_t$. Its selected-policy upper bound is $e_t + E_t$, where

$$E_t = \sup_x \{y_{i(x)} + L_t^{\text{gr}} \|x - x_{i(x)}\|_1 - f_t(x)\}.$$

Partitioning cells by nearest-node boundaries makes the expression separately affine in each coordinate. Its maximum is therefore at a partition vertex; the implementation encloses all such vertices. Both uniform and curvature-refined FVI receive this sharper one-sided account, not an artificially doubled baseline. Their acquired-state allowance is derived separately in (109).

APPENDIX E: COST-DIRECTED FEASIBLE IMPROVEMENT

A policy certificate answers a safety question relative to the optimum. It need not answer the comparison with an installed controller. We now make actual cost reduction a constraint on the improvement step. The continuation representation, including a neural realization, supplies proposals; a centered, policy-specific evaluation determines whether to implement them. The incumbent is always an admissible fallback.

Fix a feasible implemented policy π and bounded functions $h_0, \dots, h_T$ in the common evaluation domain of Section C. Its observation and rounding rules are part of π . Suppose verified constants satisfy

$$\ell_t \leq \mathcal{T}_t^\pi h_{t+1} - h_t \leq u_t \quad (t < T), \quad \ell_T \leq g - h_T \leq u_T. \quad (74)$$

Define $a_T = \ell_T$, $b_T = u_T$ and

$$a_t = \ell_t + \beta_t a_{t+1}, \quad b_t = u_t + \beta_t b_{t+1}, \quad w_t = b_t - a_t.$$

All inequalities in this section are uniform over the true state domain. Bounds at sampled points alone do not meet this requirement.

PROPOSITION 24—Centered reference advantage: *Under (74), $a_t \leq J_t^\pi - h_t \leq b_t$. For any feasible action v at state x , put*

$$d_t(x, v) = Q_t(h_{t+1}; x, v) - Q_t(h_{t+1}; x, \pi_t(x)).$$

If $\bar{d}_t(x, v)$ is a valid upper enclosure of this difference, then

$$Q_t(J_{t+1}^\pi; x, v) - J_t^\pi(x) \leq \bar{d}_t(x, v) + \beta_t w_{t+1}. \quad (75)$$

The width allowance and the action contrast are invariant to date-specific constant shifts of the fitted continuations.

PROOF: Monotonicity and cash invariance propagate the two error endpoints backward. Applying them once more to each of the two feasible actions bounds each continuation-evaluation error between $\beta_t a_{t+1}$ and $\beta_t b_{t+1}$. Their difference is at most $\beta_t w_{t+1}$. Date-specific shifts translate both endpoints by the same amount and cancel in the action contrast. *Q.E.D.*

The proposition is deliberately a reference-policy result. Substituting a bound on $h - V^*$ without controlling $h - J^\pi$ changes the statement. Likewise, accepting an action merely because it lowers $Q_t(h_{t+1})$ ignores the reference-evaluation allowance.

THEOREM 25—Incumbent-preserving Bellman improvement: *Suppose a measurable feasible proposal v_t and a nonnegative measurable margin κ_t are given. At a state where a verified enclosure satisfies*

$$\bar{d}_t(x, v_t(x)) + \beta_t w_{t+1} \leq -\kappa_t(x), \quad (76)$$

implement $v_t(x)$; otherwise implement the original $\pi_t(x)$. Call the resulting policy μ . For a cell-based observation rule, require feasibility and (76) for every true state in the reported cell, using the same proposed action on that cell. Then $J_t^\mu(x) \leq J_t^\pi(x)$ at every date and state.

Under conditional expectation, let D_t be the accepted-state set and let the state trajectory follow μ . Then the stronger quantitative assertion is

$$J_0^\mu(x) - J_0^\pi(x) \leq -\mathbb{E}_x^\mu \sum_{t < T} B_t \kappa_t(X_t) \mathbf{1}_{D_t}(X_t). \quad (77)$$

The qualitative assertion holds for the stipulated monotone cash-invariant conditional recursion. The occupancy formula is asserted only for expectation.

PROOF: At accepted states Proposition 24 gives $\mathcal{T}_t^\mu J_{t+1}^\pi \leq J_t^\pi - \kappa_t$. Elsewhere the two actions coincide, so the advantage is exactly zero, not $\beta_t w_{t+1}$. Backward monotone comparison, starting from a common terminal cost, gives the first assertion. Under expectation subtract $J_t^\pi(X_t)$ along the μ trajectory and telescope the conditional expectations. The resulting identity is the expected discounted sum of true reference advantages, each bounded above by the accepted negative margin or by zero. *Q.E.D.*

The improvement principle is related to approximate policy iteration with safeguarded updates, including Metelli et al. (2021). We do not claim that monotone policy improvement itself is new. Here the relevant connection is between a signed own-policy evaluation account, action-centered numerical errors, original-owner construction, and the actual information-constrained action. The test is available to neural and conventional proposals alike. An exact native encoding of a ReLU continuation does not become a second independent method.

E.1. A fully executed final-date specialization

A useful test of this result leaves every earlier decision fixed and improves only the final action. Set $h_T = g$. The continuation width is then zero, so the gate concerns the actual conditional final cost, without a fitted policy-evaluation error. For the investment economy in Section L, write

$$y_j = \frac{1}{16} + \frac{x_j}{2} + \frac{x_{j+1}}{8} + \frac{x_j(1 - x_{j+1})}{16} + \gamma_j a, \quad s_j = (-1)^{j+1}, \quad \sigma_z^2 = \frac{1}{3072},$$

where the cyclic convention and γ_j are as in that section. Put $b = 1/2 - 2\sum_j y_j/d$,
 $r = |\sum_j s_j|/(16d)$, and

$$\Psi_2(b, r) = \begin{cases} b_+^2, & r = 0, \\ ((b+r)_+^3 - (b-r)_+^3)/(6r), & r > 0. \end{cases} \quad (78)$$

Exact integration under the original continuous innovation law gives

$$\begin{aligned} \mathbb{E}g(y + sZ) &= \frac{4}{d} \sum_j (y_j - 5/8)^2 + \frac{1}{4d} \sum_j (y_j - y_{j+1})^2 + 2\Psi_2(b, r) \\ &\quad + 4\sigma_z^2 + \frac{\sigma_z^2}{4d} \sum_j (s_j - s_{j+1})^2. \end{aligned} \quad (79)$$

Thus no terminal innovation mesh is needed for this improvement step. Its computation is nevertheless charged. The formula includes the odd-dimensional shortage uncertainty; setting it to zero in dimension three would change the model.

COROLLARY 26—Acquired-cell final improvement: *In the declared investment economies, final conditional cost $H(x, a) = c(x, a) + \beta \mathbb{E}g(F(x, a, Z))$ is strictly convex in a , with second derivative at least $2p$. Fix an observation cell C and the implemented incumbent action a_C . Choose a lattice minimizer a_C^* at the reported midpoint on the cell-wide feasible set. Accept it only if an outward enclosure of $H(x, a_C^*) - H(x, a_C)$ is nonpositive for all $x \in C$. Keep the incumbent on every other cell and at all earlier dates. The resulting acquired policy weakly improves actual expected discounted cost for every initial state. This guarantee does not compare it to a different full-information incumbent.*

The lattice is $\{k/4096 : k \in \mathbb{N}_0\}$, truncated by the minimum capacity on the eight-bit cell. Convex derivative bisection identifies the two neighboring lattice actions; their actual cost contrast selects the midpoint proposal. The gate compares its cost with the incumbent on the entire cell, not just at the midpoint. Exact rational point comparison resolves a numerically ambiguous proposal tie. An uncertain true-state box may intersect several cells; the simulator retains their action hulls rather than selecting a favorable cell.

For each generator, old and repaired policies share their earlier trajectory. Their discounted cost difference therefore equals the conditional final contrast. Centering cancels state-only costs and tightens its interval enclosure. When two generators are compared, their trajectories differ and are simulated separately using common innovations. Both comparisons concern actual policies; neither is obtained by subtracting two policy-loss certificates.

APPENDIX F: FROM SAFE IMPROVEMENT TO FINITE-SWEEP ACCURACY

The preceding improvement result compares an implemented policy with its own incumbent. We now connect that comparison to the original Bellman optimum. This distinction matters economically: a policy can improve strictly while remaining more costly than a conventional alternative. The result below preserves the original objective and gives explicit sufficient conditions for closing its optimality gap; it does not infer those conditions from a small empirical residual.

Consider the finite-horizon expectation model of the paper, with bounded costs, feasible measurable Markov policies, and discount factors $\beta_t \geq 0$, $0 \leq t < T$. Write

$$Q_t(h; x, a) = c_t(x, a) + \beta_t \mathbb{E}[h(X_{t+1}) \mid x, a], \quad (T_t h)(x) = \inf_{a \in A_t(x)} Q_t(h; x, a),$$

and $V_T^* = g$, $V_t^* = T_t V_{t+1}^*$. The benchmark is the original fully observed optimum. An acquired-state policy remains feasible for this benchmark; a restriction to cell-wise actions may, however, impose a positive approximation floor. No vanishing information error is assumed.

Starting from a feasible π^0 , at pass k let h^k have certified signed policy-evaluation bands

$$a_t^k \leq J_t^{\pi^k}(x) - h_t^k(x) \leq b_t^k, \quad w_t^k = b_t^k - a_t^k. \quad (80)$$

These can be obtained from the signed-residual recursion already proved; the approximator need not have a small absolute level error. Let a feasible proposal v_t^k satisfy, uniformly in x ,

$$Q_t(h_{t+1}^k; x, v_t^k(x)) \leq (T_t h_{t+1}^k)(x) + \alpha_t^k. \quad (81)$$

For its centered comparison with the exact incumbent, define $d_t^k(x) = Q_t(h_{t+1}^k; x, v_t^k(x)) - Q_t(h_{t+1}^k; x, \pi_t^k(x))$. A quantitative upper enclosure satisfies

$$d_t^k(x) \leq \bar{d}_t^k(x) \leq d_t^k(x) + \zeta_t^k. \quad (82)$$

Here ζ_t^k includes acquisition-cell variation and integration or arithmetic enclosure width, not just the machine unit roundoff. The new policy uses the proposal exactly when $\bar{d}_t^k(x) + \beta_t w_{t+1}^k \leq 0$ and otherwise uses $\pi_t^k(x)$ unchanged. For an acquired cell, feasibility and this comparison must hold throughout that cell, with a single implementable action and acceptance decision. Bounds (81)–(82) are required on the same cells, including their boundaries.

THEOREM 27—Finite-sweep accuracy with incumbent preservation: *Suppose (80)–(82) hold at each pass, with nonnegative finite α_t^k and ζ_t^k . Put*

$$\epsilon_t^k = \alpha_t^k + \zeta_t^k + 2\beta_t w_{t+1}^k, \quad E_t^k = \|J_t^{\pi^k} - V_t^*\|_\infty.$$

Then $J_t^{\pi^{k+1}} \leq J_t^{\pi^k}$ at every date and state, and

$$E_t^{k+1} \leq \beta_t E_{t+1}^k + \epsilon_t^k, \quad E_T^k = 0. \quad (83)$$

Set $B_{t,0} = 1$ and $B_{t,j} = \prod_{r=t}^{t+j-1} \beta_r$. For $1 \leq m \leq T - t$,

$$E_t^m \leq B_{t,m} E_{t+m}^0 + \sum_{j=0}^{m-1} B_{t,j} \epsilon_{t+j}^{m-1-j}. \quad (84)$$

In particular, if $\epsilon_t^k \leq \bar{\epsilon}_t$ uniformly in k , then after T passes

$$E_t^T \leq \sum_{j=0}^{T-t-1} B_{t,j} \bar{\epsilon}_{t+j}. \quad (85)$$

With exact policy evaluation up to a datewise constant, exact greedy proposals, and exact centered comparisons, the policy is optimal after at most T passes, and no intermediate pass raises its cost.

The hypotheses separate three computational tasks. A trained continuation must be evaluated relative to the policy actually retained; the action search must ap-

proximate the original feasible infimum; and the comparison must be valid over the available information set. A guarantee for only one of these tasks does not imply the theorem. The same conclusion applies to native and neural representations when they implement the same objects. In particular, it does not turn their common compilation saving into a neural-specific advantage. Safe policy improvement is an established principle; the result here supplies the signed-width and acquisition-aware error account for the incumbent-preserving NBO construction.

PROPOSITION 28—An implementable cellwise proposal bound: *Let an acquired cell C have representative x_C and radius r_C . Suppose $Q_t(h; \cdot, a)$ is L_t -Lipschitz on the cell for every menu action, and $Q_t(h; x, \cdot)$ is D_t -Lipschitz on feasible actions. Let $M_C \subseteq \bigcap_{x \in C} A_t(x)$ be a nonempty finite menu such that every $a \in A_t(x)$, for every $x \in C$, lies within $\delta_{A,C}$ of some member of M_C . If $v_C \in M_C$ is an η_C -minimizer at x_C , then it satisfies (81) on C with*

$$\alpha_t^k \leq \eta_C + 2L_t r_C + D_t \delta_{A,C}.$$

The directed-cover assumption is substantive when constraints move with the state. It is not implied by a dense grid at the representative alone. A failed cover or a cell that cannot certify a comparison retains its incumbent; it must not be counted as a successful accuracy refinement.

PROPOSITION 29—Sharpness of the evaluation-width coefficient: *Under the theorem’s assumptions, the coefficient two on $\beta_t w_{t+1}^k$ in (83) cannot be uniformly replaced by a smaller constant, even for a one-period deterministic problem with two actions, an exact greedy proposal, and exact contrast arithmetic.*

The theorem is constructive but conditional on verified error budgets. It does not assert that a continuous neural fit uniformly approximates a discontinuous acquired-policy value, or that refining a fixed acquisition contract eliminates its information floor. The executed investment study below certifies a final-date improvement. It is not an execution of T full improvement passes, nor numerical evidence of the finite-sweep near-optimality conclusion. Exact finite-state regression tests exercise

the new identities, rejection cases, and sharpness example separately from the economic benchmark.

APPENDIX G: POLICY-RELEVANT CONTINUATION ERRORS

A continuation error matters for an investment decision through the actions it changes. The signed-width account above removes a datewise level error. It can nevertheless charge for a nonconstant error that has the same conditional expectation under every feasible action. We now obtain a sharper account without changing the Bellman optimum, the acquisition contract, or the incumbent-preserving rule.

Write $P_t^a f(x) = \mathbb{E}[f(X_{t+1}) \mid x, a]$, and at pass k put $e_t^k = J_t^{\pi^k} - h_t^k$. An *action-contrast certificate* is a measurable nonnegative bound $R_t^k(x; a, b)$ satisfying

$$|(P_t^a - P_t^b)e_{t+1}^k(x)| \leq R_t^k(x; a, b) \leq \chi_t^k \quad (a, b \in A_t(x)). \quad (86)$$

Both inequalities are required uniformly over true states and feasible actions. The scalar certificate $R_t^k = w_{t+1}^k$ always remains available. A bound computed for the selected proposal alone is sufficient for safety, but not for the accuracy conclusion, which compares with the entire feasible infimum.

THEOREM 30—Contrast-certified finite-sweep improvement: *In the expectation model of Theorem 27, retain the feasible proposal and comparison-enclosure assumptions (81)–(82). Suppose (86) holds. Accept $v_t^k(x)$ precisely when*

$$\bar{d}_t^k(x) + \beta_t R_t^k(x; v_t^k(x), \pi_t^k(x)) \leq 0; \quad (87)$$

otherwise retain the exact incumbent. On an acquired cell, the proposal, acceptance decision, feasibility, and all enclosures must be valid on the whole cell. Then every pass is nonworsening at every state and date. Moreover,

$$E_t^{k+1} \leq \beta_t E_{t+1}^k + \tilde{\epsilon}_t^k, \quad \tilde{\epsilon}_t^k = \alpha_t^k + \zeta_t^k + 2\beta_t \chi_t^k. \quad (88)$$

The finite-sweep unrolling and T-pass bound of Theorem 27 hold with ϵ replaced by $\tilde{\epsilon}$. In particular, exact greedy search and comparisons with $\chi_t^k = 0$ give an optimal policy after at most T passes, even when the global evaluation widths are positive. The coefficient two remains sharp over this class of certificates.

The theorem concerns expected discounted cost. It does not silently extend the linear action-kernel difference to a nonlinear recursive preference. The earlier monotone-recursion safety results remain in force under their own hypotheses. Nor does a small sampled action contrast establish (86). The next result constructs an upper bound using residuals and couplings with verified marginals; it never calls the unknown policy value.

PROPOSITION 31—Coupled residual certificate: *Fix a policy π and critic h . Let $r_T = g - h_T$ and $r_t = \mathcal{T}_t^\pi h_{t+1} - h_t$, and let w_t be valid signed-residual widths. For every pair of states let $\Lambda_t^\pi(x, y)$ be a measurable coupling of $P_t^{\pi_t(x)}(x, \cdot)$ and $P_t^{\pi_t(y)}(y, \cdot)$. Suppose verified nonnegative functions H_t enclose $|r_t(x) - r_t(y)|$. Define*

$$\begin{aligned} D_T(x, y) &= \min\{w_T, H_T(x, y)\}, \\ D_t(x, y) &= \min\left\{w_t, H_t(x, y) + \beta_t \int D_{t+1}(u, v) d\Lambda_t^\pi(x, y)(u, v)\right\}. \end{aligned} \quad (89)$$

Then $|e_t(x) - e_t(y)| \leq D_t(x, y)$. For any measurable coupling $\Lambda_t(x; a, b)$ of the two feasible action kernels,

$$R_t(x; a, b) = \min\left\{w_{t+1}, \int D_{t+1}(u, v) d\Lambda_t(x; a, b)(u, v)\right\} \quad (90)$$

is a valid action-contrast certificate. Replacing either integral by a verified nonnegative upper enclosure preserves validity. No optimal-transport minimizer or smooth incumbent is required.

The construction includes the work needed to enclose residual differences, transport them, and take the action supremum. It is an alternative certificate, not a free evaluation oracle. In a finite model, a common-uniform coupling gives an explicit terminating implementation. In a continuous model, an available common innovation gives a coupling directly, but integration and cellwise certification still have to be supplied. With coarse enclosures the minimum in (90) returns the original scalar guarantee.

G.1. Action-null components and neural realizations

Define the action-null space at date t by

$$\mathcal{N}_t = \{n : (P_t^a - P_t^b)n(x) = 0 \text{ for every } x \text{ and } a, b \in A_t(x)\}.$$

This is a linear space of bounded measurable functions, not a statement about the accuracy of a training procedure. Constants belong to it; exogenous payoff components can make it larger. Values differing by an element of $\mathcal{N}_t$ have identical action contrasts at date t , although their levels and state dependence need not agree.

COROLLARY 32—Certified removal of action-null error: *Suppose $e_{t+1}^k = n_{t+1}^k + u_{t+1}^k$, where membership $n_{t+1}^k \in \mathcal{N}_t$ is established and u_{t+1}^k has a verified range of width v_{t+1}^k . Then $R_t^k \leq v_{t+1}^k$ is available, independently of the amplitude of n_{t+1}^k . A constructive way to obtain this width is to evaluate the corrected critic $h^k + n^k$ by the original signed-residual recursion. The action comparisons using h_{t+1}^k and $h_{t+1}^k + n_{t+1}^k$ coincide.*

The qualification “established” is important. Arbitrarily declaring part of a fit to be null invalidates the certificate. If an exact null component is known, symbolic cancellation before interval arithmetic avoids subtracting two large levels. If it is only approximately null, its verified nonzero action contrast must be added to the allowance.

PROPOSITION 33—A nonconstant ReLU component in the original economy: *In the investment dynamics (112), for each declared dimension $d \geq 2$, every bounded measurable function of $y_1 - 2y_2$ is action-null under the common innovation. In particular,*

$$n_M(y) = M(y_1 - 2y_2 + 1)_+, \quad M \geq 0,$$

is a one-unit ReLU realization with oscillation $2M$ on the state cube. At the final date, the critic $h_T = g - n_M$ has global evaluation width $2M$ but zero action-contrast error. The contrast-aware terminal policy is therefore exactly the unperturbed policy of Corollary 26; its statewise cost guarantee is unchanged for every M .

This result supplies a representation-level robustness property of NBO, not an identified advantage over a native encoding of the same continuation. It also iden-

ifies an avoidable source of conservatism in a global critic-error account. It does not remove an economically relevant continuation error, a capacity constraint, or an information floor. Full proofs and an exact finite-state reference implementation are in the supplement.

APPENDIX H: DIRECTED CERTIFICATES AND COMPLETE POLICY IMPROVEMENT

The action-contrast result permits an error component to be large when it is irrelevant to a decision. A second source of conservatism remains: taking an absolute value before comparing two actions discards the direction of a useful continuation difference. We retain that direction throughout the verification calculation. The policy generator, the original economic primitives, and the Bellman optimum are unchanged.

H.1. Whole-cell improvement and accuracy

Fix a complete incumbent π^k and write

$$A_t^k(x, a) = c_t(x, a) + \beta_t P_t^a J_{t+1}^{\pi^k}(x) - J_t^{\pi^k}(x).$$

Thus $A_t^k(x, \pi_t^k(x)) = 0$. Let C be an observation cell. Candidate actions must be implementable and feasible for every true state in C . A finite family $\mathcal{B}_t(C)$ must cover every feasible action at every such state; a covering set may contain additional, infeasible actions if its lower enclosure remains valid on the feasible subset. Suppose verified numbers satisfy

$$\begin{aligned} A_t^k(x, a_i) &\leq u_t^k(C, a_i) & (x \in C), \\ \ell_t^k(C, B) &\leq A_t^k(x, a) & (x \in C, a \in B \cap A_t(x)). \end{aligned}$$

All numerical errors are included in these endpoints. The incumbent is an additional candidate with exact upper endpoint zero. Put

$$U_t^k(C) = \min\{0, \min_i u_t^k(C, a_i)\}, \quad L_t^k(C) = \min\{0, \min_{B \in \mathcal{B}_t(C)} \ell_t^k(C, B)\}. \quad (91)$$

Select an action attaining $U_t^k(C)$, retaining the incumbent on a tie, and implement all selected dates simultaneously. This is a finite computation; it does not presume access to the policy value in the display defining A .

THEOREM 34—Directed finite-sweep certificate: *Consider the bounded, finite-horizon expected-cost model of Theorem 30. Suppose the whole-cell feasibility and enclosure conditions above hold at every pass. Set*

$$\varepsilon_t^k = \sup_C \{U_t^k(C) - L_t^k(C)\}.$$

Then $\varepsilon_t^k \geq 0$, $J_t^{\pi^{k+1}}(x) \leq J_t^{\pi^k}(x)$ at every state and date, and

$$E_t^{k+1} \leq \beta_t E_{t+1}^k + \varepsilon_t^k, \quad E_t^k = \sup_x \{J_t^{\pi^k}(x) - J_t^*(x)\}. \quad (92)$$

Writing $B_{t,0} = 1$ and $B_{t,j} = \prod_{i=0}^{j-1} \beta_{t+i}$, for $K \geq 1$ and $m = \min\{K, T - t\}$,

$$E_t^K \leq \sum_{j=0}^{m-1} B_{t,j} \varepsilon_{t+j}^{K-1-j} + B_{t,m} E_{t+m}^{K-m}.$$

The final term vanishes when $K \geq T - t$. In particular, exact directed greedy certificates yield the original Bellman optimum within T passes. The coefficient one on ε cannot be reduced uniformly.

PROOF: For every $x \in C$, the selected action has advantage at most $U_t^k(C) \leq 0$. Starting from the common terminal value and proceeding backwards, monotonicity of conditional expectation gives $J_t^{\pi^{k+1}} \leq J_t^{\pi^k}$. The covering property also gives

$$A_t^k(x, \pi_t^{k+1}(x)) - \inf_{a \in A_t(x)} A_t^k(x, a) \leq U_t^k(C) - L_t^k(C).$$

The left side is nonnegative, so the right side is nonnegative. Replacing the new continuation by the old one using the already proved safety inequality, and then comparing with an arbitrarily near-optimal Bellman action, yields (92). Taking the limit removes any need for attainment of the continuous-action infimum. Iteration proves the displayed finite-sweep bound. For sharpness, take one decision date, an incumbent of cost $\delta > 0$, and an unproposed feasible action of cost zero. The

exact incumbent upper advantage is zero and the exact cover lower advantage is $-\delta$. Retaining the incumbent gives both $E = \delta$ and $\varepsilon = \delta$. Q.E.D.

This result does not replace Theorem 30. The earlier theorem is useful when only a symmetric error bound is available. The present certificate is useful when a numerical calculation retains signed endpoints. In either case safety for a proposed action is weaker than accuracy against the full feasible infimum. The action cover in (91) is indispensable for the latter. A small gap on a finite proposal menu alone is not a small Bellman gap.

H.2. Signed residual transport without a policy-value oracle

Fix a critic h and incumbent π . Let $e_t = J_t^\pi - h_t$ and $r_t = \mathcal{T}_t^\pi h_{t+1} - h_t$, with $r_T = g - h_T$. Suppose $\mathcal{H}_t(x, y)$ is an interval enclosing $r_t(x) - r_t(y)$. For a coupling $\Lambda_t^\pi(x, y)$ with the stated policy-kernel marginals, interval addition and integration define

$$\mathcal{D}_T(x, y) = \mathcal{H}_T(x, y), \quad \mathcal{D}_t(x, y) = \mathcal{H}_t(x, y) + \beta_t \int \mathcal{D}_{t+1}(u, v) d\Lambda_t^\pi(x, y)(u, v). \quad (93)$$

Integration means an enclosure of the lower and upper integrals, not midpoint quadrature. Independently valid signed evaluation bands may be intersected with this interval; an empty intersection signals a failed certificate and must stop acceptance.

PROPOSITION 35—Directed residual enclosure: *The interval $\mathcal{D}_t(x, y)$ encloses $e_t(x) - e_t(y)$. If $\Lambda_t(x; a, b)$ couples the two action kernels and $\mathcal{A}_t^h(x; a, b)$ encloses*

$$c_t(x, a) - c_t(x, b) + \beta_t(P_t^a - P_t^b)h_{t+1}(x),$$

then

$$\mathcal{A}_t^h(x; a, b) + \beta_t \int \mathcal{D}_{t+1}(u, v) d\Lambda_t(x; a, b)(u, v) \quad (94)$$

encloses $A_t^\pi(x, a)$ when $b = \pi_t(x)$. Replacing each integral by an outward enclosure preserves the result. Taking absolute endpoint maxima recovers an admissible symmetric action-contrast certificate.

PROOF: The identity $e_t = r_t + \beta_t P_t^\pi e_{t+1}$ and the marginal conditions imply the corresponding identity for differences at (x, y) . Backward induction gives (93). Substitute $J_{t+1}^\pi = h_{t+1} + e_{t+1}$ into the advantage and use the action-kernel marginal conditions to obtain (94). Endpoint integration is order preserving because coupling weights are nonnegative. Intersection with another valid enclosure preserves containment. Q.E.D.

The choice $h = 0$ is a valid verification gauge: $e = J^\pi$, the running residual is the incumbent's stage cost, and the terminal residual is g . It does not change the previously constructed policy or claim that a zero critic is an effective policy generator. Under this gauge, a common-innovation rollout of stage-cost differences computes the signed recursion directly. Its computational cost must therefore be charged as policy evaluation. Neither a neural critic nor a native representation receives that work for free.

H.3. A computable specialization in the original investment economy

The full-sweep implementation uses the same two-dimensional nonlinear investment primitives as the preceding study. For two coupled states x, y , retain a separate enclosure of $\Delta = x - y$ rather than recomputing it by subtracting two large state boxes. For the cyclic bilinear transition, the deterministic difference in coordinate j is

$$\Delta_j \left(\frac{1}{2} + \frac{1 - x_v}{16} \right) + \Delta_v \left(\frac{1}{8} - \frac{y_j}{16} \right) + b_j(a - b), \quad v \neq j, \quad (b_1, b_2) = \left(\frac{1}{2}, \frac{1}{4} \right).$$

A common innovation cancels in this difference, but is still integrated with its original continuous marginal law. Centered polynomial identities evaluate differences of quadratic and quartic costs. For a shortage term, the intersection

$$u_+ - v_+ \in ([u]_+ - [v]_+) \cap [\min\{0, \underline{\Delta}\}, \max\{0, \bar{\Delta}\}]$$

retains a verified interval for $\Delta = u - v$. In regions where both arguments are positive, their difference is exactly Δ ; where both are negative it is zero. These identities are evaluated outward.

The acquired actor is discontinuous at cell boundaries. Its action range on a rectangle is obtained from exact minima and maxima of integer action indices, including all intersected cells. No Lipschitz assumption is imposed on that actor. Innovation intervals are integrated by enclosures on dyadic bins with their exact probabilities. The final innovation is integrated analytically using the original polynomial and shortage identities. The resulting lower covers include every continuous feasible action, whereas only robustly feasible lattice actions can be selected.

PROPOSITION 36—Work, storage, and certified approximation: *Suppose an acquired actor has n^d cells per date and a range-minimum and range-maximum table over dyadic rectangles. Its preprocessing and stored table size are $O(Tn^d(1 + \log n)^d)$, and a rectangle query uses $O(d + 2^d)$ work. Let N be the number of verified observation cells, A the total number of candidate and lower-cover intervals per cell, and q the number of innovation bins at each unresolved date. If the last innovation is integrated analytically, one full pass visits at most*

$$AN \sum_{r=1}^T \sum_{j=0}^{r-1} q^j \quad (95)$$

pair nodes. Streaming the innovation children requires $O(Tmd)$ working state storage for a batch of m roots, in addition to actor tables and $O(TAN)$ retained endpoint records. Thus a full pair-state table is unnecessary, but tensor-state and exponential-horizon costs have not disappeared.

More generally, suppose an approximate signed pair calculation has a verified terminal endpoint error η_T and local endpoint error at most η_t at date t , inclusive of residual approximation, compression, integration, and arithmetic. Expanding each computed interval by

$$b_T = \eta_T, \quad b_t = \eta_t + \beta_t b_{t+1}$$

is sufficient to enclose the exact signed recursion. In an action-kernel integral the continuation expansion is multiplied by β_t ; its own integration and arithmetic errors must also be added.

PROOF: There are n^d base locations and at most $(1 + \log_2 n)^d$ side-length combinations. Each rectangular extremum is reconstructed from 2^d overlapping dyadic rectangles. A root with r remaining dates has at most $1 + q + \dots + q^{r-1}$ visited nodes; summing over roots, dates, and action intervals gives (95). Processing one child at a time gives the stated stack bound. Finally, integration against a probability coupling is nonexpansive in a uniform endpoint error. Backward induction then proves the error expansion. *Q.E.D.*

In the executed case, $d = 2$, $n = 32$, $A = 17$ (nine proposals and eight continuous-action cover intervals), and $q = 16$. Each actor table occupies 149,760 bytes in the recorded integer representation, hence 299,520 bytes for two dates and 449,280 bytes for three. These are table bytes, not process peak memory. Identity pairs can terminate without recursion and reduce actual work below (95). Verified compression could further reduce storage, but its error must enter the displayed recursion; unverified interpolation across policy jumps is inadmissible. The present experiment executes streamed interval recursion, not a low-rank pair approximation.

H.4. *Approximately null components and misspecification*

An exactly null direction is a special case of a quantitative bound. Let $n(y) = M\sigma(v^\top y + b)$, where σ is one-Lipschitz. For a verified coupling of the actual action kernels, suppose $s(x; a, b)$ bounds $\mathbb{E}|v^\top(Y^a - Y^b)|$. Then

$$|(P_t^a - P_t^b)n(x)| \leq |M|s(x; a, b). \quad (96)$$

This follows directly from the Lipschitz inequality and integration. For a common-innovation affine exposure $Y^a - Y^b = B(a - b)$, the bound is $|M| |v^\top B(a - b)|$. If $v = v_0 + \delta v$, $B = B_0 + \delta B$, and $v_0^\top B_0 = 0$, the non-null coefficient is

$$\delta v^\top B_0 + v_0^\top \delta B + \delta v^\top \delta B.$$

Any action-dependent innovation discrepancy under the actual coupling contributes its own expected absolute projection. These terms cannot be set to zero because a fitted component resembles an invariant.

The deposited perturbation catalogue varies the ReLU direction, action exposure, and an action-dependent innovation-mean shift. Its positive bound is substituted into the same incumbent gate. It also records the policy that would result from falsely declaring the component null. This is a controlled model-misspecification test with known perturbed primitives, not a learned-null discovery procedure or an inference guarantee for an unknown transition law. The original investment laws and the original cost comparisons remain unchanged.

APPENDIX I: ALLOCATING RESOURCES TO AN ECONOMIC TOLERANCE

A single resolution parameter imposes an economic restriction that the Bellman problem does not impose. State coverage, action search and integration have different error coefficients and different effects on work. Refining all three at the same rate can therefore spend most of the service budget on the wrong source of error. We separate these choices while retaining the continuation, feasible-witness interpretation and all-state comparison developed above.

I.1. *A separated primitive budget*

Consider the investment family specified in Section L, with $d \in \{2, 3, 4\}$ state coordinates, discount $\beta = 15/16$, and uniform coordinate grids with N intervals. At a state node use $K + 1$ feasible action fractions. Integrate the common scalar innovation using M equal-probability bins and their midpoints. The three integers need not agree. Let L_t be the primitive graph modulus of the constructed witness continuation, and let D_t be its action modulus. A downward numerical displacement of the action grid is denoted by δ_a .

The state and action covering radii satisfy $h_t = d/(2N)$ and $k_t \leq 1/(8K) + \delta_a$. The innovation lies in an interval of length $1/16$. Within a bin its expected absolute distance from the midpoint is $1/(64M)$. Since changing the scalar innovation by $z - z'$ changes the next state by at most $d|z - z'|$ in the ℓ^1 norm, a certified numerical query has an error allowance

$$e_t \leq \frac{\beta d L_{t+1}}{64M} + \xi_t. \quad (97)$$

Here ξ_t covers arithmetic, interval evaluation and label rounding. It does not suppress the integration error. The terminal label allowance is e_T .

PROPOSITION 37—Separated-resource policy account: *Under the feasible-witness hypotheses, the returned policy satisfies*

$$0 \leq J^\pi - V^* \leq \frac{a}{N} + \frac{b}{K} + \frac{q}{M} + \eta, \quad (98)$$

where

$$\begin{aligned} a &= d \sum_{t < T} \beta^t L_t + d \beta^T L_g, & b &= \frac{1}{8} \sum_{t < T} \beta^t D_t, \\ q &= \frac{d}{32} \sum_{t < T} \beta^{t+1} L_{t+1}, & \eta &= \sum_{t < T} \beta^t (2\xi_t + D_t \delta_a) + 2\beta^T e_T. \end{aligned} \quad (99)$$

The primitive coefficients are independent of the computed labels. The leading number of innovation evaluations in target formation is

$$T(N+1)^d(K+1)M. \quad (100)$$

The compiled continuation additionally incurs its corner-evaluation and preprocessing costs. Pilot selection, policy extraction, verification, arithmetic bit lengths and output are separate charged components.

PROOF: Insert the state and action radii and (97) into the feasible-witness loss bound. The selected-policy upper account contributes the second copy of e_t , not a second action-search allowance. Collecting the coefficients of N^{-1}, K^{-1}, M^{-1} gives (99). There are $(N+1)^d$ state nodes, $K+1$ actions per node and M innovation representatives per action at each nonterminal date. No independence of fitting errors is used. Q.E.D.

The account explains why a common further state rung need not be a diagonal refinement. In the $T=3, p=4, d=2$ cell, the exact primitive coefficients are approximately $a=51.1569, b=2.3650$ and $q=1.0486$. The choice $(N, K, M) = (128, 16, 8)$ has account about $0.6786 + \eta$, below target one when the numerical remainder fits the remaining budget. The diagonal choice $(128, 128, 128)$ is more accurate, but

its target-generation innovation count at that same state grid is $129 \cdot 128 / (17 \cdot 8)$ times larger. This is an arithmetic comparison of two declared allocations, not an observation of the unexecuted diagonal service or a runtime ratio. Both witness and conventional FVI receive the separated allocation in the experiment.

I.2. The optimal continuous allocation and its boundary cases

A useful planning problem minimizes the leading product $N^d KM$ subject to (98). This proxy isolates the primitive evaluation term; it is not the entire operation count or the measured wall clock. Write $E = \varepsilon - \eta > 0$ and suppose first that a, b, q are positive and all optimal resources exceed their required lower bounds.

PROPOSITION 38—Error shares and clipped allocation: *The unique interior minimizer of the continuous proxy has*

$$\frac{a}{N} = \frac{d}{d+2}E, \quad \frac{b}{K} = \frac{1}{d+2}E, \quad \frac{q}{M} = \frac{1}{d+2}E. \quad (101)$$

For lower resource bounds $r_i \geq \underline{r}_i > 0$, put $c = (a, b, q)$, $p = (d, 1, 1)$ and $\bar{x}_i = c_i / \underline{r}_i$. If $E < \sum_i \bar{x}_i$, there is a unique positive λ such that

$$x_i = \min\{\bar{x}_i, p_i \lambda\}, \quad \sum_i x_i = E; \quad r_i = c_i / x_i. \quad (102)$$

These resources minimize the constrained continuous proxy. If $E \geq \sum_i \bar{x}_i$, all lower resource bounds suffice. Coordinates with zero error coefficient remain at their resource lower bound.

When dyadic lower bounds are used, rounding each optimal resource upward to a power of two preserves the error guarantee and multiplies this proxy by less than 2^{d+2} .

PROOF: Set $x_i = c_i / r_i$. Minimizing the resource product is equivalent to maximizing $\sum_i p_i \log x_i$ over $0 < x_i \leq \bar{x}_i$ and $\sum_i x_i \leq E$. Strict concavity and the multiplier conditions give p_i / x_i equal on unconstrained coordinates and larger on capped coordinates. This is precisely (102). The sum of its right sides is strictly increasing until all caps bind. The unconstrained case yields (101). Upward dyadic rounding never enlarges an error term and multiplies $N^d KM$ by less than the stated factor. Neither argument bounds rational-arithmetic bit cost or computer time. *Q.E.D.*

The optimization is a standard separable resource-allocation calculation. Its role here is to expose the error coefficients of the actual own-future policy construction, including integration and implementation, rather than to claim novelty for the multiplier argument. A finite experimental menu need not contain the continuous optimum. The prototype uses a common, frozen dyadic menu for all generators and records all of it. It does not describe this menu as a globally optimal tuning policy for either FVI or the witness method.

I.3. *A complete tolerance frontier*

For method m , let a prospectively fixed ladder return certificates $g_{m,r}$ and cumulative costs $w_{m,r}$, including failed earlier rungs. Define

$$r_m(\varepsilon) = \min\{r : g_{m,r} \leq \varepsilon\}, \quad W_m(\varepsilon) = w_{m,r_m(\varepsilon)}. \quad (103)$$

An empty set is a capped failure, not an interpolated success. Sort the finite union of all certificate values and zero. On each left-closed, right-open interval between successive values, every method's first crossing and complete work are constant. Evaluating the first-crossing rule at the left endpoint therefore reconstructs the entire positive-tolerance frontier exactly, even when successive certificates are not monotone. An unbounded final interval is retained. Repetitions have separate costs, while their mathematical checkpoints must be checked for identity.

This finite partition supplies the appropriate interpretation of a threshold advantage. At a fixed resolution the witness method may cost more. At a tolerance where it crosses one rung earlier, the smaller certificate may more than offset that cost. Whether this happens, and over which tolerance interval, is determined by (103); it is not determined by a favorable single threshold. Under an additional homogeneous model $g_m(n) = A_m/n$ and $w_m(n) = C_m n^\kappa$, the continuous frontier ratio is $(C_W/C_F)(A_W/A_F)^\kappa$. This last expression is a conditional explanation, not a fitted scaling law for the experiments. The release supplies the exact finite partition, sensitivity around one, complete prefixes and separate fixed-resolution observations.

APPENDIX J: FROM RETURNED POLICIES TO ECONOMIC RESOURCE DECISIONS

An all-state upper bound and a comparison of actual policy costs answer different economic questions. We now connect the policies returned by a construction frontier to direct cost inference and then to a choice of implementation resources. The construction data select the policies; independently generated evaluation paths assess them. This separation permits a valid comparison of first-crossing policies with different resolutions.

J.1. Simultaneous inference after construction

Let $\mathcal{H}$ denote all construction information: primitive inputs, numerical labels, fitted continuations, returned policies and complete certificate ladders. Conditional on $\mathcal{H}$, fix a finite catalogue of implemented-policy contrasts. Each implementation includes its observation rule, original witness or nearest-node selector, true-state feasible repair and action quantization. For a contrast i , let D_i be the difference in expected discounted costs under a specified initial law. A common-innovation path score evaluates the two policies on their own endogenous trajectories. It is not the difference between two regret bounds.

Suppose the initial and innovation uniforms are partitioned into equal bins. Independent uniform bin indices are drawn, and interval simulation encloses the score for every continuous realization inside the drawn bins. All unresolved state, selector and repair branches remain in the enclosure. If $[A_i, B_i]$ is a deterministic support interval, clip the path endpoints to that support. Write $\underline{m}_i, \overline{m}_i$ for outward sample-mean bounds on the lower and upper endpoint samples and $v_{i,L}, v_{i,U}$ for upper bounds on their unbiased sample variances. For $n \geq 2$ define

$$R_i(v) = \sqrt{2v\ell/n} + \frac{7(B_i - A_i)\ell}{3(n-1)}, \quad \ell \geq \log(4K/\alpha), \quad (104)$$

where K bounds the number of sampled contrasts.

THEOREM 39—Construction-conditional economic intervals: *Under the stated independent-bin model, conditional on any valid construction history $\mathcal{H}$, with probability*

at least $1 - \alpha$ all catalogue contrasts satisfy

$$D_i \in [l_i, u_i] := [\max\{A_i, \underline{m}_i - R_i(v_{i,L})\}, \min\{B_i, \overline{m}_i + R_i(v_{i,U})\}]. \quad (105)$$

The same coverage holds unconditionally. Subsequent selection among these covered contrasts, or evaluation at any deterministic implementation-price vector, needs no additional union bound. This conclusion applies only to policies and initial laws included in the catalogue; it does not create direct-cost observations for every unstudied tolerance or initial state.

PROOF: For each drawn bin sequence, conditional integration over its continuous realizations places the exact path expectation between the interval endpoints. Uniform mixing over the equal bins recovers the original continuous law. The lower endpoint expectation is thus no larger than D_i , and the upper endpoint expectation is no smaller. Apply the bounded empirical-Bernstein inequality of [Maurer and Pontil \(2009\)](#) to the endpoint samples and take the union bound, with outward arithmetic included in the mean, variance and radius. Conditional validity for every $\mathcal{H}$ implies unconditional validity by integration. On this one simultaneous event, all ensuing deterministic comparisons of the covered quantities hold together. *Q.E.D.*

In particular, selecting the first policy to meet a specified certificate is permitted when that selection uses construction records rather than the subsequent evaluation paths. Repeated construction clocks are not independent draws from a neural-optimizer population. The statistical randomness here belongs to the evaluation design. A fixed pseudorandom stream implements a reproducible experiment; it is not itself a proof of the independent-bin assumption.

J.2. Replacement-fee identification

Write $D = J^W - J^F$. If witness is installed, replacing it by FVI changes total cost by $k - D$, where $k \geq 0$ is the replacement fee. If FVI is installed, the change is $k + D$. Thus the two directional gains before fees are D and $-D$.

COROLLARY 40—The break-even fee frontier: *On an event where $D \in [l, u]$, the two directional gains belong to $[l, u]$ and $[-u, -l]$. The unknown absolute cost difference*

satisfies

$$|D| \in [\text{dist}(0, [l, u]), \max\{|l|, |u|\}]. \quad (106)$$

For every $k \geq \max\{0, u, -l\}$ neither replacement strictly lowers expected cost. By contrast, $k \geq \max\{0, l, -u\}$ only removes a certified strictly profitable replacement. The two thresholds generally differ. Strict inequalities in the first criterion certify that both replacements strictly increase cost.

PROOF: Subtract the gain intervals from k and compare the resulting interval endpoints with zero. Taking the image of $[l, u]$ under the absolute-value map gives (106). The assertions hold simultaneously for every nonnegative fee on the original coverage event. *Q.E.D.*

The fee frontier is more informative than a conclusion at one generous charge. The earlier 1/64 replacement decision is retained with its original interpretation and confidence account. The new analysis reports the two directional gain intervals and both thresholds instead of selecting a favorable fee after seeing the data. There is no welfare calibration hidden in the units: a decision maker must supply the economically relevant charge, or use the full frontier to see what remains unidentified.

J.3. Observation costs on the same bit grid

Let $\mathcal{B} = \{4, \dots, 16\}$ be the available bit precisions. For one fixed returned controller, write J_b for the expected cost of its actual b -bit, robustly repaired and quantized implementation. Use $b_0 = 16$ as a common reference and obtain simultaneous intervals for $D_b = J_b - J_{b_0}$. The identity interval at $b = b_0$ is exactly $[0, 0]$. Let $s_b = db \sum_{t < T} B_t$ be the discounted number of coordinate-bits purchased; the net cost is $J_b + \lambda s_b$.

THEOREM 41—Certified choice of actual net implementation cost: Suppose all $D_b \in [l_b, u_b]$ simultaneously. For any $\lambda \geq 0$ define

$$L(\lambda) = \min_b \{l_b + \lambda s_b\}, \quad U(\lambda) = \min_b \{u_b + \lambda s_b\}, \quad (107)$$

$$\widehat{b}(\lambda) = \min \arg \min_b \{u_b + \lambda s_b\}.$$

Then, simultaneously for all $\lambda \geq 0$,

$$0 \leq J_{\widehat{b}} + \lambda s_{\widehat{b}} - \min_{b \in \mathcal{B}} (J_b + \lambda s_b) \leq U(\lambda) - L(\lambda). \quad (108)$$

The lower and upper functions are finite piecewise-affine envelopes. All changes in their active lines occur at nonnegative intersections of catalogue lines, with exact comparisons deciding ties. The smallest selected precision $\widehat{b}(\lambda)$ is weakly decreasing in λ .

PROOF: Subtract the common J_{b_0} . The selected implementation has net difference at most $U(\lambda)$, and the best catalogue implementation has net difference at least $L(\lambda)$. This proves the regret bound pointwise on the one simultaneous event and hence for every λ . A minimum of finitely many affine functions changes its active affine expression only at an intersection. For $\lambda_2 > \lambda_1$, sum the two minimizing inequalities to obtain $(\lambda_2 - \lambda_1)(s_{\widehat{b}(\lambda_2)} - s_{\widehat{b}(\lambda_1)}) \leq 0$. Since s_b increases strictly with b , the precision comparative static follows. *Q.E.D.*

This result makes a statement about *actual* net cost within the declared bit catalogue. It differs from minimizing a uniform all-state loss bound. Both criteria are useful: the latter gives distribution-free protection relative to the full-information optimum, while (108) uses a specified initial law to compare feasible implementations of one returned policy. Neither criterion gives an optimum over arbitrary sensor technologies or all policies.

For a fair comparison the conventional actor receives its own acquisition allowance. Suppose its nearest-node policy has a certified selected-policy account at a reported state, obtained from the graph modulus L_t^{gr} , and its actual interpolated critic has modulus L_t^f . Transporting the same selected node action to the true state, then using cell-wide capacity repair, adds at most

$$(L_t^{\text{gr}} + L_t^f + 2D_t K_t^A) r(b) + D_t \Delta_t. \quad (109)$$

Indeed, transport the feasible node pair directly to the true state, bound its node distance by the reported-state distance plus $r(b)$, and bound $f_t(\widehat{x}) - f_t(x)$ by $L_t^f r(b)$.

Robust repair adds $D_t(2K_t^A r(b) + \Delta_t)$. This proves (109) without assuming actor continuity. For a witness critic $L_t^f = L_t^{gr}$, it reduces to the preceding witness allowance. The empirical comparison therefore evaluates the all-state-account minimizer, the interval upper-net-cost selector, and a descriptive midpoint selector on the *same* integer bit and price grid. The midpoint selector is not called the true optimum.

APPENDIX K: JOINT POLICY-COST AND RESOURCE DECISIONS

Let a finite catalogue $\mathcal{P}$ contain implemented policies, including incumbents, repaired policies and conventional alternatives. Construction, verification and direct evaluation are distinct charged activities. Denote their measured deployment-relevant resource vector by r_j ; the published timing convention must specify whether an evaluation shared by several candidates is a common fixed charge or an attributable marginal charge. In particular, a joint paired evaluation cannot be counted as zero or silently charged only to the preferred generator.

Conditional on construction information, suppose a simultaneous event satisfies $J_j \in [L_j, U_j]$ for every $j \in \mathcal{P}$. For nonnegative prices θ , let $k_j(\theta)$ be a known implementation charge, for example $\theta \cdot r_j$ plus an information fee. Define

$$\hat{j}(\theta) \in \arg \min_{j \in \mathcal{P}} \{U_j + k_j(\theta)\}, \quad R(\theta) = U_{\hat{j}} + k_{\hat{j}}(\theta) - \min_{j \in \mathcal{P}} \{L_j + k_j(\theta)\}. \quad (110)$$

A fixed lexicographic tie rule makes the selection measurable.

PROPOSITION 42—Actual net-cost frontier: *On the simultaneous event, for every nonnegative price vector, including one selected after observing the intervals,*

$$0 \leq J_{\hat{j}} + k_{\hat{j}}(\theta) - \min_j \{J_j + k_j(\theta)\} \leq R(\theta).$$

A proposed replacement of incumbent i by j is certified nonworsening in net cost whenever a simultaneous upper bound on $J_j - J_i$, plus $k_j(\theta) - k_i(\theta)$ and the replacement fee, is nonpositive. Otherwise a rule retaining i remains safe. Policy-specific certificate widths can be displayed alongside this frontier, but do not determine its ordinate.

PROOF: Bound the selected true cost above by its upper line and the best catalogue cost below by the minimum lower line. All inequalities hold on the same finite

event for every price, so a new price query does not consume another probability allocation. The replacement assertion applies the paired upper bound directly. *Q.E.D.*

Absolute intervals can be considerably wider than centered paired differences. For a decision between an incumbent and its final repair, the centered interval is therefore the appropriate replacement input. A statistically unresolved interval is not evidence of equivalence. In contrast, Corollary 26 supplies a deterministic nonincrease, although it does not necessarily certify a strictly positive improvement large enough to cover an implementation fee.

K.1. *Information technology and the price of conservatism*

An explicit theoretical observation technology gives the bit price units. Suppose a device reports a dyadic cell in each coordinate at date t . With b bits per coordinate it transmits db bits and requires τdb resource units, where τ is a technological conversion factor. At unit resource price w , and with a setup cost K_b , its present cost is

$$C_b(w, \tau) = K_b + w\tau db \sum_{t < T} B_t. \quad (111)$$

Each reported cell entails its own feasible action contract. In the retained sensor experiment, $K_b = 0$ and the normalized price $\lambda = w\tau$ takes the declared values. This supplies a structural interpretation and a sensitivity variable, not an estimate of actual device costs. The dataset contains no empirical measurements from which w or τ could be calibrated.

Theorem 41 and Proposition 42 apply to the finite bit catalogue with charges (111). Their upper-minus-lower envelope is a valid bound on the cost of conservative choice relative to the unknown true catalogue minimizer. Replacing that minimizer by the lowest observed sample mean and calling it the ex post true optimum is not justified. The retained exact price envelopes and their regret bounds answer the identified decision question without making this substitution. A richer technology with fixed charges, nonlinear bandwidth costs or precision-dependent delays can be assessed by changing the known charge functions, but new unobserved policies require new cost bounds.

A similarly explicit interpretation applies to terminal repair. A controller that incurs cost K to install the additional gate, and q per deployed final action, is worth adopting only when its expected discounted policy gain covers those charges. The published centered gain interval defines the entire break-even range. Its mathematical safety does not by itself imply economic adoption at positive computation prices.

APPENDIX L: CERTIFIED COMPUTATION AND ECONOMIC DECISIONS

L.1. *The controlled investment family*

There are $d \in \{2, 3, 4\}$ capital coordinates on $[0, 1]^d$, one investment action, and a common continuous shock. The feasible set is

$$0 \leq a \leq \kappa(x) = \frac{1}{8} + \frac{1}{8d} \sum_{j=1}^d x_j.$$

With cyclic coordinate indices, $\gamma_j = 1/2$ for odd j and $1/4$ for even j , and $z \sim \text{Unif}[-1/32, 1/32]$, the dynamics are

$$F_j(x, a, z) = \frac{1}{16} + \frac{x_j}{2} + \frac{x_{j+1}}{8} + \frac{x_j(1 - x_{j+1})}{16} + \gamma_j a + (-1)^{j-1} z. \quad (112)$$

Current costs are

$$c(x, a) = \frac{2}{d} \sum_j (x_j - 5/8)^2 + \frac{1}{4d} \sum_j (x_j - x_{j+1})^2 + 2 \left(\frac{1}{2} - \frac{2}{d} \sum_j x_j \right)_+^2 + pa^2 + 4a^4. \quad (113)$$

Terminal cost doubles the first coefficient and retains the imbalance and shortage terms, with no action cost. Capital enters both adjustment incentives and feasible capacity. The nonlinear transition and shortage penalty couple the coordinates; the model is not separable linear-quadratic control. Prices $p \in \{1, 4\}$ and horizons $T \in \{2, 3\}$ are theoretical primitives, not calibrated estimates.

For the declared dimensions, valid primitive moduli in ℓ^1 are

$$C_x = \frac{13}{2d}, \quad L_g = \frac{9}{d}, \quad F_x = \frac{11}{16}, \quad F_a = \sum_j \gamma_j, \quad C_a = \frac{p}{2} + \frac{1}{4}, \quad K^A = \frac{1}{8d}.$$

The transition is invariant by direct interval bounds. The state column sums of its Jacobian are at most $11/16$. The costs have piecewise-affine gradients; the cube vertices and the shortage-boundary vertices give the stated bounds. The supplemental proof restricts this explicit formula to the declared dimensions rather than presuming it for an arbitrary number of coordinates. The general constructive theorem uses whichever valid primitive moduli the economy supplies.

L.2. *Matched services and realized adaptation*

The execution-before protocol fixes three generators: compiled witness, uniform multilinear FVI, and curvature-refined coordinate FVI. Every generator forms its own future labels and receives the same state-node count, feasible action fractions and integration bins. For each two-state price-horizon cell the common resource ladder is

$$(N, K, M) = (4, 2, 1), (8, 2, 1), (16, 4, 2), (32, 4, 2), (64, 8, 4), (128, 16, 8).$$

The entire ladder is run in three isolated processes per cell and method, rotating method order. This yields 36 two-state services and 216 rungs. State dimension three uses four declared rungs through $(32, 8, 4)$; dimension four uses three through $(16, 8, 4)$, with the witness and uniform generators, both horizons, price one and three repetitions. Those additional 24 services contain 84 rungs. The comparison across dimensions uses the same target five, not a dimension-specific certificate threshold. This is a finite common-accuracy experiment at a coarse tolerance, not evidence of an asymptotic or dimension-free rate.

The curvature comparator obtains an own-future pilot on a nine-node coordinate grid. For each axis it takes maxima of adjacent exact second differences across the other coordinates. It then bisects the interval maximizing squared length times the largest pilot curvature over that interval, with deterministic leftmost ties. A negligible positive floor permits a defined rule when the pilot is affine; the dominant unit baseline of the earlier heuristic is absent. Pilots, allocation and the realized axes are stored and charged. Realized nonuniformity, rather than the method's name, determines whether the experiment supplies an adaptive comparison. This

remains a coordinate tensor algorithm; it is not presented as a sparse-grid or general adaptive-cell method.

Each rung starts from primitives. The internal complete-prefix clock includes own-future integration, all earlier failed rungs, pilots, exact compilation, actor construction, all-state certification, deployment checks and durable checkpoints. Separate process clocks include startup and the identical non-economic warm-up. A single CPU is selected, numerical-library threads are fixed to one, and CPU frequency is not controlled. Peak resident memory includes the process. Primitive query counts are component counts, not complete FLOPs or bit-operation counts. Independent policy evaluation has a separate joint service and is never added to an unrelated historical clock.

TABLE L.I: Block A separated-resource attainment

Method	4	2	1	1/2
W	12/12	12/12	12/12	6/12
F	12/12	12/12	12/12	3/12
A	12/12	12/12	12/12	3/12

W: compiled witness; F: uniform FVI; A: pilot-curvature bisection FVI. Twelve services per method are four economies, each repeated three times. They are not twelve independent economic tasks or optimizer draws.

TABLE L.II: Complete positive-tolerance frontiers in two states

(T, p)	Tolerance interval	W: N ; sec.	F: N ; sec.	A: N ; sec.
(2,1)	[0,0.423176)	–	–	–
(2,1)	[0.423176,0.483393)	128; 10.658	–	–
(2,1)	[0.483393,0.846352)	128; 10.658	128; 8.458	128; 9.817
(2,1)	[0.846352,0.962257)	64; 1.458	128; 8.458	128; 9.817
(2,1)	[0.962257,1.6927)	64; 1.458	64; 1.134	64; 1.514
(2,1)	[1.6927,1.90651)	32; 0.290	64; 1.134	64; 1.514

	(T, p)	Tolerance interval	W: N ; sec.	F: N ; sec.	A: N ; sec.	
1						1
2						2
3	(2,1)	[1.90651, 2.79584)	32; 0.290	32; 0.269	32; 0.399	3
4	(2,1)	[2.79584, 3.21615)	16; 0.092	32; 0.269	32; 0.399	4
5	(2,1)	[3.21615, 5.59167)	16; 0.092	16; 0.119	16; 0.176	5
6	(2,1)	[5.59167, 6.13235)	8; 0.034	16; 0.119	16; 0.176	6
7	(2,1)	[6.13235, 10.0042)	8; 0.034	8; 0.061	8; 0.090	7
8	(2,1)	[10.0042, 10.6876)	4; 0.014	8; 0.061	8; 0.090	8
9	(2,1)	[10.6876, ∞)	4; 0.014	4; 0.029	4; 0.042	9
10	(2,4)	[0, 0.45093)	—	—	—	10
11	(2,4)	[0.45093, 0.506837)	128; 10.689	—	—	11
12	(2,4)	[0.506837, 0.901859)	128; 10.689	128; 8.453	128; 9.797	12
13	(2,4)	[0.901859, 1.00914)	64; 1.452	128; 8.453	128; 9.797	13
14	(2,4)	[1.00914, 1.80372)	64; 1.452	64; 1.134	64; 1.509	14
15	(2,4)	[1.80372, 2.00023)	32; 0.290	64; 1.134	64; 1.509	15
16	(2,4)	[2.00023, 2.92224)	32; 0.290	32; 0.268	32; 0.397	16
17	(2,4)	[2.92224, 3.31506)	16; 0.092	32; 0.268	32; 0.397	17
18	(2,4)	[3.31506, 5.84447)	16; 0.092	16; 0.118	16; 0.177	18
19	(2,4)	[5.84447, 6.33036)	8; 0.034	16; 0.118	16; 0.177	19
20	(2,4)	[6.33036, 10.3185)	8; 0.034	8; 0.061	8; 0.091	20
21	(2,4)	[10.3185, 10.9232)	4; 0.014	8; 0.061	8; 0.091	21
22	(2,4)	[10.9232, ∞)	4; 0.014	4; 0.029	4; 0.043	22
23	(3,1)	[0, 0.63562)	—	—	—	23
24	(3,1)	[0.63562, 0.690808)	128; 15.635	—	—	24
25	(3,1)	[0.690808, 1.27124)	128; 15.635	128; 12.420	128; 14.176	25
26	(3,1)	[1.27124, 1.37488)	64; 2.070	128; 12.420	128; 14.176	26
27	(3,1)	[1.37488, 2.54248)	64; 2.070	64; 1.651	64; 2.148	27
28	(3,1)	[2.54248, 2.72357)	32; 0.401	64; 1.651	64; 2.148	28
29	(3,1)	[2.72357, 4.11404)	32; 0.401	32; 0.388	32; 0.568	29
30	(3,1)	[4.11404, 4.53322)	16; 0.126	32; 0.388	32; 0.568	30
31	(3,1)	[4.53322, 8.22807)	16; 0.126	16; 0.171	16; 0.252	31
32	(3,1)	[8.22807, 8.6154)	8; 0.046	16; 0.171	16; 0.252	32
	(3,1)	[8.6154, 14.5143)	8; 0.046	8; 0.088	8; 0.130	
	(3,1)	[14.5143, 14.9762)	4; 0.019	8; 0.088	8; 0.130	

(T, p)	Tolerance interval	W: N ; sec.	F: N ; sec.	A: N ; sec.
(3,1)	$[14.9762, \infty)$	4; 0.019	4; 0.041	4; 0.061
(3,4)	$[0, 0.678554)$	—	—	—
(3,4)	$[0.678554, 0.723386)$	128; 15.815	—	—
(3,4)	$[0.723386, 1.35711)$	128; 15.815	128; 12.435	128; 14.257
(3,4)	$[1.35711, 1.44001)$	64; 2.085	128; 12.435	128; 14.257
(3,4)	$[1.44001, 2.71422)$	64; 2.085	64; 1.643	64; 2.169
(3,4)	$[2.71422, 2.85349)$	32; 0.402	64; 1.643	64; 2.169
(3,4)	$[2.85349, 4.31287)$	32; 0.402	32; 0.390	32; 0.569
(3,4)	$[4.31287, 4.66727)$	16; 0.127	32; 0.390	32; 0.569
(3,4)	$[4.66727, 8.62574)$	16; 0.127	16; 0.170	16; 0.253
(3,4)	$[8.62574, 8.88407)$	8; 0.046	16; 0.170	16; 0.253
(3,4)	$[8.88407, 15.0203)$	8; 0.046	8; 0.088	8; 0.131
(3,4)	$[15.0203, 15.306)$	4; 0.019	8; 0.088	8; 0.131
(3,4)	$[15.306, \infty)$	4; 0.019	4; 0.042	4; 0.061

Each time is the median complete prefix over three isolated repetitions. A dash is a capped failure. Adjacent bands with identical returned rungs are merged. Printed cutoffs are rounded; exact rational endpoints in the deposited frontier determine membership, including equality. The action and integration budgets are uniquely fixed by N in the common ladder. Minima, maxima and counts are deposited, not replaced by these medians.

TABLE L.III: Repeated common-accuracy comparison: target five, price one

d	T	W pass	N	sec.	F pass	N	sec.
2	2	3/3	16	0.092	3/3	16	0.119
2	3	3/3	16	0.126	3/3	16	0.171
3	2	3/3	16	2.424	3/3	16	2.074
3	3	3/3	16	3.508	3/3	16	3.032
4	2	3/3	16	58.569	3/3	16	60.809
4	3	3/3	16	82.156	3/3	16	90.099

All dimensions face the same coarse target five; caps and allocations remain explicit. Two-state observations are reused from the main catalogue and not counted twice. This finite comparison does not identify an asymptotic scaling law.

The original block A contains 60 services and 300 rungs. Repeated mathematical checkpoints are byte-identical. The pilot-curvature bisection generator realizes nonuniform axes in 0 of 252 date models; its name alone does not establish effective adaptation. At common successful declared target-by-repetition comparisons, witness is faster than uniform FVI in 13 of 39 and slower in 26. Against pilot-curvature bisection FVI, the corresponding counts are 30 faster and 9 slower out of 39. These counts summarize selected targets; Table L.II and the exact deposited partition, rather than a binary attainment headline, are the primary work comparison.

The new first-crossing direct-cost intervals identify lower witness cost in 0 cases, higher cost in 93, and leave 3 unresolved. For the separately evaluated favorable R48 work pair, the corresponding counts are 0, 0 and 8. A cheaper first crossing and a sharper bound therefore receive an actual economic comparison rather than an inferred ranking. The detailed intervals and initial actions are retained, including every adverse outcome.

Across the 24 declared controller-price decisions, the upper-net-cost selector is certified strictly cheaper than the all-state-bound minimizer in 24. Its cost-regret upper bounds relative to the best evaluated bit range from 0.00238686 to 0.00332242 in the normalized economic units. Where the comparison does not identify a strict improvement it remains unresolved; midpoint choices are descriptive, not true optima. This is a finite expected-net-cost conclusion under the uniform initial law, not a calibrated welfare result or a universal preference for coarse observation.

The direct and sensor study contains 200 sampled contrasts, totaling 52,428,800 interval paths, with a separately recorded joint evaluation time of 923.64 seconds. Identity rows use no simulated paths. This evaluation work is not allocated to one favored generator or hidden inside another method's construction clock.

L.3. *Actual first-crossing policies and replacement fees*

For the four two-state economic cells and each target 4,2,1, the protocol selects repetition zero's first-crossing witness and uniform FVI policies when both attain.

It also selects the particular favorable R48 target-two pair: $T = 3, p = 4$, witness at $N = 32$ and FVI at $N = 64$. Those old frozen policies are evaluated as an additional pair; their historical clocks are not reconstructed as part of the new service. Each selected pair is compared under the uniform initial law and point laws at $(1/8, 1/8)$, $(1/2, 1/2)$ and $(7/8, 7/8)$, with full information or eight-bit observations. Finite observations use robust cell-wide capacity repair and downward action quantum $1/4096$.

Every sampled contrast uses 262,144 common interval paths, 40-bit equal initial and innovation bins, and outward mean and variance bounds. All new policy-pair and sensor contrasts share one family of at most 300 with error $1/100$ and log upper bound 13. The independent-bin model supplies the statistical assertion; the fixed PCG64 streams supply reproducibility. Old R48 intervals retain their separate original confidence event. Simultaneous interpretation of both old and new families is at least 98 percent by a union bound, not automatically 99 percent.

TABLE L.IV: Actual costs of the returned policies and fee sensitivity

Pair	ϵ	Bits	n	W lower	W higher	Unres.	Sufficient fee range
R49	4	E	16	0	14	2	[0.00156,0.1042]
R49	4	8	16	0	15	1	[0.001565,0.1047]
R49	2	E	16	0	16	0	[0.007903,0.05741]
R49	2	8	16	0	16	0	[0.007853,0.05892]
R49	1	E	16	0	16	0	[0.003374,0.02665]
R49	1	8	16	0	16	0	[0.003379,0.02688]
R48	2	E	4	0	0	4	[0.001746,0.001875]
R48	2	8	4	0	0	4	[0.001746,0.001886]

E denotes full information. Counts cover every declared initial law and economic cell in the group; they do not rank all initial states. The fee is sufficient to certify that neither directional switch pays. It is not the smaller threshold that merely removes identified profitability. All new intervals share one 99 percent family conditional on the independent-bin model.

The complete numerical supplement reports the directional replacement gains for every pair and the fee sufficient to certify that neither switch pays. It also reports the lower threshold that merely removes an identified profitable switch. The old $1/64$ decision is preserved, but no conclusion depends on treating that fee as calibrated. The deposited exact fee frontiers permit any nonnegative fee to be assessed on the same confidence event. Lower regret certificates are never described as observed policy-cost improvements.

L.4. *Purchasing information on the evaluated grid*

For each first-target-one controller, the protocol evaluates every integer precision $b = 4, \dots, 16$ under the uniform initial law. Each policy is compared with its own 16-bit implementation; the reference self-comparison is exactly zero. Prices per coordinate-bit are $1/16384, 1/4096, 1/1024$. The deterministic all-state-account minimizer and the upper-net-cost selector of Theorem 41 use precisely that same grid. Neither is compared with an unobserved continuous optimum.

TABLE L.V: Information choice on the same evaluated bit grid

$T;p$	Method	Bound-minimizing bits	Net-selected bits	Net-regret bound range
2;1	W	16,14,12	6,6,4	[0.002563,0.002596]
2;1	F	16,14,12	7,4,4	[0.002387,0.002426]
2;4	W	16,14,12	6,6,4	[0.002691,0.00272]
2;4	F	16,14,12	4,4,4	[0.002564,0.002564]
3;1	W	16,14,12	7,4,4	[0.003092,0.003107]
3;1	F	16,14,12	6,4,4	[0.002992,0.003036]
3;4	W	16,15,13	4,4,4	[0.003322,0.003322]
3;4	F	16,14,12	4,4,4	[0.003237,0.003237]

Bit triples correspond, in order, to prices $1/16384, 1/4096, 1/1024$ per coordinate-bit. The net selector minimizes the upper actual-net-cost interval. Its regret bound is against the best bit in the finite catalogue, not an estimated unconstrained optimum. Every price-specific value and all nonnegative-price envelopes are deposited.

For each economic cell, the record additionally gives all nonnegative information prices by exact lower and upper affine envelopes. A retrospective price query within that family is allowed by simultaneous coverage and is not a new experiment. Initial-law dependence remains substantive: a uniform-state expected-cost decision need not be optimal at a particular initial state, and a conservatively optimal all-state bound may rationally purchase more information.

L.5. *The retained evidence and its interpretation*

The new service does not rewrite the original nonlinear training studies. The R42 tighter-target neural attainments, six of twelve scalar and five of six coupled services, remain capped-service outcomes; the R44 residual refinements of seven frozen failures remain diagnostics, not retraining successes. The 24 direct neural-minus-ridge intervals remain unresolved. The R45 constructive comparison still has nine of twelve neural and twelve of twelve spline attainments at its tightest target. Fixed64 and adaptive precision still certify all 36 historical precision services, with adaptive slower in 26 and about 0.98 percent slower in aggregate.

The R48 target-one count, twelve of twelve for witness versus six of twelve for each FVI method, also remains unchanged. At a relaxed tolerance 1.014 the old FVI near misses all attain; the release supplies a complete sensitivity table rather than reclassifying those original failures. The new common $N = 128$ state rung uses separately declared action and integration budgets. It is not the unexecuted diagonal (128,128,128) experiment. The original 48 constrained-policy intervals, including six favoring FVI and none favoring witness, remain available with newly exposed fee frontiers. Their adverse point-state pattern motivates inspection of actual policies; it is not removed by a favorable all-state certificate.

The complete development edition, unchanged historical companions and machine-readable preservation map retain all of these results and their proofs. This revision adds a controlled resource allocation and actual implementation decision, not a selectively favorable replacement record.

APPENDIX M: COST-DIRECTED CONSTRUCTION AND IMPLEMENTATION EVIDENCE

We now execute the distinction between constructing a safe policy and improving its actual economic performance. The new policies use the same nonlinear investment primitives as Section L; neither the price parameter nor the terminal penalty is changed to favor the proposed action rule. All earlier continuation weights, owners, policies, clocks and adverse contrasts retain their historical identities. The recovered R50 main experiment and common-accuracy amendment had separate local pre-execution freezes. R51 is an explicitly known-design reproduction, with its own Git source identity and runner clocks; the same draws do not provide additional independent observations.

M.1. *Prospective resource selection and a residual-driven comparator*

For the primary block, select powers of two $N \in \{4, 8, 16, 32, 64\}$ and $K, M \in \{1, 2, 4, 8, 16\}$ before constructing a policy. Among allocations satisfying the primitive witness budget $a/N + b/K + q/M \leq \varepsilon - 10^{-8}$, minimize the exact integer query proxy $T(N+1)^d(K+1)M$, with lexicographic ties. The reserve is checked against the realized certificate; it is not a universal fixed-precision theorem. This rule supplies an executed version of the separated allocation, rather than a retrospective proxy calculation. Conventional FVI receives the same allocation but its own actual-modulus certificate. The witness planning formula alone does not guarantee FVI attainment.

The catalogue contains seven economic cells: $d = 2, 3, 4$, $T = 2$, $p = 1$, $\varepsilon = 5$, and $d = 2$, $T = 2, 3$, $p = 1, 4$, $\varepsilon = 2$. Witness and uniform FVI are each run on all cells. Two additional services choose the smallest isotropic power-of-two allocation $N = K = M$ satisfying the same primitive target-five budget in $d = 2$, $T = 2$, $p = 1$. A further service in that cell uses residual-driven FVI. Each of these seventeen specifications is executed three times in separate single-CPU processes, with identical warm-up and rotated service order. Repetitions estimate local timing dispersion, not a distribution of economic tasks or neural initializations.

The residual-driven comparator queries its own future Bellman objective at interval endpoints and midpoints. It inserts coordinate knots where the largest interpolation surplus across the other coordinates is greatest, then equidistributes square-root surplus mass along each axis, with dyadic rounding and explicit monotone-knot checks. This last step follows the local quadratic interpolation-error scaling, rather than the curvature-mass rule of the earlier graded block. All pilot queries, residual comparisons, compilation and output are charged. Recorded geometry, not the method name, determines whether adaptation occurred. This is a coordinate-tensor comparator, not a claim to have executed the whole adaptive sparse-grid literature.

The original R50 design and source were frozen locally before their measured execution. Its early bisection-only development test and subsequent equidistribution correction are retained with their original records. R51 freezes the recovered algorithm before a known-design reproduction; it is not presented as a new blind experiment. The present clocks belong only to this execution, and the earlier R49 evidence is unchanged.

TABLE M.I: Primary prospective allocations and realized certificates

$d;T;p$	Method	N,K,M	ϵ	Bound	Pass	Sec.	Sec. range
2;2;1	W	32,2,2	2	1.9713	3/3	0.0847	[0.0841,0.0852]
2;2;1	W I	8,8,8	5	4.6296	3/3	0.0195	[0.0193,0.0200]
2;2;1	W	16,1,1	5	3.9427	3/3	0.0240	[0.0238,0.0268]
2;2;4	W	32,8,1	2	1.9317	3/3	0.1086	[0.1072,0.1105]
2;3;1	W	64,4,2	2	1.7567	3/3	0.5059	[0.5001,0.5075]
2;3;4	W	64,4,2	2	1.9149	3/3	0.5058	[0.4962,0.5095]
3;2;1	W	16,1,1	5	4.0581	3/3	0.4424	[0.4335,0.4610]
4;2;1	W	8,8,2	5	4.8628	3/3	1.8441	[1.8348,1.9009]
2;2;1	A	16,1,1	5	4.2586	3/3	1.0709	[1.0636,1.0941]
2;2;1	F	32,2,2	2	2.1589	0/3	0.0644	[0.0631,0.0646]
2;2;1	F I	8,8,8	5	5.3703	0/3	0.0232	[0.0227,0.0233]
2;2;1	F	16,1,1	5	4.2199	3/3	0.0243	[0.0242,0.0246]
2;2;4	F	32,8,1	2	2.0923	0/3	0.0900	[0.0869,0.0918]

$d;T;p$	Method	N,K,M	ε	Bound	Pass	Sec.	Sec. range
2;3;1	F	64,4,2	2	1.7761	3/3	0.3650	[0.3495,0.3661]
2;3;4	F	64,4,2	2	1.9038	3/3	0.3606	[0.3590,0.3723]
3;2;1	F	16,1,1	5	4.3316	3/3	0.3890	[0.3831,0.4258]
4;2;1	F	8,8,2	5	5.653	0/3	2.4360	[2.4275,2.4645]

W: native witness with equivalent ReLU realization; F: uniform FVI; A: residual-driven FVI; I: isotropic allocation. Each specification has three identical mathematical checkpoints and three distinct clocks. The primitive planning guarantee is witness-specific; FVI uses its actual critic modulus and certificate. Decimal rounding never determines attainment.

The primary block certifies 24 of 24 witness services, 12 of 24 uniform FVI services and 3 of three residual-driven services. These counts include timing repetitions, not independent policy draws. The residual comparator has nonuniform axes in all 3 date models of its repeated checkpoint. Its bound is 4.2586, compared with 4.2199 for uniform FVI at the same quota. Its pilot work is retained.

In the isotropic comparison cell, the planned witness uses 1,156 target innovation evaluations, versus 11,664 for the isotropic service. The latter has fewer state nodes; query savings are therefore not a claim of less work in every category. The table reports measured complete clocks rather than inferring a speed ratio from this count.

M.2. A separate common-accuracy completion

Inspection of the primary block showed that sharing a witness-sufficient allocation does not make the FVI certificate pass. We therefore froze a separate, explicit service: in $d = 2, 3, 4$ at $T = 2, p = 1, \varepsilon = 5$, begin with the same planner; if a method's actual certificate fails, double only N , reconstruct from primitives and retain the failed attempt. Stop at first success or at $N = 64$. Both methods face this rule. There are three new isolated repetitions per method and dimension. The source freeze and results are separate from the primary catalogue, and old failures remain failures.

TABLE M.II: Separate common-target completion: horizon two, price one, target five

d	Method	N range	Tries	Bound	Pass	Sec.	Sec. range	MiB
2	W	16–16	1	3.9427	3/3	0.0242	[0.0223,0.0281]	41.4
2	F	16–16	1	4.2199	3/3	0.0244	[0.0236,0.0275]	41.7
3	W	16–16	1	4.0581	3/3	0.4383	[0.4155,0.4386]	44.3
3	F	16–16	1	4.3316	3/3	0.3913	[0.3840,0.3935]	42.3
4	W	8–8	1	4.8628	3/3	1.7418	[1.7139,1.8030]	50.5
4	F	8–16	2	3.1181	3/3	29.0946	[28.9063,29.3447]	169.8

All failed state refinements are included in the complete prefix. MiB is the maximum process peak over three repetitions, not just model storage. The amendment followed the primary outcomes and was frozen before its own execution. Neither block overwrites the other. Host exclusivity and CPU frequency were not controlled.

The returned policies receive new direct actual-cost comparisons below. These observations provide a common target in every reported dimension. They remain a finite low-dimensional tensor experiment. The construction and corner factors $(N + 1)^d$ and 2^d are not removed by the compiler, the terminal repair or this allocation rule.

M.3. Direct policy costs and the source of removable loss

The primary cost experiment fixes twenty-five groups. Sixteen use the original first-target-two R49 witness and FVI policies in all four price–horizon cells, evaluated under the uniform law and each of the three declared point laws. Seven use the new planned services under uniform initial states; one compares residual-driven and uniform FVI; one uses the isotropic allocations. A separate three-group family evaluates the common-accuracy amendment’s returned policies. Every group contains both original and final-repaired policies, with the same eight-bit observation and $1/4096$ action contract. Both generators receive the same improvement.

Each group uses 131,072 continuous-law interval paths and analytic final-innovation integration. Four absolute costs and all six pair contrasts are recorded.

The primary 250 estimands and supplemental 30 estimands are covered by the conservative endpoint union-bound budget of 512 and log upper bound 13. The supplemental groups have their own conditional independent-bin model within the declared simultaneous family; the policies may depend on earlier construction information. Fixed streams establish reproducibility only. The older R49 family retains its own separate error account.

TABLE M.III: Actual policy differences and removable final-action loss

Case	$J_L - J_R$	$J_{L^+} - J_{R^+}$	$J_L - J_{L^+}$
C2;2;1;U	[0.023033,0.028988]	[0.012051,0.017775]	[0.011448,0.014567]
C3;2;1;U	[0.01796,0.023776]	[0.010129,0.015794]	[0.0077535,0.010757]
C4;2;1;U	[0.033636,0.039473]	[0.015912,0.021433]	[0.016432,0.019542]
I2;2;1;U	[0.042713,0.048868]	[0.018209,0.023856]	[0.023541,0.026849]
O2;2;1;M	[0.0064146,0.011415]	[-0.0023843,0.0023843]	[0.007622,0.010239]
O2;2;1;L	[0.054972,0.060123]	[0.044007,0.048792]	[0.0099411,0.012722]
O2;2;1;H	[0.0042173,0.0090667]	[0.0022567,0.007049]	[0.0014934,0.003946]
O2;2;1;U	[0.013161,0.018416]	[0.0064562,0.011605]	[0.0056474,0.0083165]
O2;2;4;M	[0.035287,0.040496]	[0.030381,0.035526]	[0.0036933,0.0063836]
O2;2;4;L	[0.020345,0.025506]	[0.017976,0.023095]	[0.0010828,0.0036974]
O2;2;4;H	[0.039236,0.044497]	[0.034549,0.039692]	[0.0034416,0.006155]
O2;2;4;U	[0.011181,0.016658]	[0.0058295,0.011256]	[0.004188,0.0069254]
O2;3;1;M	[0.015697,0.021868]	[0.012485,0.018641]	[0.0016956,0.0047777]
O2;3;1;L	[0.026367,0.032516]	[0.022928,0.029055]	[0.0020031,0.005116]
O2;3;1;H	[0.011221,0.017268]	[0.0095515,0.015586]	[0.00030976,0.0033545]
O2;3;1;U	[0.0095127,0.015773]	[0.0065302,0.012761]	[0.0014967,0.0045926]
O2;3;4;M	[0.010433,0.016948]	[0.0089337,0.015429]	[9.4671e-05,0.0033779]
O2;3;4;L	[0.0095446,0.01608]	[0.0074164,0.013916]	[0.0005048,0.0037993]
O2;3;4;H	[0.0073415,0.013898]	[0.0046083,0.011119]	[0.0011154,0.0044219]
O2;3;4;U	[0.0056055,0.012273]	[0.0039105,0.010565]	[0.00021625,0.0035063]
Q2;2;1;U	[0.01186,0.017187]	[0.0056983,0.010907]	[0.0055893,0.0083122]
P2;2;1;U	[0.022852,0.028801]	[0.011835,0.01755]	[0.011497,0.014615]

Case	$J_L - J_R$	$J_{L^+} - J_{R^+}$	$J_L - J_{L^+}$
Q2;2;4;U	[0.0086864,0.014074]	[0.004022,0.0093656]	[0.003535,0.0062382]
Q2;3;1;U	[0.01136,0.017753]	[0.0081953,0.014553]	[0.0017403,0.0048791]
Q2;3;4;U	[0.010559,0.017349]	[0.0074835,0.014248]	[0.0020091,0.0053568]
P3;2;1;U	[0.017874,0.023687]	[0.0099988,0.015657]	[0.0078066,0.010814]
P4;2;1;U	[0.033196,0.039029]	[0.015676,0.021193]	[0.016374,0.019482]
A2;2;1;U	[-0.0024517,0.0023493]	[-0.0024125,0.0023985]	[0.00062241,0.0030972]

A case is block/dimension; horizon; price; initial law. O: original R49 target-two policies; P: primary target-five plan; Q: primary target-two plan; C: common-accuracy amendment; I: isotropic; A: surplus FVI versus uniform FVI. Except A, left is witness and right is uniform FVI. U is uniform; L,M,H are point states 1/8,1/2,7/8. Superscript + denotes the identical terminal safety rule. Five-significant-digit display is not an exact endpoint; the rational records govern all conclusions.

All 27 declared witness comparison groups identify a strictly positive reduction of the witness incumbent's actual expected cost. The simultaneous gain lower endpoints range from 9.4671e-05 to 0.023541; upper endpoints range from 0.0033545 to 0.026849. This is a law-specific statement for the stated policies, not strict improvement at every initial state or 27 independent training successes.

Before repair, witness cost is identified as higher in 27 groups. After both methods receive the same repair, it remains higher in 26, is lower in 0, and is unresolved in 1. The residual-driven FVI comparison is unresolved both before and after repair. Thus the terminal construction repairs a genuine economic loss without manufacturing a reversal of the strong conventional comparison.

The active evaluator records 3,670,016 interval paths, 7,340,032 cell-gate queries and 5,447,560 accepted changes across the two incumbents. No acquisition ambiguity or broad-hull fallback is recorded. Its joint direct-evaluation time before durable record output is 127.95 seconds. That total is evaluation work, not a speed comparison between independently executed policies.

For each incumbent the direct reduction is the centered final contrast. The remaining witness-FVI difference after both repairs is an actual comparison of the two

changed policies. The identity

$$J_W - J_F = (J_W - J_{W+}) + (J_{W+} - J_{F+}) - (J_F - J_{F+})$$

separates removable final-action losses from the difference that remains. It does not attribute every remaining discrepancy to a single representation mechanism. The deposited date-level action means, terminal state means, accepted and blocked proposals, and capacity/quantum repair counts permit this distinction to be checked. A finite-sample positive gain can justify a strictly positive replacement charge up to its simultaneous lower endpoint; a deterministic nonincrease alone does not justify an arbitrary implementation fee.

M.4. *Actual-cost frontiers and complete resource boundaries*

The supplement reports every group's four absolute cost intervals and every direct contrast. For each new construction, its resource coordinate includes the full own-future service and, separately, the joint evaluation charge. The reader can use either a shared fixed evaluation cost for the catalogue decision or the complete joint cost required to certify a particular comparison. We do not claim to have measured independent marginal costs for four separately simulated policies when the implementation shares their earlier trajectories. The pair clock is an explicitly charged object, not an unpriced convenience.

TABLE M.IV: Construction work and actual cost of new implementations

Case	Method	Build sec.	Original cost	Repaired cost	Pair sec.
C2;2;1;U	W	0.0242	[0.6242,0.6378]	[0.61117,0.62482]	3.799
C2;2;1;U	F	0.0244	[0.59821,0.61177]	[0.59628,0.60988]	3.799
C3;2;1;U	W	0.4383	[0.58992,0.60068]	[0.58062,0.59146]	6.102
C3;2;1;U	F	0.3913	[0.56904,0.57982]	[0.56768,0.57849]	6.102
C4;2;1;U	W	1.7418	[0.62188,0.63199]	[0.60399,0.61391]	9.790
C4;2;1;U	F	29.0946	[0.5855,0.59525]	[0.5854,0.59515]	9.790

Case	Method	Build sec.	Original cost	Repaired cost	Pair sec.
I2;2;1;U	W	0.0195	[0.63959,0.65386]	[0.61456,0.6285]	3.744
I2;2;1;U	F	0.0232	[0.59407,0.60781]	[0.59363,0.60737]	3.744
Q2;2;1;U	W	0.0847	[0.60972,0.62332]	[0.60276,0.61638]	3.870
Q2;2;1;U	F	0.0644	[0.59518,0.60882]	[0.59444,0.6081]	3.870
P2;2;1;U	W	0.0240	[0.6222,0.63577]	[0.60912,0.62274]	3.901
P2;2;1;U	F	0.0243	[0.59639,0.60993]	[0.59445,0.60803]	3.901
Q2;2;4;U	W	0.1086	[0.69243,0.70702]	[0.68754,0.70213]	3.919
Q2;2;4;U	F	0.0900	[0.68106,0.69562]	[0.68086,0.69542]	3.919
Q2;3;1;U	W	0.5059	[0.72237,0.73696]	[0.71906,0.73365]	4.406
Q2;3;1;U	F	0.3650	[0.70782,0.7224]	[0.70769,0.72228]	4.406
Q2;3;4;U	W	0.5058	[0.86908,0.88468]	[0.86539,0.881]	4.513
Q2;3;4;U	F	0.3606	[0.85511,0.87074]	[0.85451,0.87015]	4.513
P3;2;1;U	W	0.4424	[0.58981,0.60055]	[0.58046,0.59128]	6.294
P3;2;1;U	F	0.3890	[0.56902,0.57978]	[0.56764,0.57844]	6.294
P4;2;1;U	W	1.8441	[0.61951,0.62957]	[0.60168,0.61155]	7.713
P4;2;1;U	F	2.4360	[0.58357,0.59328]	[0.58332,0.59303]	7.713
A2;2;1;U	A	1.0709	[0.59772,0.61129]	[0.59585,0.60945]	3.797
A2;2;1;U	F	0.0243	[0.59777,0.61134]	[0.59585,0.60946]	3.797

Pair sec. is a shared cost of the whole four-policy comparison, repeated in the display for clarity, not independently attributable to each row. Construction medians use the same cell and fresh records; no historical clocks are spliced. Absolute-cost intervals are wider than centered pair intervals. Charges enter Proposition 42 with their explicitly chosen shared or marginal accounting convention.

All stored coefficients, original owners, action scalars, model evaluation categories, transform relaxations, interpolation/corner operations, terminal derivative queries, whole-cell tests, exact fallbacks, numerical interval widths, serialized bytes and process peak resident memory remain in the records. The counters are primitive categories, not a complete FLOP or bit-operation count. The latter would require an additional cost model. Direct comparison time is not added to clocks from another host or to a historical failure in order to manufacture a new complete service.

The retained sensor catalogue has a structural interpretation through equation (111). Its prices remain normalized primitives, and its confidence-envelope

regret is relative to the best observed implementation in the finite catalogue. This revision neither estimates a sensor technology nor substitutes a sample-mean winner for the true expected-net-cost minimizer.

M.5. *Historical findings that the new evidence does not replace*

The earlier 93 of 96 strictly higher witness-cost contrasts remain visible in Table L.IV. Original curvature-FVI never changed its grid in the realized block. The later genuinely graded method also remains an adverse conventional-comparator experiment: it was less precise in all twelve matched final services and slower in all twelve. These two findings concern the two specified heuristics, not all adaptive dynamic programming. The present surplus comparator adds an independently recorded realized partition, while uniform FVI remains the strong baseline.

Likewise, a successful final-action repair does not establish better earlier continuation learning, universal neural superiority, a guarantee for unrestricted non-convex training, or a dimension-free policy construction. What is established is a cost-directed feasible improvement with an explicit safety condition, its executed original-economy specialization, and source-bound resource and actual-cost comparisons. These results extend the NBO policy-evaluation and improvement program without replacing its economic subject or erasing its previous evidence.

APPENDIX N: EXHAUSTIVE ROBUSTNESS IN THE ORIGINAL INVESTMENT ECONOMY

The action-null result permits a diagnostic in which the true policy and its economic cost are held fixed while a critic's irrelevant component changes. This avoids confusing a new representation with a new economic policy. Before execution, we fixed all eight R49 target-two checkpoint policies retained in the R51 package: witness and uniform FVI for every two-state price-horizon cell. We enumerate every one of the 256^2 closed eight-bit observation cells for each policy. There is no selection of cells by favorable cost or by the new results.

For each cell we retain the original robust capacity repair, the action lattice $1/4096$, and the convex final-date proposal. The centered outward interval for its actual

conditional cost difference is the same object used in the inherited terminal gate. We compare the global-width and contrast-aware gates after replacing g by $g - n_M$, for $M \in \{0, 1, 16, 4096\}$. Proposition 33 establishes the zero action-error allowance algebraically; it is not estimated from this grid. Closed-cell enclosures cover boundary states, while execution retains the original observation tie convention. Symbolic cancellation precedes interval evaluation.

TABLE N.I

EXHAUSTIVE ACTION-NUL ROBUSTNESS ON THE ORIGINAL INVESTMENT ECONOMY

Incumbent	T	p	$M = 0$	$M = 1$	$M = 16$	$M = 4096$
Witness	2	1	48872	0	0	0
Witness	2	4	60574	0	0	0
Witness	3	1	48634	0	0	0
Witness	3	4	60114	0	0	0
FVI	2	1	42198	0	0	0
FVI	2	4	54249	0	0	0
FVI	3	1	37574	0	0	0
FVI	3	4	54249	0	0	0

Each row enumerates all 65,536 closed eight-bit cells. Entries count strict changes accepted by the global-width gate. The contrast-aware gate accepts exactly the $M = 0$ column at every amplitude, with identical returned actions. Earlier policy dates are unchanged. Counts concern cells, not additional expected-cost observations.

The records retain the incumbent and proposed action indices, the returned contrast-aware index, the exact binary64 endpoint encodings, and every global-width decision for every cell. An independent record audit checks cell completeness, cell-wide capacity, all gates, counts and accepted upper bounds. The contrast-aware actions are also compared with a fresh call to the unmodified inherited terminal implementation on every cell. Thus the reported invariance concerns the actual stored policy, not equality of two rounded average costs. The exact finite-state suite separately exercises nonzero action-error bounds, rejection, and all passes of the finite-sweep theorem.

The experiment measures resistance to irrelevant continuation error. It is not a new estimate of $J^{\text{witness}} - J^{\text{FVI}}$. The contrast-aware policies are identical to the unperturbed terminal-repaired incumbents, so repeating their cost simulation would not create additional independent economic observations. The prior actual-cost intervals, their confidence events, and the 93 witness-higher, zero witness-lower and three unresolved R49 comparisons remain unchanged. Safe improvement of either incumbent remains different from superiority over the competing policy.

Every service includes input loading, proposals, true-cost interval comparisons, independent inherited-policy replay, amplitude decisions and durable cell output. Machine-independent query counters, serialization time and peak memory are retained separately. The supplementary work table is the cost of this exhaustive diagnostic, not the construction cost of an alternative learned model. A full T -pass investment-economy execution is not inferred from a final-date ablation. The general finite-sweep result has explicit verified-error hypotheses; the continuous economic experiment executes only the declared specialization.

APPENDIX O: FULL POLICY SWEEPS IN THE NONLINEAR INVESTMENT ECONOMY

The directed certificate is now evaluated beyond an exact-null identity and beyond the final decision. This section uses the original continuous-law two-state investment economy, with state-dependent capacity, nonlinear costs and transitions, and acquired-state implementation. It does not substitute a finite-state model for that economy. The construction and direct-cost records were executed in the separately frozen R53 primary and adaptive studies; the present revision integrates them into the paper and replays their full deposited outputs. That replay adds neither new policies nor new independent samples.

O.1. *Design and the economic endpoint*

The primary design has horizons $T = 2, 3$, price parameter $p = 1$, and two generators: the compiled witness construction and tensor FVI. Each service starts from the primitives and executes every rung of the original prefix up to the declared checkpoint. It checks that its reconstructed original policy has the same identity as the

historical checkpoint. A five-bit observation contract in each coordinate then gives 32^2 closed cells per date, with downward action spacing $1/4096$. Both methods use the same acquisition and robust-feasibility rules.

Starting from that complete acquired incumbent, each service executes exactly T all-date passes. At every pass all new dates are computed against the same old complete incumbent and then adopted together. Each cell has nine robustly feasible proposed actions and eight interval covers of the full continuous-action range. Sixteen dyadic common-innovation bins enclose every unresolved expectation; the last innovation is integrated analytically. Rectangle extrema enclose discontinuous future actors. These calculations implement the directed certificate of Section H with the verification gauge $h = 0$. They are not a claim that a zero critic generates an effective policy, and their policy-evaluation cost is recorded.

The primary numerical endpoint is the actual expected cost of each complete returned policy under the original uniform initial law and continuous innovations. Conditional on the frozen policies, 65,536 common-path rows per horizon produce outward path-cost enclosures using forty-bit continuous-bin cells. The bounded empirical-Bernstein account applies to the endpoint samples under the declared independent-bin model. Dependence between policies through common paths is retained. The primary protocol allocates a simultaneous family error of $1/100$; the adaptive extension has a separate $1/100$ family. Combining the two families requires the corresponding union bound, rather than counting their repeated policies as independent economic objects.

The distinction between cumulative gain and comparative ranking remains important. In the three-date experiment, the witness policy's gain after three full sweeps lies in $[0.00344390, 0.01578829]$, with the displayed endpoints rounded outward. Its final witness-minus-FVI interval is $[-0.00287355, 0.00923410]$. Thus the cumulative own-incumbent gain is identified, while the final cross-method difference is unresolved. The latter is not evidence of equality and not an identified reversal of the original ranking. Individual step gains can be statistically unresolved even though the deterministic certificate guarantees nonworsening expected cost at every true state.

TABLE O.I

ACTUAL EXPECTED POLICY COSTS BEFORE AND AFTER FULL SWEEPS

T	Method	Initial cost	Final cost	Initial minus final
2	W	[0.60258, 0.62351]	[0.59092, 0.61168]	[0.00665, 0.01684]
2	F	[0.58999, 0.61078]	[0.58984, 0.61063]	[-0.00463, 0.00494]
3	W	[0.71853, 0.74116]	[0.70898, 0.73147]	[0.00344, 0.01579]
3	F	[0.70576, 0.72833]	[0.70576, 0.72833]	[-0.00596, 0.00597]

W is the compiled witness construction; F is tensor FVI. Every cost is for a complete acquired policy under the original continuous initial and innovation laws. Endpoints are rounded outward for display.

Paired gain intervals, not subtraction of the displayed marginal intervals, determine gain signs. The policies and observations were frozen and executed in the R53 primary study.

TABLE O.II

PAIRED WITNESS-MINUS-FVI COST AFTER EVERY PASS

T	Pass	Expected-cost difference
2	0	[0.00757, 0.01775]
2	1	[0.00151, 0.01148]
2	2	[-0.00376, 0.00590]
3	0	[0.00663, 0.01896]
3	1	[0.00285, 0.01513]
3	2	[-0.00105, 0.01112]
3	3	[-0.00288, 0.00924]

A positive interval identifies higher witness cost. An interval containing zero is unresolved, not evidence of equality or witness superiority. The same paths couple policies within a horizon; no independence between these estimands is assumed.

For an installed witness controller with a replacement charge $\tau \geq 0$, let $[l, u]$ be the corresponding cumulative gain interval. The net benefit of replacement is in $[l - \tau, u - \tau]$ on the same simultaneous event. Hence every $0 \leq \tau < l$ is supported as strictly beneficial when $l > 0$, and every $\tau > u$ is ruled out as beneficial. The intermediate region is unresolved. This gives a full adoption-fee account for an actual improved policy, rather than only a reduced certificate for an unchanged

policy. These are normalized model units, not an estimated installation charge or a calibrated welfare exercise.

O.2. All-date acceptance and the original Bellman gap

The improvement account records each proposal, the exact retained incumbent, the interval endpoints, the selected action, the continuous-action lower bound and the induced directed gap. Table O.III gives every primary pass, including nonterminal changes. The supplementary tables report every date's positive contrast allowance and directed gap, the number of blocked and admitted candidates, and the comparison with scalar-width and symmetric-contrast gates. No failed or blocked cell is omitted.

TABLE O.III

ALL-DATE POLICY CHANGES AND CERTIFIED GAP RECURSION

<i>T</i>	Method	Pass	Changed	Nonterminal	Gap bound	Prefix (s)
2	W	1	590	17	3.72986	3.462
2	W	2	265	265	0.08149	6.608
2	F	1	653	1	3.59553	3.477
2	F	2	0	0	0.06317	6.639
3	W	1	612	44	5.11709	51.769
3	W	2	282	282	3.65096	101.038
3	W	3	135	135	0.22556	149.583
3	F	1	64	0	4.87775	50.846
3	F	2	0	0	3.40006	99.948
3	F	3	0	0	0.09350	149.173

Changed counts refer to distinct deployed cell-date actions in that pass, not to candidate comparisons. The gap is the directed all-state upper bound against the original continuous-action Bellman optimum.

It need not decrease monotonically. Prefix work includes primitive construction and previous passes through their recorded array outputs; the final service and inference records are charged separately in the catalogue account.

There are two distinct guarantees. A nonpositive directed upper advantage gives safety for an accepted action. The lower interval cover of all feasible actions gives the additional accuracy quantity $\varepsilon_t^k = \sup_C (U_t^k - L_t^k)$. The computed recursion is then $G_t^{k+1} = \beta_t G_{t+1}^k + \varepsilon_t^k$, initialized by a valid support bound. Its terms include cell uncertainty, actor ambiguity, action covering, integration and arithmetic through their actual endpoint propagation. They are not separately estimated from residual samples and then presumed negligible.

A coarse continuous-action cover or a wide future-actor range can make G conservative. The theorem guarantees that the returned policy does not worsen, not that a particular numerical upper bound must decrease at every pass. Executing T passes at nonzero tolerances also does not assert exact optimality: the remaining gap is the displayed sum of propagated directed errors. The exact T -pass optimality statement and its sharp coefficient have their own rational finite-model regression tests.

O.3. *A genuinely nonuniform conventional comparison*

The extension adds surplus-driven FVI at the same two-date primitive checkpoint. Its adaptation history and the number of nonuniform date models are recorded, together with every pilot and every earlier construction rung. It is not the earlier curvature rule that left its grid unchanged. Compiled witness and tensor FVI are reconstructed on the same extension runner, receive the same two complete passes, and reproduce their primary policy identities. A fresh path stream compares all three methods and all passes within a separately frozen statistical family.

This comparison tests an actual changed conventional representation, rather than merely a new label for the tensor baseline. Its adaptive grid nevertheless remains a tensor grid with nonuniform coordinates. It is therefore not presented as a sparse-grid or non-tensor high-dimensional demonstration. The earlier failed curvature rule, the adverse graded refinement, and the original direct-cost comparisons remain in their original tables and records.

TABLE O.IV

FRESH ACTUAL-COST COMPARISON WITH A NONUNIFORM CONVENTIONAL METHOD

Method	Pass	Actual expected cost	Gain from initial policy
W	0	[0.60337, 0.62425]	–
W	1	[0.59714, 0.61789]	[0.00129, 0.01130]
W	2	[0.59173, 0.61245]	[0.00662, 0.01681]
F	0	[0.59081, 0.61156]	–
F	1	[0.59065, 0.61141]	[-0.00462, 0.00494]
F	2	[0.59065, 0.61141]	[-0.00462, 0.00494]
A	0	[0.59068, 0.61143]	–
A	1	[0.59060, 0.61135]	[-0.00470, 0.00486]
A	2	[0.59060, 0.61135]	[-0.00470, 0.00486]

A is surplus-driven FVI; W and F denote the same generators as above. All three are reconstructed in the same separately frozen cohort and receive two full improvement passes. W and F policy identities reproduce the primary cohort, while the cost paths are a fresh family. Nonuniformity and pilot work are recorded explicitly.

O.4. Complete catalogue work at common actual-cost targets

A finite catalogue of returned policies supplies a precise work account. For each method we select the pass with the lowest simultaneous actual-cost upper endpoint. We charge that method's entire observed construction-and-improvement service, including all executed passes and durable records, plus the entire shared cost-inference service. Charging the full shared service to each eligible method is conservative and avoids assigning an unobserved marginal inference time to a method evaluated only jointly. The separate publication audit reports the additional replay and document-production overhead.

For a target u , the eligible methods are exactly those whose selected actual-cost upper bound is at most u . Taking the least recorded complete-catalogue work among eligible methods defines the observed finite frontier. Its breakpoints and method choices are deposited without rounding; the printed table is descriptive. This is a frontier for releasing a policy from an already executed finite catalogue. It is not a claim about the minimum work of an unexecuted sequential stopping

TABLE O.V

COMPLETE-CATALOGUE WORK TO A VERIFIED ACTUAL-COST UPPER BOUND

Cohort	Method	Pass	Cost upper	Service (s)	Inference (s)	Sum (s)
P2	W	2	0.61168	6.612	1.850	8.463
P2	F	1	0.61063	6.643	1.850	8.493
P3	W	3	0.73147	149.587	3.182	152.769
P3	F	1	0.72833	149.178	3.182	152.359
A2	W	2	0.61245	6.646	2.667	9.312
A2	F	1	0.61141	6.579	2.667	9.245
A2	A	1	0.61135	18.521	2.667	21.188

For each method the selected pass minimizes its actual-cost upper endpoint. Work conservatively charges every own executed construction and sweep, plus the entire shared inference service. It is not a first-crossing or minimum-work sequential stopping estimate. Replay and document-production overhead are separately timed in the release audit. The exact finite threshold partition, eligible methods and least-recorded-work choices are deposited in the machine-readable audit.

algorithm, an optimized resource allocation, or a universal fastest generator. In particular, improving the witness incumbent does not by itself establish a better work-to-cost frontier than FVI.

The work tables separately identify construction time, pair-node counts, coupling work, actor queries, actor-table bytes, process peak memory, and the cumulative clocks. A common-innovation certificate avoids a stored table over all state pairs, but the work bound in Proposition 36 makes its horizon dependence explicit. These records turn a constructive mathematical certificate into an executed and costed verification method in the stated economy. They do not remove the tensor observation cover or establish a representation-specific scaling advantage.

O.5. *Controlled approximate-null and false-null tests*

The extension varies three departures from the exact action-null example: the ReLU direction, the deterministic action exposure, and a known action-dependent innovation-mean coefficient. Each departure ranges over 0, $1/4096$, $1/256$, and $1/16$; the amplitudes are 1, 16, and 4096. Two original incumbents yield 384 controlled

cases, each evaluated on all 32^2 closed cells. A verified positive allowance replaces the zero certificate whenever the altered coefficient is nonzero.

TABLE O.VI
CONTROLLED DEVIATIONS FROM EXACT ACTION-NULL STRUCTURE

Quantity	Recorded value
Controlled cases	384
Closed cells per case	1024
Cell–case combinations	393216
Positive-bound cases	378
Safe changed cell–cases	170374
Harmful false-null cell–cases	116185
Maximum contrast allowance	315.00000
New statistical observations	0

The catalogue changes known ReLU directions, action exposures and action-dependent innovation means. Harm means a strictly positive lower true-cost difference for an action selected by the deliberately false-null rule. Corrected decisions pass the nonpositive true-upper check. Cell–cases are deterministic repeated diagnostics, not independent economic observations.

The corrected gate and the deliberately false-null gate are both replayed. Every selected corrected action is feasible and has a nonpositive true-advantage upper endpoint. For the false-null policy, a strictly positive true-advantage lower endpoint identifies a harmful decision, not merely an uncertain one. The full catalogue retains accepted, blocked and unchanged actions, including cases where large non-nullity makes the verified gate conservative. These are deterministic diagnostics for known perturbed primitives. They are not additional cost-inference samples, an automatic discovery of learned null structure, or a guarantee under an unknown misspecified transition law.

APPENDIX P: ECONOMIC SCOPE AND CONCLUSION

The economic contribution is an explicit connection between a policy’s construction resources and its implementable cost. A primitive error budget can determine

whether to refine the state cover, action search or integration. A complete finite tolerance frontier states when one generator returns a certified policy earlier. Independent comparisons of the returned policies then determine what can be said about controller replacement and purchased information. These are different steps; none can be replaced by the favorable outcome of another.

The constructive backend and exact compiler make positive deterministic statements under explicit moduli and effective-computability conditions. Their guarantees are not success probabilities for unrelated nonconvex training algorithms. The matched experiments and direct-cost intervals make finite, source-bound claims about the declared economies, resource ladders and initial laws. Common-accuracy dimension cells do not establish an asymptotic scaling law, and endogenous precision within a finite sensor catalogue is not an optimum over all information technologies. The original broad theory, applications and adverse experimental record remain part of the complete NBO development.

The cost-directed result adds a further positive statement: a feasible acquired policy can be improved without sacrificing its incumbent cost guarantee, provided its centered advantage is verified. In the executed final-stage specialization every studied witness incumbent has a strictly positive identified reduction under the declared initial law. Most witness–FVI differences remain unfavorable after both methods are repaired. The improvement theorem and this comparative observation are consistent: safe improvement relative to an incumbent is not dominance over another method.

Neural Bellman Operators therefore provide a way to organize continuation construction, feasible implementation, certification and economic resource choice around the policy actually returned. Whether a particular representation or generator is preferable is decided by its full certified frontier and actual economic consequences, rather than by the name of its function class.

APPENDIX Q: SUBMISSION AND PRESERVATION MAP

This article and its technical supplement are the active R54 exposition of the same paper. Every R49 main-article and supplementary label is retained in its

corresponding current document. The complete development editions retain the R48 main and supplementary bodies and add the current theory and evidence, with a reading notice distinguishing historical wording from the active empirical account. The original R49 sources remain unchanged under `preserved/R49`; R48 sources and earlier theory/application companions are likewise retained. Preservation is not a new mathematical validation of every historical application.

The local revision contains ordinary sources, both immutable R49 evidence archives, the R48 publication input, newly executed records, generated tables, a point-by-point response and compiled documents. The direct numerical correction retains its original observations and replays the same streams with explicitly outward cross-policy subtraction. The corrected intervals, not their raw predecessors, govern the active direct-cost tables. Source and evidence identities are checked before a clean offline Git-archive rebuild. The local delivery record distinguishes a completed build from a remote publication: no remote branch or commit is claimed without a confirmed push. This limitation is administrative, not a substitute for the substantive revised manuscript.

APPENDIX R: THE NONLINEAR INVESTMENT FAMILY

The economic law is the original dimension-normalized investment family. For $x \in [0, 1]^d$, scalar action $0 \leq a \leq \kappa(x)$, and cyclic coordinate $v = j + 1$,

$$\kappa(x) = \frac{1}{8} + \frac{1}{8d} \sum_j x_j, \quad (114)$$

$$F_j(x, a, Z) = \frac{1}{16} + \frac{x_j}{2} + \frac{x_v}{8} + \frac{x_j(1 - x_v)}{16} + \gamma_j a + s_j Z,$$

where (γ_j, s_j) alternates between $(1/2, 1)$ and $(1/4, -1)$ and Z is uniform on $[-1/32, 1/32]$. The initial state is uniform on the cube and $\beta = 15/16$. The stage and terminal costs are

$$c(x, a) = \frac{2}{d} \sum_j (x_j - 5/8)^2 + \frac{1}{4d} \sum_j (x_j - x_{j+1})^2 + 2 \left(\frac{1}{2} - \frac{2}{d} \sum_j x_j \right)_+^2 + a^2 + 4a^4, \quad (115)$$

$$g(x) = \frac{4}{d} \sum_j (x_j - 5/8)^2 + \frac{1}{4d} \sum_j (x_j - x_{j+1})^2 + 2 \left(\frac{1}{2} - \frac{2}{d} \sum_j x_j \right)_+^2.$$

The original coefficients and continuous laws are unchanged. The domain is invariant: the drift lies in $[1/16, 11/16]$, $a \leq 1/4$, $\gamma_j \leq 1/2$, and the innovation magnitude is at most $1/32$. The results are for this theoretical economy, not calibrated installation fees or estimated welfare magnitudes.

APPENDIX S: TRAINED NEURAL CONTINUATIONS AND EXACT INNOVATION INTEGRATION

The constructive backend supplies a feasible reference without relying on the success of nonconvex optimization. A trained backend addresses a different computational question: whether features fitted to the operator's own continuation targets produce useful candidate actions at a prescribed resource budget. These two backends are not identified with one another. In particular, fitting the hidden weights below is not compilation of the min-plus construction in Theorem 22.

Let the transition have the one-factor form

$$F_t(x, a, Z) = \mu_t(x, a) + s_t(x, a)Z, \quad Z \sim \text{Unif}[-r, r].$$

The deterministic map μ_t may be nonlinear and the state may be multidimensional. Consider a fitted continuation with m hidden units,

$$h(y) = c + v^\top y + \sum_{j=1}^m u_j (w_j^\top y + b_j)_+. \quad (116)$$

Both hidden and output coefficients may be data dependent. Once a candidate has been constructed, its stored coefficients specify an exact numerical function; uncertainty about whether training found a good function is not an arithmetic uncertainty about the function that was stored.

LEMMA 43—Integrated neural continuation: Put $z_j = w_j^\top \mu_t(x, a) + b_j$ and $R_j = r|w_j^\top s_t(x, a)|$. Then

$$P_t^a h(x) = c + v^\top \mu_t(x, a) + \sum_{j=1}^m u_j \Phi(z_j, R_j), \quad (117)$$

where $\Phi(z, 0) = z_+$ and, for $R > 0$,

$$\Phi(z, R) = \begin{cases} 0, & z \leq -R, \\ (z + R)^2 / (4R), & -R < z < R, \\ z, & z \geq R. \end{cases}$$

For fixed stored coefficients and evaluable primitives, a point expectation uses $O(md)$ arithmetic operations in addition to evaluating μ_t and s_t . Outward evaluation on a state-action box gives a valid enclosure. There is no innovation-grid discretization remainder in this analytic expectation.

PROOF: The scalar random variable $w_j^\top s_t Z$ is uniform on $[-R_j, R_j]$, or is identically zero. Integrating its positive part gives the displayed three cases. Linearity of expectation gives (117); the affine innovation term has mean zero. Each feature requires two inner products and a bounded number of scalar operations. The function Φ is nondecreasing in each argument on $R \geq 0$: in the interior its radius derivative is $(R^2 - z^2)/(4R^2) \geq 0$, and outside the interior it is constant in the radius. Endpoint evaluation with directed arithmetic, enclosing an uncertain branch rather than choosing a favorable branch, therefore preserves containment. Negative output weights must reverse interval endpoints in the usual way. *Q.E.D.*

The formula is a property of this neural class under the stated innovation law, not a general exact-integration result for arbitrary deep networks or unknown kernels. Quadratic critics also have analytic expectations here; tree critics use valid bin enclosures. Thus eliminating quadrature for the neural backend is a specific computational facility, not a claim that conventional approximation necessarily needs quadrature. On the active hinge region the equivalent expression $(z + R)^2 / (4R)$

avoids subtracting two nearly equal squares. At $R = 0$ the separate branch is essential.

The fitted construction proceeds backward. At each date it draws training states, forms a finite feasible action menu, evaluates current cost plus the integrated *own* subsequent critic, minimizes these labels, and fits the next critic. Its training labels are not evaluations of the optimal Bellman continuation. Fitted value iteration with general function classes has a substantial literature; see [Munos and Szepesvári \(2008\)](#) and the tree-based construction of [Ernst et al. \(2005\)](#). Our use of training is accompanied by a separate whole-cell verification obligation. No training mean-square error is substituted for that obligation.

In the executed neural construction the hidden layer has width 32. All hidden and output coefficients receive 120 full-batch Adam updates, followed by a ridge fit of the affine and output coefficients on the learned features. Coefficients are rounded to 24 binary fractional bits. The verifier evaluates this quantized function, not a hypothetical unrounded optimizer. The fitted hidden weights, targets, states, seeds and quantized output are retained. A fitted nonlinear network, a quadratic continuation, an ExtraTrees continuation, and the two original constructive frontends are therefore identifiable, distinct candidate generators. Whether their additional proposals change the deployed policy is an empirical question that we report separately from their fitting errors.

S.1. *Certified action search for the fitted neural continuation*

The same integrated representation also permits a constructive action search, rather than a second unconverged optimization problem. Fix a reported state in the original economy, so the deterministic next state is $\mu(x, 0) + \gamma a$. Let the feasible deployed actions be $\{0, \Delta, \dots, M\Delta\}$, with $M\Delta$ no greater than the whole-cell lower capacity. Interpret stored coefficients and the reported state as exact dyadic numbers.

THEOREM 44—Algebraic neural action witnesses: *For a continuation of the form (116), its integrated Bellman objective in the original scalar investment action has at most*

2m interior breakpoints. Between consecutive breakpoints it is a quartic polynomial

$$4a^4 + q_2a^2 + q_1a + q_0,$$

with rational coefficients. An exact minimizer on the deployed action lattice can be found among at most $22m + 13$ candidate indices, independently of the number of lattice points. Construction requires $O(md + m^2 + m \log m)$ rational arithmetic and comparison operations, plus $O(m)$ cubic-root isolation problems at the action precision. Coefficient bit lengths and root-isolation work must also be charged. The selected action is an original fitted-critic action witness, not a min-plus encoding.

PROOF: For feature j , write its preactivation mean as $z_j + s_j a$ and its innovation radius as R_j . When $s_j \neq 0$, branch changes occur only at $(\pm R_j - z_j)/s_j$; for $R_j = 0$ these coincide. Retain those breakpoints inside the feasible interval and its two endpoints. Lemma 43 makes each feature quadratic, affine or zero between breakpoints. Adding $a^2 + 4a^4$ and the affine part of the critic gives the displayed polynomial. Its derivative is $16a^3 + 2q_2a + q_1$ and has at most three real roots in each of at most $2m + 1$ pieces.

Between stationary points and breakpoints the objective is monotone. Hence a lattice minimizer lies immediately to either side of a breakpoint or a stationary point, or at a feasible endpoint. Isolate each stationary root in a rational interval of width at most $\Delta/4$ and include every index from the floor of its lower endpoint divided by Δ through the ceiling of its upper endpoint divided by Δ . This contributes at most three indices per isolated root and contains the required neighboring lattice points. The breakpoints and endpoints contribute at most $2(2m + 2)$ indices. Thus $9(2m + 1) + 4m + 4 = 22m + 13$ is a valid bound. Extra indices and duplicate candidates do not invalidate minimization. Exact rational evaluation and lexicographic value-action comparison settle ties.

Computing feature coefficients costs $O(md)$, sorting the breakpoints costs $O(m \log m)$ comparisons, and literal formation of the polynomial on each piece costs $O(m^2)$ operations. Rational isolation of the finitely many cubic roots terminates at positive action precision, including multiple roots. This proves the construction and

the stated separation between arithmetic counts and bit-dependent root-isolation work. Q.E.D.

The solver is implemented with rational root isolation and is checked against exhaustive exact lattice minimization, including saved fitted critics from each of the three state dimensions. It minimizes the stored critic’s objective at the reported state. Whole-cell safety still requires the true-advantage certificate; critic optimality alone does not justify deployment. The recorded economic experiment used its prospectively fixed 17-action construction menu. Its policies and measured clocks are not retroactively replaced by this new action solver, and its cost intervals are not attributed to untested replacement policies.

APPENDIX T: A COMPLETE-POLICY CERTIFICATE ON NON-TENSOR ACQUIRED STATES

For the remainder of the new verification and stopping results, the recursion is the bounded finite-horizon expected-cost model. Let $\mathcal{C}$ be a finite measurable partition, with closed boxes used only to verify the cells. Boundary ties follow a specified observation rule. A deployed action is constant on each reported cell and feasible at every state in its closure. A verification box meeting several cells must retain every possible future action. The argument does not require a tensor partition or a continuous actor.

Fix a complete incumbent π^k and define its true advantage

$$A_t^k(x, a) = c_t(x, a) + \beta_t P_t^a J_{t+1}^{\pi^k}(x) - J_t^{\pi^k}(x).$$

The incumbent has advantage zero. For each cell C , let a_i be robustly feasible candidates and let $\mathcal{B}_t(C)$ be a finite interval cover containing every action feasible at every $x \in C$. The cover may include additional actions; its bounds must remain valid on the feasible subset. Suppose verified endpoints satisfy

$$A_t^k(x, a_i) \leq u_i(C), \quad \ell_B(C) \leq A_t^k(x, a) \quad (a \in B \cap A_t(x)).$$

Set $U_t^k(C) = \min\{0, \min_i u_i(C)\}$ and $L_t^k(C) = \min\{0, \min_B \ell_B(C)\}$. Select the candidate attaining U , retaining the incumbent on a tie. All dates are computed against the same old policy and adopted simultaneously.

THEOREM 45—Directed policy improvement: *Under these conditions the new policy is feasible and*

$$J_t^{\pi^{k+1}}(x) \leq J_t^{\pi^k}(x) \quad (118)$$

for every state and date. With $E_t^k = \sup_x (J_t^{\pi^k}(x) - J_t^*(x))$ and $\epsilon_t^k = \max_C (U_t^k(C) - L_t^k(C))$,

$$E_t^{k+1} \leq \beta_t E_{t+1}^k + \epsilon_t^k. \quad (119)$$

For $B_{t,j} = \prod_{i=0}^{j-1} \beta_{t+i}$, $B_{t,0} = 1$, and $m = \min\{K, T - t\}$,

$$E_t^K \leq \sum_{j=0}^{m-1} B_{t,j} \epsilon_{t+j}^{K-1-j} + B_{t,m} E_{t+m}^{K-m}.$$

The final term vanishes when $K \geq T - t$. Exact full-action greedy certificates thus recover the original Bellman optimum within T passes. The coefficient one multiplying the directed error cannot be reduced uniformly.

PROOF: The selected advantage is at most $U \leq 0$. Starting with the common terminal value, backward monotonicity proves (118). The full-action lower cover gives

$$A_t^k(x, \pi_t^{k+1}(x)) - \inf_{a \in A_t(x)} A_t^k(x, a) \leq U_t^k(C) - L_t^k(C).$$

Replacing the new continuation by the old one using safety, and comparing with an arbitrarily near-optimal Bellman action, yields (119). Taking the infimum limit avoids an attainment assumption. Iterating the inequality gives the displayed finite-horizon expansion. For sharpness take one decision date, an incumbent costing $\delta > 0$, and a feasible unproposed action costing zero. An exact incumbent upper endpoint is zero and an exact covering lower endpoint is $-\delta$. Retaining the incumbent gives both error and directed allowance equal to δ . *Q.E.D.*

Theorem 45 restates the original directed result in the partition notation used by the new implementation. It is deliberately representation neutral. Its premise is stronger than checking a finite candidate menu: without the full-action cover, a safe improvement need not have a small gap from the continuous-action optimum. A change of acquisition partition between passes is permitted when the old policy is transferred exactly to its child cells; the complete incumbent remains the same function.

T.1. A one-state residual and value cache

Let $h_T = g$ and h_t be any stored critics, including fitted networks. Write $e_t = J_t^\pi - h_t$ and

$$r_t = c_t(\cdot, \pi_t) + \beta_t P_t^{\pi_t} h_{t+1} - h_t, \quad e_t = r_t + \beta_t P_t^{\pi_t} e_{t+1}.$$

A cellwise residual interval $\mathcal{R}_t(C)$ encloses r_t throughout the cell. If $\mathcal{E}_{t+1}(C')$ encloses the next error, each innovation bin maps a current box and action into a next-state box. The hull of $\mathcal{E}_{t+1}(C')$ over all cells intersecting that box encloses every possible next error. Denote the probability-weighted sum of those hulls by $\mathfrak{P}_t^a \mathcal{E}_{t+1}(C)$. Bin probabilities refer to the original continuous law, not an empirical law of midpoints.

PROPOSITION 46—Cached residual verification: *Starting from $\mathcal{E}_T = \{0\}$, the backward recursion*

$$\mathcal{E}_t(C) = \mathcal{R}_t(C) + \beta_t \mathfrak{P}_t^{\pi_t} \mathcal{E}_{t+1}(C) \tag{120}$$

encloses e_t on every cell. An independently computed direct value interval $\mathcal{V}_t(C)$ may be intersected with $h_t(C) + \mathcal{E}_t(C)$, and $\mathcal{E}_t(C)$ with $\mathcal{V}_t(C) - h_t(C)$. These intersections preserve containment. For feasible a, b the sum

$$c_t(C, a) - c_t(C, b) + \beta_t \{ (P_t^a - P_t^b) h_{t+1}(C) + \mathfrak{P}_t^a \mathcal{E}_{t+1}(C) - \mathfrak{P}_t^b \mathcal{E}_{t+1}(C) \}$$

encloses the fixed-incumbent action difference. Intersecting it with the analogous direct-value difference gives a valid directed interval. Identical actions have the exact interval $\{0\}$.

PROOF: For every innovation in a bin, the next state belongs to one of the intersected cells; its error lies in the corresponding stored interval and hence in their hull. Integration against the true bin probabilities preserves order. Backward induction in the residual identity proves (120). Each displayed intersection combines two sets known to contain the same real value. Substituting $J^\pi = h + e$ in the action difference proves the final statement. An empty intersection indicates an invalid numerical calculation and must stop acceptance, rather than licensing a selected midpoint. Q.E.D.

This construction stores one-state bands rather than a table over state pairs. It sacrifices dependence between two trajectories when taking separate hulls; in particular, actor ambiguity and broad residual ranges can make it conservative. For locally Lipschitz critics, a centered residual evaluation can be intersected with the natural interval. A valid generalized-gradient enclosure suffices for the line-segment mean-value bound of a ReLU critic. This is not permission to differentiate a discontinuous acquired actor. On one cell its action is fixed; across cells the cache explicitly retains the alternative actions.

T.2. Signed sensitivity of the installed reference

The original investment family in Section R has a smooth zero-investment reference. Its stage and terminal state costs are continuously differentiable, including the squared shortage term. Let μ denote the zero-noise drift and γ the action exposure. For a state box X_0 , propagate boxes X_{j+1} containing $\mu(X_j, 0) + sZ$ for the complete innovation support. Starting from a terminal gradient enclosure $\mathcal{G}_r \supset \nabla g(X_r)$, put

$$\mathcal{G}_j \supset \nabla c(X_j, 0) + \beta D_x \mu(X_j, 0)^\top \mathcal{G}_{j+1}. \quad (121)$$

The inclusion means outward interval multiplication and addition. It is evaluated backward along forward state tubes, with the true sparse Jacobian.

PROPOSITION 47—Signed reference certificate: *The interval $\mathcal{G}_0$ encloses ∇J_0^0 throughout X_0 . For a nonnegative proposed action a , enclose the gradient of the remaining zero policy at every state $\mu(x, \theta a) + sZ$, $0 \leq \theta \leq 1$. If $[v_-, v_+]$ encloses the average*

projection $\mathbb{E}[\gamma^\top \nabla J_{t+1}^0]$ over that segment, then

$$A_t^0(x, a) \in a^2(1 + 4a^2) + \beta a[v_-, v_+]. \quad (122)$$

The signed interval may be intersected with Proposition 46. The construction is valid for the complete zero reference; it does not differentiate a subsequently changed, cellwise constant policy.

PROOF: Differentiate the finite fixed-policy recursion. On the invariant compact domain, the derivatives are bounded, so differentiation and the finite conditional expectations may be interchanged. For every shock path the chain rule has the enclosed Jacobian form in (121); backward induction and expectation preserve the enclosure. The fundamental theorem of calculus along the first action segment gives

$$(P_t^a - P_t^0)J_{t+1}^0(x) = a \int_0^1 \mathbb{E}[\gamma^\top \nabla J_{t+1}^0(\mu(x, \theta a) + sZ)] d\theta.$$

The state-only current costs cancel and the remaining current difference is $a^2 + 4a^4$. This proves (122). Partitioning the first innovation into bins and integrating valid projected gradient hulls with their exact probabilities preserves the same argument. *Q.E.D.*

PROPOSITION 48—Reference stability and computational work: *In the dimension-normalized investment family, for every $d \geq 2$,*

$$\|D_x \mu\|_\infty \leq \kappa := 11/16, \quad \|\nabla c(\cdot, 0)\|_1 \leq 15/2, \quad \|\nabla g\|_1 \leq 10.$$

Consequently, with r stages remaining,

$$\text{Lip}_\infty(J_r^0) \leq \frac{15}{2} \sum_{j=0}^{r-1} (\beta\kappa)^j + 10(\beta\kappa)^r \leq \frac{1920}{91}, \quad \beta = 15/16. \quad (123)$$

For a partition with N leaves, A candidate and covering intervals per leaf, and q innovation bins, the literal cached range-query implementation uses

$$O(TAqN^2d)$$

cell-membership work, in addition to primitive and critic evaluation. Signed reference tubes over a full sweep use $O(ANqdT^2)$ work and can stream their innovation children. Their working state storage is $O(NdT)$ in addition to stored critics, cached bands and retained $O(TAN)$ endpoint records. No q^T recursion or full state-pair table is required.

PROOF: In row j , the only state derivatives of the drift are $1/2 + (1 - x_{j+1})/16$ and $1/8 - x_j/16$, both nonnegative. Their sum is at most $11/16$. The target, cyclic imbalance and shortage terms have gradient ℓ^1 bounds $5/2$, 1 and 4 for current cost; the terminal target bound is 5 . This gives the stated constants. The chain rule yields $L_r \leq 15/2 + \beta\kappa L_{r-1}$ with $L_0 \leq 10$. Its geometric expansion and $\beta\kappa = 165/256$ prove (123).

A queried next-state box is tested against N leaves in d coordinates. There are $O(TANq)$ such root-bin queries. The backward cache therefore has the displayed membership cost. Each gradient tube with r remaining dates visits $O(Nr)$ state rows and $O(Ndr)$ Jacobian coordinates for one first-innovation bin. Summing r across current dates gives $O(ANqdT^2)$. Streaming bins requires storage for one forward tube and its backward gradients, not all future shock combinations. *Q.E.D.*

The stability constant is a bound for this smooth reference, not a dimension-free complexity theorem for solving arbitrary dynamic programs. Resolving a desired policy or all-state accuracy may require the leaf count, network width and action cover to grow. The membership bound explicitly charges a quadratic leaf-count dependence. The executed locally split partition removes a tensor-product representation; it does not prove that a fixed 32-leaf partition resolves every eight-dimensional economic policy.

T.3. What a tighter interval establishes

For fixed candidates, let $C_t(C)$ be the greater of $L_t(C)$ and the minimum of their lower endpoints, including the incumbent's zero. Then

$$U - L = (U - C) + (C - L), \quad U - C \geq 0, \quad C - L \geq 0. \quad (124)$$

The first term measures uncertainty remaining between candidate lower information and the selected upper bound; the second measures the additional continuous-cover allowance. Their maxima need not add to the maximum of their sum. Actor ambiguity, acquisition width, innovation enclosures and arithmetic can enter both terms, so these are not separately estimated causal error components.

More detailed refinement can be recorded without such a causal interpretation. Keep the candidate actions fixed, replace each candidate upper bound by the minimum of its old and refined bounds, and replace the global lower bound by the maximum of its old and refined bounds. Then $U_r \leq U_{r-1}$, $L_r \geq L_{r-1}$, and the certificate improvement at step r is exactly

$$(U_{r-1} - L_{r-1}) - (U_r - L_r) = (U_{r-1} - U_r) + (L_r - L_{r-1}) \geq 0.$$

State subdivision, action-cover subdivision, innovation refinement and higher-precision arithmetic can each supply an independent valid refinement. The present execution records the cache-to-signed-tube refinement and both terms of (124); it does not label an unexecuted six-way refinement experiment as measured evidence.

APPENDIX U: PROSPECTIVE ECONOMIC STOPPING AND LEARNED STRUCTURE

A policy certificate determines whether a proposed change is safe at every state. An economic stopping rule asks when the returned complete policy has attained a specified expected-cost target under the initial law. Neither question is answered by choosing the best recorded policy after all computation has finished.

U.1. *A first-attainment service*

Fix an installed policy π^0 , a target $q \in (0, 1)$, a finite resource ladder, and cumulative sample sizes at each stage. At a stage, construct a critic from primitives, transfer the old policy exactly to any refined partition, and execute a verified simultaneous sweep. Freeze the resulting complete policy before drawing that stage's fresh evaluation stream. At each declared look form a confidence interval for

$$D_{k,q} = \mathbb{E} J_0^{\pi^k} - q \mathbb{E} J_0^{\pi^0}. \quad (125)$$

Return the policy at the first nonpositive upper endpoint. A strictly positive lower endpoint ends the current stage and advances to the next resource stage. An unresolved interval leads to the next declared look. Exhausting the last stage returns the safe incumbent and an explicit budget-exhaustion status. It is not a conclusion that the target is unattainable.

THEOREM 49—Validity at prospective return: *Suppose each accepted sweep satisfies Theorem 45. At each declared stage and look, assume the interval for (125) has conditional failure probability at most $\alpha_{k,l}$ given the information preceding the stage's fresh evaluation stream. If $\sum_{k,l} \alpha_{k,l} \leq \alpha$, then, with probability at least $1 - \alpha$, every successful return satisfies its stated expected-cost target. Every returned incumbent, including one returned after budget exhaustion, is statewise no worse than the original installed policy. No independence between different methods or different looks is required for the union bound. Each within-stage endpoint sample must satisfy its stated conditional sampling assumptions.*

PROOF: Condition on the history preceding each executed stage. Its frozen policy and fresh independent-bin stream satisfy the fixed-policy interval guarantee. Integrating conditional failure probabilities and taking the union over the prespecified potential looks gives a simultaneous event of probability at least $1 - \alpha$. The first-crossing rule selects one already covered interval on this event, so its nonpositive upper endpoint implies (125) is nonpositive. Unexecuted later looks need not be computed. Statewise safety follows by composition of the deterministic sweep inequalities and exact policy transfer, independently of the economic inference event. *Q.E.D.*

The theorem also permits a countable predetermined ladder with summable failure allowances. The experiment uses the finite ladder specified below, not an unimplemented anytime extension. A process clock begins before construction and ends after its durable output. Failed stages, own-policy evaluation and inference are charged to that service. A method reaching a target early must not be charged for hypothetical later stages, and a method failing early must not have its completed work omitted.

For the implemented bounded-sample account, each path is enclosed over its complete forty-bit initial and innovation cells. The last innovation is integrated analytically. The interval in Theorem 39 applies at each declared look, using the known support and outward sample moments. The new family limit is 2,048 intervals with logarithm upper bound 14 and error 1/100; the planned maximum is 1,188 intervals. Earlier looks reuse prefixes within a stage and are covered by the same finite union. Distinct stages use fresh streams. A fixed pseudorandom seed establishes reproducibility under the declared independent-bin model, not mathematical independence of a deterministic sequence.

COROLLARY 50—Economic adoption at return: *Suppose the returned policy has cumulative gain $G = \mathbb{E}J_0^{\pi^0} - \mathbb{E}J_0^{\pi} \in [l, u]$ on the simultaneous event. For every replacement charge $\tau \geq 0$, computation price $\lambda \geq 0$, and recorded complete work $W \geq 0$, its net gain belongs to*

$$[l - \tau - \lambda W, u - \tau - \lambda W].$$

Thus $\tau + \lambda W < l$ certifies profitable adoption and $\tau + \lambda W > u$ rules out profitable adoption. These comparisons hold simultaneously over the entire nonnegative price plane on the original event.

PROOF: Subtract the same known quantity $\tau + \lambda W$ from both endpoints of the gain enclosure. The assertion is pathwise on the simultaneous event, so choosing prices after observing recorded work requires no extra confidence allocation. *Q.E.D.*

This is a decision at return, not a claim that the executed relative-cost stopping rule optimizes all possible prices or future computation schedules. The units of λ must convert recorded work into the model's cost units. In an uncalibrated theoretical economy the price plane is an explicit economic sensitivity analysis, not an estimated welfare magnitude.

U.2. Learning an approximately null direction

Suppose the transition is $\mu_t(x, 0) + Ba + sZ$ with unknown $B \in [0, 1]^d$, known nonlinear drift and known bounded additive action-invariant innovation law. Fit

a direction $\hat{v}$ on training observations. Independent validation observations Y_i estimate B coordinatewise; each coordinate has mean B_j and range length at most w . Let the recorded mean enclosure be $[\underline{m}_j, \overline{m}_j]$ and let δ_{bin} cover numerical bin displacement. Form

$$\mathcal{K}_N = \prod_{j=1}^d \left[\max\{0, \underline{m}_j - r_N\}, \min\{1, \overline{m}_j + r_N\} \right], \quad r_N = w\sqrt{\ell/(2N)} + \delta_{\text{bin}}. \quad (126)$$

Only validation data, the fitted direction and the declared parameter prior enter the gate.

THEOREM 51—Confidence-qualified action-null correction: *For a finite collection of H validation boxes, each with at most d_{max} coordinates, $2Hd_{\text{max}}e^{-\ell} \leq \alpha_B$ gives simultaneous coverage of their true exposures with probability at least $1 - \alpha_B$. On this event, for every fitted direction, amplitude M , and nonnegative feasible action pair,*

$$|(P_B^a - P_B^b)M(\hat{v}^\top y + c)_+| \leq |M| \sup_{B' \in \mathcal{K}_N} |\hat{v}^\top B'| |a - b|. \quad (127)$$

A whole-cell critic-advantage upper bound evaluated uniformly over the same box, plus the discounted right-hand side of (127), is a valid upper bound for the corresponding true advantage with the nuisance removed. Accepting only when that bound is nonpositive preserves safety simultaneously for all tested amplitudes and actions. The simplification is rejected whenever its uncertainty allowance prevents the gate from passing.

PROOF: The bounded-coordinate inequality of [Hoeffding \(1963\)](#) and a union bound give coverage of (126); numerical mean and bin allowances preserve containment. On that event, use a common innovation to couple the two actual kernels. Their next-state difference is $B(a - b)$. The positive-part map is one-Lipschitz, so its projected difference is bounded by $|\hat{v}^\top B| |a - b|$ for every realization. Integration proves (127). The box event controls all linear projections, not merely a preselected direction. Uniformly enclosing the rest of the critic advantage over that same parameter set prevents an unobserved true-parameter oracle from entering the remaining calculation. Subtracting a nuisance difference cannot increase the upper bound by more than its absolute allowance. *Q.E.D.*

The experiment uses $w = 1/2$, $\ell = 14$, 27 validation boxes and at most eight coordinates, giving error at most $1/200$. The direction is fitted on 256 observations and then quantized; the exact stored direction is used in (127). Validation sizes are 512, 4,096 and 32,768. The guarantee concerns an unknown action-exposure vector in a maintained structural model. It does not infer an arbitrary unknown transition law, and coverage failures must remain in the evidence rather than being discarded by reference to the simulator’s known truth.

An empty validation box invalidates the parameter qualification and rejects the simplified certificate. It cannot authorize a policy change.

APPENDIX V: PROSPECTIVE RETURNS IN THE NONLINEAR INVESTMENT ECONOMY

V.1. *Design*

Tasks are $(d, T) = (2, 2), (4, 4), (8, 6)$. Each method/task/target combination is a separate from-primitives service. Targets $q = 9/10$ and $49/50$ require ten- and two-percent expected-cost reductions relative to installed zero investment. The stopping statistic is the signed contrast $J(\pi) - qJ(0)$, not a ratio of plug-in mean estimates. Two resource stages use respectively 256 and 1,024 training states, 32 and 128 non-tensor leaves, and four and eight innovation bins. Actions are downward quantized at $1/4096$ using the lower capacity over the complete leaf. Partition geometry is chosen before the generator and is nested between stages.

Trained ReLU, quadratic and ExtraTrees fitted-value methods use their own future targets. The neural architecture and training are described in Section S. ExtraTrees uses 16 trees, depth at most eight and minimum leaf size three; quadratic FVI fits a ridge-regularized quadratic. The original two-dimensional frontends are reconstructed at $N = 8, 16$, $K = 4$ and $M = 4$. Every method receives nine common robust candidate actions and one fitted proposal, as well as eight intervals covering the full continuous feasible action set. The new algebraic neural action solver is evaluated separately on saved critics; it was not substituted for the frozen experimental construction menu.

After one complete simultaneous sweep at a stage, inference uses cumulative looks of 4,096, 16,384 and 65,536 interval paths. It stops on the first covered target crossing, advances a stage at a strictly positive lower target contrast or exhausted look budget, and never executes a later stage after success. All computation, including failed stages and each method's own inference, is retained. Three isolated sequential processes repeat identical training and inference seeds, with rotated method order, a fixed available CPU core and one numerical-library thread. These measure reproducibility and clock variation, not a population of independently trained networks. Clock frequency is recorded where available but not controlled.

V.2. *The refinement and the complete returned policies*

The first cohort uses the residual/value cache. Its complete 66-process record contains 39 target attainments and 27 budget exhaustions. In the eight-dimensional task it makes no policy change. The separate sensitivity extension was motivated by that completed result and frozen before its own production and fresh inference. It retains the same targets, fitted constructions and common actions, while intersecting the old cache comparison with the signed reference interval of Proposition 47. This chronology matters: the extension is not described as prespecified before observing the primary cohort.

All 66 signed-reference services attain their targets at the first resource stage. This means one complete all-date sweep on 32 leaves, not T exact sweeps and not attainment of the Bellman optimum. For the neural eight-dimensional, six-date controller, the changed leaf counts by date are (26, 24, 24, 20, 14, 0). Its identified cost reduction is generated by nonterminal decisions; it is not a terminal-repair experiment relabeled as a long-horizon result. At the ten-percent target its cumulative gain interval is $[0.21452658, 0.29774911]$, with endpoints rounded outward. The complete expected cost interval is $[1.42162445, 1.47873990]$.

By Corollary 50, the eight-dimensional return supports every replacement-and-computation charge below its lower gain endpoint, in the model's normalized units. A charge above its upper gain endpoint rules out profitable adoption. The interval between them remains unresolved. This translates the returned policy into a full

TABLE V.I

OUTCOMES BEFORE THE SIGNED REFERENCE REFINEMENT

d	T	q	Stages	Paths	Status	Last-pass changes
2	2	49/50	1	16384	Hit	27
2	2	9/10	2	65536	Hit	93
4	4	49/50	2	16384	Hit	90
4	4	9/10	2	16384	Exhausted	90
8	6	49/50	2	65536	Exhausted	0
8	6	9/10	2	4096	Exhausted	0

This primary-cache cohort is not dropped after its limitations become visible. These are the neural rows; the candidate generators return policy-identical final arrays within each primary task and target. Exhausted means the declared finite service did not certify its target, not that improvement is impossible. All own construction, inference and failed-stage work remains charged.

TABLE V.II

ACTUAL COSTS AT THE PROSPECTIVE NEURAL-POLICY RETURN

d	T	q	Paths	Expected cost	Installed minus returned
2	2	49/50	4096	[0.5338, 0.6834]	[0.0622, 0.2408]
2	2	9/10	16384	[0.5950, 0.6503]	[0.1294, 0.1777]
4	4	49/50	4096	[0.8975, 1.0658]	[0.1199, 0.3796]
4	4	9/10	4096	[0.8974, 1.0658]	[0.1177, 0.3775]
8	6	49/50	4096	[1.3427, 1.5356]	[0.0918, 0.4154]
8	6	9/10	16384	[1.4216, 1.4788]	[0.2145, 0.2978]

N is the trained ReLU generator. Every row is a separately executed target service with the signed reference certificate. Endpoints are rounded outward. The interval covers the complete implemented policy under the original continuous law. Timing repetitions have identical policy and inference identities and are not additional economic observations. Other generators and every earlier-stage outcome are retained in the supplement.

economic price sensitivity, rather than choosing a favorable fee after inspecting one estimate.

The directed comparison, all-state optimality account and initial-distribution cost interval continue to answer different questions. At date zero of the eight-

TABLE V.III

WHERE THE SIGNED REFERENCE CERTIFICATE CHANGES DECISIONS

Date	Changed	Base width	Refined width	$U - C$	$C - L$	$U - L$
0	26	14.094	0.610	0.604	0.418	1.021
1	24	12.158	0.594	0.588	0.407	0.994
2	24	10.094	0.569	0.562	0.388	0.950
3	20	7.892	0.530	0.521	0.358	0.879
4	14	5.543	0.466	0.453	0.310	0.762
5	0	1.055	1.055	0.640	0.074	0.713

Eight-dimensional state, six dates, trained ReLU, first stage, 32 non-tensor leaves. Widths are maxima over candidate enclosures; decomposition entries are maxima over leaves, rounded upward. The last row uses analytic terminal differences rather than differentiating an acquired actor. A tighter comparison certificate is not itself a new policy-cost observation.

dimensional neural service, the maximum candidate interval width falls from more than 14 to less than 0.610 after the signed refinement. The full-action greedy allowance remains about 1.021, and the inherited finite-sweep all-state bound is much looser than the observed expected cost gain. The experiment does not claim that 32 leaves or one sweep solve the continuous-action Bellman problem accurately at every state. It does establish the declared prospective expected-cost improvement for the complete implemented policy, together with deterministic statewise incumbent protection.

V.3. Method comparisons and complete work

The trained ReLU proposal changes no deployed cell beyond the nine common actions in either cohort. The only tenth-proposal effect is one compiled-witness cell in the two-dimensional sensitivity task, repeated across targets and timing runs. These identities are checked explicitly rather than inferred from similar cost intervals.

The conventional generators receive the same acquired states, common candidate actions, signed reference information and economic stopping rule. The audit reconstructs the trained proposal and then removes the tenth candidate from the

TABLE V.IV

COMPLETE PROCESS WORK TO THE TEN-PERCENT COST TARGET

d	T	Method	Min (s)	Median (s)	Max (s)	Output KiB	RSS MiB
2	2	W	0.781	0.808	1.345	739	42.0
2	2	E	1.678	1.737	1.740	1100	110.0
2	2	Q	0.695	0.701	0.745	781	42.3
2	2	N	1.389	1.409	1.431	788	50.6
2	2	F	0.807	0.809	0.853	731	42.0
4	4	E	11.515	11.684	11.781	1286	107.3
4	4	Q	5.576	5.582	5.583	489	39.0
4	4	N	9.073	9.087	9.238	512	50.7
8	6	E	31.833	31.886	35.303	2588	112.3
8	6	Q	21.027	21.243	21.882	1350	43.8
8	6	N	26.714	26.771	30.501	1400	51.8

N: trained ReLU; Q: quadratic FVI; E: ExtraTrees FVI; W: compiled witness; F: tensor FVI. Three isolated sequential processes use identical seeds within a method. Each clock includes its own imports, construction, verification, inference and durable output. Compare methods within a task runner, not across runners. Affinity and numerical-library threads are fixed; clock frequency is not controlled. Output excludes prior-stage archives outside the service; RSS is process peak memory. Both targets and all repetitions appear in the supplement.

recorded gate to determine whether that proposal changes the selected action. This is a deterministic candidate-attribution diagnostic, not a fresh policy-cost experiment. Exact deployed-policy identities are checked across generators and across timing repetitions.

In the four- and eight-dimensional sensitivity tasks, the generators return policy-identical arrays within each target. In the two-dimensional task the compiled-witness return differs slightly, while the trained ReLU, quadratic, ExtraTrees and tensor-FVI returns coincide. Thus the policy improvement is not evidence that the neural generator dominates these alternatives. At the ten-percent target, quadratic FVI has the lowest recorded median complete-process work in each task. The neural backend supplies a genuinely trained nonlinear continuation and analytic

integration, but those facilities must not be confused with an identified economic ranking.

The clocks include imports, primitive construction, critic fitting, acquired-policy preparation, verification, own inference, serialization and durable outputs. Method comparisons are within each task's controlled runner; clocks from different task runners are not a hardware-controlled scaling regression.

V.4. Neural action witnesses and learned null directions

TABLE V.V

EXACT LATTICE SEARCH ON SAVED FITTED NEURAL CRITICS

d	T	Full lattice	Candidates	Old index	Exact index	Critic decrease
2	2	896	10	279	303	0.0000800
4	4	768	4	191	169	0.0000289
8	6	640	6	559	555	0.0000012

One declared cell center and its saved next-date critic per task. Rational isolation and comparison agree with exhaustive rational lattice minimization. The reported decrease is in that fixed approximate Bellman objective, not actual economic policy cost. These solver regressions neither retrain a critic nor produce new policy-cost observations.

The algebraic action solver gives exact fitted-objective minimizers on the full deployed lattice in its retained regressions, using only a small subset of lattice indices. This is a constructive neural-search result. Its objective decreases concern the saved approximate continuation at a specified cell center, not the expected cost of a new deployed controller. All historical inference remains attached to the policies that were actually evaluated.

The sample-split study fits nine directions across dimensions two, four and eight and three independent direction-fitting seeds. Three validation sizes give 27 confidence boxes; three amplitudes on 64 acquired leaves give 81 cases and 5,184 cell-case decisions. All 27 recorded boxes contain the simulator exposure. The corrected gate makes 165 safe changes. No strictly harmful plug-in decision is identified in this learned-direction catalogue. That zero count is retained; it is not

TABLE V.VI

VALIDATION-QUALIFIED LEARNED DIRECTIONS

d	Validation rows	Projection-bound range	Corrected changes	Harmful plug-in
2	512	[0.0660, 0.0670]	36	0
2	4096	[0.0241, 0.0271]	54	0
2	32768	[0.0121, 0.0141]	75	0
4	512	[0.0608, 0.0636]	0	0
4	4096	[0.0217, 0.0232]	0	0
4	32768	[0.0081, 0.0102]	0	0
8	512	[0.0586, 0.0622]	0	0
8	4096	[0.0207, 0.0236]	0	0
8	32768	[0.0078, 0.0101]	0	0

Each row aggregates three independently fitted directions and three amplitudes on 64 leaves per case.

The projection range reports the minimum and maximum robust upper allowance across fitted directions, not a confidence interval for a population mean. All 27 exposure boxes cover the simulator exposure in this record. No strictly harmful plug-in choice is identified here; the earlier known-perturbation catalogue is a different diagnostic and remains unchanged.

replaced by the large harmful count from the earlier, different known-perturbation experiment.

The primary policy-cost, sensitivity policy-cost and learned-exposure families have errors 1/100, 1/100 and 1/200, respectively. Their joint union error is at most 1/40. This addition combines covered statements; it does not make identical policies, timing repetitions, amplitudes or reused prefixes independent observations. The learned-direction conclusion remains conditional on the known drift and maintained additive action-invariant innovation structure.

V.5. *Relation to the original evidence*

The earlier witness, FVI, acquisition, common-repair, exact-null and adaptive-grid results remain in the complete development article and its supplement. In particular, earlier adverse witness-cost comparisons are not erased by the improvement of a newly installed zero controller. The new first-attainment services replace retrospective catalogue work as the primary work-to-target endpoint for this experiment; the

historical catalogue remains a valid account of a different executed procedure. The full source and result archive preserves both and states which economic policy and law each interval actually covers.

APPENDIX W: FROM FITTED ACTION WITNESSES TO ECONOMIC POLICY ACCURACY

An exact action witness has a definite but limited meaning: it removes optimization error for the stored continuation at a reported state. The next result identifies the further obligations needed to turn that statement into an error bound for the implemented economic policy. It also isolates a comparison that can be made without attributing the common verifier's contribution to a particular representation.

Fix a complete reference policy π , a date t , and an acquired cell C with representative x_C . Write

$$Q^h(x, a) = c_t(x, a) + \beta_t P_t^a h_{t+1}(x), \quad Q^\pi(x, a) = c_t(x, a) + \beta_t P_t^a J_{t+1}^\pi(x).$$

The superscripts refer to the continuation used, not to different economic primitives. Let G_C be a finite action lattice feasible throughout C , containing the reference action. The reference is assumed implementable on the cell; all selected actions are measurable under the same observation contract. For $g_C \in G_C$, define its reported-state optimization loss by

$$\kappa_C = Q^h(x_C, g_C) - \min_{b \in G_C} Q^h(x_C, b). \quad (128)$$

This is a nonnegative, finite quantity. Theorem 44 makes it exactly zero for its returned lattice witness, regardless of the training loss or the signs of the learned output weights.

For all $x \in C$ and $a, b \in G_C$, suppose the following nonnegative bounds are available:

$$|[Q^h(x, a) - Q^h(x, b)] - [Q^h(x_C, a) - Q^h(x_C, b)]| \leq \eta_C, \quad (129)$$

$$|[Q^\pi - Q^h](x, a) - [Q^\pi - Q^h](x, b)| \leq \chi_C. \quad (130)$$

Suppose also that every $a \in A_t(x)$ has a $b \in G_C$ with

$$Q^\pi(x, b) \leq Q^\pi(x, a) + \delta_C. \quad (131)$$

These are distinct requirements. The first concerns acquisition within a cell, the second the action-relevant continuation error, and the third the difference between the robust deployed lattice and the original state-dependent continuous feasible set.

Let $[l_C(g_C), u_C(g_C)]$ enclose the true reference advantage $Q^\pi(x, g_C) - J_t^\pi(x)$ throughout C . Let w_C be a verified upper bound on this interval's width. The adoption rule of Theorem 45 is applied to a candidate set containing g_C and the reference, with selected upper endpoint $U_C \leq \min\{0, u_C(g_C)\}$. Denote its complete-action lower endpoint by L_C .

THEOREM 52—Fitted witness transfer: *Under (128)–(131) and the valid whole-cell endpoints above, the selected action a_C satisfies, at every $x \in C$,*

$$Q^\pi(x, a_C) \leq J_t^\pi(x), \quad (132)$$

$$Q^\pi(x, a_C) - \inf_{a \in A_t(x)} Q^\pi(x, a) \leq \min\{U_C - L_C, \rho_C\}, \quad (133)$$

$$\rho_C := \kappa_C + \eta_C + \chi_C + \delta_C + w_C.$$

Consequently the directed finite-sweep recursion remains valid with

$$\varepsilon_t^k = \sup_C \min\{U_t^k(C) - L_t^k(C), \rho_t^k(C)\}.$$

In particular, exact neural action search removes κ_C from this bound, not the other four terms. The coefficient one on the sum ρ_C is sharp without additional restrictions.

PROOF: Fix $x \in C$, an arbitrary feasible action a , and its covering lattice point b . Equations (128)–(130) imply

$$Q^\pi(x, g_C) - Q^\pi(x, b) \leq \kappa_C + \eta_C + \chi_C.$$

Adding (131) gives a bound against $Q^\pi(x, a)$ with the additional term δ_C . The endpoint rule gives

$$Q^\pi(x, a_C) - J_t^\pi(x) \leq U_C \leq u_C(g_C) \leq Q^\pi(x, g_C) - J_t^\pi(x) + w_C.$$

Combining these inequalities and taking the infimum over a proves the ρ_C bound. The full-action lower cover gives the independent $U_C - L_C$ bound. Also $U_C \leq 0$, proving (132). Backward induction from the common terminal value gives safety for simultaneous adoption of all dates. The proof of Theorem 45 then applies with the smaller valid local error; no optimizer attainment for the continuous infimum is needed.

For sharpness, take one decision date, one observation cell containing one state, exact fitted values, and two actions with costs zero and $\delta > 0$. Let the incumbent and only proposed action be the action of cost δ , but let the lattice and full lower cover contain both actions. Then $\kappa_C = \delta$, the other four terms are zero, and the retained policy has loss exactly δ . Thus the coefficient multiplying the unrestricted sum cannot be smaller than one. *Q.E.D.*

The width term records the conservatism of certifying the proposed action and then choosing among upper endpoints. It is necessary: the action with the smallest upper bound need not have the smallest true advantage. A separately verified bound on $u_C(g_C) - [Q^\pi(x, g_C) - J_t^\pi(x)]$ can replace the full interval width. Likewise, sharper contrast bounds may replace (129) and (130). Their validity must be established before they are used in the policy account.

W.1. A primitive continuous-action allowance

The continuous-action condition is explicit in the investment family. Let $\bar{a}(x)$ be its capacity, Δ the action spacing, and

$$m_C \Delta = \Delta \left\lceil \frac{\inf_{x \in C} \bar{a}(x)}{\Delta} \right\rceil, \quad G_C = \{0, \Delta, \dots, m_C \Delta\}.$$

If $Q^\pi(x, \cdot)$ is uniformly L_a -Lipschitz, then

$$\delta_C = L_a \{\text{osc}_C \bar{a} + \Delta\} \tag{134}$$

is admissible. For the smooth installed zero policy in the original normalized family, one may take

$$L_a = \frac{3}{4} + \frac{900}{91}. \quad (135)$$

PROOF: For a feasible a , use the smaller of its downward lattice rounding and $m_C \Delta$. Its distance from a is at most $\text{osc}_C \bar{a} + \Delta$, proving (134). The derivative of $a^2 + 4a^4$ has absolute value at most $3/4$ on $[0, 1/4]$. Proposition 48 bounds the zero-policy value's Lipschitz constant in the maximum norm by $1920/91$. The action exposure has maximum norm $1/2$ and the discount is $15/16$. Therefore the continuation's action Lipschitz constant is at most $900/91$, proving (135). *Q.E.D.*

The primitive allowance is valid but can be conservative. The executed policy bound uses the direct continuous-action lower cover, not a claim that the primitive sum is always tighter. The transfer theorem permits both accounts and shows exactly where an improved action search can enter the original Bellman-accuracy objective.

W.2. Nested proposals and identifiable verification gains

Consider two candidate sets $M_0 \subseteq M_1$ for the same reference, acquired cells, complete-action lower cover and valid upper-endpoint procedure. The endpoints of shared actions must be identical; this condition prevents changes in a fitted residual cache from being confused with changes in the proposed actions. Both sets retain the reference with exact upper advantage zero.

COROLLARY 53—Nested-candidate certificate: *Let $U_i(C)$ be the minimum candidate upper endpoint for M_i . Then*

$$U_1(C) \leq U_0(C) \leq 0, \quad 0 \leq U_1(C) - L(C) \leq U_0(C) - L(C).$$

The improvement is strict precisely when an added candidate has upper endpoint strictly below $U_0(C)$. Both returned complete policies are statewise safe relative to the reference. These inequalities alone do not order their actual expected costs.

PROOF: Taking a minimum over a larger finite set proves the first inequalities. The valid full-action lower bound is no greater than every feasible action's advantage and hence no greater than either selected upper bound. This proves the nonnegative gap inequalities and the strictness criterion. Safety follows from Theorem 45. To see the last distinction, consider a single state and date, an incumbent of cost 2, an old candidate of cost 1 with advantage enclosure $[-1, -1/2]$, and a new candidate of cost $5/4$ with exact advantage $-3/4$. The new candidate has the smaller upper endpoint and is adopted, although its cost is higher than that of the old candidate. Both improve the incumbent. A lower cover of -1 makes both certificates valid. Q.E.D.

In the new experiment the ReLU-exact mode retains every ReLU-menu action and adds the algebraic witness. Both use the same fitted coefficients at a given construction attempt and a critic-independent smooth-reference verifier. Thus the corollary is an observed, testable comparison rather than a representation claim. Exact center-objective regret, stricter upper endpoints, changed deployed actions and independent expected-cost differences are reported separately.

W.3. *A paid decision between returned policies*

Let $[l, u]$ cover $J(\pi_1) - J(\pi_0)$ on a simultaneous event after both policies have been selected. If adopting π_1 rather than π_0 incurs an additional replacement charge τ and measured computational work difference ΔW , its net cost difference at computation price $\lambda \geq 0$ lies in

$$[l + \tau + \lambda\Delta W, u + \tau + \lambda\Delta W]. \quad (136)$$

A negative upper endpoint supports adoption; a positive lower endpoint rules it out on the same event. Between those regions the comparison is unresolved. This affine transformation is valid for every declared nonnegative price on that event. It does not turn a normalized benchmark into an estimated installation fee. It links a particular fitted-action computation to the same economic policy and resource question as the original NBO construction.

W.4. Explicit moduli for the trained neural representation

The transfer bound need not leave acquisition error as an unspecified constant. In the original investment family, write the learned continuation as

$$h(y) = v^\top y + c + \sum_{j=1}^m u_j \sigma(w_j^\top y + b_j).$$

Let $s_j = w_j^\top \gamma$, let R_j be the feature's scalar-innovation radius from Lemma 43, and set $D_C = \text{diam } G_C$ and $r_C = \sup_{x \in C} \|x - x_C\|_\infty$. Define

$$d_j(D) = \begin{cases} 0, & s_j = 0, \\ \min\{1, |s_j|D/(2R_j)\}, & s_j \neq 0, R_j > 0, \\ 1, & s_j \neq 0, R_j = 0. \end{cases}$$

All these quantities are determined by the stored coefficients and the declared primitive law.

PROPOSITION 54—Action-sensitive neural moduli: *For the original separable current action cost and drift with maximum-norm Lipschitz constant $11/16$, the acquisition allowance*

$$\eta_C = \beta_t \frac{11}{16} r_C \sum_{j=1}^m |u_j| \|w_j\|_1 d_j(D_C) \quad (137)$$

satisfies (129). In particular, the affine part and every exactly action-null ridge contribute zero to this allowance. For the installed zero reference, a valid continuation-error allowance is

$$\chi_C = \beta_t D_C \left\{ \frac{960}{91} + |v^\top \gamma| + \sum_{j=1}^m |u_j s_j| \right\}. \quad (138)$$

Either allowance may be replaced by the minimum of itself and any independently valid sharper enclosure. Computing the displayed bounds requires $O(md)$ arithmetic, with outward rounding when evaluated numerically.

PROOF: Write the integrated ridge as $\Phi_R(z) = \mathbb{E}\sigma(z + RZ)$ for Z uniform on $[-1, 1]$. For $R > 0$, its derivative is zero, then affine with slope $1/(2R)$, then one.

Hence the function

$$z \mapsto \Phi_R(z + s_j a) - \Phi_R(z + s_j b)$$

is Lipschitz with constant at most $\min\{1, |s_j(a - b)|/(2R)\}$. For $R = 0$ the difference of the two hinge derivatives is bounded in absolute value by one almost everywhere, giving the same conclusion with the declared convention. For $s_j = 0$ the function is identically zero. The preactivation mean varies by at most $\|w_j\|_1(11/16)\|x - x_C\|_\infty$. The affine continuation contrast is $v^\top \gamma(a - b)$, which does not depend on x . The state-only current cost cancels, and the current action-cost difference also does not depend on x . Summing the ridge contributions and discounting proves (137).

Under a common innovation, the zero-policy continuation contrast is bounded by $(1920/91)\|\gamma\|_\infty|a - b| = (960/91)|a - b|$. The one-Lipschitz hinge gives a fitted continuation contrast bound $\{|v^\top \gamma| + \sum_j |u_j s_j|\}|a - b|$. The triangle inequality and discount give (138). Independently valid upper bounds may be minimized because each bounds the same nonnegative quantity. The arithmetic count follows from computing the feature projections and norms. *Q.E.D.*

These formulas connect feature geometry to acquired-state accuracy without assuming that training discovers a null direction. They complement the confidence-qualified learned-direction result: estimated primitive coefficients must still range over the corresponding confidence set. In the present known-law action-search experiment, the direct signed reference enclosure determines adoption and the full-action lower cover determines the displayed Bellman error. The primitive moduli supply an additional certified account, not a retroactive claim that their conservative sum was the stopping criterion.

APPENDIX X: EXECUTING FITTED NEURAL ACTION WITNESSES

The preceding study establishes a prospective economic return and a non-tensor verification procedure. Its candidate-attribution audit also identifies a necessary next comparison: the common action menu often determines the same policy regardless of the fitted generator. The algebraic solver in Theorem 44 was tested there as a solver, not used to create the policies whose costs were reported. We now

execute that solver in new prospective services and evaluate the resulting policies on new streams. The original investment laws, Bellman optimum and NBO research object are unchanged.

X.1. *A controlled action-search comparison*

The new protocol fixes the three tasks $(d, T) = (2, 2), (4, 4), (8, 6)$ before production. It compares five modes. Common-only uses the nine robust actions without fitting a critic. ReLU-menu uses a genuinely fitted 32-feature continuation and adds its 17-menu proposal. ReLU-exact retains every one of those candidates and adds the exact fitted-critic lattice minimizer. Quadratic-exact supplies the analogous exact action search for its fitted quadratic continuation. ExtraTrees-menu supplies its own future tree-based proposal. All modes receive the same exact terminal-action search, so an analytic terminal improvement is not assigned to neural representation.

The nonterminal verification rule is identical across generators: signed sensitivity of the installed zero-investment policy from Proposition 47. Its state boxes, innovation bins and primitive arithmetic do not depend on the fitted critic. Shared candidate actions therefore have identical upper endpoints. This design implements the premise of Corollary 53, rather than changing the candidate generator and its error certificate simultaneously.

Each mode is reconstructed independently for each task, target and one of three distinct training seeds. ReLU-menu and ReLU-exact use the same seed within an attempt, but neither loads the other's fit free of charge. Their fitted-parameter identities are checked. Different seeds produce different fitted neural objects. The three seeds are not a population of economic tasks, and their timing variation is not described as a controlled timing-only repetition. The preceding fixed-seed timing repetitions remain a separate experiment.

X.2. *Prospective economic targets*

The two targets require ten- and twenty-percent expected-cost reductions from the installed zero policy. Each target service starts from primitives. It has two declared construction attempts, with training rows 256, 1024, acquired leaves 32, 64,

and innovation bins 4, 8. The same locally split non-tensor observation geometry and downward action spacing $1/4096$ apply to every mode.

At each attempt a complete candidate is certified against the installed policy and frozen before inference. Intermediate attempts are not deployed. In particular, refinement does not apply the smooth-reference derivative formula to an already changed discontinuous actor. Every attempt is safe relative to the installed zero policy; no ordering between the policies of different attempts is assumed.

The own-service stopping estimand is $J(\pi) - qJ(0)$, with $q = 9/10$ or $4/5$, rather than a ratio of two estimated means. Its cumulative inference looks have 4096, 16384 and 65536 paths. A nonpositive upper endpoint returns immediately. A positive lower endpoint advances the construction attempt; an unresolved interval advances the declared look. After the final unsuccessful attempt the service returns a budget-exhausted safe policy, not a claim that the target is impossible. Every attempted fit, exact search, verification, inference and durable output belongs to the service's surrounding process clock.

There are 90 separately executed target services. After each set of five modes has returned, a separate 65,536-path stream compares ReLU-exact with each other returned policy. This stream is independent of construction and stopping information under the declared independent-bin model. It evaluates the original continuous initial and innovation laws using forty-bit bin enclosures and analytic final integration. Negative path differences are retained. Full policy identities receive exact zero contrasts; other overlapping intervals are not treated as identities.

The new finite confidence family has error at most $1/100$. Its maximum possible interval count is

$$90 \cdot 2 \cdot 3 \cdot 3 + 18 \cdot 4 = 1692 < 2048,$$

covering own cost, target contrast and gain at every possible look, together with the independent validation contrasts. The retained logarithm allowance is checked by exact rational arithmetic. The first-attainment argument of Theorem 49 applies conditionally at every fresh stage. This finite-look account is not an undeclared confidence sequence valid at arbitrary future sample sizes.

The new cohort executes 90 separate prospective target services. Of these, 75 attain their declared target and 15 return a budget-exhausted safe candidate. The audit regenerates 78 distinct fitted-model certificate and action-search calculations and checks 330 stopping looks. In the trained ReLU-exact services, 1354 nonterminal cell-attempt records have strictly positive menu optimization loss, while the additional algebraic candidate changes 9 deployed cell-attempt decisions. These counts refer to the declared service records, not independent samples. Among the 72 fresh paired comparisons with ReLU-exact, 0 identify lower ReLU-exact cost, 0 identify higher cost, 42 are exact policy identities and 30 remain unresolved.

TABLE X.I
PROSPECTIVE TWENTY-PERCENT TARGET SERVICES

d/T	Mode	Reached	Cost upper	Median (s)	Range (s)
2/2	C	3/3	0.61949	3.776	3.763–3.793
2/2	N	3/3	0.61949	4.511	4.497–4.536
2/2	X	3/3	0.61948	5.216	5.179–5.238
2/2	Q	3/3	0.61949	4.083	4.043–4.107
2/2	E	3/3	0.61962	5.185	5.168–5.202
4/4	C	3/3	0.97342	12.881	12.439–12.983
4/4	N	3/3	0.97342	15.077	15.004–15.240
4/4	X	3/3	0.97342	17.951	17.786–18.369
4/4	Q	3/3	0.97342	13.780	13.475–13.991
4/4	E	3/3	0.97342	15.984	15.872–16.513
8/6	C	0/3	1.36544	35.153	35.049–35.571
8/6	N	0/3	1.36544	39.239	39.190–39.458
8/6	X	0/3	1.36544	44.189	44.183–45.418
8/6	Q	0/3	1.36544	38.682	38.257–38.954
8/6	E	0/3	1.36544	40.766	40.704–41.016

C uses common candidates without fitting. N uses the trained ReLU menu; X adds its exact algebraic witness; Q uses quadratic exact search; E uses ExtraTrees menu proposals. All modes receive the same exact terminal action. Three different training seeds are reconstructed separately. Cost upper is the largest returned-policy upper endpoint across those seeds, not a confidence bound for a population mean. Complete process times include every own failed attempt and inference look.

X.3. What the neural action computation changes

Table X.II separates positive center-objective optimization loss, strict improvement of a whole-cell upper endpoint and an actual change of the deployed action. These events are not interchangeable. The center optimization loss is evaluated exactly from the same stored critic; it is zero for the algebraic minimizer. Removing only the added algebraic candidate reconstructs the within-attempt menu policy with the same verifier. Thus every reported stricter endpoint is a genuine nested-menu improvement.

TABLE X.II
OBJECTIVE SEARCH AND WHOLE-CELL DEPLOYMENT ATTRIBUTION

d	Mode	$\kappa > 0$	Changed	Stricter	Roots	Values
2	X	163	7	7	337	3632
2	Q	145	0	0	321	2276
4	X	451	2	2	626	10849
4	Q	224	0	0	409	3965
8	X	740	0	0	840	14603
8	Q	37	0	0	127	4866

Counts are cell-attempt records across all own target services, including repeated reconstruction at the other target; they are not independent economic observations. Positive optimization loss refers to the nonterminal menu witness compared with the exact lattice witness. Changed and stricter counts remove only the extra algebraic proposal while retaining the shared verifier. Root and value counts include the common terminal solver, which is not credited as a neural-specific contribution.

The new execution supplies the previously missing connection between the trained representation's exact action witness and a deployed policy. It also measures the extent of that connection. A large number of positive optimization losses can coexist with few additional deployed actions because the common safe menu already determines most selected upper bounds. Even a strictly improved upper bound does not prove that the resulting policy costs less than its menu counterpart. The independent paired contrasts in Table X.III, and not the fitted objective or

TABLE X.III

FRESH RETURNED-POLICY COST COMPARISONS

<i>d</i>	X versus	X lower	X higher	Identity	Unresolved
2	C	0	0	2	4
2	N	0	0	2	4
2	Q	0	0	2	4
2	E	0	0	2	4
4	C	0	0	3	3
4	N	0	0	4	2
4	Q	0	0	3	3
4	E	0	0	2	4
8	C	0	0	6	0
8	N	0	0	6	0
8	Q	0	0	6	0
8	E	0	0	4	2

Each row covers both targets and all three training seeds. A negative paired difference favors X. Identity means equal complete policy arrays on the same observation partition, not merely overlapping intervals. Other unresolved intervals are not equivalence findings. The fresh comparison stream is independent of the corresponding stopping information under the sampling model.

interval width, determine that economic comparison. Every individual interval is reported in the supplement.

The common-only control pays no fitting cost and is a substantive comparator. Quadratic-exact receives the same algebraic-search opportunity, while ExtraTrees retains its declared fitted-value construction. This prevents the exact solver's availability from being conflated with a special verification privilege. The experiment compares the complete returned procedures and identifies their components; it does not presume that fitting a neural continuation must beat a common-action or conventional-value alternative.

X.4. *Work, uniform accuracy and adoption*

The main table reports complete own-target process work at a matched economic target. It includes failed intermediate attempts and all inference actually needed to return. The supplementary work ledger separately records fitting, root isolation, candidate evaluation, signed verification, path evaluation, serialized bytes and resident memory. Different tasks run on different task runners; within each task processes are sequential, use one available core and one numerical-library thread, and rotate mode order. Available governor information is recorded, but clock frequency is not controlled.

The fresh paired validation is additional study work. Its recorded timer ends after endpoint evaluation and serialization, before the final comparison JSON is committed. It is therefore reported as a partial validation-service clock, not a complete process clock, and is not concealed inside the own-target time. The protocol's intended complete comparison-process timing was not implemented; this timing boundary does not change the prospective stopping rule or its fully recorded own-service work. The raw timing records are retained without retrospective adjustment.

Every returned policy also receives a one-sweep all-state loss account. At date t it includes the old-reference continuation term $\beta_t E_{t+1}^0$ as well as the directed greedy allowance. The tables use a valid support bound for that reference error and cap the result by the original total-cost support. They do not substitute the local $U - L$ gap for loss against the Bellman optimum. A successful expected-cost target under the initial law is thus not mislabeled as exact Bellman convergence.

The gain from replacing installed zero investment remains an actual economic gain, subject to the stated confidence event and normalized primitives. Choosing between two returned methods requires the additional incremental-cost calculation (136). It charges the measured work difference and any replacement fee, and uses the fresh paired policy-cost interval. The complete fee-computation-price regions follow from that interval without choosing a favorable fee after seeing the outcome. The experiment is a theoretical investment exercise, not an externally calibrated welfare estimate.

APPENDIX Y: LOSSLESS SEARCH FOR TRAINED NEURAL ACTION WITNESSES

The fitted-witness transfer result separates an optimization error from the errors in acquisition, continuation evaluation and full-action verification. An exact solver eliminates the first error, but exactness alone does not make the solver economical. We now exploit the learned features to reduce the work of returning precisely the same action witness. The continuation, economic primitives and original Bellman-accuracy target remain unchanged.

Y.1. *Exact reduction of integrated neural features*

Fix a state and a stored affine-ReLU continuation. Interpret its finite binary coefficients as rational numbers. With a scalar uniform innovation, the expectation of a ridge is

$$\phi_r(z) = \begin{cases} ((z+r)_+^2 - (z-r)_+^2)/(4r), & r > 0, \\ z_+, & r = 0. \end{cases}$$

The radius r is the absolute projection of the innovation support on the ridge direction. After removing a constant independent of the action, the fitted objective in the investment model takes the form

$$F(a) = q_1 a + q_2 a^2 + 4a^4 + \sum_{j=1}^m \omega_j \phi_{r_j}(z_j + s_j a). \quad (139)$$

All coefficients are rational. Negative fitted weights are permitted. No convexity of the fitted continuation is assumed. The robust feasible action lattice is $\mathcal{L} = \{k\delta : 0 \leq k \leq K\}$, where δ and $K\delta$ are fixed by the acquisition contract.

A feature can be simplified before any lattice search. Let $[v_j^-, v_j^+]$ be the exact range of $z_j + s_j a$ on $[0, K\delta]$. If $v_j^+ \leq -r_j$, its contribution vanishes. If $v_j^- \geq r_j$, it is affine. If $r_j > 0$ and $[v_j^-, v_j^+] \subseteq [-r_j, r_j]$, it is quadratic. Absorb these contributions into a rational constant c_0 and the coefficients of a and a^2 . Only features crossing an activation-region boundary remain in the ridge sum. Consequently

$$F(a) = c_0 + \tilde{q}_1 a + \tilde{q}_2 a^2 + 4a^4 + \sum_{j \in \mathcal{C}} \omega_j \phi_{r_j}(z_j + s_j a) \quad (0 \leq a \leq K\delta). \quad (140)$$

This is an identity, not pruning by a small fitted weight. A feature with a large weight may be eliminated exactly, whereas a small crossing feature remains unless its whole contribution is enclosed.

Y.2. An interval screen with an exact return

Write $f_k = F(k\delta) - c_0$. Compute outward enclosures $[\ell_k, u_k]$ of all f_k and put

$$b = \min_{0 \leq k \leq K} u_k, \quad S = \{k : \ell_k \leq b\}. \quad (141)$$

Compare the original objective exactly on S , choosing the smallest index on a tie. The weak inequality in the definition of S is essential. Floating-point endpoints reject a point only when they prove that it cannot minimize the original objective; they never supply an approximate answer in place of the exact comparison.

THEOREM 55—Lossless neural action search: *Suppose the intervals in (141) are finite, ordered and valid. Then S contains every exact lattice minimizer. Exact rational comparison on S , with the same smallest-index convention, returns exactly the witness of exhaustive lattice minimization and of the terminating algebraic solver.*

Let $\eta = \max_k (u_k - \ell_k)$. Every survivor satisfies

$$F(k\delta) - \min_{a \in \mathcal{L}} F(a) \leq 2\eta. \quad (142)$$

In particular, a unique minimizer is the only survivor whenever its gap to the second-best lattice point exceeds 2η .

More generally, let $\mathcal{M} \subseteq \mathcal{L}$ be the exact minimizer set. If, for some $\mu > 0$,

$$F(a) - \min_{\mathcal{L}} F \geq \frac{\mu}{2} \text{dist}(a, \mathcal{M})^2 \quad (a \in \mathcal{L}),$$

then

$$|S| \leq \min \left\{ K + 1, |\mathcal{M}| \left(1 + 2 \left\lfloor \frac{2\sqrt{\eta/\mu}}{\delta} \right\rfloor \right) \right\}.$$

The growth condition is an additional assumption, not a property inferred from a fitted ReLU representation.

PROOF: Let k_* minimize f_k . For every j , $f_{k_*} \leq f_j \leq u_j$, hence $\ell_{k_*} \leq f_{k_*} \leq b$ and $k_* \in S$. This argument retains all minimizers, including ties. Exact comparison with the same tie rule therefore returns the same index. The algebraic solver enumerates lattice neighbors of every activation boundary and isolated stationary point on each polynomial piece; its established exhaustive-search equivalence then gives the same witness.

For a survivor k , validity and the width bound imply

$$f_k \leq \ell_k + \eta \leq b + \eta \leq u_{k_*} + \eta \leq f_{k_*} + 2\eta.$$

This proves (142) and the unique-gap statement. Under the additional growth condition a survivor is within $2\sqrt{\eta/\mu}$ of some element of $\mathcal{M}$. There are at most $1 + 2\lfloor 2\sqrt{\eta/\mu}/\delta \rfloor$ lattice points within that distance of one lattice minimizer. Taking a union, and then the total lattice cardinality, proves the last bound. *Q.E.D.*

The generic screening implication uses only interval order. The implementation contribution here is the exact reduction of integrated trained ridges, outward evaluation of their reduced objective, and exact recovery of the original NBO witness. On a failed arithmetic check or a lattice exceeding the declared screening limit, the implementation calls the original terminating algebraic solver. A fallback is charged as work and cannot be recorded as an approximate successful screen. A wide but valid enclosure may retain the entire lattice; the theorem promises correctness in that case, not a speed gain.

PROPOSITION 56—Arithmetic work and working storage: *Let $N = K + 1$, let $m_c = |C|$, and let $s = |S|$. Forming the state-specific neural coefficients and performing exact region reduction requires $O(dm)$ arithmetic operations for an m -ridge continuation in dimension d . The screen uses $O(N(1 + m_c))$ fixed-precision interval operations, followed by $O(sm)$ rational operations to compare survivors using the original objective. The vectorized implementation uses $O(N(1 + m_c))$ interval working storage, in addition to the fitted coefficients, and returns $O(s)$ retained indices. These are arithmetic counts; the cost of rational arithmetic also depends on operand bit lengths.*

PROOF: Each ridge requires state and innovation projections and a constant number of exact range tests. Each remaining ridge contributes one scalar interval expression at each lattice point. Pairwise outward summation of the ridge contributions is linear in their number. The exact return evaluates the original m -ridge objective at each survivor. The implementation retains ridge-by-lattice batches, giving the stated storage rather than a streaming bound. Integer index comparison resolves equal rational values. None of these counts eliminates the dependence on d , lattice resolution or fitted width. *Q.E.D.*

Terminal optimization is unchanged in the experiment. The screened intervention concerns only nonterminal trained-ReLU continuations. In particular, time saved by an analytic terminal formula is not credited to learned-feature reduction.

Y.3. *Preservation of a prospective economic return*

A solver is used inside a procedure, not in isolation. Consider a deterministic prospective controller conditional on its construction data and random streams. At each attempted resolution it constructs an acquired policy, verifies whole-cell advantages and continuous-action lower covers, and then runs the declared sequence of cost-inference looks. It returns at the first certified economic target, advances resolution when required by its rule, or reports budget exhaustion. Suppose its branching decisions depend on these mathematical outputs, not on elapsed time or solver-specific diagnostic counters.

THEOREM 57—Prospective transcript invariance: *Run the same controller with the algebraic solver and the lossless screened solver. Suppose corresponding construction steps produce identical stored continuations, acquisition partitions and other candidate actions, and use identical inference streams. Then both controllers return the same policy, have identical whole-cell verification endpoints and cost-inference endpoints, and stop at the same resolution and inference look with the same status. Their exact fitted-optimization loss is the same, namely zero on the declared action lattice. Their all-state Bellman-error certificate is unchanged.*

The returned policies have exactly equal expected economic cost under every admissible initial distribution. This equality does not require statistical inference. It does not assert equality of computation cost.

PROOF: Condition on the common construction inputs and streams. Theorem 55, or the exact fallback, gives identical proposed action indices at every leaf and date. The remaining candidate menu and its ordering are the same. The same verifier therefore evaluates the same intervals and resolves the same incumbent-preserving comparisons, producing an identical acquired policy and continuous-action error account. Evaluating that policy on the same interval paths gives identical path endpoints, hence identical inference bounds and the same next controller decision. Induction over inference looks and attempted resolutions proves the entire transcript assertion. No additional attempt occurs after a successful return. The inherited fitted-witness transfer and directed policy-improvement theorems apply to the identical witnesses and certificates. Finally, identical policies have the same controlled law and expected cost for every initial law. *Q.E.D.*

The theorem is conditional on actual parameter identities, not merely on a shared training seed. The empirical audit checks those identities. It does not apply to a wall-clock deadline rule: a faster solver can reach a different stage before a real-time deadline. Nor does a target that is not attained under the declared finite budgets become impossible in the underlying economic model.

COROLLARY 58—Computation prices at an unchanged policy: *Suppose both implementations return π and face specified computation charges C_A, C_S . At a nonnegative price λ , their total expected policy-plus-computation cost difference is*

$$\{J(\pi) + \lambda C_S\} - \{J(\pi) + \lambda C_A\} = \lambda(C_S - C_A).$$

If an installed controller has a verified replacement-gain interval $[l, u]$, a fixed installation charge τ gives net gain in $[l - \tau - \lambda C_i, u - \tau - \lambda C_i]$ for implementation i .

PROOF: Cancel the identical policy costs in the first display and subtract the specified deterministic charges from the gain interval in the second. *Q.E.D.*

This comparative statement holds for specified resource charges. Three measured process repetitions do not estimate a distribution-free expected runtime, and the study does not calibrate λ or τ . They do, however, permit a directly interpretable work comparison at exactly the same economic and statistical accuracy, without attributing different welfare to identical neural policies.

APPENDIX Z: EQUAL-POLICY WORK AT PROSPECTIVE ECONOMIC TARGETS

The preceding experiment asks whether an exact trained-neural witness changes a returned policy. A different, necessary question is how economically that same witness can be computed. The lossless screen permits a controlled comparison in which economic accuracy, fitted parameters, whole-cell verification and statistical stopping remain identical. We integrate the separately frozen R58 execution into this revision and independently reconstruct its solver, certificate and path calculations. The integration creates no new training run or independent statistical observation.

Z.1. *A matched intervention within Neural Bellman Operators*

The experiment retains the original nonlinear cyclic investment laws. Its tasks are $(d, T) = (2, 2), (4, 4), (8, 6)$ and its targets are ten- and twenty-percent expected-cost reductions relative to installed zero investment. Each task–target has two solver modes and three sequential complete-process repetitions, for 36 services. Both modes independently train the same 32-feature own-future ReLU continuation. Exact stored-parameter hashes establish matching; the common seed alone is not used as proof of matching.

The sole intervention is the nonterminal fitted-action solver. Algebraic search partitions the hinge objective and isolates its cubic stationary roots. Screened search performs exact activation-region reduction, outward lattice elimination and rational comparison of survivors. Both use the same robust action spacing $1/4096$, the same smallest-index tie rule and the same exact terminal solver. They retain the nine common actions and the critic-menu proposal; the exact action is an additional candidate. The same signed installed-reference verifier supplies all whole-cell upper endpoints and continuous-action lower covers.

The acquired observation geometry is the existing locally split binary partition, not a tensor grid. The two resource attempts use $(256, 32, 4)$ and $(1024, 64, 8)$ for training rows, acquired leaves and innovation bins. All attempts are checked against installed zero investment. Intermediate candidates are not deployed, and no derivative argument is applied to a changed discontinuous acquired actor. This experiment is a prospective refinement service, not a claim that two refinement attempts are two policy-iteration sweeps.

Z.2. First attainment, including nonattainment

At each attempted resolution the complete policy is frozen before inference. The target estimand is

$$D_q = J(\pi) - qJ(0), \quad q \in \{9/10, 4/5\}.$$

Cumulative sample sizes are 4096, 16384 and 65536. A nonpositive upper endpoint of D_q returns immediately. A strictly positive lower endpoint advances the construction attempt. An unresolved interval advances the inference look; an exhausted look budget advances the attempt. Failure to certify after the last attempt returns a safe policy with status *budget exhausted*. The protocol specifies these rules before its production execution, rather than selecting a favorable pass from a completed catalogue.

The R58 stopping stream is distinct from the R57 stopping and comparison streams. Within a task, target and attempt, the two solver modes and timing repetitions deliberately regenerate the same stream. Repeated computation of that stream does not multiply sample size. The confidence family has error at most $1/100$ under the maintained independent-bin model. Its conservative maximum possible number of interval evaluations is

$$36 \cdot 2 \cdot 3 \cdot 3 = 648 < 2048,$$

including cost, target contrast and gain at every possible look. The frozen empirical-Bernstein implementation retains its verified logarithm allowance. The union bound

requires no independence between policies or looks. It does not authorize unplanned sample sizes or an arbitrary future stopping rule.

TABLE Z.I

COMPLETE SERVICE WORK FOR EXACTLY MATCHING POLICIES

(d, T)	Reduction	Return	Algebraic (s)	Screened (s)	Ratio
(2,2)	10%	Met	1.212	1.101	1.101
(2,2)	20%	Met	3.577	3.409	1.049
(4,4)	10%	Met	5.837	5.487	1.064
(4,4)	20%	Met	20.554	19.373	1.061
(8,6)	10%	Met	20.229	19.270	1.050
(8,6)	20%	Exhausted	63.827	62.075	1.028

Each time is the median of all three separately reconstructed fixed-work process repetitions. The ratio is algebraic divided by screened. Corresponding policies, fitted parameters, verification endpoints and stopping histories agree exactly. Exhaustion is not successful target attainment. Different tasks use different runners; the comparison is within task and target.

TABLE Z.II

ORIGINAL-LAW ECONOMIC COSTS AT THE COMMON RETURN

(d, T)	Reduction	Policy cost	Gain from installed	Target contrast
(2,2)	10%	[0.5964, 0.6517]	[0.1302, 0.1786]	[-0.0983, -0.0548]
(2,2)	20%	[0.5984, 0.6203]	[0.1595, 0.1740]	[-0.0171, -0.0060]
(4,4)	10%	[0.8995, 1.0672]	[0.1205, 0.3800]	[-0.2482, -0.0056]
(4,4)	20%	[0.9527, 0.9718]	[0.2616, 0.2809]	[-0.0324, -0.0167]
(8,6)	10%	[1.4177, 1.4748]	[0.2130, 0.2963]	[-0.1235, -0.0456]
(8,6)	20%	[1.3464, 1.3649]	[0.3334, 0.3558]	[-0.0141, 0.0050]

The two modes have identical intervals, not independent estimates. Endpoints are rounded outward.

The signed target contrast, rather than a ratio of estimated means, determines attainment. The twenty-percent target in eight dimensions remains unresolved at the final budget.

Thirty services attain their target. Six exhaust their budgets: the two modes, each repeated three times, at $(d, T) = (8, 6)$ and the twenty-percent target. Their final target-contrast interval contains zero. The returned policy nevertheless has a positive gain interval relative to installation. Thus failure to certify the stronger target is not failure of policy safety, evidence of zero gain, or a proof that the economic target is impossible.

Every matched pair has identical fitted parameters, acquired actions, certificates, original-law path endpoints and stopping decisions. Consequently its policy-cost difference is exactly zero under every initial law, by Theorem 57. No separate paired statistical experiment is needed to establish this equality. In contrast, the distinct neural and conventional policies in Section X retain their own paired cost intervals and must not be assigned this identity conclusion.

Z.3. What the computational intervention removes

TABLE Z.III
NONTERMINAL NEURAL SEARCH WORK, INCLUDING REJECTED LATTICE POINTS

(d, T)	Reduction	Queries	Exact A	Exact S	Roots A	Screened points
(2,2)	10%	32	189	32	12	23040
(2,2)	20%	96	612	96	43	70144
(4,4)	10%	96	723	96	41	62976
(4,4)	20%	288	2283	288	112	192000
(8,6)	10%	160	1033	160	27	94720
(8,6)	20%	480	3467	480	126	289280

A denotes algebraic search and S screened search. Counts sum all attempted resolutions in one service; all three repetitions reproduce them. The screened mode retains one point per nonterminal neural query and performs zero root isolations and zero fallbacks in this catalogue. Terminal work is common and excluded from these neural counts.

The machine-independent records identify the mechanism. Across the frozen workloads, the screen removes noncrossing integrated ridges exactly and retains one lattice point per nonterminal neural query. That point is compared rationally

against the original objective. No production call uses the algebraic fallback. The number of original-objective rational evaluations therefore falls to the number of neural queries, whereas the algebraic mode compares several candidate points per query and performs cubic-root isolation. The table charges all lattice points evaluated by the interval screen; it does not count only survivors as total work.

The screen is not an approximate optimizer with a favorable tolerance. The audit recomputes both solvers from every distinct saved model, compares every recorded witness and endpoint hash, regenerates each whole-cell certificate, and reconstructs the original interval paths. Distinct computations are memoized only after their input identities are checked. Every service file and stopping transcript is still compared. Separate exact fixtures verify all lattice endpoints for disjoint small problems, signed weights, activation boundaries, ties, non-dyadic lattices and fallback cases.

Median action-search times decrease in all six task–target groups, as do the complete process medians. The full-service algebraic-to-screened ratios range from approximately 1.028 to 1.101. These gains are smaller than the action-search gains because training, common terminal search, verification and inference are unchanged. The eight-dimensional unsuccessful target is included in the work comparison: faster exhaustion is not described as faster successful attainment.

The effect is a measured implementation gain within the trained NBO procedure at exactly matching policies and statistical accuracy. It is not a ranking of neural methods above common-only, quadratic or tree-based alternatives. The stronger conventional comparisons, the common-menu attribution results and all unresolved differences in the preceding studies remain part of the evidence.

Z.4. *Complete work and the economic decision*

Every displayed service clock surrounds a separate process from invocation through exit and log flush. It includes imports, source checks, independently repeated fitting, action search, signed verification, own path evaluation and inference, all attempted resolutions and durable records. Available frequency-governor information, processor affinity, numerical-library versions, output bytes and peak

resident memory are retained. The parent driver's own checkpoint-writing overhead is outside that measured interval and is identified as controller overhead, not assigned a zero cost.

Within a task, processes run sequentially on one available core with one numerical-library thread, and solver order rotates by repetition and target. All three measurements are printed in the supplement. Frequencies are observed where available but not controlled. Different task runners do not supply a hardware-controlled dimension experiment. Three fixed-work repetitions establish neither a distribution-free runtime advantage nor a population of independently fitted policies. The three distinct fitted seeds in the preceding experiment answer a different question and remain separately identified.

For a specified computation price, Corollary 58 converts an equal-policy resource difference directly into a total-cost difference. Relative to the installed controller, the same actual gain interval supports a larger region of installation fees and computation prices when the specified implementation charge is lower. This is a substantive difference in the cost of obtaining an unchanged economic answer, not a claim that two representations of the same policy produce different welfare. The model's units remain normalized; no empirical installation fee, computation price or welfare calibration is inferred from the timings.

The original accuracy objective remains visible throughout. Every returned policy retains its one-sweep all-state loss account, including the installed-reference continuation bound as well as its full-action directed gap. Lossless search changes neither that account nor the amount of acquisition information. A successful initial-law cost-reduction target is not substituted for convergence to the original Bellman optimum. The constructive resource theory, fitted-witness transfer and directed finite-sweep theorem continue to state the conditions under which that stronger objective is controlled.

APPENDIX: CURRENT EXACT-SEARCH AND ACCURACY ADDITIONS

APPENDIX R61.1: EXACT ACTION RECOVERY AND THE BELLMAN ACCURACY CONTRACT

The preceding lossless screen makes the number of exact comparisons depend on the separation of the best lattice actions. That dependence is useful when separation is substantial, but it is not a reason to assume a favorable regime. We now exploit the integrated objective directly. The resulting exact solver does not require a growth constant, a unique minimizer, or isolation of an algebraic root. We also separate this local optimization question from verification against the original Bellman optimum.

R61.1.1. *Finite differences instead of root isolation*

Fix a state and a stored continuation. After the exact activation reductions, a scalar-action objective under a finite mixture of scalar uniform innovations has the form

$$F(a) = c_1 a + c_2 a^2 + c_4 a^4 + \sum_{j=1}^m w_j \Phi_{r_j}(z_j + s_j a) + \sum_{j=1}^h v_j (y_j + t_j a)_+^2, \quad a \in [0, \bar{a}], \quad (143)$$

up to a common constant. Here $c_4 \geq 0$, all stored coefficients are rational, and $\Phi_r(z) = \mathbb{E}(z + U)_+$ for U uniform on $[-r, r]$, with $\Phi_0(z) = z_+$. The original investment model has $c_4 = 4$. Mixture probabilities and shifts can be absorbed into the ridge weights and intercepts. Negative weights are allowed. Breakpoints are rational and divide the interval into at most $K \leq 2m + h + 1$ polynomial pieces. On each closed piece, continuity gives the polynomial $P(a) = d_0 + d_1 a + d_2 a^2 + c_4 a^4$ at the endpoints as well.

Let $p_j = P(j/Q)$ on the consecutive feasible indices $j = l, \dots, u$ of one piece. Define $d_j = p_{j+1} - p_j$. Direct expansion gives

$$d_j = \frac{d_1}{Q} + \frac{d_2(2j+1)}{Q^2} + \frac{c_4(4j^3 + 6j^2 + 4j + 1)}{Q^4}, \quad (144)$$

$$d_{j+1} - d_j = \frac{2d_2}{Q^2} + \frac{c_4(12j^2 + 24j + 14)}{Q^4}. \quad (145)$$

The last expression is nondecreasing for $j \geq 0$, including when $d_2 < 0$. This is a discrete curvature property, not global convexity of the fitted objective.

LEMMA 59—Exact minimization of one piece: Suppose $0 \leq l \leq u$ and $c_4 \geq 0$. The smallest minimizing index of $p_l, \dots, p_u$ can be recovered using $O(1 + \log(u - l + 1))$ exact rational sign queries and at most four exact objective comparisons. If every second difference is negative, endpoints suffice. Otherwise let j_0 be the first index with nonnegative second difference. In addition to the endpoints, it suffices to retain j_0 and the first $k \geq j_0$ with $d_k \geq 0$, when that index exists. If it does not exist, the right endpoint is the minimum of the remaining tail.

PROOF: A sequence with decreasing first differences is discretely concave and has a minimum at an endpoint. If j_0 exists, the prefix through p_{j_0+1} is concave, whereas the tail beginning at p_{j_0} is convex. A minimum of the full sequence is therefore either the left endpoint or a minimum of that convex tail. The latter is at the first nonnegative first difference, or at the right endpoint if all its first differences are negative. The first such index is the smallest tail minimizer, including a run of exact ties. Monotonicity of (145) locates j_0 by integer bisection, and monotonicity of d_j on the tail locates k in the same way. Intervals with fewer than three indices are handled by their endpoints. Exact comparisons of the retained values complete the argument. Q.E.D.

THEOREM 60—Root-free exact action witness: For (143) on $\{j/Q : 0 \leq j \leq \bar{j}\}$, with $\bar{a} = \bar{j}/Q$, form the rational activation partition, intersect each piece with the integer lattice, and apply Lemma 59. Evaluate the original objective at the union of the retained indices, choosing the smallest index on an exact tie. The result is precisely the original lattice minimizer, without a separation assumption or an algebraic root calculation.

A direct implementation uses

$$O((m + h)^2 + (m + h) \log(m + h + 1) + K \log(\bar{j} + 2)) \quad (146)$$

rational operations and comparisons, and $O(m + h + K)$ stored rational numbers and indices. These are operation counts, not unit-cost bit bounds. If b bounds the bit lengths of all operands and $C_{\text{rat}}(b)$ bounds one rational operation or comparison, multiplying (146)

by $C_{\text{rat}}(b)$ is a valid bit-cost account. For input coefficient bit length at most B , fixed degree
and denominator Q , a conservative bound is

$$b = O((m + h + 1)(B + \log(Q + 1) + \log(\bar{j} + 2) + \log(m + h + 2))).$$

PROOF: The activation formulas are exact on every piece. Its lattice intersection is
the consecutive interval from the ceiling of its left endpoint times Q to the floor of its
right endpoint times Q ; empty intersections are skipped. These intersections cover
every lattice action. The lemma retains a minimizing index for every nonempty in-
tersection and its smallest minimizer. Exact comparison under the original objective
is unaffected by the piecewise constants or the earlier common-constant reduction,
and gives the original tie rule globally.

Sorting breakpoints costs $O((m + h) \log(m + h + 1))$ comparisons. Constructing
the polynomial coefficients independently in each piece costs $O(K(m + h))$ opera-
tions. The integer bisections cost $O(K \log(\bar{j} + 2))$. There are $O(K)$ final candidate
indices, and direct evaluation of the original objective costs $O(K(m + h))$. The im-
plementation stores the features, breakpoints and candidate indices, rather than an
array indexed by every lattice point. Finally, products of a fixed number of input ra-
tionals have bit length proportional to the input lengths. Summing at most $O(m + h)$
such terms admits a common denominator formed by their product; this gives the
stated conservative operand bound, including powers of Q , the action-index bound,
and the rational breakpoint calculations. Q.E.D.

The theorem is useful in precisely the case where interval separation is inconclu-
sive. A flat piece, a tie between consecutive actions, or a large signed cancellation
does not prevent exact bisection. It does not assert that bisection is always fastest:
when the lattice is short relative to the activation partition, native exhaustive eval-
uation can be cheaper. The ablation below therefore includes both, as well as the
original symbolic solver after the same exact reduction. Discrete curvature and veri-
fied interval exclusion are established numerical principles (Murota, 2003, Moore
et al., 2009). The result here is their explicit, root-free realization for the integrated
NBO objective, with an exact original-witness contract and a bit-aware implementa-
tion account.

R61.1.2. *Several controls and uniform-box shocks*

The lossless witness contract itself need not be scalar. Let $a \in \mathbb{R}^p$ belong to a bounded rational polytope and to a declared rational action lattice. For a ridge of a one-hidden-layer ReLU continuation, affine exposure gives an argument $z + s^\top a + \sum_{i=1}^q U_i$, where the U_i are independent and uniform on $[-r_i, r_i]$. Zero radii are removed. Finite mixtures of such boxes are allowed, with nonnegative rational probabilities summing to one.

PROPOSITION 61—Exact integrated ridge: *For $q \geq 1$ and $r_i > 0$,*

$$\mathbb{E}\left(z + s^\top a + \sum_{i=1}^q U_i\right)_+ = \frac{\sum_{\epsilon \in \{-1,1\}^q} \left(\prod_{i=1}^q \epsilon_i\right) \left(z + s^\top a + \sum_{i=1}^q \epsilon_i r_i\right)_+^{q+1}}{2^q (q+1)! \prod_{i=1}^q r_i}. \quad (147)$$

Consequently each fitted objective is rational at every feasible lattice action. Its value and valid box enclosures can be computed with signed output weights and exact ties. Evaluation costs include the 2^q terms per ridge; the formula does not remove shock-dimension dependence.

PROOF: An iterated antiderivative of x_+ of order q is $x_+^{q+1}/(q+1)!$. Apply the fundamental theorem of calculus in each integration coordinate and divide by the uniform-box volume. Each boundary choice supplies the corresponding sign in (147). Rationality follows from the rational primitives and lattice. The expected hinge is monotone in its scalar argument, so evaluating it at the exact minimum and maximum of $z + s^\top a$ on a rectangle supplies a valid enclosure. A negative output weight reverses the endpoints. Addition and nonnegative mixture weighting preserve inclusion. Q.E.D.

The executed two-control extension uses shared capacity $a_1 + a_2 \leq 1/8$, exposure columns $(1/2, 1/4)$ and $(1/4, 1/2)$, and the original first shock together with an optional independent second shock. The extension is stated rather than conflated with the scalar experiment: setting $a_2 = 0$ and disabling the second shock recovers the original action-dependent objective. The quartic lemma is used only in its

1 stated scalar, piecewise-quartic subclass. The two-shock and two-control calculation 1
 2 instead uses verified box subdivision and exact feasible-point comparison. 2

3
 4 THEOREM 62—Finite constrained recovery with a work budget: *Suppose a finite* 3
 5 *rational action lattice is intersected with a bounded rational polytope, and every visited* 4
 6 *integer rectangle has a valid lower objective bound. Maintain an exactly evaluated feasible* 5
 7 *incumbent and a fixed lexicographic tie rule. Discard a rectangle only if it is infeasible, its* 6
 8 *lower bound is strictly greater than the incumbent value, or equality holds and its smallest* 7
 9 *possible index is not earlier than the incumbent. Split any other nonsingleton rectangle* 8
 10 *into two disjoint integer rectangles, and evaluate singleton actions exactly. Then exhaustive* 9
 11 *subdivision returns the original smallest minimizing feasible action. Switching after any* 10
 12 *declared node budget to an exact exhaustive solver preserves this conclusion and charges the* 11
 13 *preceding work in addition to the fallback.* 12

14 For side cardinalities $N_1, \dots, N_p$, a depth-first implementation has depth at most 13
 15 $\sum_i \lceil \log_2 N_i \rceil$ under midpoint splitting. Its rectangle stack stores $O(p \sum_i \log(N_i + 1))$ inte- 14
 16 gers. Without pruning it can still visit $O(\prod_i N_i)$ nodes. Objective arithmetic, rational bit 15
 17 lengths and exact-polytope feasibility work must be added to these combinatorial counts. 16

18 PROOF: Each split partitions the remaining integer indices without omission. 18
 19 A discarded infeasible rectangle contains no admissible competitor. Strict lower- 19
 20 bound exclusion removes no minimizer that could improve the incumbent. In the 20
 21 equality case no point can improve its value, and the index condition excludes 21
 22 an earlier tied minimizer. Every unpruned feasible singleton is compared exactly. 22
 23 Finite subdivision therefore retains the globally best value and the specified first 23
 24 minimizer. An exact terminating fallback solves the same remaining mathematical 24
 25 problem, whether or not it reexamines previously visited points. Its repeated work 25
 26 affects cost, not validity. Midpoint splitting halves the active side range, yielding 26
 27 the depth and stack bounds; a complete binary partition has fewer than twice as 27
 28 many nodes as its full rectangular lattice has points. Q.E.D. 28

29
 30 Replacing a solver by either theorem preserves the earlier prospective transcript 30
 31 whenever that controller depends on the mathematical witness rather than elapsed 31
 32 time or solver diagnostics. This transfers exactness to the acquired policy and its 32

certificates; it does not imply that a fitted-optimal action is optimal for the true continuation. That latter error remains in the existing transfer theorem and is examined by an independent all-state calculation next.

R61.1.3. *A direct certificate against the original Bellman optimum*

Let $\mathcal{C}_t$ be a finite cover of the original state domain with an unambiguous ownership rule for shared boundaries. Starting from terminal bounds on g , maintain a piecewise lower function L_{t+1} for the Bellman value and a piecewise upper function U_{t+1} for the value of an explicitly constructed continuation policy. On a cell C , let $\mathcal{B}_t(C)$ cover every action feasible at every state in C , and let $\mathcal{A}_t(C)$ contain only actions feasible throughout C . Suppose verified endpoints satisfy

$$\begin{aligned} \ell_t(C, B) &\leq c_t(x, a) + \beta_t P_t^a L_{t+1}(x) & (x \in C, a \in B \cap A_t(x)), \\ u_t(C, b) &\geq c_t(x, b) + \beta_t P_t^b U_{t+1}(x) & (x \in C, b \in \mathcal{A}_t(C)). \end{aligned}$$

A lower cover may include additional infeasible actions, but every evaluated extension must retain a valid lower enclosure on its feasible subset. All reached state cells and the full original innovation law are included in the expectations. Define

$$L_t(C) = \min_{B \in \mathcal{B}_t(C)} \ell_t(C, B), \quad U_t(C) = \min_{b \in \mathcal{A}_t(C)} u_t(C, b),$$

and implement the minimizing upper candidate on C .

THEOREM 63—Whole-domain Bellman bracket: *The constructed complete policy π is feasible and satisfies*

$$L_t(C) \leq J_t^*(x) \leq J_t^\pi(x) \leq U_t(C), \quad x \in C. \quad (148)$$

Hence $\max_C \{U_t(C) - L_t(C)\}$ is an all-state upper bound on loss against the original continuous-action Bellman optimum. A declared tolerance is attained when its outward bound passes the corresponding comparison; no initial distribution or reference-zero gain is needed.

For a refining family, suppose the exact lower and upper endpoint maps at each fixed approximation level are order preserving and β_t -nonexpansive in the supremum norm.

Suppose their one-step errors at the true next Bellman value, together with uniformly bounded numerical endpoint errors on the relevant input range, are at most η_t , inclusive of robust-action approximation and state, action and innovation covering. If the terminal lower and upper errors are each at most η_T , then

$$\sup_C \{U_t(C) - L_t(C)\} \leq 2 \sum_{j=t}^T \left(\prod_{i=t}^{j-1} \beta_i \right) \eta_j.$$

In particular, uniformly vanishing such allowances give arbitrarily small Bellman brackets in finite horizon. The runtime certificate is (148); the consistency statement does not require the algorithm to observe J^* .

PROOF: Assume the next-date inequalities. Monotonicity of each action kernel gives $P_t^a L_{t+1} \leq P_t^a J_{t+1}^*$. The complete action cover and the lower enclosure therefore imply $L_t(C) \leq \inf_{a \in A_t(x)} \{c_t(x, a) + \beta_t P_t^a J_{t+1}^*(x)\} = J_t^*(x)$. The chosen upper action is feasible at every state in C . Its continuation policy has value at most U_{t+1} , so its current value is at most the chosen upper endpoint $U_t(C)$. Feasibility gives $J_t^* \leq J_t^\pi$. Backward induction proves the bracket.

For consistency, let e_t^- bound $J_t^* - L_t$ and let e_t^+ bound $U_t - J_t^*$. Nonexpansivity of the fixed-cover endpoint maps bounds the propagated next-date discrepancy by β_t times its supremum norm. Adding the uniformly budgeted local and numerical endpoint errors consequently gives $e_t^- \leq \eta_t + \beta_t e_{t+1}^-$ and $e_t^+ \leq \eta_t + \beta_t e_{t+1}^+$, with terminal errors at most η_T . Adding the two unrolled inequalities proves the bound. Primitive uniform continuity and consistent enclosure rules supply vanishing allowances on compact domains when the robust feasible menus and full-action covers converge. These are hypotheses for consistency, not an uncharged value-function oracle. Q.E.D.

The numerical implementation uses exact rectangle extrema of stored outward bounds. At the last date it integrates the original terminal cost analytically. At earlier dates each innovation bin retains its exact original probability, and no midpoint is substituted for the bin's conditional expectation. This calculation is a conventional constructive benchmark in the same NBO economy. It puts an attained all-state accuracy and the work paid for it beside the learned prospective services, rather than

presenting improvement relative to zero investment as a solution of the Bellman problem.

R61.1.4. *The economic value of an exact implementation change*

For two implementations returning the identical policy, let w_0, w_1 be their measured per-service work, k_0, k_1 their one-time compilation or installation work, and $\lambda \geq 0$ the price of the recorded work unit. Over H uses, with an additional implementation charge I , the new-minus-old total cost is

$$\lambda\{H(w_1 - w_0) + (k_1 - k_0)\} + I.$$

When $w_1 < w_0$, the recorded computation saving pays for a positive setup difference only beyond the corresponding break-even use count. This condition is reported separately from policy-cost comparisons and from the cold single-use boundary. No electricity tariff, processor frequency, installation price or external welfare calibration is inferred from a hosted-runner clock.

R61.1.5. *Common-reference cancellation at the original optimum*

The width of an absolute value box can be dominated by state variation common to all actions. Subtracting the same reference function from both Bellman endpoints removes that variation without changing the decision problem. This is a use of the original policy-relevant error principle, not a replacement of the Bellman objective by a zero-policy comparison.

Fix any feasible complete reference π^0 and write $h_t = J_t^{\pi^0}$. Define its true reference advantage by

$$A_t^0(x, a) = c_t(x, a) + \beta_t P_t^a h_{t+1}(x) - h_t(x).$$

Set $e_t^*(x) = J_t^*(x) - h_t(x)$ and $e_t^\pi(x) = J_t^\pi(x) - h_t(x)$. Their exact recursions are

$$e_t^*(x) = \inf_{a \in A_t(x)} \{A_t^0(x, a) + \beta_t P_t^a e_{t+1}^*(x)\},$$

$$e_t^\pi(x) = A_t^0(x, \pi_t(x)) + \beta_t P_t^{\pi_t(x)} e_{t+1}^\pi(x), \quad e_T^* = e_T^\pi = 0.$$

The reference need not be fitted, and a numerical routine need not evaluate h_t separately if it encloses its signed advantages directly from the primitives.

COROLLARY 64—Reference-centered Bellman certificate: *Apply the lower-cover and robust upper-candidate construction of Theorem 63 to the excess recursions above, using verified enclosures of A_t^0 and the original probability kernels. Denote the resulting piecewise excess bounds by ℓ_t and u_t . Then, at every true state,*

$$\ell_t(x) \leq J_t^*(x) - h_t(x) \leq J_t^\pi(x) - h_t(x) \leq u_t(x).$$

Consequently $\sup_x \{u_t(x) - \ell_t(x)\}$ is an all-state loss certificate for the original constructed policy against the original continuous-action Bellman optimum. There is no separate reference-approximation allowance when the advantage enclosure is obtained directly from the exact reference law.

For an independently fixed feasible policy $\hat{\pi}$, an upper-only excess recursion using its full acquired-actor action range gives $\hat{u}_t \geq J_t^{\hat{\pi}} - h_t$. The quantity $\sup_x \{\hat{u}_t(x) - \ell_t(x)\}$ therefore certifies that particular returned policy. It does not require the policy to be regenerated on the verification grid.

PROOF: Subtract h_t from the Bellman equation and from the policy-value equation to obtain the two recursions. The induction in Theorem 63 applies verbatim to these transformed equations: kernel monotonicity and the complete feasible-action cover give the lower excess bound, and implementability and the upper continuation bound give the upper excess bound. Both endpoints at a given true state contain the identical $h_t(x)$, which cancels in their difference. The same upper induction applies to any fixed feasible actor when its complete action range is enclosed on every verification cell. Subtracting the lower optimum bound proves the final assertion. Q.E.D.

In the two-date investment economy, choose the original zero-investment policy as reference. Its final-date advantage is the original analytic final-cost difference. At the initial date, use a common innovation to propagate the state difference caused by the proposed action and integrate the difference of subsequent zero-policy costs. Centered polynomial identities evaluate that difference; the final innovation's com-

mon variance cancels, while its correct marginal law is retained. The construction therefore does not assume a policy-value oracle, a differentiable acquired policy, or a learned invariant.

The centered experiment is explicitly subsequent to the uncentered one. It retains the unsuccessful original tolerance tests and applies a separately frozen five-rung verification design. The same centered lower table can also certify the previously returned learned and conventional policies. Such revalidation has its own cost and does not retroactively alter the original prospective stopping criterion or its measured clock. This connects the original Bellman-accuracy question to the actual returned objects, rather than only to an alternative constructed policy.

APPENDIX R61.2: STABLE ENCLOSURES AND CERTIFICATION OF A RETURNED POLICY

A valid one-step interval and a stable recursion are different requirements. The first establishes a runtime bracket; the second controls how an approximation allowance propagates when the next-date input changes. We make the distinction explicit and then apply the common-reference construction to a policy that has already been returned by a prospective service.

R61.2.1. *Fixed-cover stability*

At a fixed approximation level, hold the state cells, action covers, candidate actions, innovation bins and their probabilities fixed. An endpoint map may take minima or maxima over the cells intersecting a transition box, average those endpoints with nonnegative weights of total mass one, add a stage-cost endpoint and finally minimize over a fixed finite family of action intervals or feasible candidates. A reference-shifted endpoint map may additionally add a fixed signed reference advantage. These operations describe the exact endpoint maps beneath the outward arithmetic used here.

LEMMA 65—Stability of the executable endpoint maps: *Every such exact endpoint map $\mathcal{E}_t$ is order preserving and satisfies*

$$\|\mathcal{E}_t f - \mathcal{E}_t g\|_\infty \leq \beta_t \|f - g\|_\infty. \quad (149)$$

Before intersection with an independent bound, it also satisfies $\mathcal{E}_t(f + c) = \mathcal{E}_t f + \beta_t c$ for constants c . Intersecting a lower endpoint with an independently valid lower bound by taking a maximum, or an upper endpoint with an independently valid upper bound by taking a minimum, preserves order and the nonexpansive estimate when the independent bound is fixed. A numerical endpoint perturbation with verified uniform magnitude δ_t enters the propagated error as an additional local allowance, rather than being presumed nonexpansive because it is outward.

PROOF: For any nonempty finite index set I , both $\min_{i \in I} f_i$ and $\max_{i \in I} f_i$ are order preserving and one-Lipschitz in the supremum norm. A nonnegative weighted average of total mass one has the same properties. A fixed additive term does not alter a difference, and multiplication by $\beta_t \geq 0$ supplies the factor in (149). The final minimum over actions is again one-Lipschitz. Each of these operations commutes with addition of a constant in the stated way. Finally, $x \mapsto \min\{x, c\}$ and $x \mapsto \max\{x, c\}$ are monotone and one-Lipschitz. A separate numerical perturbation is bounded by the triangle inequality. Q.E.D.

Consequently, the consistency clause of Theorem 63 applies to the fixed-cover maps actually used in the computation. If the ideal lower and upper maps at the true next Bellman value have errors at most α_t , and their numerical evaluation has uniformly bounded endpoint error δ_t on the relevant range of inputs, put $\eta_t = \alpha_t + \delta_t$. Then the two one-sided errors satisfy

$$e_t^- \leq \eta_t + \beta_t e_{t+1}^-, \quad e_t^+ \leq \eta_t + \beta_t e_{t+1}^+.$$

Their sum gives the displayed two-sided finite-horizon bound. Different approximation levels may use different covers; the stability comparison is within each fixed level. Arbitrarily small tolerances require all nonvanishing sources of error to be

refined, including the action quantum when it is a binding discretization. A finite ladder at quantum $1/4096$ is not, by itself, an executed zero-error limit.

Inclusion alone would not justify this step. In a one-state example with continuation $v \geq 0$ and discount β , the endpoint $\beta v + v^2$ is a valid upper bound on βv and is exact at $v = 0$. Its variation from 0 to 1 nevertheless exceeds β . Thus an ideal error statement only at the true continuation cannot replace (149). The present proof supplies the needed property for the actual min-max and probability-average maps; it does not assume stability of every conceivable enclosing procedure.

R61.2.2. *The original optimum and the actual returned actor*

Let $h_t = J_t^0$ denote the value of the installed zero-investment policy in the original model. A common-reference Bellman calculation supplies a function ℓ_t satisfying

$$\ell_t(x) \leq J_t^*(x) - h_t(x).$$

Now fix a complete, previously returned acquired policy π . Its observation partition and quantum action indices are part of its identity and are not changed by verification. On a verification cell C , let $R_t(C)$ contain every action actually taken by π_t at a state in C . Such a rectangle is an outer enclosure of the graph of the fixed policy, not a proposal to implement every action in that rectangle. In particular, its endpoints need not be feasible throughout C ; the actual policy must remain feasible at each actual state, and the primitive formulas used on the outer enclosure must remain valid there.

Suppose $\mathcal{V}_t(C)$ is an outward enclosure, for all actual graph points $(x, \pi_t(x))$ with $x \in C$, of

$$A_t^0(x, \pi_t(x)) + \beta_t P_t^{\pi_t(x)} u_{t+1}(x), \quad (150)$$

where $A_t^0(x, a) = c_t(x, a) + \beta_t P_t^a h_{t+1}(x) - h_t(x)$ and $u_T = 0$. Define $\tilde{u}_t(C)$ as the upper endpoint of $\mathcal{V}_t(C)$.

PROPOSITION 66—Paid revalidation of an unchanged policy: *If the stored actor is feasible and the enclosures in (150) include every reached state and the original innovation*

law, the recursion $u_t = \tilde{u}_t$ gives

$$\ell_t(x) \leq J_t^*(x) - h_t(x) \leq J_t^\pi(x) - h_t(x) \leq u_t(x).$$

If an independent whole-domain policy certificate has established $J_t^\pi \leq h_t$ at every date, the recursion may instead use

$$u_t(C) = \min\{\tilde{u}_t(C), 0\}. \quad (151)$$

In either case, an outward maximum of $u_t(C) - \ell_t(C)$ bounds the loss of this particular returned actor against the original continuous-action Bellman optimum. Verification does not change the actor or require a policy-cost sampling event.

PROOF: The identity

$$J_t^\pi - h_t = A_t^0(\cdot, \pi_t) + \beta_t P_t^{\pi_t}(J_{t+1}^\pi - h_{t+1})$$

holds under the original kernels. Starting from the common terminal value, monotonicity of expectation and the whole-graph enclosure prove the upper inequality by backward induction. The lower inequality is the previously established Bellman lower bound; feasibility gives the middle inequality. When $J_t^\pi - h_t \leq 0$ has been independently proved, intersection with zero preserves the upper bound at each induction step. Subtracting the two bounds cancels the same $h_t(x)$ pointwise. No difference between extrema at unrelated states is substituted for this cancellation. Q.E.D.

The independent property used in (151) is obtained here by reconstructing the original reference-advantage gate at every acquired cell and date. Its nonpositive upper endpoints imply $A_t^0(x, \pi_t(x)) \leq 0$ throughout the domain, and backward policy improvement gives $J_t^\pi \leq J_t^0$. The revalidation also retains the entirely unclipped upper recursion and its final gaps. This permits the reader to distinguish a justified independent bound from a numerically favorable truncation. A record of nonpositive values at sampled centers would not supply the required property.

The verification grid is finer than the actor's acquisition partition. Actor ranges include all leaves intersecting each closed verification cell, including both sides of a

discontinuity. Rectangle extrema and original-law innovation bins then enclose the continuation. These calculations have a real cost: the common lower-table construction, each fixed-actor upper recursion and the whole-process return overhead are separate measured services. They do not retroactively reduce a prospective service's original stopping cost or certify an unexecuted high-dimensional case.

APPENDIX R61.3: EXACT COMPUTATION AND THE ECONOMIC ACCURACY CONTRACT

The preceding experiments establish distinct parts of the NBO procedure. We now examine how the exact implementation behaves away from one-survivor screening, what it costs inside a complete economic service, and whether the policies actually returned can be certified against the original Bellman optimum. The original nonlinear investment family and its continuous laws are retained. The additional two-control calculation is identified explicitly as a nested extension rather than being counted as another scalar-action policy service.

R61.3.1. *Design, information sets, and accounting*

The completed comparison has four generators: a common-action controller with no fitted continuation, a quadratic continuation with native exact witnesses, a trained ReLU continuation with lossless screening, and the same trained ReLU construction with native finite-difference witnesses. The complete services use $(d, T) = (2, 2), (4, 4), (8, 6)$, training seeds 60103 and 60209, and targets $q = 9/10$ and $q = 4/5$. The target requires expected policy cost at most q times the expected cost of installed zero investment under the declared original initial law. These targets are not a claim about the optimum at every state.

Every service starts from the primitives. Its declared construction ladder uses $(256, 32, 4)$ and $(1024, 64, 8)$ for training samples, acquired leaves and innovation bins. The original action quantum is $1/4096$. Verification uses the same signed reference-advantage calculation, complete-cell feasibility rule and continuous-action lower cover for all four generators. Sequential inference examines 4096, 16,384 and 65,536 original-law interval paths, stopping at the first declared condition. The

current execution has two hosted-worker repetitions of every mathematical design. Repeating identical seeds, policies and paths on a second worker does not double the number of independent economic observations.

The source freeze precedes the scientific services. A first execution failed in the parent recorder after an exclusive-write file was opened a second time. Its completed child outputs and failure logs remain in the attempt archive. A separately frozen correction changes that recorder to immutable numbered checkpoints and a single final ledger; the mathematical child service, seeds, targets and stopping rules are unchanged. The full corrected catalogue is used below, with no mixture of partial first-run clocks and completed rerun clocks. The separate centered-Bellman amendment was written after the uncentered frontier was observed. The fixed-policy revalidation and native factorial protocols likewise identify which predecessor results were known before their own frozen computations.

R61.3.2. *Complete services and actual policy costs*

Table [R61.3.I](#) reports every task–target–generator group on the complete return boundary. It includes imports, fitting, the executed construction and search attempts, whole-cell verification, sequential inference, durable scientific outputs, the process log and the first durable parent receipt. Cold use additionally charges native compilation; warm use starts with the compiled binary available. The source registry gives every one of the 96 services, including the 16 exhausted budgets. In total, 80 services meet their declared economic target.

The screened and native ReLU configurations return the same fitted parameters, action witnesses, acquired policy, verification arrays, stopping stage, path counts and cost intervals. Their difference is therefore an implementation-work difference at the same mathematical answer. The full-service comparison charges the native request and response as well as the exact computation. The native query benchmark below is not substituted for this complete return boundary, and a reduction in exact comparison counts is not presumed to translate proportionally into service time.

A different generator requires a different comparison. Table [R61.3.III](#) reports native-ReLU-minus-common and native-ReLU-minus-quadratic cost intervals for

TABLE R61.3.I

COMPLETE RETURN WORK AT THE SAME ECONOMIC TARGETS

d/T	q	Mode	Met 0/1	Warm 0 (s)	Warm 1 (s)	Cold 0 (s)	Cold 1 (s)
2/2	9/10	C	2/2	1.244	1.247	1.244	1.247
2/2	9/10	Q	2/2	1.310	1.310	8.615	6.046
2/2	9/10	S	2/2	1.562	1.573	1.562	1.573
2/2	9/10	N	2/2	1.532	1.537	8.837	6.273
2/2	4/5	C	2/2	3.792	3.793	3.792	3.793
2/2	4/5	Q	2/2	3.981	3.956	11.286	8.692
2/2	4/5	S	2/2	4.820	4.818	4.820	4.818
2/2	4/5	N	2/2	4.781	4.730	12.086	9.466
4/4	9/10	C	2/2	4.347	4.366	4.347	4.366
4/4	9/10	Q	2/2	4.738	4.712	12.044	9.448
4/4	9/10	S	2/2	5.523	5.558	5.523	5.558
4/4	9/10	N	2/2	5.455	5.464	12.760	10.200
4/4	4/5	C	2/2	15.835	15.887	15.835	15.887
4/4	4/5	Q	2/2	17.012	16.904	24.317	21.640
4/4	4/5	S	2/2	19.796	19.709	19.796	19.709
4/4	4/5	N	2/2	19.524	19.858	26.830	24.594
8/6	9/10	C	2/2	16.858	16.787	16.858	16.787
8/6	9/10	Q	2/2	18.371	18.271	25.676	23.007
8/6	9/10	S	2/2	19.675	19.631	19.675	19.631
8/6	9/10	N	2/2	19.575	19.365	26.880	24.101
8/6	4/5	C	0/0	52.701	52.132	52.701	52.132
8/6	4/5	Q	0/0	57.282	57.480	64.588	62.216
8/6	4/5	S	0/0	62.443	62.636	62.443	62.636
8/6	4/5	N	0/0	61.769	61.469	69.075	66.206

C uses the common verified action family without a fitted continuation; Q uses native quadratic witnesses; S and N use the screened and native ReLU witnesses. The target is expected policy cost at most q times zero-policy cost. Each worker has two declared seeds; Met 0/1 gives the attained counts separately, each out of two. Warm and cold entries are within-worker two-seed medians of the complete durable-return clock. Cold work additionally charges that worker's measured native compilation for Q and N. Worker repetitions are not new independent economic samples. Failed targets remain in the table.

TABLE R61.3.II
COMPLETE-SERVICE EFFECT OF REPLACING SCREENED WITH NATIVE ReLU SEARCH

d/T	q	Saving 0 (%)	Saving 1 (%)	Faster 0/1
2/2	9/10	1.95	2.29	2/2
2/2	4/5	0.81	1.83	2/2
4/4	9/10	1.24	1.68	2/2
4/4	4/5	1.37	-0.75	2/1
8/6	9/10	0.51	1.35	1/2
8/6	4/5	1.08	1.86	2/2

Savings are paired relative reductions in complete warm return time, summarized by the median of the two seeds within each worker. Negative values mean slower native service. Faster counts are out of two per worker. The exact returned policy, fitted parameters, stopping stage, inference looks and intervals coincide. These are descriptive timings, not a frequency-controlled or cross-platform speed theorem.

each task, target and seed on a fresh final-policy comparison stream. Positive intervals identify a more costly native-ReLU policy; exact identities are reported as identities; other intervals containing zero remain unresolved. The matched native-minus-screened contrast is exactly zero because their complete actor identities coincide. Conventional methods receive the same verifier and economic criterion. Their lower construction or return work is not removed from the ranking merely because they do not use learned ReLU features.

Across the twelve mathematical designs, native ReLU and each of the common-action and quadratic comparators have seven exact policy identities and five unresolved cost differences. No remaining comparison identifies a strictly lower native-ReLU cost. The common-action service has the smallest within-worker median warm return time in every task–target group; in a group whose target is exhausted by all generators this is a statement about observed work, not the work required to attain the target. These outcomes distinguish the implemented NBO contract from a claim of universal neural superiority.

The additional exact witness is available only after the common action family and the finite fitted menu have been evaluated. Comparing the selected policy immediately before and after that addition identifies the decisions attributable

TABLE R61.3.III

ACTUAL RETURNED-POLICY COST: NATIVE ReLU MINUS CONVENTIONAL COMPARATORS

d/T	q	Seed	N minus C	N minus Q
2/2	9/10	60103	[0.00000, 0.00000]	[0.00000, 0.00000]
2/2	9/10	60209	[-0.00514, 0.00515]	[-0.00514, 0.00515]
2/2	4/5	60103	[-0.00517, 0.00515]	[-0.00517, 0.00515]
2/2	4/5	60209	[-0.00545, 0.00488]	[-0.00545, 0.00488]
4/4	9/10	60103	[-0.00781, 0.00751]	[-0.00781, 0.00751]
4/4	9/10	60209	[0.00000, 0.00000]	[0.00000, 0.00000]
4/4	4/5	60103	[-0.00769, 0.00760]	[-0.00769, 0.00760]
4/4	4/5	60209	[0.00000, 0.00000]	[0.00000, 0.00000]
8/6	9/10	60103	[0.00000, 0.00000]	[0.00000, 0.00000]
8/6	9/10	60209	[0.00000, 0.00000]	[0.00000, 0.00000]
8/6	4/5	60103	[0.00000, 0.00000]	[0.00000, 0.00000]
8/6	4/5	60209	[0.00000, 0.00000]	[0.00000, 0.00000]

Intervals use the fresh final-policy comparison stream, conditional on the original selection. Endpoints are rounded outward. Positive intervals identify higher native-ReLU cost; an interval containing zero is unresolved. Exact policy identities have an exact zero interval. There are twelve mathematical designs, each repeated on two workers with the same policies and stream, not twenty-four independent comparisons. N minus S is identically zero.

to the extra witness under a fixed verifier. The complete reconstruction checks every such cell, including unchanged and rejected proposals. Across the distinct reconstructed verification objects, it records 12 additional-witness cell changes. This count includes the tested fitted generators; it is not an independent sample size, and it does not imply an identified average welfare gain for each changed cell. Actual economic consequences are assessed by the full-policy intervals rather than by assigning a policy value to a local reduction in a certificate.

R61.3.3. *A language- and representation-matched action-search experiment*

The eight-method stress experiment varies three training-style fixture seeds, two feature widths, three action-lattice cardinalities and five regimes. The regimes

include mixed activation, activation crossing, large signed cancellation, an exact tie between neighboring actions and a flat objective. Every rational objective is checked against exhaustive Fraction evaluation with the original smallest-index tie rule. An additional fixed subset receives seven shuffled repetitions per regime and worker. Coarser endpoint bounds and large-lattice guard cases are retained whether or not pruning is effective.

The eight methods separate symbolic reduction, vectorized enclosure, screening, exhaustive exact evaluation, native finite differences and adaptive subdivision. However, one direct contrast in that catalogue combines two changes: native exhaustive search uses the unreduced objective, whereas the native piecewise solver first reduces activation regions. We therefore execute a further two-by-two comparison before drawing an algorithm-specific timing conclusion. All four cells use the identical C++ arbitrary-size rational engine, objective, lattice, exact tie rule and original-value postcheck. They cross activation reduction, off or on, with exhaustive or finite-difference search. Reduction is charged to both reduced implementations. The same ninety fixed rational instances, rather than an outcome-selected favorable subset, are used on two newly allocated worker jobs.

TABLE R61.3.IV
FINITE-DIFFERENCE VERSUS EXHAUSTIVE SEARCH AT THE SAME REDUCED REPRESENTATION

Width	Lattice	Faster 0	Faster 1	Saving 0 (%)	Saving 1 (%)
8	33	13/15	8/15	2.51	0.40
8	129	15/15	15/15	23.05	25.20
8	513	15/15	15/15	60.51	65.07
32	33	1/15	0/15	-17.79	-18.93
32	129	15/15	15/15	26.77	30.69
32	513	15/15	15/15	70.43	72.56

Both methods use the same C++ arbitrary-size rational engine, identical reduced objective and original-value postcheck. Each row includes three fixed seeds and all five regimes. Savings are medians of paired per-instance relative timings, not ratios of unrelated medians. Reduction, process creation, serialization and parsing are charged to both methods. The unreduced pair and all four-cell repeated-query results are retained in the supplement.

At a 513-point lattice, finite differences save a median of approximately 61–73 percent relative to native exhaustive search at the same reduced representation, depending on width and worker. Every instance in those cells is faster. At width 32 and lattice size 33, the median saving is negative on both workers, approximately –18 and –19 percent. Thus the operation-count advantage at large lattices has an observed small-lattice boundary; a root-free exact solver is not uniformly the cheapest implementation. The separate width-eight, 33-point cell is close to the process-overhead floor and is more sensitive to the worker and instance. The complete four-cell repetition table retains these effects rather than attributing a language change or activation simplification to screening alone.

Batching supplies another measured distinction. The complete ninety-query catalogue is evaluated at batch sizes one, eight and thirty-two for each native implementation. Preprocessing, reduction where applicable, serialization, process creation, parsing and exact original-objective checks are all included. Every batch returns the same ninety exact answers. The batch study changes only the boundary of an implementation service; it does not create a new fitted policy or a new economic stopping experiment. Its full timings and byte counts are reported in the supplement, separately from the complete policy-service clocks.

Flatness and cancellation also clarify what a lossless result means. An interval screen is allowed to retain every minimizer and many nonminimizers; it need not produce one survivor. The finite-difference construction instead uses the exact polynomial piece, retaining the first minimizer even on a flat segment. Adaptive subdivision can reach its node budget and call an exact fallback. A fallback preserves correctness but charges its prefix work in addition to the subsequent solver. The full stress tables report all fallback events, endpoint resolutions, recorded operand lengths and interval-array memory. Endpoint coarsening after an exact rational bound is not described as a change in internal floating-point precision.

R61.3.4. *Two controls and an additional innovation*

The multi-action calculation uses saved, genuinely trained two-state ReLU continuations at two specified states. The two controls share capacity $a_1 + a_2 \leq 1/8$.

Their exposure columns are $(1/2, 1/4)$ and $(1/4, 1/2)$, and the optional second independent shock has a different exposure. Quantum denominators 64 and 128 give 45 and 153 feasible lattice points. Setting the second action to zero and disabling the second shock recovers the original action-dependent objective.

TABLE R61.3.V
CONSTRAINED TWO-CONTROL EXACT WITNESS RECOVERY

Cap index	Shocks	Designs	Exhaustive	Adaptive	Max nodes	Fallback	Time ratio
8	1	4	45	7.0	107	0	3.24
8	2	4	45	7.0	107	0	3.19
16	1	4	153	80.0	256	4	2.70
16	2	4	153	80.0	256	4	2.68

Each of four mathematical designs per row uses a saved genuinely fitted ReLU continuation, two state points and two training seeds, repeated on two workers. Shared capacity is $1/8$, with quantum denominator eight times Cap index. Adaptive counts are selected exact evaluations; node-bound computation is additional work. Time ratio is the median paired adaptive/exhaustive query time over the eight records. A ratio above one means slower adaptive execution despite fewer exact evaluations.

These are action-search results, not complete two-control welfare comparisons.

Both exact exhaustive evaluation and verified subdivision return the original constrained minimizer in every recorded query. The analytic uniform-box formula includes all shock dimensions actually used, with signed fitted output weights. Fewer exact point evaluations do not eliminate the cost of computing rectangle bounds: the adaptive/exhaustive time ratio is therefore reported along with evaluations and visited nodes. These experiments validate and cost a non-scalar witness extension. They do not constitute a complete two-control policy-cost or welfare study, and the scalar quartic search theorem is not invoked for a two-shock objective outside its stated degree structure.

R61.3.5. *Attaining an original Bellman-accuracy target*

The direct Bellman calculation retains the original two-state, two-date economy. Lower intervals cover every continuous feasible action, whereas upper actions are

quantum-valued and feasible throughout each acquired cell. Terminal innovation integration is analytic; unresolved earlier expectations use enclosures over the full original innovation bins. The uncentered procedure gives a valid but wide bracket. Centering at the same zero-policy continuation on both sides preserves cancellation before the gap is formed.

TABLE R61.3.VI
ORIGINAL-MODEL WHOLE-DOMAIN BELLMAN ACCURACY

N	A	q	Uncentered gap	Centered gap	Centered work (s)
8	8	4	2.18231	0.62981	0.478
16	16	8	1.12578	0.28126	2.290
32	32	16	0.59109	0.14409	11.288
64	64	32	0.30286	0.07289	80.081
128	64	32	–	0.03979	289.742

The model remains the original two-state, two-date investment economy with its continuous feasible actions and innovation law. N is the number of verification cells per state coordinate. The common-reference lower and upper endpoints cancel the same reference value pointwise. Centered work includes every preceding declared rung and endpoint serialization on that execution. The uncentered gap is the corresponding separately executed numerical bracket, not a same-host timing comparator. Tolerances $1/2$, $1/4$, $1/8$ and $1/16$ are attained; $1/32$ is not.

The centered ladder reaches the all-state tolerances $1/2$, $1/4$, $1/8$ and $1/16$. At $N = 128$, its maximum datewise gap is at most 0.03979, compared with approximately 0.30330 at the final uncentered rung. The centered construction does not attain $1/32$ on its declared ladder. Its cumulative recorded work includes every earlier centered rung and serialized endpoint array. These are constructive bounds against the original continuous-action optimum, not confidence intervals for gains relative to zero investment. The state partition here is a tensor verification grid; the non-tensor prospective actor and its high-dimensional service evidence remain distinct objects.

The bracket by itself certifies the policy constructed by that Bellman calculation. To assess the policies actually returned by the prospective services, we separately

execute Proposition 66. The complete two-date service catalogue has 32 returned records and 5 distinct policy-plus-partition identities. Every identity is included, whether its policy came from a neural, quadratic or common-action generator. The lower table is fixed in advance at $N = 128$, $A = 64$, $q = 32$. The upper recursion is recomputed for each original actor through all intersecting acquired leaves. Neither training nor action selection is repeated.

TABLE R61.3.VII

REVALIDATION OF EVERY DISTINCT ACTUALLY RETURNED TWO-DATE ACTOR

Actor	Leaves	Modes	Aliases	Date 0 gap	All-date gap	Unclipped	Own work (s)
2cc41f05	64	N/S	4	0.07176	0.07176	0.07176	2.254
3f2668fe	64	C/Q	8	0.07176	0.07176	0.07176	2.278
4e5b6e26	32	N/S	4	0.23801	0.30330	0.30330	2.228
718597a7	64	N/S	4	0.06778	0.06778	0.06778	2.279
a06e05ba	32	C/N/Q/S	12	0.23801	0.30330	0.30330	2.223

All thirty-two returned service records are represented by five complete policy-plus-partition identities.

No actor is retrained, improved or replaced. Verification uses $N=128$ and $q=32$ with the deposited centered lower Bellman table. Closed actor ranges include every intersecting leaf. The independent global nonworsening bound is intersected with upper excess endpoints, but the maximum gaps here coincide with the retained unclipped calculation. Own work covers each fixed-actor recursion through its endpoint archive; the shared lower-table reconstruction and process-return overhead are separate audit records.

Three of these returned actors have all-state bounds between approximately 0.06778 and 0.07176, and attain tolerance $1/8$ but not $1/16$. The other two retain bounds around 0.30330. The policy distinction is economically important: a service that meets a ten-percent initial-law improvement target can still return an actor with a materially wider worst-state optimality account. Conversely, the tighter fixed-policy bounds establish original-optimum accuracy for actual deployed actors rather than assigning the Bellman comparator's bound to an unrelated neural policy. The maximum gap is unchanged by the independently justified zero intersection in these results; the unclipped endpoints and gaps are also retained.

The shared lower-table reconstruction, each actor's additional upper recursion and the whole-process receipt are charged separately. These post-selection verification costs are not spliced into a fictitious earlier stopping run, and no old service clock is replaced by a replay timing. Deterministic complete-domain verification adds no independent policy-cost observation.

R61.3.6. *Economic interpretation and computational scope*

For two implementations returning the same policy, the gain is a reduction in computation cost, conditional on the measured execution environment. The setup-and-use formula above gives the full range of implementation charges and use counts supported at any stated price of the work unit. For different policies, a returned-policy gain interval $[l, u]$, a replacement fee τ and an incremental measured-work price $\lambda\Delta w$ give the net-benefit interval

$$[l - \tau - \lambda\Delta w, u - \tau - \lambda\Delta w].$$

A positive lower endpoint supports replacement under that price specification; a negative upper endpoint rules it out; the remaining cases are unresolved. This algebra does not estimate an installation fee, an electricity price or a welfare normalization from a hosted-runner clock.

The experiments specify CPU affinity and record the available processor, compiler and runtime metadata. Processor frequency is not controlled and hardware counters were unavailable under the recorded permissions. Separate worker allocations do not establish independent physical architectures. Rational operation counts, observed operand-size diagnostics, interval-array bytes, process resident memory, cold compilation and complete return times therefore remain separate quantities. The new evidence supports a terminating, costed exact witness and an executable original-optimum certificate within the original NBO development. It does not require suppressing adverse conventional comparisons or inferring an unexecuted scaling result.

APPENDIX: REFERENCES

- DUFFIE, DARRELL AND LARRY G. EPSTEIN (1992): "Stochastic Differential Utility," *Econometrica*, 60 (2), 353–394. [7]
- EPSTEIN, LARRY G. AND STANLEY E. ZIN (1989): "Substitution, Risk Aversion, and the Temporal Behavior of Consumption and Asset Returns: A Theoretical Framework," *Econometrica*, 57 (4), 937–969. [7]
- ERNST, DAMIEN, PIERRE GEURTS, AND LOUIS WEHENKEL (2005): "Tree-Based Batch Mode Reinforcement Learning," *Journal of Machine Learning Research*, 6, 503–556. [7, 133]
- FELZENSZWALB, PEDRO F. AND DANIEL P. HUTTENLOCHER (2012): "Distance Transforms of Sampled Functions," *Theory of Computing*, 8 (19), 415–428. [42, 69, 75]
- HAN, JIEQUN, ARNULF JENTZEN, AND WEINAN E (2018): "Solving High-Dimensional Partial Differential Equations Using Deep Learning," *Proceedings of the National Academy of Sciences*, 115 (34), 8505–8510. [7]
- HOEFFDING, WASSILY (1963): "Probability Inequalities for Sums of Bounded Random Variables," *Journal of the American Statistical Association*, 58 (301), 13–30. [144]
- HORNIK, KURT, MAXWELL STINCHCOMBE, AND HALBERT WHITE (1990): "Universal Approximation of an Unknown Mapping and Its Derivatives Using Multilayer Feedforward Networks," *Neural Networks*, 3 (5), 551–560. [7]
- JUDD, KENNETH L. (1998): *Numerical Methods in Economics*, Cambridge, MA: MIT Press. [7]
- KAKADE, S., AND J. LANGFORD (2002): "Approximately Optimal Approximate Reinforcement Learning," in *Proceedings of the Nineteenth International Conference on Machine Learning*, 267–274. [7]
- KONDA, VIJAY R. AND JOHN N. TSITSIKLIS (1999): "Actor-Critic Algorithms," in *Advances in Neural Information Processing Systems*, MIT Press, vol. 12. [7]
- LAKSHMINARAYANAN, CHANDRASHEKAR AND SHALABH BHATNAGAR (2014): "Approximate Dynamic Programming based on Projection onto the $(\min, +)$ Subsemimodule," arXiv:1403.4175. [69]
- LONGSTAFF, FRANCIS A. AND EDUARDO S. SCHWARTZ (2001): "Valuing American Options by Simulation: A Simple Least-Squares Approach," *Review of Financial Studies*, 14 (1), 113–147. [7]
- MAURER, ANDREAS AND MASSIMILIANO PONTIL (2009): "Empirical Bernstein Bounds and Sample Variance Penalization," in *Proceedings of the 22nd Annual Conference on Learning Theory*, arXiv:0907.3740. [7, 51, 57, 70, 97]
- MC SHANE, EDWARD J. (1934): "Extension of Range of Functions," *Bulletin of the American Mathematical Society*, 40 (12), 837–842. [25, 40, 69, 74]
- METELLI, ALBERTO MARIA, MATTEO PIROTTA, DANIELE CALANDRIELLO, AND MARCELLO RESTELLI (2021): "Safe Policy Iteration: A Monotonically Improving Approximate Policy Iteration Approach," *Journal of Machine Learning Research*, 22 (97), 1–83. [78]
- MOORE, RAMON E., R. BAKER KEARFOTT, AND MICHAEL J. CLOUD (2009): *Introduction to Interval Analysis*. Philadelphia: Society for Industrial and Applied Mathematics. [178]

- 1 MUNOS, RÉMI, AND CSABA SZEPESVÁRI (2008): “Finite-Time Bounds for Fitted Value Iteration,” *Journal* 1
2 *of Machine Learning Research*, 9 (27), 815–857. [133] 2
- 3 MUROTA, KAZUO (2003): *Discrete Convex Analysis*. Philadelphia: Society for Industrial and Applied 3
4 Mathematics. [178] 4
- 5 PIROTTA, MATTEO, MARCELLO RESTELLI, ALESSIO PECORINO, AND DANIELE CALANDRIELLO (2013): 5
6 “Safe Policy Iteration,” in *Proceedings of the 30th International Conference on Machine Learning*, vol. 28 of 6
7 *Proceedings of Machine Learning Research*, 307–315. [7] 7
- 8 POWELL, WARREN B. (2011): *Approximate Dynamic Programming: Solving the Curses of Dimensionality*, 8
9 Hoboken, NJ: Wiley, 2 ed. [7] 9
- 10 SCHERRER, BRUNO, MOHAMMAD GHAVAMZADEH, VICTOR GABILLON, BORIS LESNER, AND 10
11 MATTHIEU GEIST (2015): “Approximate Modified Policy Iteration and its Application to the Game of 11
12 Tetris,” *Journal of Machine Learning Research*, 16 (49), 1629–1676. [7] 12
- 13 SIRIGNANO, JUSTIN AND KONSTANTINOS SPILIOPOULOS (2018): “DGM: A Deep Learning Algorithm 13
14 for Solving Partial Differential Equations,” *Journal of Computational Physics*, 375, 1339–1364. [7] 14
- 15 WANG, YANG, BRENDAN O’DONOGHUE, AND STEPHEN BOYD (2015): “Approximate Dynamic Pro- 15
16 gramming via Iterated Bellman Inequalities,” *International Journal of Robust and Nonlinear Control*, 25 16
17 (10), 1472–1496. [69] 17
18 18
19 19
20 20
21 21
22 22
23 23
24 24
25 25
26 26
27 27
28 28
29 29
30 30
31 31
32 32